Advanced Metaprogramming in Classic C++

Davide Di Gennaro

Advanced Metaprogramming in Classic C++

Davide Di Gennaro

apress[®]

Advanced Metaprogramming in Classic C++

Copyright © 2015 by Davide Di Gennaro

This work is subject to copyright. All rights are reserved by the Publisher, whether the whole or part of the material is concerned, specifically the rights of translation, reprinting, reuse of illustrations, recitation, broadcasting, reproduction on microfilms or in any other physical way, and transmission or information storage and retrieval, electronic adaptation, computer software, or by similar or dissimilar methodology now known or hereafter developed. Exempted from this legal reservation are brief excerpts in connection with reviews or scholarly analysis or material supplied specifically for the purpose of being entered and executed on a computer system, for exclusive use by the purchaser of the work. Duplication of this publication or parts thereof is permitted only under the provisions of the Copyright Law of the Publisher's location, in its current version, and permission for use must always be obtained from Springer. Permissions for use may be obtained through RightsLink at the Copyright Clearance Center. Violations are liable to prosecution under the respective Copyright Law.

ISBN-13 (pbk): 978-1-4842-1011-6

ISBN-13 (electronic): 978-1-4842-1010-9

Trademarked names, logos, and images may appear in this book. Rather than use a trademark symbol with every occurrence of a trademarked name, logo, or image we use the names, logos, and images only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

The use in this publication of trade names, trademarks, service marks, and similar terms, even if they are not identified as such, is not to be taken as an expression of opinion as to whether or not they are subject to proprietary rights.

While the advice and information in this book are believed to be true and accurate at the date of publication, neither the authors nor the editors nor the publisher can accept any legal responsibility for any errors or omissions that may be made. The publisher makes no warranty, express or implied, with respect to the material contained herein.

Managing Director: Welmoed Spahr

Lead Editor: Steve Anglin

Technical Reviewer: Sverrir Sigmundarson

Editorial Board: Steve Anglin, Louise Corrigan, Jonathan Gennick, Robert Hutchinson, Michelle Lowman, James Markham, Susan McDermott, Matthew Moodie, Jeffrey Pepper, Douglas Pundick, Ben Renow-Clarke, Gwenan

Spearing, Steve Weiss

Coordinating Editor: Mark Powers Copy Editor: Kezia Endsley Compositor: SPi Global Indexer: SPi Global

Artist: SPi Global

Cover Designer: Anna Ishchenko

Distributed to the book trade worldwide by Springer Science+Business Media New York, 233 Spring Street, 6th Floor, New York, NY 10013. Phone 1-800-SPRINGER, fax (201) 348-4505, e-mail orders-ny@springer-sbm.com, or visit www.springeronline.com. Apress Media, LLC is a California LLC and the sole member (owner) is Springer Science + Business Media Finance Inc (SSBM Finance Inc). SSBM Finance Inc is a Delaware corporation.

For information on translations, please e-mail rights@apress.com, or visit www.apress.com.

Apress and friends of ED books may be purchased in bulk for academic, corporate, or promotional use. eBook versions and licenses are also available for most titles. For more information, reference our Special Bulk Sales—eBook Licensing web page at www.apress.com/bulk-sales.

Any source code or other supplementary material referenced by the author in this text is available to readers at www.apress.com/9781484210116. For detailed information about how to locate your book's source code, go to www.apress.com/source-code/.

Template metaprogramming and expression templates are not techniques for novice programmers, but an advanced practitioner can use them to good effect.

Technical Report on C++ Performance, ISO/IEC TR 18015:2006(E)

Nothing described in this report involves magic.

Technical Report on C++ Performance, ISO/IEC TR 18015:2006(E)

People should not be asked to do things "just because we say so". At least we can try to explain the reasons behind the rules.

An Interview with Bjarne Stroustrup - Dr. Dobb's Journal

I hope Tibet will find this book as interesting as all the others he reads



Contents at a Glance

About the Author

About the Technical Reviewer

Acknowledgments

Preface

Chapter 1: Templates

Chapter 2: Small Object Toolkit

■ #include <techniques>

Chapter 3: Static Programming

Chapter 4: Overload Resolution

Chapter 5: Interfaces

Chapter 6: Algorithms

Chapter 7: Code Generators

Chapter 8: Functors

Chapter 9: The Opaque Type Principle

■ #include <applications>

Chapter 10: Refactoring

Chapter 11: Debugging Templates

Chapter 12: C++0x

Appendix A: Exercises

Appendix B: Bibliography

Index

Contents

About the Author

About the Technical Reviewer

Acknowledgments

Preface

Chapter 1: Templates

1.1. C++ Templates

- 1.1.1. Typename
- 1.1.2. Angle Brackets
- 1.1.3. Universal Constructors
- 1.1.4. Function Types and Function Pointers
- 1.1.5. Non-Template Base Classes
- 1.1.6. Template Position

1.2. Specialization and Argument Deduction

- 1.2.1. Deduction
- 1.2.2. Specializations
- 1.2.3. Inner Class Templates

1.3. Style Conventions

- 1.3.1. Comments
- 1.3.2. Macros
- 1.3.3. Symbols
- 1.3.4. Generality
- 1.3.5. Template Parameters
- 1.3.6. Metafunctions
- 1.3.7. Namespaces and Using Declarations

1.4. Classic Patterns

- 1.4.1. size_t and ptrdiff_t
- 1.4.2. void T::swap(T&)
- 1.4.3. bool T::empty() const; void T::clear()
- 1.4.4. X T::get() const; X T::base() const
- 1.4.5. X T::property() const; void T::property(X)
- 1.4.6. Action(Value); Action(Range)

1.4.7. Manipulators
1.4.8. Position of Operators
1.4.9. Secret Inheritance
1.4.10. Literal Zero

1 4 11 D 1 ...

1.4.11. Boolean Type

1.4.12. Default and Value Initialization

1.5. Code Safety

1.6. Compiler Assumptions

1.6.1. Inline

1.6.2. Error Messages

1.6.3. Miscellaneous Tips

1.7. Preprocessor

1.7.1. Include Guards

1.7.2. Macro Expansion Rules

Chapter 2: Small Object Toolkit

2.1. Hollow Types

2.1.1. instance_of

2.1.2. Selector

2.1.3. Static Value

2.1.4. Size of Constraints

2.2. Static Assertions

2.2.1. Boolean Assertions

2.2.2. Assert Legal

2.2.3. Assertions with Overloaded Operators

2.2.4. Modeling Concepts with Function Pointers

2.2.5. Not Implemented

2.3. Tagging Techniques

2.3.1. Type Tags

2.3.2. Tagging with Functions

2.3.3. Tag Iteration

2.3.4. Tags and Inheritance

■ #include <techniques>

Chapter 3: Static Programming

- 3.1. Static Programming with the Preprocessor
- 3.2. Compilation Complexity
- 3.3. Classic Metaprogramming Idioms

3.4. Hidden Template Parameters

- 3.4.1. Static Recursion on Hidden Parameters
- 3.4.2. Accessing the Primary Template
- 3.4.3. Disambiguation

3.5. Traits

- 3.5.1. Type Traits
- 3.5.2. Type Dismantling

3.6. Type Containers

- 3.6.1. typeat
- 3.6.2. Returning an Error
- 3.6.3. Depth
- 3.6.4. Front and Back
- 3.6.5. Find
- 3.6.6. Push and Pop
- 3.6.7. More on Template Rotation
- 3.6.8. Agglomerates
- 3.6.9. Conversions
- 3.6.10. Metafunctors

3.7. A Summary of Styles

Chapter 4: Overload Resolution

4.1. Groups

- 4.1.1. From Overload to Groups
- 4.1.2. Runtime Decay

4.2. More Traits

- 4.2.1. A Function Set for Strings
- 4.2.2. Concept Traits
- 4.2.3. Platform-Specific Traits
- 4.2.4. Merging Traits

4.3. SFINAE

- 4.3.1. SFINAE Metafunctions
- 4.3.2. Multiple Decisions
- 4.3.3. Only If
- 4.3.4. SFINAE and Returned Functors
- 4.3.5. SFINAE and Software Updates
- 4.3.6. Limitations and Workarounds
- 4.3.7. SFINAE with Partial Specializations

4.4. Other Classic Metafunctions with Sizeof

4.5. Overload on Function Pointers

- 4.5.1. Erase
- 4.5.2. Swap
- 4.5.2. Argument Dominance

Chapter 5: Interfaces

5.1. Wrapping References

5.2. Static Interfaces

- 5.2.1. Static Interfaces
- 5.2.2. Common Errors
- 5.2.3. A Static_Interface Implementation
- 5.2.4. The Memberspace Problem
- 5.2.5. Member Selection

5.3. Type Hiding

- 5.3.1. Trampolines
- 5.3.2. Typeinfo Wrapper
- 5.3.3. Option_Map
- 5.3.4. Option_Parser
- 5.3.5. Final Additions
- 5.3.6. Boundary Crossing with Trampolines

5.4. Variant

- 5.4.1. Parameter Deletion with Virtual Calls
- 5.4.2. Variant with Visitors

5.5. Wrapping Containers

Chapter 6: Algorithms

6.1. Algorithm I/O

- 6.1.1. Swap-Based or Copy-Based
- 6.1.2. Classification of Algorithms
- 6.1.3. Iterator Requirements
- 6.1.4. An Example: Set Partitioning
- 6.1.5. Identifying Iterators
- 6.1.6. Selection by Iterator Value Type

6.2. Generalizations

- 6.2.1. Properties and Accessors
- 6.2.2. Mimesis
- 6.2.3. End of Range

6.3. Iterator Wrapping

- 6.3.1. Iterator Expander
- 6.3.2. Fake Pairs

6.6. The Barton-Nackman Trick **Chapter 7: Code Generators** 7.1. Static Code Generators 7.2. Double checked Stop 7.3. Static and Dynamic Hashing 7.3.1. A Function Set for Characters 7.3.2. Changing Case 7.3.3. Mimesis Techniques 7.3.4. Ambiguous Overloads 7.3.5. Algorithm I/O 7.3.6. Mimesis Interface 7.4. Nth Minimum 7.5. The Template Factory Pattern 7.6. Automatic Enumeration of Types 7.7. If-Less Code 7.7.1. Smart Constants 7.7.2. Converting Enum to String 7.7.3. Self-Modifying Function Tables **Chapter 8: Functors** 8.1. Strong and Weak Functors

8.2. Functor Composition Tools

8.3.1. Conversion of Functions to Functors8.3.2. Conversion of Members to Functors

8.3.3. More on the Double Wrapper Technique

8.7. Forwarding and Reference Wrappers

8.3. Inner Template Functors

8.4.1. A Step-by-Step Implementation

8.4. Accumulation

8.5. Drivers

8.6. Algors

6.4. Receipts

6.5.1. Less and NaN

6.5. Algebraic Requirements

Chapter 9: The Opaque Type Principle

9.1. Polymorphic Results

9.2. Classic Lambda Expressions

- 9.2.1. Elementary Lambda Objects
- 9.2.2. Lambda Functions and Operators
- 9.2.3. Refinements
- 9.2.4. Argument and Result Deduction
- 9.2.5. Deducing Argument Type
- 9.2.6. Deducing Result Type
- 9.2.7. Static Cast
- 9.2.8. Arrays

9.3. Creative Syntax

9.3.1. Argument Chains with () and []

9.4. The Growing Object Concept

- 9.4.1. String Concatenation
- 9.4.2. Mutable Growing Objects
- 9.4.3. More Growing Objects
- 9.4.4. Chain Destruction
- 9.4.5. Variations of the Growing Object

9.5. Streams

- 9.5.1. Custom Manipulators and Stream Insertion
- 9.5.2. Range Insertion with a Growing Object

9.6. Comma Chains

9.7. Simulating an Infix

■ #include <applications>

Chapter 10: Refactoring

10.1. Backward Compatibility

10.2. Refactoring Strategies

- 10.2.1. Refactoring with Interfaces
- 10.2.2. Refactoring with Trampolines
- 10.2.3. Refactoring with Accessors

10.3. Placeholders

- 10.3.1. Switch-Off
- 10.3.2. The Ghost

Chapter 11: Debugging Templates

11.1. Identify Types 11.1.1. Trapping Types 11.1.2. Incomplete Types 11.1.3. Tag Global Variables

11.2. Integer Computing

- 11.2.1. Signed and Unsigned Types
- 11.2.2. References to Numeric Constants

11.3. Common Workarounds

- 11.3.1. Debugging SFINAE
- 11.3.2. Trampolines
- 11.3.3. Compiler Bugs

Chapter 12: C++0x

- 12.1. Type Traits
- 12.2. Decltype
- 12.3. Auto
- 12.4. Lambdas
- 12.5. Initializers
- 12.6. Template Typedefs

12.7. Extern Template

- 12.7.1. Linking Templates
- 12.7.2. Extern Template

12.9. Variadic Templates

Appendix A: Exercises

A.1. Exercises

- A.1.1. Extension
- A.1.2. Integer
- A.1.3. Date Format
- A.1.4. Specialization
- A.1.5. Bit Counting
- A.1.6. Prime Numbers
- A.1.7. Typeinfo without RTTI
- A.1.8. Hints and Partial Solutions

Appendix B: Bibliography

Index

About the Author

I'm like a dog, when you whistle: yep, yep. Template? Good, good...

—Andrei Alexandrescu, build 2012

Davide loves to introduce himself as a mathematician, but a better definition would be a philosopher. After studying history of art and functional analysis, he switched to algorithm design and C++. He has been showing the marvels of metaprogramming techniques since the late 90s. As nobody could really understand him, he was eventually nicknamed "the professor". He works for big companies, where his real identity is ignored, and he spends his free time as a photographer.

Someone said that, "he makes the impossible possible".



Tibet was born on September, 6th 1998, just close to the C++ standard. He immediately showed an unusual intelligence, learning more than 100 keywords in our natural language.

Active and proud of his C++-related work, during 2014 his health started to decline. Readers often ask when he will write another book. He knows, but he simply smiles.

About the Technical Reviewer



Sverrir Sigmundarson has over 15 years of industry experience building high performance, mission-critical software for the finance and software industries. He holds an MSc degree in Computer Science from Reykjavik University in Iceland. He is currently on a special assignment as a stay-at-home-dad living in Strasbourg, France with his wife and son. He can be contacted through his website coruscantconsulting.co.uk or via linkedin.com/in/sverrirs.

Acknowledgments

Learning C++ is a process that never ends. As someone wrote, C++ is the only language whose features get discovered as though they were unexplored lands.

While a book may be the work of a single person, discovery always comes from teamwork.

The author would like to thank all the teams that made possible his journey through C++. They all had something to teach, and their contributions—direct or indirect—led to this book. His family, Carla, Alberto, Tibet and Asia; the people at Logikos, especially Max; the Natam core team, Alberto L., Alberto T., Bibo, Fabio, Graziano, Marco, Roberto, Rocco; the friends at Brainpower, in particular Alberto, Andrea, Davide, Fabio, Giacomo, Giancarlo, Luca, Marco D., Marco M., Matteo, Paolo, Pino, Vincenzo; and all the others.

I would like thank the many people at Apress who worked on the book, including Steve Anglin who talked me into it, Mark Powers for managing the project, Sverrir Sigmundarson for a fine job as technical reviewer, Jeff Pepper and Kezia Endsley for clarifying the words.

Many Googlers kindly reviewed the first chapters of the draft and provided suggestions, fixes, constructive criticism, exercises, or simply appreciation.

A very special thank goes to Attilio Meucci, who proved that writing a book is not impossible, and it's always worthwhile.

Preface

Template Metaprogramming (TMP from here on) is a new way of using C++:

- It has a *scope*: a known set of problems where it proves useful.
- It has a *philosophy*: a peculiar way of thinking about problems.
- It has a *language*: idioms and patterns.

This book, according to the 80-20 law, aims to be an introduction to the first 20% of metaprogramming—its philosophy, scope, and language—that can improve 80% of daily programming activities. All the chapters are driven by some simple ideas:

- With modern compilers, most practical benefits come from simple techniques, when correctly applied.
- TMP indeed produces better software. "Better" is simply a placeholder for faster, safer, more maintainable, more expressive, or a combination of these features.
- State-of-the-art TMP libraries usually offer a huge set of features. Unfortunately, documentation is either too large or too small. While reuse is a long-term winning strategy, mastering the basic principles may suffice.
- Getting gradually accustomed with elementary techniques, the reader will develop a deeper comprehension of the problems and eventually, if necessary, look for more advanced tools.

The reader is assumed at ease with classic C++ programming, including STL concepts and conventions.

A systematic study of TMP exceeds the capacity (in C++ sense) of any single book. With over five years invested in creating this book, I hope you will find it more than a useful starting point. For comprehensive and robust training, the interested reader may want to see the bibliography.

Source Code

This book is not focused on results, but on the path—the steps and motivations that lead to a project's implementation. Many examples derive from production code. However, in a book, problems must look as easy and evident as possible, sometimes even more. In practice, they are never this way.

So for illustration purposes, the source code is unquestionably sub-optimal and oversimplified. Oversimplification means partial or full omission of implementation details, special cases, namespaces, system headers, compiler bugs, and so on. The most advanced programming technique is

hardly an advantage if it crashes the company's official compiler.

In short, these details are important, as they make the difference between a curious prototype and a useful implementation.

In addition, code has been streamlined to satisfy visual constraints. In particular, indentation is systematically inconsistent, some function bodies have been removed, names may be shorter than necessary, and macros have been introduced for the sole purpose of shortening the text.

Readers are asked to be patient and review the Errata section that follows.

Finally, I admit that results are rarely supported with experimental data. TMP techniques give a compiler the opportunity to create optimized code, and as a rule, this book doesn't verify that it is indeed the case.

Classic and Modern C++

The C++ standard is being updated with lots of new features. The first edition of the document in 1998 had fewer than 800 pages. A 200-page technical report was published in 2003 and revised in 2006. In March 2010, the committee released the FCD, a milestone draft more than 1,300 pages long. In August 2014, the vote to approve the C++14 standard was completed. Some of the new language additions have already been implemented in compilers.

This book deals with a very small part of "C++0x" (yes, I use the familiar nickname of the new standard) and "C++14". More precisely, it discusses what has a serious impact on TMP code and is also available in the major compilers. The focus of the book remains on classic C++, which can be utilized in any implementation of C++. The so-called "modern C++" constituting the revisions incorporated in C++11 and C++14 is the topic of discussion in Chapter 12 and is referenced accordingly in other parts of this book.

Book Structure

The book is divided into three sections, and chapters are designed to be read in order. Each chapter starts with its own rationale, or a summary of the motivations for previous arguments.

The first section deals with the basics, and in particular Chapter 2 is a prerequisite for most of the source code contained in the book. Chapter 2 contains a description of the basic class templates that will be constantly and silently reused without further comments.

The second part of the book develops some techniques for writing software, in the approximate order of increasing complexity.

The third part contains some practical advice for real-world issues, so it has been pretentiously labeled "applications".

I refer to some compilers with abbreviations, followed by a version number: MSVC for Microsoft Visual C++ and GCC for GNU G++.

From time to time, I show the output of some compiler, without mentioning explicitly which one, to emphasize what a "generic" compiler would emit.

This is a note. The following text contains a sample of the typographic conventions used in this book.

```
// filename.cpp
this->is(source*code);
```

This is the resulting compiler output.

The same format denotes an algorithm description in pseudo-code.

Odd for a book that emphasizes readability, fragments of source code have no syntax highlighting, so they will look scarier than they actually are.

Errata

Readers are encouraged to send their feedback to the book's page on Apress.com (www.apress.com/9781484210116).

Errata are published regularly on http://acppmp.blogspot.com.

Note to the Third Revision

This book was born in 2005, when C++11 was yet to come, and finished just before the new standard was published. On purpose, most of the new techniques on the way were ignored, simply because they were not widely available, not finalized, or just not completely understood. None of the revisions of this book changed this view, which is still essentially correct. So, while vendors are still releasing C++11 compilers, no herculean attempt was made to upgrade the book contents.

Nonetheless, this should not be considered a limitation in any way. Starting TMP at a low level and with simpler language tools means that your code will run on existing compilers, and is a powerful educational experience, and it will lead to a stronger appreciation of all the "syntactic sugar" that modern C++ offers.

PART 1

#include #include <techniques>
 #include <applications>

CHAPTER 1

Templates

```
"C++ supports a variety of styles."
```

Bjarne Stroustrup, A Perspective on ISO C++

Programming is the process of teaching something to a computer by talking to the machine in one of its common languages. The closer to the machine idiom you go, the less natural the words become.

Each language carries its own expressive power. For any given concept, there is a language where its description is simpler, more concise, and more detailed. In assembler, we have to give an extremely rich and precise description for any (possibly simple) algorithm, and this makes it very hard to read back. On the other hand, the beauty of C++ is that, while being close enough to the machine language, the language carries enough instruments to enrich itself.

C++ allows programmers to express the same concept with different *styles* and good C++ looks more natural.

First you are going to see the connection between the templates and the style, and then you will dig into the details of the C++ template system.

Given this C++ fragment:

return x*x;

```
double x = sq(3.14);
  Can you guess what sq is? It could be a macro:
#define sq(x) ((x)*(x))
  A function:
double sq(double x)
{
  return x*x;
}
  A function template:
template <typename scalar t>
```

inline scalar_t sq(const scalar_t& x)

```
A type (an unnamed instance of a class that decays to a double):
class sq
   double s ;
public:
   sq(double x)
   : s (x*x)
   { }
   operator double() const
   { return s ; }
};
   A global object:
class sq t
public:
   typedef double value type;
   value type operator()(double x) const
      return x*x;
};
const sq t sq = sq t();
   Regardless of how sq(3.14) is implemented, most humans can guess what sq(3.14) does
for example, passing a square to a function template will trigger an unexpected argument deduction:
```

}

just looking at it. However, visual equivalence does not imply interchangeableness. If sq is a class,

```
template <typename T> void f(T x);
f(cos(3.14)); // instantiates f<double>
f(sq(3.14)); // instantiates f<sq>. counterintuitive?
```

Furthermore, you would expect every possible numeric type to be squared as efficiently as possible, but different implementations may perform differently in different situations:

```
std::vector<double> v;
std::transform(v.begin(), v.end(), v.begin(), sq);
```

If you need to transform a sequence, most compilers will get a performance boost from the last implementation of sq (and an error if sq is a macro).

The purpose of TMP is to write code that is:

- Visually clear to human users so that nobody needs to look underneath.
- Efficient in most/all situations from the point of view of the compiler.
- Self-adapting to the rest of the program.¹

Self-adapting means "portable" (independent of any particular compiler) and "not imposing constraints". An implementation of sq that requires its argument to derive from some abstract base class would not qualify as self-adapting.

The true power of C++ templates is *style*. Compare the following equivalent lines:

```
double x1 = (-b + sqrt(b*b-4*a*c))/(2*a);
double x2 = (-b + sqrt(sq(b)-4*a*c))/(2*a);
```

All template argument computations and deductions are performed at compile time, so they impose no runtime overhead. If the function sq is properly written, line 2 is at least as efficient as line 1 and easier to read at the same time.

Using sq is elegant:

- It makes code readable or self-evident
- It carries no speed penalty
- It leaves the program open to future optimizations

In fact, after the concept of squaring has been isolated from plain multiplication, you can easily plug in specializations:

```
template <typename scalar_t>
inline scalar_t sq(const scalar_t& x)
{
   return x*x;
}

template <>
inline double sq(const double& x)
{
   // here, use any special algorithm you have!
}
```

1.1. C++ Templates

The classic C++ language admits two basic types of templates—*function templates* and *class templates*²:

Here is a function template:

// ...

When you supply suitable values to all its parameters, a template generates entities during compilation. A function template will produce functions and a class template will produce classes. The most important ideas from the TMP viewpoint can be summarized as follows:

- You can exploit class templates to perform computations at compile time.
- Function templates can auto-deduce their parameters from arguments. If you call sq(3.14), the compiler will automatically figure out that scalar_t is double, generate the function sq<double>, and insert it at the call site.

Both kinds of template entities start declaring a *parameter list* in angle brackets. Parameters can include *types* (declared with the keyword typename or class) and non-types: integers and pointers.³

Note that, when the parameter list is long or when you simply want to comment each parameter separately, you may want to indent it as if it were a block of code within curly brackets.

Parameters can in fact have a default value:

A template can be seen as a metafunction that maps a tuple of parameters to a function or a class. For example, the sq template

```
template <typename scalar_t>
scalar_t sq(const scalar_t& x);
```

maps a type T to a function:

```
T \rightarrow T (*) (const T_{\&})
```

In other words, sq<double> is a function with signature double (*) (const double&). Note that double is the value of the parameter scalar t.

Conversely, the class template

```
template <typename char_t = char>
class basic_string;
```

maps a type T to a class:

```
T \rightarrow basic\_string<T>
```

With classes, *explicit specialization* can limit the domain of the metafunction. You have a general template and then some specializations; each of these may or may not have a body.

char_t and scalar_t are called *template parameters*. When basic_string<char> and sq<double> are used, char and double are called *template arguments*, even if there may be some confusion between double (the template argument of sq) and x (the argument of the function sq<double>).

When you supply template arguments (both types and non-types) to the template, seen as a metafunction, the template is *instantiated*, so if necessary the compiler produces machine code for the entity that the template produces.

Note that different arguments yield different instances, even when instances themselves are identical: sq<double> and sq<const double> are two unrelated functions.⁴

When using function templates, the compiler will usually figure out the parameters. We say that an argument *binds* to a template parameter.

```
template <typename scalar_t>
scalar_t sq(const scalar_t& x) { return x*x; }
```

```
double pi = 3.14; sq(pi); // the compiler "binds" double to scalar_t  

double x = sq(3.14); // ok: the compiler deduces that scalar_t is double double x = sq<double>(3.14); // this is legal, but less than ideal
```

All template arguments must be compile-time constants.

- Type parameters will accept everything known to be a type.
- Non-type parameters work according to the most automatic casting/promotion rule.⁵

Here are some typical errors:

The best syntax for a compile-time constant in classic C++ is static const [[integer type]] name = value;

The static prefix could be omitted if the constant is local, in the body of a function, as shown previously. However, it's both harmless and clear (you can find all the compile-time constants in a project by searching for "static const" rather than "const" alone).

The arguments passed to the template can be the result of a (compile-time) computation. Every valid integer operation can be evaluated on compile-time constants:

- Division by zero causes a compiler error.
- Function calls are forbidden.⁷
- Code that produces an intermediate object of non-integer/non-pointer type is non-

```
portable, except when inside sizeof: (int) (N*1.2), which is illegal.
Instead use (N+N/5). static_cast<void*>(0) is fine too.
SomeClass<(27+56*5) % 4> s1;
```

Division by zero will cause a compiler error only if the computation is entirely static. To see the difference, note that this program compiles (but it won't run).

SomeClass<sizeof(void*)*CHAR BIT> s1;

On the other hand, compare the preceding listing with the following two, where the division by zero happens during compilation (in two different contexts):

And with:

```
tricky<0/0> t;
test.cpp(12) : error C2975: 'N' : invalid template argument for
'tricky',
```

expected compile-time constant expression

More precisely, compile-time constants can be:

- Integer literals, for example, 27, CHAR BIT, and 0x05
- sizeof and similar non-standard language operators with an integer result (for example, alignof where present)
- Non-type template parameters (in the context of an "outer" template)

```
template <int N>
class AnotherClass
{
   SomeClass<N> myMember_;
};
```

• Static constants of integer type

```
template <int N, int K>
struct MyTemplate
{
   static const int PRODUCT = N*K;
};
SomeClass< MyTemplate<10,12>::PRODUCT > s1;
```

• Some standard macros, such as __LINE__ (There is actually some degree of freedom; as a rule they are constants with type long, except in implementation-dependent "edit and continue" debug builds, where the compiler must use references. In this case, using the macro will cause a compilation error.)⁹

```
SomeClass< LINE > s1; // usually works...
```

A parameter can depend on a previous parameter:

template

```
class Y
};
Y<int, 7> y1; // fine
Y<double, 3> y2; // error: the constant '3' cannot have type
'double'
  Classes (and class templates) may also have template member functions:
// normal class with template member function
struct mathematics
  template <typename scalar t>
  scalar t sq(scalar t x) const
     return x*x;
};
// class template with template member function
template <typename scalar t>
struct more mathematics
  template \langle typename other t \rangle^{10}
  static scalar t product(scalar t x, other t y)
     return x*y;
};
double A = mathematics().sq(3.14);
double B = more mathematics<double>().product(3.14, 5);
```

1.1.1. Typename

best

The keyword typename is used:

- As a synonym of class, when declaring a type template parameter
- Whenever it's not evident to the compiler that an identifier is a type name

For an example of "not evident" think about MyClass<T>:: Y in the following fragment:

```
template <typename T>
struct MyClass
  typedef double Y;
                                             // Y may or may not be
a type
  typedef T Type;
                                             // Type is always a type
};
template < >
struct MyClass<int>
  static const int Y = 314;
                                             // Y may or may not be
a type
  typedef int Type;
                                             // Type is always a type
};
int Q = 8;
template <typename T>
void SomeFunc()
 MyClass<T>::Y * Q; // what is this line? it may be:
                    // the declaration of local pointer-to-double
named Q;
                    // or the product of the constant 314, times the
global variable Q
};
  Y is a dependent name, since its meaning depends on T, which is an unknown parameter.
  Everything that depends directly or indirectly on unknown template parameters is a dependent
name. If a dependent name refers to a type, then it must be introduced with the typename keyword.
template <typename X>
class AnotherClass
  MyClass<X>::Type t1 ;
                                       // error: 'Type' is a dependent
name
  typename MyClass<X>::Type t2 ; // ok
  MyClass<double>::Type t3; // ok: 'Type' is independent of
Χ
};
  Note that typename is required in the first case and forbidden in the last:
template <typename X>
class AnotherClass
```

// ok, but it won't

typename MyClass<X>::Y member1 ;

```
typename may introduce a dependent type when declaring a non-type template parameter:
template <typename T, typename T::type N>
struct SomeClass
{
};
struct S1
{
   typedef int type;
};
SomeClass<S1, 3> x; // ok: N=3 has type 'int'
```

typename MyClass<double>::Y member2; // error

As a curiosity, the classic C++ standard specifies that if the syntax typename T1::T2 yields a non-type during instantiation, then the program is ill-formed. However, it doesn't specify the converse: if T1::T2 has a valid meaning as a non-type, then it could be re-interpreted later as a type, if necessary. For example:

```
template <typename T>
struct B
{
   static const int N = sizeof(A<T>::X);
   // should be: sizeof(typename A...)
};
```

Until instantiation, B "thinks" it's going to call sizeof on a non-type; in particular, sizeof is a valid operator on non-types, so the code is legal. However, X could later resolve to a type, and the code would be legal anyway:

```
template <typename T>
struct A
{
   static const int X = 7;
};

template <>
struct A<char>
{
   typedef double X;
};
```

compile if X is 'int'.

Although the intent of typename is to forbid all such ambiguities, it may not cover all corner cases. 11

1.1.2. Angle Brackets

Even if all parameters have a default value, you cannot entirely omit the angle brackets:

Template parameters may carry different meanings:

- Sometimes they are really meant to be generic, for example, std::vector<T> or std::set<T>. There may be some conceptual assumptions about T—say constructible, comparable...—that do not compromise the generality.
- Sometimes parameters are assumed to belong to a fixed set. In this case, the class template is simply the common implementation for two or more similar classes. 12

In the latter case, you may want to provide a set of regular classes that are used without angle brackets, so you can either derive them from a template base or just use typedef¹³:

```
template <typename char_t = char>
class basic_string
{
    // this code compiles only when char_t is either 'char' or
'wchar_t'
    // ...
};
class my_string : public basic_string<>
{
    // empty or minimal body
    // note: no virtual destructor!
};
typedef basic string<wchar t> your string;
```

A popular compiler extension (officially part of C++0x) is that two or more adjacent "close angle brackets" will be parsed as "end of template," not as an "extraction operator". Anyway, with older compilers, it's good practice to add extra spaces:

1.1.3. Universal Constructors

A template copy constructor and an assignment are not called when dealing with two objects of the very same kind:

```
template <typename T>
class something
public:
  // not called when S == T
  template <typename S>
  something(const something<S>& that)
  }
  // not called when S == T
  template <typename S>
  something& operator=(const something<S>& that)
     return *this;
};
  something<int> s0;
  something <double > s1, s2;
  s0 = s1; // calls user defined operator=
  s1 = s2;
              // calls the compiler generated assignment
```

The user-defined template members are sometimes called *universal copy constructors* and *universal assignments*. Note that universal operators take something<X>, not X.

The C++ Standard 12.8 says:

- "Because a template constructor is never a copy constructor, the presence of such a template does not suppress the implicit declaration of a copy constructor."
- "Template constructors participate in overload resolution with other constructors, including copy constructors, and a template constructor may be used to copy an object if it provides a better match than other constructors."

In fact, having very generic template operators in base classes can introduce bugs, as this example shows:

```
struct base
```

```
template <typename T>
  base(T x) {}

;;

struct derived : base
{
  derived() {}

  derived(const derived& that)
    : base(that) {}
};

  derived d1;
  derived d2 = d1;
```

The assignment d2 = d1 causes a stack overflow.

An implicit copy constructor must invoke the copy constructor of the base class, so by 12.8 above it can never call the universal constructor. Had the compiler generated a copy constructor for derived, it would have called the base copy constructor (which is implicit). Unfortunately, a copy constructor for derived is given, and it contains an explicit function call, namely base (that). Hence, following the usual overload resolution rules, it matches the universal constructor with T=derived. Since this function takes x by value, it needs to perform a copy of that, and hence the call is recursive. ¹⁴

1.1.4. Function Types and Function Pointers

Mind the difference between a function type and a pointer-to-function type:

```
template <double F(int)>
struct A
{
};

template <double (*F)(int)>
struct B
{
};

They are mostly equivalent:

double f(int)
{
   return 3.14;
```

```
A<f> t1; // ok
B<f> t2; // ok
```

Usually a function decays to a function pointer exactly as an array decays to a pointer. But a function type cannot be constructed, so it will cause failures in code that look harmless:

This problem is mostly evident in functions that return a functor (the reader can think about std::not1 or see Section **4.3.4**). In C++, function templates that get parameters by reference prevent the decay:

For what concerns pointers, function templates with explicit parameters behave like ordinary functions:

```
double f(double x)
{
   return x+1;
}

template <typename T>
T g(T x)
{
   return x+1;
}

typedef double (*FUNC_T) (double);

FUNC_T f1 = f;
FUNC_T f2 = g<double>;
```

However, if they are members of class templates and their context depends on a yet unspecified parameter, they require an extra template keyword before their name¹⁵:

```
template <typename X>
struct outer
{
   template <typename T>
   static T g(T x)
   {
      return x+1;
   }
};

template <typename X>
void do_it()
{
   FUNC_T f1 = outer<X>::g<double>;
   FUNC_T f2 = outer<X>::template g<double>;
   // correct
}
```

Both typename and template are required for inner template classes:

```
template <typename X>
struct outer
{
   template <typename T>
   struct inner {};
};

template <typename X>
```

```
void do_it()
{
   typename outer<X>::template inner<double> I;
}
```

Some compilers are not rigorous at this.

1.1.5. Non-Template Base Classes

If a class template has members that do not depend on its parameters, it may be convenient to move them into a plain class:

```
template <typename T>
class MyClass
  double value ;
  std::string name ;
  std::vector<T> data ;
public:
  std::string getName() const;
};
should become:
class MyBaseClass
protected:
  ~MyBaseClass() {}
  double value ;
  std::string name ;
public:
  std::string getName() const;
};
template <typename T>
class MyClass : MyBaseClass
  std::vector<T> data_;
public:
  using MyBaseClass::getName;
};
```

The derivation may be public, private, or even protected. ¹⁶ This will reduce the compilation

complexity and potentially the size of the binary code. Of course, this optimization is most effective if the template is instantiated many times.

1.1.6. Template Position

The body of a class/function template must be available to the compiler at every point of instantiation, so the usual header/cpp file separation does not hold, and everything is packaged in a single file, with the hpp extension.

If only a declaration is available, the compiler will use it, but the linker will return errors:

```
// sq.h
template <typename T>
T sq(const T& x);

// sq.cpp
template <typename T>
T sq(const T& x)
{
   return x*x;
}

// main.cpp
#include "sq.h" // note: function body not visible
int main()
{
   double x = sq(3.14); // compiles but does not link
```

A separate header file is useful if you want to publish only some instantiations of the template. For example, the author of sq might want to distribute binary files with the code for sq<int> and sq<double>, so that they are the only valid types.

In C++, it's possible to explicitly force the instantiation of a template entity in a translation unit without ever using it. This is accomplished with the special syntax:

```
template class X<double>;
template double sq<double>(const double&);
```

Adding this line to sq.cpp will "export" sq<double> as if it were an ordinary function, and the plain inclusion of sq.h will suffice to build the program.

This feature is often used with algorithm tags. Suppose you have a function template, say encrypt or compress, whose algorithmic details must be kept confidential. Template parameter T represents an option from a small set (say T=fast, normal, best); obviously, users of the algorithm are not supposed to add their own options, so you can force the instantiation of a small

number of instances—encrypt<fast>, encrypt<normal>, and encrypt<best>—and distribute just a header and a binary file.

Note C++0x adds to the language the external instantiation of templates. If the keyword extern is used before template, the compiler will skip instantiation and the linker will borrow the template body from another translation unit.

See also Section 1.6.1 below.

1.2. Specialization and Argument Deduction

By definition, we say that a name is *at namespace level*, at *class level*, or *at body level* when the name appears between the curly brackets of a namespace, class, or function body, as the following example shows:

Function templates—member or non-member—can automatically deduce the template argument looking at their argument list. Roughly speaking, ¹⁷ the compiler will pick the most specialized function that matches the arguments. An exact match, if feasible, is always preferred, but a conversion can occur.

A function F is more specialized than G if you can replace any call to F with a call to G (on the same arguments), but not vice versa. In addition, a non-template function is considered more specialized than a template with the same name.

Sometimes overload and specialization look very similar:

```
template <typename scalar_t>
inline scalar t sq(const scalar t& x); // (1) function template
```

But they are not identical; consider the following counter-example:

The basic difference between overload and specialization is that a function template acts as a single entity, regardless of how many specializations it has. For example, the call sq(y) just after (3) would force the compiler to select between entities (1) and (2). If y is double, then (2) is preferred, because it's a normal function; otherwise, (1) is instantiated based on the type of y: only at this point, if y happens to be int, the compiler notices that sq has a specialization and picks (3).

Note that two different templates may overload:

```
template <typename T>
void f(const T& x)
{
   std::cout << "I am f(reference)";
}
or:
template <typename T>
void f(const T* x)
{
   std::cout << "I am f(pointer)";
}</pre>
```

On the other hand, writing a specialization when overloaded templates are present may require you to specify explicitly the parameters:

Remember that template specialization is legal only at the namespace level (even if most compilers will tolerate it anyway):

```
class mathematics
  template <typename scalar t>
   inline scalar t sq(const scalar t& x) { ... }; // template
member function
  template <>
   inline int sq(const int& x) { ... };
                                                      // illegal
specialization!
};
  The standard way is to call a global function template from inside the class:
// global function template: outside
template <typename scalar t>
inline scalar t gsq(const scalar t& x) { ... };
// specialization: outside
template <>
inline int gsq(const int& x) { ... };
class mathematics
  // template member function
  template <typename scalar t>
  inline scalar t sq(const scalar t& x)
   {
     return gsq(x);
};
  Sometimes you may need to specify explicitly the template parameters because they are unrelated
to function arguments (in fact, they are called non-deducible):
class crc32 { ... };
class adler { ... };
template <typename algorithm t>
size t hash using(const char* x)
  // ...
```

In this case, you must put non-deducible types and arguments first, so the compiler can work out

size t j = hash using<crc32>("this is the string to be hashed");

all the remaining:

Argument deduction obviously holds only for function templates, not for class templates.

It's generally a bad idea to supply an argument explicitly, instead of relying on deduction, except in some special cases, described next.

• When necessary for disambiguation:

• When a type is non-deducible 18:

```
template <typename T>
T get_random()
{ ... }

double r = get_random<double>();
```

• When you want a function template to look similar to a built-in C++ cast operator:

```
template <typename X, typename T>
X sabotage_cast(T* p)
{
   return reinterpret_cast<X>(p+1);
}
std::string s = "don't try this at home";
double* p = sabotage cast<double*>(&s);
```

• To perform simultaneously a cast and a function template invocation:

```
double y = sq<int>(6.28) // casts 6.28 to int, then squares the value
```

• When an algorithm has an argument whose default value is template-dependent (usually a functor)¹⁹:

```
template <typename LESS_T>
void nonstd_sort (..., LESS_T cmp = LESS_T())
{
    // ...
}

// call function with functor passed as template argument
nonstd_sort< std::less<...> > (...);

// call function with functor passed as value argument
nonstd_sort (..., std::less<...>());
```

A template name (such as std::vector) is different from the name of the class it generates (such as std::vector<int>). At the class level, they are equivalent:

As a rule, the word something alone, without angle brackets, represents a template, which is a well-defined entity of its own. In C++, there are *template-template parameters*. You can declare a template whose parameters are not just types, but are class templates that match a given pattern:

```
template <template <typename T> class X>
class example
{
    X<int> x1_;
    X<double> x2_;
};
```

```
typedef example<something> some_example; // ok: 'something'
matches
```

Note that class and typename are not equivalent here:

```
template <template <typename T> typename X> // error
```

Class templates can be fully or partially *specialized*. After the general template, we list specialized versions:

```
// in general T is not a pointer
template <typename T>
struct is a pointer type
  static const int value = 1;
};
// 2: full specialization for void*
template <>
struct is a pointer type<void*>
  static const int value = 2;
};
// 3: partial specialization for all pointers
template <typename X>
struct is a pointer type<X*>
  static const int value = 3;
};
int b1 = is a pointer type<int*>::value; // uses 3 with X=int
int b2 = is a pointer type<void*>::value; // uses 2
int b3 = is a pointer type<float>::value; // uses the general
template
```

Partial specialization can be recursive:

```
template <typename X>
struct is_a_pointer_type<const X>
{
   static const int value = is_a_pointer_type<X>::value;
};
```

The following example is known as the pointer paradox:

```
template <typename T>
void f(const T& x)
```

#include <iostream>

```
std::cout << "My arg is a reference";
}

template <typename T>
void f(const T* x)
{
   std::cout << " My arg is a pointer";
}

   In fact, the following code prints as expected:

const char* s = "text";
f(s);
f(3.14);

My arg is a pointer
My arg is a reference
   Now write instead:

double p = 0;
f(&p);</pre>
```

You would expect to read pointer; instead you get a call to the *first* overload. The compiler is correct, since type double* matches const T* with one trivial implicit conversion (namely, adding const-ness), but it matches const T& perfectly, setting T=double*.

1.2.1. Deduction

{

Function templates can deduce their parameters, matching argument types with their signature:

Deduction also covers non-type arguments:

However, remember that deduction is done via "pattern matching" and the compiler is not required to perform any kind of algebra²⁰:

```
// this template is formally valid, but deduction will never
succeed...
template <int I>
arg<I> f(arg<I+1>)
{
    // ...
}

arg<3> a;
f(a);    // ...the compiler will not solve the equation
I+1==3
arg<2+1> b;
f(b);    // ...error again

No matching function for call to 'f'
Candidate template ignored: couldn't infer template argument 'I'
```

On the other hand, if a type is contained in a class template, then its context (the parameters of the outer class) cannot be deduced:

Note that this error does *not* depend on the particular invocation. This kind of deduction is logically not possible; T may not be unique.

```
template <typename T>
struct A
{ typedef double type; };

// if A<X>::type is double, X could be anything
```

A dummy argument can be added to enforce consistency:

```
template <typename T>
void f(std::vector<T>&, typename std::vector<T>::iterator);
```

The compiler will deduce \mathbb{T} from the first argument and then verify that the second argument has the correct type.

You could also supply explicitly a value for T when calling the function:

```
template <typename T>
void f(typename std::vector<T>::iterator);
std::vector<double> w;
f<double>(w.begin());
```

Experience shows that it's better to minimize the use of function templates with non-deduced parameters. Automatic deduction usually gives better error messages and easier function lookup; the following section lists some common cases.

First, when a function is invoked with template syntax, the compiler does not necessarily look for a template. This can produce obscure error messages.

```
struct base
  template <int I, typename X> // template, where I is non-
deduced
  void foo (X, X)
};
struct derived : public base
                                   // not a template
  void foo(int i)
    foo<314>(i, i);
                                  // line #13
};
1>error: 'derived::foo': function call missing argument list; use
'&derived::foo' to create a pointer to member
1>error: '<' : no conversion from 'int' to 'void (__cdecl
derived::* )(int)'
    There are no conversions from integral values to pointer-to-
member values
1>error: '<' : illegal, left operand has type 'void ( cdecl
derived::* )(int)'
1>warning: '>' : unsafe use of type 'bool' in operation
1>warning: '>' : operator has no effect; expected operator with
side-effect
```

When the compiler meets foo<314>, it looks for any foo. The first match, within derived, is void foo (int) and lookup stops. Hence, foo<314> is misinterpreted as (ordinary function name) (less) (314) (greater). The code should explicitly specify base::foo.

Second, if name lookup succeeds with multiple results, the explicit parameters constrain the overload resolution:

```
template <typename T>
void f();
template <int N>
void f();
f<double>(); // invokes the first f, as "double" does not match
"int N"
              // invokes the second f
f < 7 > ();
  However, this can cause unexpected trouble, because some overloads<sup>21</sup> may be silently ignored:
template <typename T>
void g(T x);
double pi = 3.14;
g<double>(pi); // ok, calls g<double>
template <typename T>
void h(T x);
void h(double x);
double pi = 3.14;
h<double>(pi);
                      // unexpected: still calls the first h
  Here's another example:
template <int I>
class X {};
template <int I, typename T>
void q(X < I >, T x);
                                 // a special 'g' for X<0>
template <typename T>
void g(X<0>, T x);
                                 // however, this is q<T>, not q<0, T>
double pi = 3.14;
X<0> x;
                                 // calls the first q
g<0>(x, pi);
```

Last but not least, old compilers used to introduce subtle linker errors (such as calling the wrong

// calls the second q

g(x, pi);

function).

1.2.2. Specializations

Template specializations are valid only at the namespace level²²:

The compiler will start using the specialized version only after it has compiled it:

```
template <typename scalar_t>
scalar_t sq(const scalar_t& x)
{ ... }

struct A
{
    A(int i = 3)
    {
        int j = sq(i); // the compiler will pick the generic template
      }
};

template <>
int sq(const int& x) // this specialization comes too late, compiler gives error
{ ... }
```

However, the compiler will give an error in such a situation (stating that *specialization comes after instantiation*). Incidentally, it can happen that a generic class template explicitly "mentions" a special case, as a parameter in some member function. The following code in fact causes the aforementioned compiler error.

```
template <typename T>
struct C
  C(C<void>)
};
template <>
struct C<void>
};
  The correct version uses a forward declaration:
template <typename T>
struct C;
template <>
struct C<void>
};
template <typename T>
struct C
{
  C(C<void>)
};
  Note that you can partially specialize (and you'll do it often) using integer template parameters:
// general template
template <typename T, int N>
class MyClass
{ ... };
// partial specialization (1) for any T with N=0
template <typename T>
class MyClass<T, 0>
{ ... };
// partial specialization (2) for pointers, any N
template <typename T, int N>
class MyClass<T*, N>
{ ... };
```

However, this approach can introduce ambiguities:

Usually you must explicitly list all the "combinations". If you specialize X<T1, T2> for all $T1 \in A$ and for all $T2 \in B$, then you must also specialize explicitly X<T1, $T2> \in A\times B$.

```
// partial specialization (3) for pointers with N=0
template <typename T>
class MyClass<T*, 0>
{ ... };
```

It's illegal to write a *partial* specialization when there are dependencies between template parameters in the general template.

```
// parameters (1) and (2) are dependent in the general template
template <typename int_t, int_t N>
class AnotherClass
{};
template <typename T>
class AnotherClass<T, 0>
{};
error: type 'int_t' of template argument '0' depends on template
parameter(s)
```

Only a full specialization is allowed:

```
template <>
class AnotherClass<int, 0>
{};
```

A class template specialization may be completely unrelated to the general template. It need not have the same members, and member functions can have different signatures.

While a gratuitous interface change is a symptom of bad style (as it inhibits any generic manipulation of the objects), the freedom can be usually exploited:

```
template <typename T, int N>
struct base_with_array
{
   T data_[N];
   void fill(const T& x)
   {
      std::fill_n(data_, N, x);
   }
};
template <typename T>
```

```
struct base_with_array<T, 0>
{
    void fill(const T& x)
    {
    }
};

template <typename T, size_t N>
class cached_vector : private base_with_array<T, N>
{
    // ...

public:
    cached_vector()
    {
        this->fill(T());
    }
};
```

1.2.3. Inner Class Templates

A class template can be a member of another template. One of the key points is syntax; the inner class has its own set of parameters, but it knows all the parameters of the outer class.

```
template <typename T>
class outer
{
public:
   template <typename X>
   class inner
   {
       // use freely both X and T
   };
};
```

The syntax for accessing inner is outer<T>::inner<X> if T is a well-defined type; if T is a template parameter, you have to write outer<T>::template inner<X>:

It's usually difficult or impossible to specialize inner class templates. Specializations should be listed outside of outer, so as a rule they require two template <...> clauses, the former for T (outer), the latter for X (inner).

```
template <typename T>
                                                                     class outer
                                                                       template <typename X>
Primary template: it defines an inner<X> which we'll call informally
inner_1.
                                                                       class inner
                                                                       };
                                                                     };
                                                                     template <>
                                                                     class outer<int>
                                                                       template <typename X>
Full specializations of outer may contain an inner<X>, which to the compiler
                                                                       class inner
is completely unrelated to inner 1; we'll call this inner 2.
                                                                         // ok
                                                                       };
                                                                     };
                                                                     template <>
                                                                     class
                                                                     outer<int>::inner<float>
inner_2 can be specialized:
                                                                       // ok
                                                                     };
                                                                     template <>
                                                                     template <typename X>
                                                                     class outer<double>::inner
specialization of inner 1 for fixed T (=double) and generic X.
                                                                       // ok
                                                                     };
                                                                     template <>
                                                                     template <>
                                                                     outer<double>::inner<char>
specialization of inner_1 for fixed T (=double) and fixed X (=char).
                                                                       // ok
                                                                     };
                                                                     template <typename T>
                                                                     template <>
```

```
It's illegal to specialize inner_1 for fixed X with any T.
```

```
class
outer<T>::inner<float>
{
    // error!
};
```

Note that, even if X is the same, inner_1<X> and inner_2<X> are completely different types:

```
template <typename T>
struct outer
  template <typename X> struct inner {};
};
template <>
struct outer<int>
  template <typename X> struct inner {};
};
int main()
{
  outer<double>::inner<void> I1;
  outer<int>::inner<void> I2;
  I1 = I2;
}
error: binary '=' : no operator found which takes a right-hand
operand of type 'outer<int>::inner<X>' (or there is no acceptable
conversion)
```

It's impossible to write a function that, say, tests any two "inner"s for equality, because given an instance of inner<X>, the compiler will not deduce its outer<T>.

```
template <typename T, typename X>
bool f(outer<T>::inner<X>);    // error: T cannot be deduced?
```

The actual type of variable I1 is not simply inner<void>, but outer<double>::inner<void>. If for any X, all inner<X> should have the same type, then inner must be promoted to a global template. If it were a plain class, it would yield simply:

```
struct basic_inner
{
};
template <typename T>
struct outer
```

```
{
  typedef basic inner inner;
};
template <>
struct outer<int>
  typedef basic inner inner;
};
  If inner does not depend on T, you could write<sup>23</sup>:
template <typename X>
struct basic inner
};
template <typename T>
struct outer
  template <typename X>
  struct inner : public basic inner<X>
     inner& operator=(const basic inner<X>& that)
        static cast<basic inner<X>&>(*this) = that;
        return *this;
     }
   };
};
template <>
struct outer<int>
  template <typename X>
  struct inner : public basic inner<X>
     inner& operator=(const basic inner<X>& that)
        static cast<basic inner<X>&>(*this) = that;
        return *this;
     }
   };
};
```

Otherwise, you have to design basic_inner's template operators that support mixed operations:

```
template <typename X, typename T>
struct basic inner
  template <typename T2>
  basic inner& operator=(const basic inner<X, T2>&)
  { /* ... */ }
};
template <typename T>
struct outer
  template <typename X>
  struct inner : public basic inner<X, T>
     template <typename ANOTHER T>
     inner& operator=(const basic inner<X, ANOTHER T>& that)
        static cast<basic inner<X, T>&>(*this) = that;
        return *this;
     }
  };
};
template <>
struct outer<int>
  template <typename X>
  struct inner: public basic inner<X, int>
     template <typename ANOTHER T>
     inner& operator=(const basic inner<X, ANOTHER T>& that)
        static cast<basic inner<X, int>&>(*this) = that;
        return *this;
     }
  };
};
int main()
  outer<double>::inner<void> I1;
  outer<int>::inner<void> I2;
  I1 = I2;  // ok: it ends up calling basic inner::operator=
}
  This is known in the C++ community as the SCARY initialization.<sup>24</sup>
```

SCARY stands for "Seemingly erroneous (constrained by conflicting template parameters), but

actually work with the right implementation". Put simply, two inner types that should be different (specifically, outer<T1>::inner and outer<T2>::inner) actually share the implementation, which means it's possible to treat them uniformly as "two inners".

As you've seen for function templates, you should never instantiate the master template before the compiler has met all the specializations. If you use only full specializations, the compiler will recognize a problem and stop. *Partial* specializations that come too late will be just ignored:

```
struct A
  template <typename X, typename Y>
  struct B
  {
     void do it() {} // line #1
  };
  void f()
     B<int,int> b; // line #2: the compiler instantiates
B<int,int>
     b.do it();
};
template <typename X>
                         // this should be a specialization of
struct A::B<X, X>
B < X, X >
                     // but it comes too late for B<int,int>
                        // line #3
  void do it() {}
} ;
A a;
a.f();
                         // calls do it on line #1
  Furthermore, adding a full specialization of B will trigger a compiler error:
template <>
struct A::B<int, int>
  void do it() {}
};
error: explicit specialization; 'A::B<X,Y>' has already been
instantiated
       with
       X=int,
```

Y=int

The obvious solution is to move the function bodies after the specializations of A::B.

1.3. Style Conventions

Style is the way code is written; this definition is so vague that it includes many different aspects of programming, from language techniques to the position of curly braces.

All the C++ objects in namespace std exhibit a common style, which makes the library more coherent.

For example, all names are lowercase²⁵ and multi-word names use underscores. Containers have a member function bool T::empty() const that tests if the object is empty and a void T::clear() that makes the container empty. These are elements of style.

A fictional STL written in pure C would possibly have a global function clear, overloaded for all possible containers. Writing code such as cont.clear() or clear (&cont) has the same net effect on cont, and might even generate the same binary file, but granted, it has a very different style.

All these aspects are important during code reviews. If style agrees with the reader forma mentis, the code will look natural and clear, and maintenance will be easier.

Some aspects of style are indeed less important, because they can be easily adjusted. For example, using beautifiers—each worker in a team might have a pre-configured beautifier on his machine, integrated with the code editor, which reformats braces, spaces, and newlines at a glance.

Note JEdit (see http://www.jedit.org) is a free multiplatform code editor that supports plugins.

AStyle (Artistic Style) is a command-line open source code beautifier (see http://astyle.sourceforge.net) whose preferences include the most common formatting option (see Figure 1-1).

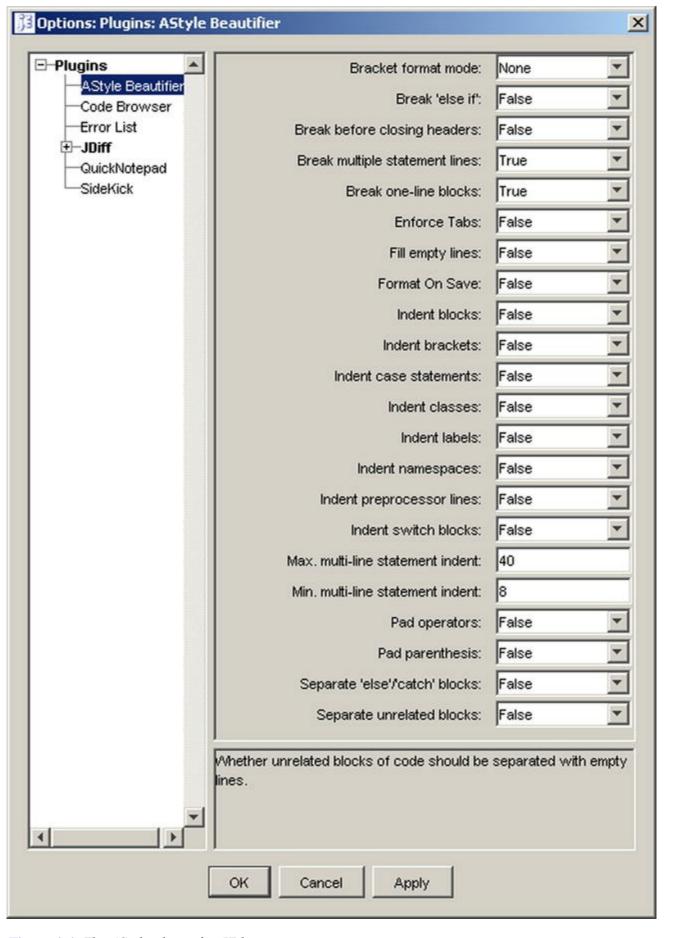


Figure 1-1. The AStyle plugin for JEdit

Most reasonable style conventions are equivalent; it's important to pick one and try to be consistent for some time. ²⁶

Ideally, if code is written according to some common behavior conventions, a reader may deduce

how it works based on the style, without looking into the details.

For example:

```
void unknown_f(multidimensional_vector<double, 3, 4>& M)
{
   if (!M.empty())
      throw std::runtime_error("failure");
}
```

Most readers will describe this fragment as, "If the multidimensional vector is not empty, then throw an exception". However, nothing in the code states that this is the intended behavior except style.

In fact, multidimensional_vector::empty could in principle make the container empty and return a non-zero error code if it does not succeed.²⁷

The naming convention is a big component of style.

The following example lists some ideas for how to convey extra meaning when building the name of an object. It is not intended as a set of axioms, and in particular no item is worse/better than its opposite, but it's a detailed example of how to assemble a style that can help you diagnose and solve problems.

Remember that the C++ standard prescribes that some identifiers are "reserved to the implementation for any use" and some are reserved for names in the global or std namespace. That means user names should never:

- Begin with an underscore (in particular, followed by a capital letter)
- Contain a double underscore
- Contain a dollar sign (it's tolerated by some compilers, but it's not portable)

1.3.1. Comments

"Many good programming practices boil down to preparing for change or expressing intent. Novices emphasize the former, experts the latter."

—John D. Cook

Remember to add lots of comments to your code. If this is valid for any programming language, it is especially true for TMP techniques, which can easily be misunderstood. The correct behavior of TMP is based on bizarre entities, like empty classes, void functions, and strange language constructs that look like errors. It's really hard for the author of the code to remember why and how these techniques work, and even harder for other people who have to maintain the code.

1.3.2. Macros

Macros play a special role in TMP. Some programmers consider them a necessary evil and indeed

they are necessary, but it's not obvious they are also evil.

Macros must:

- Allow the reader to recognize them
- Prevent name collisions

The easiest way to satisfy both requirements is to choose a unique and sufficiently ugly common prefix for all macros and play with lower/uppercase to give extra meaning to the name.

As an example, you could agree that all macros begin with MXT_. If the macro is persistent, i.e. never undefined, the prefix will be MXT. If the macro's scope is limited (it's defined and undefined later in the same file), the prefix will be mXT.

```
#ifndef MXT_filename_
#define MXT_filename_
name it MXT_*

#define mXT_MYVALUE 3
const int VALUE = mXT_MYVALUE;
#undef mXT_MYVALUE

#endif //MXT filename

// this is "exported" - let's
// this macro has limited "scope"
// let's name it mXT_*
//
#endif //MXT filename
```

A lowercase prefix mxt is reserved to remap standard/system function names in different platforms:

For better code appearance, you could decide to replace some keywords with a macro:

And/or enclose the namespace directives in an ASCII-art comment box:

It's useful to have some (integer) functions as a set of macros:

```
#define MXT_M_ABS(a) ((a) <0 ? -(a) : (a)) #define MXT M SQ(a) ((a) *(a))
```

The infix M stands for "macro" and these will be useful when working with templates:

```
template <int N>
struct SomeClass
{
   static const int value = MXT_M_SQ(N)/MXT_M_MAX(N, 1);
};
```

■ **Note** The C++11 Standard introduced a new keyword: constexpr.²⁸

A function-declared constexpr has no side effects and it always returns the same result, deterministically, from the same arguments. In particular, when such a function is called with compile-time constant arguments, its result will also be a compile-time constant:

```
constexpr int sq(int n) { return n*n; }
constexpr int max(int a, int b)
{ return a < b ? b : a; }

template < int N >
struct SomeClass
{
    static const int value = sq(N)/max(N, 1);
```

Finally, consider a special class of macros. A *macro directive* is a macro whose usage logically takes an entire line of code.

In other words, the difference between an ordinary macro and a directive is that the latter cannot coexist with anything on the same line (except possibly its arguments):

The definition of a macro directive, in general, should not end with a semicolon, so the user is forced to close the line manually (when appropriate), as if it were a standard function call.

```
// note: no trailing ';'
#define MXT_INT_I(k) int i = (k)
int main()
{
   MXT_INT_I(0); // put ';' here return 0;
}
```

Here is a more complex example. Note that the trailing semicolon is a very strong style point, so it's used even in places where, in ordinary code, a semicolon would be unnatural.

```
#define mXT C(NAME, VALUE)
static scalar t NAME()
{
  static const scalar t NAME## = (VALUE);
  return NAME## ;
}
template <typename scalar t>
struct constant
  // the final ';' at class level is legal, though uncommon
  mXT C(Pi, acos(scalar t(-1)));
  mXT C(TwoPi, 2*acos(scalar t(-1)));
  mXT C(PiHalf, acos(scalar t(0)));
  mXT C(PiQrtr, atan(scalar t(1)));
  mXT C(Log2, log(scalar t(2)));
};
#undef mXT C
double x = constant<double>::TwoPi();
```

However, special care is required when invoking macro directives, which expand to a sequence of instructions:

```
#define MXT_SORT2(a,b) if ((b)<(a)) swap((a),(b))

#define MXT_SORT3(a,b,c) \
    MXT_SORT2((a),(b)); MXT_SORT2((a),(c)); MXT_SORT2((b),(c))

int a = 5, b = 2, c = 3;

MXT_SORT3(a,b,c); // apparently ok: now a=2, b=3, c=5</pre>
```

Nevertheless, this code is broken:

```
int a = 5, b = 2, c = 3;
if (a>10)
   MXT SORT3(a,b,c);
                                  // problem here!
   Since it expands to:
if (a>10)
   MXT SORT2(a,b);
MXT SORT2(a,c);
MXT SORT2(b,c);
   More surprising is that the following fragment is clear, but incorrect:
if (a>10)
   MXT SORT2(a,b);
else
   MXT SORT2(c,d);
   Because of the way if-then-else associates in C++, the macro expands as
if (a>10)
   if (a < b)
      swap(a,b);
else
   if (c < d)
      swap(c,d);
   The indentation does not resemble the way code is executed; the block actually groups as
if (a>10)
   if (a < b)
      swap(a,b);
   else if (c < d)
      swap(c,d);
}
   To solve the problem, you can use the do {...} while (false) idiom:
#define MXT SORT3(a,b,c)
  do { MXT SORT2((a),(b)); MXT SORT2((a),(c)); MXT SORT2((b),(c));
  while (false)
   This allows both to put "local code" inside a block and to terminate the directive with a
```

Remember that this will not save you from an error like:

semicolon.

```
MXT_SORT3(a, b++, c); // error: b will be incremented more than
once
```

This is why we insist that macros are immediately recognizable by a "sufficiently ugly" prefix. To tackle the "if" macro problem, write a do-nothing else branch:

```
#define MXT SORT2(a,b) if ((b)<(a)) swap((a),(b)); else
```

Now MXT_SORT2 (a, b); expands to if (...) swap(...); else; where the last semicolon is the empty statement. Even better²⁹:

```
\#define\ MXT\_SORT2(a,b) if (!((b)<(a))) {} else swap((a),(b))
```

As a final remark, never make a direct use of types that come from macros. Always introduce a typedef. If the macro is not carefully written, the association between * and const may give unexpected results. Consider:

```
T \times = 0; const T^* p = &x; // looks correct Unless:
```

Instead, consider intercepting the macro:

1.3.3. Symbols

#define T char*

Most C++ projects contain several kinds of symbols (classes, functions, constants, and so on). A rough division line can be drawn between system/framework utilities, which are completely abstract and generic, and project specific entities, which contain specific logic and are not expected to be reused elsewhere.

This simple classification may turn out important for (human) debuggers. If any piece of code is considered a "system utility," then it's implicitly trusted and it may usually be "stepped over" during debug. On the other hand, project-specific code is possibly less tested and should be "stepped in".

We can agree that stable symbols should follow the STL naming conventions (lowercase, underscores, such as stable_sort, hash_map, and so on). This often will be the case for class templates.

The rest should be camel case (the Java convention is fine).

```
(framework header) sq.hpp
template <typename scalar t>
```

A functor is an instance of an object that implements at least one operator (), so the name of the instance behaves like a function. 30

A functor is said to be *modifying* if it takes arguments by non-const references.

A *predicate* is a non-modifying functor that takes all arguments of the same type and returns a boolean. For example, less is a binary predicate:

```
template <typename T>
struct less
{
   bool operator()(const T&, const T&) const;
};
```

Most functors contain a typedef for the return type of operator (), usually named result type or value type. 31

Functors are usually stateless or they carry few data members, so they are built on the fly. Occasionally, you may need a meaningful name for an instance, and this may not be so easy, because if the functor has a limited "scope," the only meaningful name has already been given to the class.

```
calendar myCal;
std::find_if(year.begin(), year.end(), is_holiday(myCal));
// is_holiday is a class
// how do we name an instance?
```

You can use one of the following:

• Use a lowercase functor name and convert it to uppercase for the instance:

```
calendar myCal;
is_holiday IS_HOLIDAY(myCal);
std::find_if(year.begin(), year.end(), IS_HOLIDAY);
```

• Use a lowercase functor name with a prefix/postfix and remove it in the instance:

```
calendar myCal;
is_holiday_t is_holiday(myCal);
std::find_if(year.begin(), year.end(), is_holiday);
```

1.3.4. Generality

The best way to improve generality is to reuse standard classes, such as std::pair.

This brings in well-tested code and increases interoperability; however, it may often hide some specific logic, for example the meaning of pair::first and pair::second may not be obvious at first sight. See the following paradigmatic example:

```
struct id_value
{
   int id;
   double value;
};

id_value FindIDAndValue(...);

   This may be replaced by:

std::pair<int, double> FindIDAndValue(...)
```

However, the caller of the first function can write p.id and p.value, which is easier to read than p.first and p.second. You may want to provide a *less generic* way to access pair members:

Macros

• Global functions (these are called *accessors*; see Section **6.2.1**)

```
inline int& id(std::pair<int, double>& P)
{ return P.first; }
inline int id(const std::pair<int, double>& P)
{ return P.first; }
```

• Global pointer-to-members

```
typedef std::pair<int, double> id_value;
int id_value::*ID = &id_value::first;
double id_value::*VALUE = &id_value::second;
// later
std::pair<int, double> p;
p.*ID = -5;
p.*VALUE = 3.14;
```

To make ID and VALUE constants, the syntax is:

```
int id_value::* const ID = &id_value::first;
```

1.3.5. Template Parameters

A fairly universally accepted convention is to reserve UPPERCASE names for non-type template parameters. This could cause some name conflict with macros. It's not always necessary to give a name to template parameters (as with function arguments), so when it's feasible, you'd better remove the name entirely:

```
// the following line is likely to give strange errors
// since some compilers define BIGENDIAN as a macro!

template <typename T, bool BIGENDIAN = false>
class SomeClass
{
};

template <typename T>
class SomeClass<T, true>
{
};
```

A safer declaration would be³²:

```
template <typename T, bool = false>
class SomeClass
```

Type parameters are usually denoted by a single uppercase letter—usually T (or T1, T2...) if type can be indeed anything.³³ A and R are also traditionally used for parameters that match arguments and results:

```
int foo(double x) { return 5+x; }
template <typename R, typename A>
inline R apply(R (*F)(A), A arg)
```

```
{
  return F(arg);
template <typename R, typename A1, typename A2>
inline R apply(R (*F)(A1, A2), A1 arg1, A2 arg2)
   return F(arg1, arg2);
double x = apply(\&foo 3.14);
  Otherwise, you might want to use a (meaningful) lowercase name ending with t (for example,
int t, scalar t, object t, any t, or that t).
template <typename T, int N>
class do nothing
};
template <typename int_t> // int_t should behave as an integer
type<sup>34</sup>
struct is_unsigned
   static const bool value = ...;
};
  The suffix t, which in C originally means typedef, is also widely used for (private)
typedefs standing for template instances:
template <typename scalar t>
class SomeContainer
   // informally means:
   // within this class, a pair always denotes a pair of scalars
  private:
     typedef std::pair<scalar t, scalar t> pair t;
};
  On the other hand, a public typedef name usually is composed of lowercase regular English
words (such as iterator category). In that case, type is preferred:
template <typename scalar t>
class SomeContainer
  public:
     typedef scalar t result type;
};
```

1.3.6. Metafunctions

template <typename T>

We often meet stateless class templates whose members are only enumerations (as a rule, anonymous), static constants, types (typedefs or nested classes), and static member functions.

Generalizing Section 1.1, we consider this template a *metafunction* that maps its tuple of parameters to a class, which is seen as a *set of results* (namely, its members).

```
template <typename T, int N>
struct F
{
   typedef T* pointer_type;
   typedef T& reference_type;
   static const size_t value = sizeof(T)*N;
};
```

The metafunction F maps a pair of arguments to a triple of results:

```
(T,N) \rightarrow (pointer\_type, reference\_type, value)
\{type\} \times \{int\} \rightarrow \{type\} \times \{size t\}
```

Most metafunctions return either a single type, conventionally named type, or a single numeric constant (an integer or an enumeration), conventionally named value.³⁵

```
struct largest_precision_type;

template <>
struct largest_precision_type<float>
{
   typedef double type;
};

template <>
struct largest_precision_type<double>
{
   typedef double type;
};

template <>
struct largest_precision_type<int>
{
   typedef long type;
};
```

Similarly:

```
template <unsigned int N>
struct two to
  static const unsigned int value = (1<<N);
};
template <unsigned int N>
struct another two to
  enum { value = (1 << N) };
};
                                   // invocation
unsigned int i = two to<5>::value;
largest precision<int>::type j = i + 100; // invocation
  Historically, the first metafunctions were written using enums:
template <size t A>
struct is prime
  enum { value = 0 };
};
template <>
struct is prime<2>
  enum { value = 1 };
};
template <>
struct is prime<3>
  enum { value = 1 };
};
// ...
```

The main reason was that compilers were unable to deal with static const integers (including bool). The advantage of using an enum over a static constant is that the compiler will *never* reserve storage space for the constant, as the computation is either static or it fails.

Conversely, a static constant integer could be "misused" as a normal integer, for example, taking its address (an operation that the compiler will disallow on enums).

Note According to the classic C++ standard, the use of a static constant as a normal integer is illegal (unless the constant is re-declared in the .cpp file, as any other static data member of a class). However, most compilers allow it, as long as the code does not try to take the address of the constant or bind it to a const reference. The requirement was removed in modern C++.

Furthermore, the language allows declaring a static integer constant (at function scope, not at class scope) that is *dynamically* initialized, and so *not* a compile-time constant:

```
static const int x = INT_MAX;
static const int y = std::numeric_limits<int>::max();  // dynamic
static const int z = rand();  // dynamic
double data[y];  // error
```

In practice, an enum is usually equivalent to a *small* integer. enums are in general implemented as signed int, unless their value is too large. The most important difference is that you cannot bind an unnamed enum to a template parameter without an explicit cast:

```
double data[10];
std::fill_n(data, is_prime<3>::value, 3.14); // may give error!
  The previous code is non-portable, because std::fill n may be defined.
template <..., typename integer t, ...>
void fill n(..., integer t I, ...)
 ++I; // whatever...
 --I; // whatever...
error C2675: unary '--': "does not define this operator or
a conversion to a type acceptable to the predefined operator
see reference to function template instantiation
'void std:: Fill n<double*, Diff, Ty>( OutIt, Diff,const Ty
&, std:: Range checked iterator tag) ' being compiled
       with
       Diff=,
          __Ty=double *,
          _OutIt=double **
       1
```

In practice, an enum is fine to store a small integer (for example, the logarithm of an integer in base 2). Because its type is not explicit, it should be avoided when dealing with potentially large or unsigned constants. As a workaround for the std::fill_n call, just cast the enumeration to an appropriate integer:

```
std::fill_n(..., int(is_prime<3>::value), ...); // now ok!
```

Frequently, metafunctions invoke helper classes (you'll see less trivial examples later):

```
template <int N>
struct ttnp1_helper
{
```

```
static const int value = (1<<N);
};

template <int N>
struct two_to_plus_one
{
   static const int value = ttnp1_helper<N>::value + 1;
};
```

The moral equivalent of auxiliary variables are private members. From a TMP perspective, numeric constants and type(def)s are equivalent compile-time entities.

```
template <int N>
struct two_to_plus_one
{
private:
    static const int aux = (1<<N);

public:
    static const int value = aux + 1;
};</pre>
```

The helper class is not private and not hidden,³⁶ but it should not be used, so its name is "uglified" with helper or t (or both).

1.3.7. Namespaces and Using Declarations

Usually all "public" framework objects are grouped in a common namespace and "private" objects reside in special nested namespaces.

```
namespace framework
{
   namespace undocumented_private
   {
      void handle_with_care()
      {
            // ...
      };
   }
   inline void public_documented_function()
      {
            undocumented_private::handle_with_care();
      }
}
```

It's not a good idea to multiply the number of namespaces unnecessarily, since argument-

dependent name lookup may introduce subtle problems, and friend declarations between objects in different namespaces are problematic or even impossible.

Usually, the core of a general-purpose metaprogramming framework is a set of headers (the extension * . hpp is in fact used for pure C++ headers). *Using-namespace declarations* in header files are generally considered bad practice:

```
my_framework.hpp
using namespace std;
main.cpp
#include "my_framework.hpp"

// main.cpp doesn't know, but it's now using namespace std
```

However, *using-function declarations* in header files are usually okay and even desirable (see the do_something example later in the paragraph).

A special use for using-namespace declarations is header versioning.³⁷ This is a very short example:

```
namespace X
{
   namespace version_1_0
   {
     void func1();
     void func2();
   }

   namespace version_2_0
   {
     void func1();
     void func2();
   }

#ifdef USE_1_0
   using namespace version_1_0;
#else
   using namespace version_2_0;
#endif
}
```

Thus the clients using the header always refer to X::func1.

Now we are going to describe in detail another case where using declarations can make a difference.

Function templates are often used to provide an "external interface," which is a set of global functions that allow algorithms to perform generic manipulations of objects³⁸:

The author of a fictitious framework1 provides a function is empty that works on a broad

```
class of containers and on C strings:
```

One of the good properties of this approach is the ease of extensibility. For any new type X, you can provide a specialized is_empty that will have priority over the default implementation. However, consider what happens if the function is explicitly qualified:

```
// framework2.hpp
#include "framework1.hpp"
MXT NAMESPACE BEGIN(framework2)
template <typename string t>
void do something(string t const& x)
                                         // line #2
  if (!framework1::is empty(x))
    // ...
MXT NAMESPACE END(framework2)
#include "framework2.hpp"
namespace framework3
{
  class EmptyString
  } ;
  bool is empty(const EmptyString& x)
```

Thus, you let argument-dependent lookup pick an available is _empty but ensure that framework1 can always supply a default candidate (see also the discussion in Section 1.4.2).

1.4. Classic Patterns

if (!is empty(x))

//...

};

void do something(string t const& x)

When coding a framework/library, it's typical to use and reuse a small set of names. For example, containers can be expected to have a member function [[integer type]] size() const that returns the number of elements.

Adopting a uniform style increases interoperability of objects; for more details, see Chapter 6.All the following paragraphs will try to describe the traditional meaning connected to a few common C++ names.

1.4.1. size_t and ptrdiff_t

In C++ there's no unique standard and portable way to name large integers. Modern compilers will in general pick the largest integers available for long and unsigned long. When you need a large

and fast integer quickly, the preferred choices are size t (unsigned) and ptrdiff t (signed).

size_t, being the result of sizeof and the argument of operator new, is large enough to store any amount of memory; ptrdiff_t represents the difference of two pointers. Since the length of an array of chars is end-begin, as a rule of thumb they will have the same size.

Furthermore, in the flat C++ memory model, sizeof(size_t) also will be the size of pointers, and these integers will likely have the natural size in an architecture—say, 32 bits on a 32-bit processor and 64 bits on a 64-bit processor. They will also be fast (the processor bus will perform atomic transport from registers to memory).

Given this class:

```
template <int N>
struct A
{
   char data[N];
};

sizeof(A<N>) is at least N, so it also follows that size_t is not smaller than int.<sup>39</sup>
```

1.4.2. void T::swap(T&)

This function is expected to swap *this and the argument in constant time, without throwing an exception. A practical definition of *constant* is "an amount of time depending only on T". ⁴⁰

If T has a swap member function, the user expects it to be not worse than the traditional three-copy swap (that is, X=A; A=B; B=X). Indeed, this is always possible, because a member function can invoke each member's own swap:

```
class TheClass
{
   std::vector<double> theVector_;
   std::string theString_;
   double theDouble_;

public:
   void swap(TheClass& that);
   {
     theString_.swap(that.theString_);
     theVector_.swap(that.theVector_);
     std::swap(theDouble_, that.theDouble_);
   }
};
```

The only step that could take non-fixed time is swapping dynamic arrays element by element, but this can be avoided by swapping the arrays as a whole.

The class std::tr1::array<T, N> has a swap that calls std::swap_range on an array of length N, thus taking time proportional to N and depending on T. However, N is part of the type, so

according to this definition, it is constant time. Furthermore, if T is a swappable type (e.g., std::string), swap_range will perform much better than the three copy procedure, so the member swap is definitely an advantage.

The first problem to address is how to swap objects of unspecified type T:

The explicit qualification std:: is an unnecessary constraint. You'd better introduce a using declaration, as seen in Section 1.3.7:

```
using std::swap;

template <typename T>
class TheClass
{
    T theObj_;

public:
    void swap(TheClass<T>& that) // line #1
    {
       swap(theObj_, that.theObj_); // line #2
    }
};
```

However, this results in a compiler error, because by the usual C++ name resolution rules, swap in line 2 is the swap defined in line 1, which does not take two arguments.

The solution, an idiom known as *swap with ADL*, is to introduce a global function with a different name:

```
using std::swap;

template <typename T>
inline void swap_with_ADL(T& a, T& b)
{
   swap(a, b);
}

template <typename T>
class TheClass
{
```

```
T theObj_;
public:
   void swap(TheClass<T>& that)
   {
     swap_with_ADL(theObj_, that.theObj_);
}
```

Due to lookup rules, <code>swap_with_ADL</code> forwards the call to either a <code>swap</code> function defined in the same namespace as <code>T</code> (which hopefully is <code>T</code>'s own version), or to <code>std::swap</code> if nothing else exists. Since there's no local member function with a similar name, lookup escapes class level.

The traditional argument for swap is T&; however, it may make sense to provide more overloads. If an object internally holds its data in a standard container of type X, it might be useful to provide void swap (X&), with relaxed time-complexity expectations:

```
template <typename T>
class sorted vector
  std::vector<T> data ;
public:
  void swap(sorted vector<T>& that)
     data .swap(that.data);
  }
  void swap(std::vector<T>& that)
     data .swap(that);
     std::sort(data .begin(), data .end());
};
  And even more<sup>41</sup>:
struct unchecked type t { };
inline unchecked type t unchecked() { return unchecked type t(); }
template <typename T>
class sorted vector
  // ...
  void swap(std::vector<T>& that, unchecked_type_t (*)())
  {
     assert(is sorted(that.begin(), that.end()));
     data .swap(that);
  }
```

```
sorted_vector<double> x;
std::vector<double> t;

load_numbers_into(x);
x.swap(t);

// now x is empty and t is sorted
// later...

x.swap(t, unchecked); // very fast
To sum up:
```

- Explicitly qualify std::swap with parameters of fixed native type (integers, pointers, and so on) and standard containers (including string).
- Write a using declaration for std::swap and call an unqualified swap when parameters have undefined type T in global functions.
- Call swap with ADL inside classes having a swap member function.

std::swap grants the best implementation for swapping both native and std types. swap is used in algorithms with move semantics:

```
void doSomething(X& result)
{
   X temp;
   // perform some operation on temp, then...
   swap(temp, result);
}
```

and in implementing an exception-safe assignment operator in terms of the copy constructor:

```
} ;
```

If you perform an unconditional swap, the most efficient solution is to take the argument by value:

```
X& operator=(x that)
{
   that.swap(*this);
   return *this;
}
```

On the other hand, you might want to perform additional checks before invoking the copy constructor by hand, even if it's less efficient⁴²:

```
X& operator=(const X& that)
{
   if (this != &that)
   {
      X temp(that);
      temp.swap(*this);
   }
   return *this;
}
```

The drawback is that at some point, both that and temp are alive, so you may need more free resources (e.g., more memory).

1.4.3. bool T::empty() const; void T::clear()

The former function tests whether an object is empty; the latter makes it empty. If an object has a member function size(), then a call to empty() is expected to be no slower than size() ==0.

Note that an object may be empty but still control resources. For example, an empty vector might hold a raw block of memory, where in fact no element has yet been constructed.

In particular, it's unspecified if a clear function will or won't release object resources; clear is a synonym of reset.

To enforce resource cleanup of an auto variable, the usual technique is to swap the instance with a temporary:

```
T x;
// now x holds some resources...
T().swap(x);
```

1.4.4. X T::get() const; X T::base() const

The name get is used when type T wraps a simpler type X. A smart pointer's get would thus return the internal plain pointer.

The function base instead is used to return a copy of the wrapped object, when the wrapper is just a different interface. Since a smart pointer typically adds some complexity (for example, a reference count), the name base would not be as appropriate as get. On the other hand, std::reverse_iterator is an interface that swaps ++ and -- of an underlying iterator, so it has a base().

1.4.5. X T::property() const; void T::property(X)

In this section, "property" is a symbolic name. A class can expose two overloaded member functions called "property" with two different intents.

The first form returns the current value of the property for the current instance; the second sets the property to some new value. The property-set function can also have the form:

```
X T::property(X newval)
{
   const X oldval = property();
   set_new_val(newval);
   return oldval;
}
```

This convention is elegant but not universally used; it is present in std::iostream.

1.4.6. Action(Value); Action(Range)

In this section, "action" is again a symbolic name for an overloaded function or member function.

If an object's own action—for example container.insert (value)—is likely to be invoked sequentially, an object may provide one or more range equivalents. In other words, it can provide member functions with two or more parameters that identify a series of elements at a time. Some familiar examples are:

- An element and a repeat counter
- Two iterators pointing to (begin...end)
- An array and two indexes

It's up to the implementation to take advantage of the range being known in advance. As usual, the range-equivalent function should never be worse than the trivial implementation action(range) := for (x in range) { action(x); }.

1.4.7. Manipulators

Manipulators are one of the least known and more expressive pieces of the C++ standard. They are simply functions that take a stream as an argument. Since their signature is fixed, streams have a

special insertion operator that runs them:

```
class ostream
{
public:
    ostream& operator<<(ostream& (*F) (ostream&))
    {
       return F(*this);
    }

    inline ostream& endl(ostream& os)
    {
       os << '\n';
       return os.flush();
    }
};

int main()
{
    // actually execute endl(cout << "Hello world")
    std::cout << "Hello world" << std::endl;
}</pre>
```

Some manipulators have an argument. The implementation may use a template proxy object to transport this argument to the stream:

```
struct precision_proxy_t
{
   int prec;
};

inline ostream& operator<<(ostream& o, precision_proxy_t p)
{
   o.precision(p.prec);
   return o;
}

precision_proxy_t setprecision(int p)
{
   precision_proxy_t result = { p };
   return result;
}

cout << setprecision(12) << 3.14;</pre>
```

Note that a more realistic implementation may want to *embed* a function pointer in the proxy, so as to have only one insertion operator:

```
class ostream;
```

```
template <typename T, ostream& (*FUNC)(ostream&, T)>
struct proxy
  T arg;
  proxy(const T& a)
     : arg(a)
};
class ostream
public:
  template <typename T, ostream& (*FUNC)(ostream&, T)>
  ostream& operator<<(proxy<T, FUNC> p)
     return FUNC(*this, p.arg);
};
ostream @ global setpr(ostream @ o, int prec)
  o.precision(prec);
  return o;
}
proxy<int, global setpr> setprecision(int p)
  return p;
cout << setprecision(12) << 3.14;</pre>
Note Observe that in classic C++ FUNC would just be a member:
template <typename T>
struct proxy
  T arg;
   ostream& (*FUNC)(ostream&, T);
};
class ostream
```

```
public:
    template <typename T>
    ostream& operator<<(proxy<T> p)
    {
       return p.FUNC(*this, p.arg);
    }
};
```

In principle, a function template could be used as a manipulator, such as:

```
stream << manip1;
stream << manip2(argument);
stream << manip3<N>;
stream << manip4<N>(argument);
```

But in practice this is discouraged, as many compilers won't accept manip3.

1.4.8. Position of Operators

It's important to understand the difference between member and non-member operators.

When member operators are invoked, the left side has already been statically determined, so if any adjustment is necessary, it's performed only on the right side. Alternatively, non-member operators will only match exactly or give errors.

Suppose you are rewriting std::pair:

// ...

```
template <typename T1, typename T2>
struct pair
{
   T1 first;
   T2 second;

   template <typename S1, typename S2>
   pair(const pair<S1, S2>& that)
   : first(that.first), second(that.second)
   {
   }
};

   Now add operator==. First as a member:

template <typename T1, typename T2>
struct pair
{
```

```
inline bool operator== (const pair<T1,T2>& that) const
{
    return (first == that.first) && (second == that.second);
};

Then you compile the following code:

pair<int, std::string> P(1, "abcdefghijklmnop");
pair<const int, std::string> Q(1, "qrstuvwxyz");
if (P == Q)
```

This will work and will call pair<int, string>::operator==. This function requires a constant reference to pair<int, string> and instead it was given pair<const int, string>. It will silently invoke the template copy constructor and make a copy of the object on the right, which is undesirable, as it will make a temporary copy of the string.

It is slightly better to put the operator outside the class:

{ ... }

```
template <typename T1, typename T2>
bool operator== (const pair<T1,T2>& x, const pair<T1,T2>& y)
{
   return (x.first == y.first) && (x.second == y.second);
}
```

At least, this code will now fail to compile, since equality now requires identical pairs. Explicit failure is always more desirable than a subtle problem.

Analogous to the classic C++ rule, "if you write a custom copy constructor, then you'll need a custom assignment operator," we could say that if you write a universal copy constructor, you'll *likely* need universal operators, to avoid the cost of temporary conversions. In this case, use either a template member function with two parameters or a global operator with four. Some programmers prefer global operators when it's possible to implement them using only the public interface of the class (as previously shown).

```
template <typename T1, typename T2 >
struct pair
{
    // ...

template <typename S1, typename S2>
    inline bool operator== (const pair<S1, S2>& that) const
    {
       return (first == that.first) && (second == that.second);
    }
};
```

This will work if this->first and that.first are comparable (for example, int and

const int). Note that you may still have temporary conversions, because you are delegating to an unspecified T1::operator==.⁴³

1.4.9. Secret Inheritance

Public derivation from a concrete class can be used as a sort of "strong typedef":

```
class A
{
    // concrete class
    // ...
};

class B : public A
{
};

// now B works "almost" as A, but it's a different type
```

You may need to implement one or more "forwarding constructors" in B.

This is one of the strategies to simulate template typedefs (which do not exist yet in C++; see Section 12.6):

```
template <typename T1, typename T2>
class A
{
    // ...
};

template <typename T>
class B : public A<T, T>
{
};
```

However, this is acceptable only if A is a private class whose existence is unknown or undocumented:

A secret base class is often a good container of operators that does not depend on some template

parameters. For example, it may be reasonable to test equality between two objects, ignoring all the parameters that are purely cosmetic:

```
template <typename T, int INITIAL_CAPACITY = 16>
class C;

template <typename T>
class H
{
public:
    H& operator==(const H&) const;
};

template <typename T, int INITIAL_CAPACITY>
class C : public H<T>
{
};
```

Comparisons between two containers C with a different INITIAL_CAPACITY will succeed and call their common base H::operator==.

1.4.10. Literal Zero

Sometimes you need to write a function or an operator that behaves differently when a literal zero is passed. This is often the case with smart pointers:

```
template <typename T>
class shared_ptr
{
    //...
};
shared_ptr<T> P;
T* Q;

P == 7; // should not compile
P == 0; // should compile
P == Q; // should compile
```

You can distinguish 0 from a generic int by writing an overload that accepts a pointer to member of a class that has no members:

```
class dummy {};

typedef int dummy::*literal_zero_t;

template <typename T>
class shared_ptr
```

```
// ...
bool operator==(literal_zero_t) const {
```

The user has no way of creating a literal_zero_t, because dummy has no members of type int, so the only valid argument is an implicit cast of a literal zero (unless a more specialized overload exists).

1.4.11. Boolean Type

Some types, such as std::stream, have a cast-to-boolean operator. If implemented naively, this can lead to inconsistencies:

```
class stream
  // ...
  operator bool() const
   // ...
stream s;
                       // ok, that's what we want
if (s)
  int i = s + 2; // unfortunately, this compiles
  A classic workaround was to implement cast to void*:
class stream
  operator void*() const
     // return 'this' when true or '0' when false
};
stream s;
                        // ok, that's what we want
if (s)
```

1.4.12. Default and Value Initialization

If T is a type, then the construction of a default instance does not imply anything about the initialization of the object itself. The exact effect of

```
T x;
```

 $T \times = T();$

depends heavily on T. If T is a fundamental type or a POD, then its initial value is undefined. If T is a class, it's possible that some of its members are still undefined:

```
class A
{
   std::string s_;
   int i_;

public:
   A() {} // this will default-construct s_ but leave i_
uninitialized
};

On the other hand, the line
```

will initialize T to 0, say for all fundamental types, but it may crash if T is A, because it's illegal to copy the uninitialized member i from the temporary on the right into x.

So to sum up:

```
T a();
               // error:
             // a is a function taking no argument and returning T
             // equivalent to T (*a)()
               // ok only if T is a class with default constructor
T b;
             // otherwise T is uninitialized
T c(T());
               // error: c is a function taking a function and
returning T
             // equivalent to T (*c)(T (*)())
               // ok only if T is a simple aggregate44 (e.g.
T d = \{\};
a struct
             // without user-defined constructors)
T e = T();
               // requires a non-explicit copy constructor
             // and may yield undefined behaviour at runtime
```

Value initialization (see paragraphs 8.5.1-7 of the standard) is a way to work around this problem. Since it works only for class members, you have to write:

```
template <typename T>
struct initialized_value
{
   T result;
   initialized_value()
   : result()
   {
   }
};
```

If T is a class with a default constructor, that will be used; otherwise, the storage for T will be set to 0. If T is an array, each element will be recursively initialized:

1.5. Code Safety

The spirit of TMP is "elegance first". In theory, some techniques can open vulnerabilities in source code, which a malicious programmer could exploit to crash a program.⁴⁵

```
Consider the following situations:
#include <functional>
class unary F : public std::unary function<int,float>
public:
   // ...
int main()
  unary F u;
  std::unary function<int,float>* ptr = &u; // ok, legal!
  delete ptr;
                                                    // undefined
behaviour!
   return 0;
}
  The system header <functional> could make the counter-example fail by defining a protected
destructor in the unary function:
template<class Arg, class Result>
struct unary function
{
  typedef _Arg argument_type;
  typedef Result result type;
protected:
  ~unary function()
};
  But this in general does not happen.<sup>46</sup>
  The following idea is due to Sutter ([4]):
myclass.h
class MyClass
  private:
      double x ;
      int z ;
  public:
      template <typename stream t>
```

void write x to(stream t& y)

```
y << x ;
};
   Is it possible to legally read/modify the private member MyClass::z ? Just add a
specialization somewhere after including myclass.h:
struct MyClassHACK
};
template <>
void MyClass::write x to(MyClassHACK&)
   // as a member of MyClass, you can do anything...
   z = 3;
   Finally, there are problems when declaring template friendship. First, there's no standard and
portable way to declare friendship with a template parameter (refer to [5] for more details).
template <typename T, int N>
class test
   friend class T; // uhm...
};
   Second, there is no way to make test<T, N> a friend of test<T, J> (there is nothing like
partial template friendship). A common workaround is to declare test<T, N> a friend of
test<X, J> for any other type X.
template <typename T, int N>
class test
   template <typename X, int J>
                                     // ok, but every test<X,J> has
      friend class test;
access
};
   The same malicious user, who wrote MyClassHACK, can add:
template <>
class test<MyClassHACK, 0>
   public:
```

template <typename T, int N>
void manipulate(test<T,N>& x)

{

```
// a friend can do anything!
};
```

You'll see that TMP sometimes makes use of techniques that are correctly labeled bad practice in conventional C++, including (but not limiting to):

- Lack of non-virtual protected destructor in (empty) base class
- Implementing cast operators operator T() const
- Declaring a non-explicit constructor with a single argument

1.6. Compiler Assumptions

Heavy usage of templates implies massive work for the compiler. Not all standard-conforming techniques behave identically on every *platform*.⁴⁷

You denote by *language-neutral idioms* all the language features that don't have a standard-prescribed behavior but only a reasonable expected behavior. In other words, when you use language-neutral idiom, you can expect that most compilers will converge on some (optimal) behavior, even if they are not demanded by the standard to do so.

Note For example, the C++ standard prescribes that sizeof(T) > 0 for any type T, but does not require the size of a compound type to be minimal. An empty struct could have size 64, but we expect it to have size 1 (or at worst, a size not larger than a pointer).

A standard-conforming compiler can legally violate the optimality condition, but in practice, such a situation is rare. In other words, a language-neutral idiom is a language construction that does not make a program worse, but gives a nice opportunity of optimization to a good compiler.

Several possible problems can arise from a perfect standard-conforming code fragment:

- Unexpected compiler errors
- Failures at runtime (access violations, core dumps, blue screens, and panic reactions)
- Huge compilation/link time
- Suboptimal runtime speed

The first two issues are due to compiler bugs and involve finding a language workaround (but the second one is usually met when it's too late).

The third problem mostly depends on poorly written template code.

The fourth problem involves finding language-neutral idioms that are not recognized by the optimizer and therefore unnecessarily slow down the program execution.

An example of expected behavior we do care about is the addition of an empty destructor to a base class.

```
class base
{
   public:
      void do_something() {}
   protected:
      ~base() {}
};
class derived : public base
{
};
```

Since the empty destructor adds no code, we expect the executable to be identical both with and without it.⁴⁸

The compiler will be assumed able to understand and deal optimally with the situations listed in the next paragraphs.

1.6.1. Inline

The compiler must be able to manage function inlining by itself, ignoring the inline directives and the code positioning (where the body of the member functions is written).

The all-inline style places definitions and declarations inside the body of the class; every member function is implicitly inline:

```
template <typename T>
class vector
{
public:
   bool empty() const
   {
       // definition and declaration
   }
};
```

The merged header style splits definitions and declarations of non-inline member functions, but keeps them in the same file:

```
template <typename T>
class vector
{
public:
  bool empty() const; // declaration, non inline
};
```

```
template <typename T>
bool vector <T>::empty() const
{
    // definition
}
```

In any case, whether you explicitly write it or not, the inline directive is just more than a hint. Some popular compilers indeed have an option to inline any function at the compiler's discretion. Specifically, we assume that

• A sequence of inline functions is always "optimal" if the functions are simple enough, no matter how long the sequence is:

```
template <typename T, int N>
class recursive
{
   recursive<T,N-1> r_;

public:
   int size() const
   {
     return 1 + r_.size();
   }
};

template <typename T>
class recursive<T, 0>
{
   public:
   int size() const
   {
     return 0;
   }
};
```

In the previous construction, recursive<T, N>::size() will be inlined and the optimizer will simplify the call down to return N. 49

• The compiler can optimize a call to a (const) member function of a stateless object, the typical case being binary relation's operator().

It's a common STL idiom to let a class hold a copy of a functor as a private member:

```
template <typename T>
struct less
{
  bool operator()(const T& x, const T& y) const
```

```
{
     return x<y;
};
template < typename T, typename less t = std::less<T> >
class set
                                // the less functor is a member
  less t less ;
public:
  set(const less t& less = less t())
    less (less)
  }
  void insert(const T& x)
     // ...
     if (less (x,y)) // invoking less t::operator()
     // ...
  }
};
  If the functor is indeed stateless and operator () is const, the previous code should be
equivalent to:
template <typename T>
struct less
  static bool apply(const T& x, const T& y)
     return x<y;
};
template < typename T, typename less t = std::less<T> >
class set
public:
  void insert(const T& x)
     // ...
     if (less t::apply(x,y))
     { }
  }
};
```

However, you pay for the greater generality since the <code>less_</code> member will consume at least one byte of space. You can solve both issues if the compiler implements the EBO (*empty base optimization*).

```
class stateless_base
{
};

class derived : public stateless_base
{
    // ...
};
```

In other words, any derivation from a stateless base will not make the derived class larger. ⁵⁰ If less is actually a stateless structure, the EBO will not add extra bytes to the layout of set.

```
template <typename T>
struct less
  bool operator()(const T& x, const T& y) const
     return x<y;
};
template < typename T, typename less t = std::less<T> >
class set : private less t
  inline bool less(const T& x, const T& y) const
     return static cast<const less t&>(*this)(x,y);
public:
  set(const less t& l = less t())
  : less t(1)
  {
  }
  void insert(const T& x)
     // ...
     if (less(x,y)) // invoking less t::operator() through
*this
     { }
};
```

Note the auxiliary member function less, which is intended to prevent conflicts with any other set::operator().

1.6.2. Error Messages

You would like a compiler to give precise and useful error diagnostics, especially when dealing with templates. Unfortunately, the meaning of "precise" and "useful" may not be the same for a human and a compiler.

Sometimes TMP techniques specifically induce the compiler to output a hint in the error message. The user, on the other hand, should be ready to figure out the exact error from some keywords contained in the compiler log, ignoring all the noise. Here's an example of noise:

```
\include\algorithm(21) : error 'void DivideBy10<T>::operator ()(T &)
const' : cannot convert parameter 1 from 'const int' to 'int &'
       with
       Γ
           T=int
       Conversion loses qualifiers
       iterator.cpp(41) : see reference to function template
instantiation ' Fn1
       std::for each<XT::pair iterator<iterator t,N>,DivideBy10<T>>
(_InIt,_InIt,_Fn1)'
        being compiled
       with
           Fn1= DivideBy10<int>,
          iterator t=std:: Tree<std:: Tmap traits<int, double, std::1</pre>
             <std::pair<const int,double>>,false>>::iterator,
          N=1,
           T=int,
          InIt=xT::pair_iterator<std:: Tree<std:: Tmap traits<int,dou</pre>
             std::allocator<std::pair<const</pre>
int,double>>,false>>::iterator,1>
  Here's what the user should see:
iterator.cpp(41): error in 'std::for each (iterator, iterator,
DivideBy10<int>)'
     with
```

iterator = XT::pair iterator<std::map<int,</pre>

'void DivideBy10<T>::operator ()(T &) const' : cannot convert parameter

double>::const iterator, 1>

1 from 'const int' to 'int &'

This means that the caller of for_each wants to alter (maybe divide by 10?) the (constant) keys of a std::map, which is illegal. While the original error points to <header>, the true problem is in iterator.cpp.

Unfriendly entries in error messages happen because the "bare bones error" that the compiler sees may be "distant" from the semantic error.

Long Template Stack

As shown previously, a function template can report an error, due to a parameter passed from its callers. Modern compilers will list the whole chain of template instantiations. Since function templates usually rely on template frameworks, these errors are often several levels deep in the stack of function calls.

Implementation Details

In the previous example, the compiler shows std::_Tree instead of std::map because map::iterator happens to be defined in a separate base class (named _Tree). std::map has a public typedef that borrows an iterator from its base class:

```
typedef typename Tree<...>::iterator iterator;
```

These implementation details, which are usually hidden from the user of std::map, may leak in the error log.

Expanded Typedefs

An error with std::string may show up as std::basic_string<char, ...> because some compilers will replace typedefs with their definition. The substitution may introduce a type that's unknown to the user.

However, it is truly impossible for the compiler to decide whether it's convenient or not to perform these substitutions.

Suppose there are two metafunctions called F<T1>::type and G<T2>::type:

```
typedef typename G<T>::type GT;
typedef typename F<GT>::type FGT;
```

An error may occur

• When T is not a valid argument for G, and in this case you'd like to read:

```
error "F<GT> [where GT=G<int>::type]...".
```

• Because G<T>::type (which is defined but unknown to the user) is rejected by F, so it may be more useful:

```
error "F<GT> [where GT=double]...".
```

However, if you don't know the result of G, a log entry such as F<X> [where X=double]... can be misleading (you may not even be aware that you are invoking F<double>).

Incomplete Types

If wisely used, an incomplete type can cause a specific error (see Section 2.2). However, there are situations where a type is *not yet* complete and this may cause bizarre errors. A long, instructive example is in Appendix A.

As a rule, when a compiler says that "a constant is not a constant" or that "a type is not a type," this usually means that you are either defining a constant recursively or are using a not-yet-complete class template.

1.6.3. Miscellaneous Tips

Regardless of assumptions, real compilers can do any sort of things, so this section outlines a few generic tips.

Don't Blame the Compiler

Bugs can lie:

- In the code, with probability $(100-\epsilon)\%$
- In the optimizer, with probability slightly greater than $(\epsilon/2)\%$
- In the compiler, with probability less than $(\epsilon/2)\%$

Even problems that show up only in release builds are rarely due to optimizer bugs. There are some natural differences between debug and release builds, and this may hide some errors in the program. Common factors are #ifdef sections, uninitialized variables, zero-filled heap memory returned by debug allocators, and so on.

Compilers do have bugs, but a common misconception is that they show up only in release builds. The following code, compiled by MSVC7.1, produces the right values in release and not in debug:

```
#include <iostream>
int main()
{
  unsigned __int64 x = 47;
  int y = -1;
  bool test1 = (x+y)<0;
  x += y;</pre>
```

```
bool test2 = (x<0);
bool test3 = (x<0);
std::cout << test1 << test2 << test3; // it should print 000
return 0;</pre>
```

GCC4 in Mac OSX in debug builds does not warn the user that there are multiple main functions in a console program and it silently produces a do-nothing executable.⁵¹

Keep Warnings at the Default Level

A warning is just a guess. All compilers can recognize "idioms" that can be, with some probability, a symptom of human errors. The higher the probability is, the lower the warning level. Displaying top-level warnings is very unlikely to reveal an error, but it will flood the compiler log with innocuous messages.⁵²

Do Not Silence Warnings with "Dirty" Code Modifications

If some particular warning is annoying, legitimate, and probably not an error, don't modify the code. Place a compiler-specific #pragma disable-warning directive around the line. This will be useful to future code reviewers.

However, this solution should be used with care (a warning in a deeply-nested function template might generate many long, spurious entries in the compiler log).

One of the most dangerous warnings that should *not* be fixed is the "signed/unsigned comparison".

Many binary operations between mixed operands involve the promotion of both to unsigned, and negative numbers become positive and very large.⁵³ Compilers will warn in some—not all—of these situations.

```
bool f(int a)
{
   unsigned int c = 10;
   return ((a+5)<c);
}

test01.cpp(4): warning C4018: '<': signed/unsigned mismatch</pre>
```

The function returns true for $a \in \{-5,-4,...,4\}$. If you change c to int, the warning disappears, but the function will behave differently.

The same code in a metafunction produces no warning at all:

```
template <int A>
class BizarreMF
{
```

In real code, two situations are likely vulnerable to "signedness bugs":

• Updating a metafunction return type from enum to a static unsigned constant:

```
static const bool value = (A+5) < OtherMF<B>::value;
// unpredictable result: the type of OtherMF is
unknown / may vary
```

• Changing a container:

The C++ standard does not define univocally an integer type for array indices. If p has type T^* , then p[i] == *(p+i), so i should have type $ptrdiff_t$, which is signed. vector<T>::operator[] however takes an unsigned index.

To sum up, warnings are:

- Compiler specific
- Not related to code correctness (there exist both correct code that produces warnings and incorrect code that compiles cleanly)

Write code that produces the least warnings possible, but not less.

Maintain a Catalog of Compiler Bugs

This will be most useful when upgrading the compiler.

Avoid Non-Standard Behavior

This advice is in every book about C++, but we repeat it here. Programmers⁵⁴ tend to use their favorite compiler as the main tool to decide if a program is correct, instead of the C++ standard. A reasonable empirical criterion is to use two or more compilers, and if they disagree, check the standard.

Don't Be Afraid of Language Features

Whenever there's a native C++ keyword, function, or std:: object, you can assume that it's impossible to do better, unless by trading some features.⁵⁵

It's usually true that serious bottlenecks in C++ programs are related to a *misuse* of language features (and some features are more easily misused than others; candidates are virtual functions and dynamic memory allocation), but this does not imply that these features should be avoided.

Any operating system can allocate heap memory fast enough that a reasonable number of calls to operator new will go unnoticed. 56

Some compilers allow you to take a little memory from the stack via a function named alloca; in principle, alloca followed by a placement new (and an explicit destructor call) is roughly equivalent to new, but it incurs alignment problems. While the standard grants that heap memory is suitably aligned for any type, this does not hold for a stack. Even worse, building objects on unaligned memory may work by chance on some platforms and, totally unobserved, may slow down all data operations.⁵⁷

The opposite case is trading features. It is sometimes possible to do better than new under strong extra hypotheses; for example, in a single threaded program where the allocation/deallocation pattern is known:

Here you may hope to outperform the default new-based allocator, since two deletions are always followed by a single allocation. Roughly speaking, when this is handled by system new/delete, the operating system has to be notified that more memory is available in line #2, but line #3 immediately reclaims the same amount of memory back.⁵⁸

Think About What Users of Your Code Would Do

Human memory is not as durable as computer memory. Some things that may look obvious or easily deducible in classic C++ may be more difficult in TMP.

Consider a simple function such as:

```
size t find number in string(std::string s, int t);
```

You can easily guess that the function looks for an occurrence of the second argument within the first. Now consider:

```
template <typename T, typename S>
size_t find_number_in_string(S s, T t);
```

While this may look natural to the author (S stands for string, after all), we should consider some memory-helping tricks.

• Any IDE with code completion will show the argument names:

```
template <typename T, typename S>
size_t find_number_in_string(S str, T number);

template <typename NUMBER_T, typename STRING_T>
size_t find_number_in_string(STRING_T str, NUMBER_T number);
```

- Insert one line of comment in the code before the function; an IDE could pick it up and show a tooltip.
- Adopt some convention for the order of arguments, or the result type (like C's memcpy).

1.7. Preprocessor

1.7.1. Include Guards

As already mentioned, a project is usually spread across many source files. Each file must be organized such that all dependencies and prerequisites are checked by the included file, not by the caller. In particular, header inclusion should never depend on the order of #include statements.

```
// internally uses std::vector...
};
```

Most frameworks end up having a sort of root file that takes care of preparing the environment:

- Detection of the current platform
- Translation of compiler-specific macros to framework macros
- Definition of general macros (such as MXT NAMESPACE BEGIN)
- Inclusion of STL headers
- Definition of lightweight structures, typedefs, and constants

All other headers begin by including the root file, which is rarely modified. This will often decrease compilation time, since compilers can be instructed to distill a pre-compiled header from the root file.

An example follows:

```
// platform detection
#if defined( MSC VER)
#define MXT INT64
            int64
#elif defined( GNUC )
#define MXT IN\overline{164} long long
#else
// ...
#endif
// macro translation
// the framework will rely on MXT DEBUG and MXT RELEASE
#if defined(DEBUG) || defined(DEBUG) || !defined(NDEBUG)
#define MXT DEBUG
#else
#define MXT RELEASE
#endif
// general framework macros
#define MXT NAMESPACE BEGIN(x)
                     namespace x {
#define MXT NAMESPACE END(x)
                      }
// STL
#include <complex>
```

According to the basic *include guard* idiom, you should enclose each header in preprocessor directives, which will prevent multiple inclusions in the same translation unit:

```
#ifndef MXT_filename_
#define MXT_filename_

// put code here
#endif //MXT filename
```

As a small variation of this technique, you can assign a value to MXT_filename_. After all, the whole point of this book is storing information in unusual places:

```
#ifndef MXT_filename_
#define MXT_filename_ 0x1020  // version number

// put code here

#endif //MXT_filename_
#include "filename.hpp"

#if MXT_filename_ < 0x1010
#error You are including an old version!
#endif</pre>
```

Anyway, such a protection is ineffective against inclusion loops. Loops happen more frequently in TMP, where there are only headers and no *.cpp file, so declarations and definitions either coincide or lie in the same file.

Suppose A. hpp is self-contained, B. hpp includes A. hpp, and C. hpp includes B. hpp.

```
// file "A.hpp"
```

```
#ifndef MXT A
#define MXT A 0x1010
template <typename T> class A {};
#endif
// file "B.hpp"
#ifndef MXT B
#define MXT B 0x2020
#include "A.hpp"
#endif
  Later, a developer modifies A. hpp so that it includes C. hpp.
// file "A.hpp"
#ifndef MXT A
#define MXT A 0x1020
#include "C.hpp"
  Now unfortunately, the preprocessor will produce a file that contains a copy of B before A:
// MXT A is not defined, enter the #ifdef
#define MXT A 0x1020
// A.hpp requires including "C.hpp"
  // MXT C is not defined, enter the #ifdef
  #define MXT C 0x3030
  // C.hpp requires including "B.hpp"
     // MXT B is not defined, enter the #ifdef
     #define MXT B 0x2020
     // B.hpp requires including A.hpp
     // however MXT A is already defined, so do nothing!
     template <typename T> class B {};
     // end of include "B.hpp"
  template <typename T> class C {};
```

```
// end of include "C.hpp"
template <typename T> class A {};
```

This usually gives bizarre error messages.

To sum up, you should detect circular inclusion problems where a file includes (indirectly) a copy of itself before it has been fully compiled.

The following skeleton header helps (indentation is for illustration purposes only).

```
#ifndef MXT filename
#define MXT filename 0x0000 // first, set version to "null"
 #include "other_header.hpp"
 MXT NAMESPACE BEGIN (framework)
 // write code here
 MXT NAMESPACE END(framework)
 // finished! remove the null guard
 #undef MXT filename
 // define actual version number and quit
 #define MXT filename 0x1000
#else
                      // if guard is defined...
 #if MXT filename == 0x0000
                      // ...but version is null
 #error Circular Inclusion
                      // ...then something is wrong!
 #endif
#endif //MXT filename
```

Such a header won't *solve* circular inclusion (which is a design problem), but the compiler will *diagnose* it as soon as possible. Anyway, sometimes it might suffice to replace the #error statement with some forward declarations:

```
#ifndef MXT_my_vector_
#define MXT_my_vector_ 0x0000

template <typename T>
  class my_vector
{
    public:
    // ...
```

```
# #undef MXT_my_vector_
# #define MXT_my_vector_ 0x1000

#else

#if MXT_my_vector_ == 0x0000

template <typename T>
class my_vector;
#endif

#endif //MXT my vector
```

1.7.2. Macro Expansion Rules

A smart use of macros can simplify metaprogramming tasks, such as automation of member function generation. We briefly mention the non-obvious preprocessor rules here⁵⁹:

• The token-concatenation operator ## produces one single token from the concatenation of two strings. It's not just a "whitespace elimination" operator. If the result is not a single C++ token, it's illegal:

```
#define M(a,b,c) a ## b ## c
int I = M(3,+,2); // error, illegal: 3+2 is not
a single token
int J = M(0,x,2); // ok, gives 0x2
```

- The stringizer prefix # converts text⁶⁰ into a *valid* corresponding C++ string, thus it will insert the right backslashes, and so on.
- Generally macro expansion is recursive. First, arguments are completely expanded, then they are substituted in the macro definition, then the final result is again checked and possibly expanded again:

• The two operators # and ##, however, *inhibit macro expansion* on their arguments, so:

```
Z(B, A1); // expands to BA1, not to B100
```

• To make sure everything is expanded, you can add an additional level of indirection that apparently does nothing:

```
#define Y(a,b) a ## b
#define Z(a,b) Y(a,b)

Z(B,A1);
// expands first to Y(B,A1). Since neither B nor A1
is an operand
// of # or ##, they are expanded, so we get
Y(B,100),
// which in turn becomes B100
```

• Macros cannot be recursive, so while expanding Z, any direct or indirect reference to Z is not considered:

```
#define X Z
#define Z X+Z

Z;
// expands first as X+Z. The second Z is ignored;
then the first X
// is replaced by Z, and the process stops,
// so the final result is "Z+Z"
```

• A popular trick is to define a macro as itself. This is practically equivalent to an #undef, except that the macro is still defined (so #ifdef and similar directives don't change behavior).

#define A A

As a final recap:

Observe that, in this code, X may look just as a convenience shortcut for X2, but *it's not*. Normally you cannot observe the difference, but before X expands to X2, argument expansion occurs, something that direct invocation of X2 could have prevented.

How safe is it to replace a macro that defines an integer with a constant (either enum or static const int)? The answer is in the previous code snippet. After the change, preprocessor tricks will break:

```
//#define A 1
static const int A = 1;

// ...

X(A); // const char* cA = "A";
Y(A); // const char* cA = "A";
```

But if A is not *guaranteed* to be a macro, the replacement should be transparent.⁶¹

One more rule that is worth mentioning is that the preprocessor respects the distinction between macros with and without arguments. In particular it will not try to expand A followed by an open bracket, and similarly for X not followed by a bracket. This rule is exploited in a popular idiom that prevents construction of unnamed instances of a type C^{62} :

Finally, since many template types contain a comma, it's not generally possible to pass them safely through macros:

There are several workarounds for this:

• Extra brackets (as a rule, this is *unlikely* to work, as in C++ there's not much use for a type in brackets):

```
DECLARE_x_OF_TYPE((std::map<int, double>));
// > (std::map<int, double>) x; > error
```

• A typedef will work, unless the type depends on other macro arguments:

```
typedef std::map<int, double> map_int_double;
DECLARE_x_OF_TYPE(map_int_double);
```

• Another macro:

¹Loosely speaking, that's the reason for the "meta" prefix in "metaprogramming".

²In modern C++ there are more, but you can consider them extensions; the ones described here are metaprogramming first-class citizens. Chapter 12 has more details.

³Usually any integer type is accepted, including named/anonymous enum, bool, typedefs (like ptrdiff_t and size_t), and even compiler-specific types (for example, __int64 in MSVC). Pointers to member/global functions are allowed with no restriction; a pointer to a variable (having external linkage) is legal, but it cannot be dereferenced *at compile time*, so this has very limited use in practice. See Chapter 11.

⁴The linker may eventually collapse them, as they will likely produce identical machine code, but from a language perspective they are different.

⁵An exception being that literal 0 may not be a valid pointer.

⁶See Sections 1.3.6 and **11.2.2** for more complete discussions.

⁷See the note in Section 1.3.2.

⁸You can cast a floating-point literal to integer, so strictly speaking, (int) (1.2) is allowed. Not all compilers are rigorous in regard to this rule.

⁹The use of __LINE__ as a parameter in practice occurs rarely; it's popular in automatic type enumerations (see Section 7.6) and in some implementation of custom assertions.

 $^{^{10}}$ We have to choose a different name, to avoid shadowing the outer template parameter scalar t.

¹¹ See also http://www.open-std.org/jtc1/sc22/wg21/docs/cwg_defects.html#666.

¹² Even if it's not a correct example, an open-minded reader may want to consider the relationship between std::string, std::wstring, and std::basic_string<T>.

¹³See 1.4.9.

¹⁴As a side note, this shows once more that in TMP, the less code you write, the better.

¹⁵ Compare with the use of typename described in Section 1.1.1.

 $^{^{16}}$ See the "brittle base class problem" mentioned by Bjarne Stroustrup in his "C++ Style and Technique FAQ" at http://www.research.att.com/~bs/.

¹⁷The exact rules are documented and explained in [2]. You're invited to refer to this book for a detailed explanation of what's summarized here in a few paragraphs.

¹⁸See the next section.

¹⁹This example is taken from [2].

 $^{^{20}}$ In particular, the compiler is not required to notice that void f(arg<2*N>) and void f(arg<N+N>) are the same template

- function, and such a double definition would make a program ill-formed. In practice, however, most compilers will recognize an ambiguity and emit an appropriate error.
- ²¹Template functions cannot be partially specialized, but only overloaded.
- ²²Unfortunately, some popular compilers tolerate this.
- ²³Consider the simpler case when outer<T> is a container, inner1 is an "iterator," inner2 is "const_iterator," and they both derive from an external common base, basic outer iterator.
- ²⁴The extra "Y" is little more than poetic license. Refer to the excellent article from Danny Kalev at http://www.informit.com/guides/content.aspx?g=cplusplus&seqNum=454.
- 25Except std::numeric limits<T>::quiet NaN().
- ²⁶Even source code has a lifecycle and eventually it's going to "die," i.e., it will be rewritten from scratch. However the more robust the design, the longer its life will be, and style is part of the design. See also [5].
- As discussed in [5], usually member function names should be actions. Thus empty should be a synonym for make_empty and not for is_empty. However, STL convention is established and universally understood. When in doubt, do as std::vector does.
- ²⁸See http://en.cppreference.com/w/cpp/language/constexpr for the exact requirements and specifications.
- ²⁹The difference between the last two implementations is largely how they react to an invalid syntax. As an exercise, consider some malicious code like MXT_SORT2(x, y) if (true) throw an_exception;
- ³⁰The reader might want to review the simple example early in this chapter.
- ³¹See Section **6.2.1**.
- ³²Some compilers, such as MSVC71, used to have problems with unnamed parameters; refer to paragraph 11.3.3 for a detailed example.
- 33 Some authors reserve the keyword typename for this purpose. In other words, they declare template <typename T> to mean that T is "any type" and template <class T> to suggest that T is indeed a class as opposed to a native type. However, this distinction is rather artificial.
- ³⁴Note that this is not a formal requirement; it's just a name! The name reflects how we think the type should be; later we will enforce this, if necessary.
- ³⁵The mathematically inclined reader should consider the latter as a special case of the former. The constant 5' can be replaced by a type named five or static_value<int, 5>. This leads to greater generality. See [3] for more information.
- ³⁶It should reside in an anonymous namespace, but this does not make it inaccessible.
- ³⁷The advantages are described extensively in Apple Technical Note TN2185; refer to the following page: http://developer.apple.com/technotes/tn2007/tn2185.html.
- ³⁸Such functions are denoted shims in [5].
- ³⁹If a is an array of T of length 2, then (char*) (&a[1]) (char*) (&a[0]) is a ptrdiff_t, which is at least as large as sizeof(T). That means ptrdiff_t is at least as large as int as well. This argument actually shows that every result of sizeof can be stored in a ptrdiff_t. A generic size_t may not be stored in a ptrdiff_t, because sizeof is not necessarily surjective—there may be a size_t value that is larger than every possible sizeof.
- ⁴⁰For example, to create a copy of std::string takes time proportional to the length of the string itself, so this depends not only on the type, but also on the *instance*; alternatively, copying a double is a constant-time operation. Mathematically speaking, the notion of "constant time" is not well defined in C++; the issue is too complex for a footnote, but we'll sketch the idea. An algorithm is O(1) if its execution time is bounded by a constant K, for any possible input. If the number of possible inputs is finite, even if it's huge, the algorithm is automatically O(1). For example, in C++ the sum of two int is O(1). In general, the C++ memory model has a finite addressable space (because all objects have a fixed size, and an "address" is an object) and this implies that the number of possible inputs to some algorithms is finite. Quicksort complexity is O(N*log(N)), but std::sort may be formally considered O(1), where—loosely speaking—the constant K is the time required to sort the largest possible array.
- ⁴¹Compare with Section **2.3.1**.
- ⁴²Some objects may want to check in advance if overwrite is feasible. For example, if T is std::string whose

- size() == that.size() then it might be able to perform a safe memcpy.
- ⁴³Note that the best option is to demand that the paired objects provide suitable operators, so we delegate the comparison. For example, pair<const char*, int> and pair<std::string, int> are unlikely to trigger the construction of temporary strings, because we expect the STL to supply an operator== (const char*, const std::string&).
- ⁴⁴The definition of "aggregate" changed in C++11, where uniform initialization was introduced. As the issue is quite complex and detailed, readers may want to see the bibliography.
- ⁴⁵There may be a huge cost in increased complexity that comes from writing code "totally bulletproof". Sometimes this complexity will also inhibit some compiler optimizations. As a rule, programmers should always reason pragmatically and accept the fact that code will not handle every possible corner case.
- ⁴⁶See Section 1.6.
- ⁴⁷By platform, usually we mean the set { processor, operating system, compiler, linker }.
- ⁴⁸From empirical analysis, it looks like sometimes a protected empty destructor inhibits optimizations. Some measurements have been published in [3].
- 49 Note that recursive<T, -1> will not compile.
- ⁵⁰Most compilers implement this optimization, at least in the case of single inheritance.
- ⁵¹Mac OS X 10.4.8, XCode 2.4.1, GCC 4.01.
- ⁵²Set warnings at maximum level only once, in the very last development phase or when hunting for mysterious bugs.
- ⁵³See 3.7.2 in the standard.
- ⁵⁴Including the author of this book.
- ⁵⁵Of course, there are known exceptions to this rule: some C runtime functions (sprintf, floor) and even a few STL functions (string::operator+).
- ⁵⁶Releasing memory may be a totally different matter, anyway.
- ⁵⁷On AMD processors, double should be aligned to an 8-byte boundary; otherwise, the CPU will perform multiple unnecessary load operations. On different processors, accessing an unaligned double may instantly crash the program.
- ⁵⁸In an empirical test on a similar algorithm, a map with a custom allocator improved the whole program by 25%. A general strategy is to reserve memory in chunks and free them with some degree of laziness.
- ⁵⁹For a complete reference, consider the GNU manual http://gcc.gnu.org/onlinedocs/cpp.pdf.
- ⁶⁰It can be applied only to a macro argument, not to arbitrary text.
- 61 Some C libraries, for example, list all the possible error codes without specifying the exact nature of these constants. In this case, they should be used as enums. In particular, it should be safe to undefine them, if they happen to be macros, and to replace them with real enumerations.
- ⁶²This example will actually be clear only after reading Section 2,2.

CHAPTER 2

Small Object Toolkit

The previous chapter focused on the connection between template programming and style. In short, templates are elegant, as they allow you to write efficient code that looks simple because they hide the underlying complexity.

If you recall the introductory example of sq from Chapter 1, it's clear that the first problem of TMP is choosing the best C++ entity that models a concept and makes the code look clear at the point of instantiation.

Most classic functions use internally temporary variables and return a result. Temporary variables are cheap, so you must give the intermediate results a name to increase the readability of the algorithm:

```
int n_dogs = GetNumberOfDogs();
int n_cats = GetNumberOfCats();
int n_food_portions = n_dogs + n_cats;
BuyFood(n food portions);
```

In TMP, the equivalent of a temporary variable is an auxiliary type.

To model a concept, we will freely use lots of different types. Most of them do nothing except "carry a meaning in their name," as in n food portions in the previous example.

This is the main topic of Section 2.3.

The following paragraphs list some extremely simple objects that naturally come up as building blocks of complex patterns. These are called "hollow," because they carry no data (they may have no members at all). The code presented in this chapter is freely reused in the rest of the book.

2.1. Hollow Types

2.1.1. instance_of

One of the most versatile tools in metaprogramming is instance_of:

```
template <typename T>
```

```
struct instance of
   typedef T type;
   instance of (int = 0)
};
   The constructor allows you to declare global constants and quickly initialize them.
const instance of<int> I INT = instance of<int>(); // ok but
cumbersome
const instance of<double> I DOUBLE = 0;
                                                               // also fine.
Note Remember that a const object must either be explicitly initialized or have a user-defined
default constructor. If you simply write
struct empty
   empty() {}
};
const empty EMPTY;
the compiler may warn that EMPTY is unused. A nice workaround to suppress the warning is in fact:
struct empty
```

2.1.2. Selector

};

A traditional code in classic C++ stores information in variables. For example, a bool can store two different values. In metaprogramming, all the information is contained in the type itself, so the equivalent of a bool is a (template) type that can be instantiated in two different ways. This is called a selector:

```
template <bool PARAMETER>
struct selector
```

 $empty(int = 0) {}$

const empty EMPTY = 0;

```
typedef selector<true> true_type1;
typedef selector<false> false_type;
```

Note that all instances of selector<true> convey the same information. Since their construction is inexpensive, instance_of and selector are both useful to replace explicit template parameter invocation:

```
template <bool B, typename T>
void f(const T& x)
}
int main()
  double d = 3.14;
  f<true>(d);
                                       // force B=true and deduce
T=double
};
  Or equivalently:
template <typename T, bool B>
void f(const T& x, selector<B>)
int main()
  double d = 3.14;
  f(d, selector<true>());
                                      // deduce B=true and T=double
};
```

One of the advantages of the latter implementation is that you can give a meaningful name to the second parameter, using a (cheap) constant:

```
const selector<true> TURN_ON_DEBUG_LOGGING;
// ...
double d = 3.14;
f(d, TURN_ON_DEBUG_LOGGING); // deduce B=true and T=double
```

2.1.3. Static Value

The generalization of a selector is a static value:

```
template <typename T, T VALUE>
```

```
struct static_parameter
{
};

template <typename T, T VALUE>
struct static_value : static_parameter<T, VALUE>
{
   static const T value = VALUE;
};
```

Note that you could replace selector with static_value<bool, B>. In fact from now on, you can assume that the implementation of the latter is the same.²

In a static_value, T must be an integer type; otherwise, the static const initialization becomes illegal. Instead, in static_parameter, T can be a pointer (and VALUE can be a literal zero).

A member cast operator may be added to allow switching from a static constant to a runtime integer³:

```
template <typename T, T VALUE>
struct static_value : static_parameter<T, VALUE>
{
   static const T value = VALUE;
   operator T () const
   {
      return VALUE;
   }
   static_value(int = 0)
   {
   }
};
```

So you can pass an instance of static_value<int, 3> to a function that requires int. However, it's usually safer to write an external function:

```
template <typename T, T VALUE>
inline T static_value_cast(static_value<T, VALUE>)
{
   return VALUE;
};
```

2.1.4. Size of Constraints

The C++ standard does not impose strict requirements on the size of elementary types⁴ and compound types can have internal padding anywhere between members.

Given a type T, say you want to obtain another type, T2, whose sizeof is different. A very simple solution is:

```
template <typename T>
class larger_than
{
   T body_[2]; // private, not meant to be used
};
```

It must hold that $sizeof(T) < 2*sizeof(T) \le sizeof(larger_than < T>)$. However, the second inequality can be indeed strict, if the compiler adds padding (suppose T is char and any struct has a minimum size of four bytes).

The most important use of this class is to define two types (see Section 4.2.1):

```
typedef char no_type;
typedef larger than<no type> yes type;
```

Warning These definitions are not compatible with C++0x std::false_type and std::true_type, which instead are equivalent to static_value<bool, false> and static_value<bool, true>.

In practice, you can safely use char, whose size is 1 by definition, and ptrdiff_t (in most platforms a pointer is larger than one byte).

It is possible to declare a type having exactly size N (with N>0):

```
template <size_t N>
struct fixed_size
{
  typedef char type[N];
};
```

So that sizeof(fixed size<N>::type) == N.

Note that fixed size<N> itself can have any size (at least N, but possibly larger).

Remember that it's illegal to declare a function that returns an array, but a *reference* to an array is fine and has the same size⁵:

```
fixed_size<3>::type f();  // error: illegal
int three = sizeof(f());
fixed_size<3>::type& f();  // ok
int three = sizeof(f());  // ok, three == 3
```

2.2. Static Assertions

Static assertions are simple statements whose purpose is to induce a (compiler) error when a template parameter does not meet some specification.

I illustrate here only the most elementary variations on the theme.

The simplest form of assertion just *tries to use* what you require. If you need to ensure that a type T indeed contains a constant named value or a type named type, you can simply write:

```
template <typename T>
void myfunc()
{
   typedef typename T::type ERROR_T_DOES_NOT_CONTAIN_type;
   const int ASSERT_T_MUST_HAVE_STATIC_CONSTANT_value(T::value);
};
```

If T is not conformant, you will get an error pointing to a sort of "descriptive" line.

For more complex assertions, you can exploit the fact that an incomplete type cannot be constructed, or that sizeof(T) causes a compiler error if T is incomplete.

2.2.1. Boolean Assertions

The easiest way to verify a statement is to use a selector-like class whose body is not present if the condition is false:

```
template <bool STATEMENT>
struct static_assertion
{
};

template <>
struct static_assertion<false>;

int main()
{
    static_assertion<sizeof(int) == 314> ASSERT_LARGE_INT;
    return 0;
}

error C2079: 'ASSERT_LARGE_INT' uses undefined struct
'static_assertion<false>'
```

All variations on the idiom try to trick the compiler into emitting more user-friendly error messages. Andrei Alexandrescu has proposed some enhancements. Here's an example.

```
template <bool STATEMENT>
```

```
struct static assertion;
template <>
struct static assertion<true>
   static assertion()
   { }
   template <typename T>
   static assertion(T)
   { }
};
template <> struct static assertion<false>;
struct error CHAR IS UNSIGNED { };
int main()
  const static assertion<sizeof(double)!=8> ASSERT1("invalid
double");
  const static assertion<(char(255)>0)>
ASSERT2 (error CHAR IS UNSIGNED());
   If the condition is false, the compiler will report something like, "cannot build
static assertion<false> from error CHAR IS UNSIGNED".
   Each assertion wastes some bytes on the stack, but it can be wrapped in a macro directive using
sizeof:
#define MXT ASSERT(statement)
sizeof(static assertion<(statement)>)
   The invocation
   MXT ASSERT (sizeof (double) !=8);
will translate to [[some integer]] if successful and to an error otherwise. Since a statement
like 1 is a no-op, the optimizer will ignore it.
   The very problem with macro assertions is the comma:
   MXT ASSERT(is well defined< std::map<int, double> >::value);
   //
   //
                                             comma here
   //
   // warning or error! MXT ASSERT does not take 2 parameters
   The argument of the macro in this case is probably the string up to the first comma
(is well defined < std::map<int), so even if the code compiles, it won't behave as
```

intended.

Two workarounds are possible—you can either typedef away the comma or put extra brackets around the argument:

```
typedef std::map<int, double> map_type;
MXT_ASSERT( is_well_defined<map_type>::value );
or:
MXT_ASSERT(( is_well_defined< std::map<int, double> >::value ));
```

The C++ preprocessor will be confused only by commas that are at the same level⁶ as the argument of the macro:

```
assert( f(x,y) == 4 ); // comma at level 2: ok
assert( f(x),y==4 ); // comma at level 1: error

static_assertion can be used to make assertions in classes using private inheritance:

template <typename T>
class small_object_allocator : static_assertion<(sizeof(T)<64)>
{
};
```

■ **Note** static_assert is a keyword in the modern C++ Standard. Here, I use a similar name for a class for illustration purposes. C++0x static_assert behaves like a function that takes a constant Boolean expression and a string literal (an error message that the compiler will print):

```
static assert(sizeof(T) < 64, "T is too large");</pre>
```

Similarly to the private inheritance described previously, C++0x static_assert can also be a class member.

2.2.2. Assert Legal

A different way of making assertions is to require that some C++ expression represents valid code for type T, returning non-void (most often, to state that a constructor or an assignment is possible).

If void is allowed instead, just put a comma operator inside sizeof:

For example:

```
template <typename T>
```

```
void do_something(T& x)
{
    MXT_ASSERT_LEGAL(static_cast<bool>(x.empty()));
    If (x.empty())
    {
        // ...
    }
}
```

This example will compile, and thus it will not reject T if x.empty(), whatever it means, returns (anything convertible to) bool. T could have a member function named empty that returns int or a member named empty whose operator() takes no argument and returns bool.

Here's another application:

```
#define MXT CONST REF TO(T)
                                 (*static cast<const T*>(0))
                                  (*static castT^*>(0))
#define MXT REF TO(T)
template <typename obj t, typename iter t>
class assert iterator
 enum
    verify construction =
      MXT ASSERT LEGAL(obj t(*MXT CONST REF TO(iter t))),
    verify assignment =
      MXT ASSERT LEGAL (MXT REF TO (obj t)
= *MXT CONST REF TO(iter t)),
    verify preincr =
       MXT ASSERT LEGAL (++MXT REF TO (iter t)),
    verify postincr =
       MXT ASSERT LEGAL (MXT REF TO (iter t) ++)
  };
};
```

A human programmer should read, "I assert it's legal to construct an instance of obj_t from the result of dereferencing a (const) instance of iter t" and similarly for the remaining constants.

Note Observe that some standard iterators may fail the first test. For example, a back_insert_iterator may return itself when dereferenced (a special assignment operator will take care of making *i = x equivalent to i = x).

The assert_iterator<T, I> will compile only if I acts like an iterator having a value type (convertible to) T. For example, if I does not support post-increment, the compiler will stop and

```
report an error in assert iterator<T, I>::verify postincr.
```

Remember that, with the usual restrictions on comma characters in macros, $\texttt{MXT_ASSERT_LEGAL}$ never instantiates objects. This is because sizeof performs only a dimensional check on its arguments⁷.

Also, note the special use of a macro directive. MXT_ASSERT_LEGAL should take the whole line, but since it resolves to a compile-time integer constant, you can use enums to "label" all the different assertions about a class (as in assert iterator) and make the code more friendly.

The compiler might also emit useful warnings pointing to these assertions. If obj_t is int and iter_t is double*, the compiler will refer to the verify_assignment enumerator and emit a message similar to:

Using the very same technique, you can mix static assertions of different kinds:

As an exercise, I list some more heuristic assertions on iterators.

As is, class assert_iterator validates forward const_iterators. We can remove the const-ness:

```
template <typename obj_t, typename iter_t>
class assert_nonconst_iterator : public assert_iterator<obj_t,
iter_t>
{
   enum
```

Sometimes, an algorithm that works on iterators does not need to know the actual type of the underlying objects, which makes the code even more general. For example, std::count could look like this:

```
template <typename iter_t, typename object_t>
int count(iter_t begin, const iter_t end, const object_t& x)
{
  int result = 0;
  while (begin != end)
  {
    if (*begin == x)
          ++result;
  }
  return result;
}
```

You don't need to know if *begin has the same type as x. Regardless of what exactly *begin is, you can assume that it defines an operator== suitable for comparing against an object t.

Suppose instead you have to store the result of *begin before comparison.

You may require the iterator type to follow the STL conventions, which means that object_t and iterator::value type must somehow be compatible⁸:

```
template <typename obj_t, typename iter_t>
class assert_stl_iterator
{
  typedef typename std::iterator_traits<iter_t>::value_type
  value_type;
  enum
  {
    assign1 =
        MXT_ASSERT_LEGAL(MXT_REF_TO(obj_t))
        = MXT_CONST_REF_TO(value_type)),
        assign2 =
            MXT_ASSERT_LEGAL(MXT_REF_TO(value_type))
        = MXT_CONST_REF_TO(obj_t))
        };
};
```

Finally, you can perform a rough check on the iterator type, using indicator traits to get

```
enum
{
    random_access =
        MXT_ASSERT_LEGAL(
        MXT_CONST_REF_TO(iter_t) + int() == MXT_CONST_REF_TO(iter_t))
```

2.2.3. Assertions with Overloaded Operators

its tag or writing operations with MXT ASSERT LEGAL:

sizeof can evaluate the size of an arbitrary expression. You can thus create assertions of the form sizeof(f(x)), where f is an overloaded function, which may return an incomplete type.

Here, I just present an example, but the technique is explained in Section 4.2.1.

Suppose you want to put some checks on the length of an array:

```
T arr[] = { ... };
// later, assert that length of(arr) is some constant
```

Since static assertions need a compile-time constant, you cannot define length_of as a function.

```
template <typename T, size_t N>
size_t length_of(T (&)[N])
{
   return N;
}

MXT_ASSERT(length_of(arr) == 7); // error: not a compile-time constant
```

A macro would work:

};

```
#define length of(a) sizeof(a)/sizeof(a[0])
```

But it's risky, because it can be invoked on an unrelated type that supports operator[] (such as std::vector or a pointer), with nasty implications.

However, you can write:

```
class incomplete_type;
class complete_type {};

template <size_t N>
struct compile_time_const
{
    complete_type& operator==(compile_time_const<N>) const;
```

```
template <size_t K>
  incomplete_type& operator==(compile_time_const<K>) const;
};

template <typename T>
compile_time_const<0> length_of(T)
{
    return compile_time_const<0>();
}

template <typename T, size_t N>
compile_time_const<N> length_of(T (&)[N])
{
    return compile_time_const<N>();
}
```

This works, but unfortunately the syntax of the assertion is not completely natural:

```
MXT ASSERT LEGAL(length of (arr) == compile time const<7>());
```

You can combine these techniques and the use of fixed_size<N>::type from Section 2.1.4, wrapping in an additional macro:

```
template <typename T, size_t N>
typename fixed_size<N>::type& not_an_array(T (&)[N]); // note: no
body

#define length of(X) sizeof(not an array(X))
```

Now length_of is again a compile-time constant, with some additional type-safety checks. The name not_an_array was chosen on purpose; it is usually hidden from the user, but it will usually be printed when the argument is incorrect:

```
class AA {};
int a[5];
int b = length_of(a);

AA aa;
int c = length_of(aa);
error: no matching function for call to 'not an array(AA&)'
```

2.2.4. Modeling Concepts with Function Pointers

The following idea has been documented by Bjarne Stroustrup.

A *concept* is a set of logical requirements on a type that can be translated to syntactic requirements.

For example, a "less-than comparable" type must implement operator < in some form. The exact signature of a < b doesn't matter as long as it can be used as a Boolean.

Complex concepts may require several syntactic constraints at once. To impose a complex constraint on a tuple of template parameters, you simply write a static member function, where all code lines together model the concept (in other words, if all the lines compile successfully, the constraint is satisfied). Then, you induce the compiler to emit the corresponding code simply by initializing a dummy function pointer in the constructor of a dedicated assertion class (the concept function never runs):

The concept check can be triggered when you're either building an instance on the stack or deriving from it:

```
template <typename T>
T sqrt(T x)
{
    static_assert_can_copy_T1_to_T2<T, double> CHECK1;
}
template <typename T>
class math_operations : static_assert_can_copy_T1_to_T2<T, double>
{};
```

2.2.5. Not Implemented

While C++0x allows you to "delete" member functions from a class, in classic C++, you'll sometimes want to express the fact that an operator should not be provided:

```
template <typename T>
class X
{
    // ...
```

```
X<T>& operator= (X<T>& that) { NOT_IMPLEMENTED; }
};
```

where the last statement is a macro for a static assertion that fails. For example:

The rationale for this idiom is that the member operator will be compiler-only on first use, which is exactly what you want to prevent.

However, this technique is risky and non-portable. The amount of diagnostics that a compiler can emit on unused template member function varies. In particular, if an expression does not depend on T, the compiler may legitimately try to instantiate it, so MXT ASSERT (false) may trigger anytime.

At least, the return type should be correct:

```
X<T>& operator= (X<T>& that) { NOT_IMPLEMENTED; return *this; }
```

A second choice is to make the assertion dependent on T:

```
#define NOT IMPLEMENTED MXT ASSERT(sizeof(T) == 0)
```

Finally, a portable technique is to cause a *linker* error with a fake annotation. This is less desirable than a compiler error, because linker errors usually do not point back to a line in source code. This means they are not easy to trace back.

```
#define NOT_IMPLEMENTED

X<T>& operator= (X<T>& that) NOT IMPLEMENTED;
```

2.3. Tagging Techniques

Assume you have a class with a member function called swap and you need to add a similar one called unsafe swap. In other words, you are adding a function that's a variation of an existing one. You can:

• Write a different function with a similar name and (hopefully) a similar signature:

```
public:
   void swap(T& that);
   void unsafe_swap(T& that);
```

• Add (one or more) overloads of the original function with an extra runtime argument:

```
private:
    void unsafe_swap(T& that);
public:
```

```
void swap(T& that);
enum swap_style { SWAP_SAFE, SWAP_UNSAFE };

void swap(T& that, swap_style s)
{
  if (s == SWAP_SAFE)
    this->swap(that);
  else
    this->unsafe_swap(that);
}
```

• Add an overload of the original function with an extra static *useless* argument:

```
public:
   void swap(T& that);
   void swap(T& that, int);   // unsafe swap: call as
x.swap(y, 0)
```

None of these options is completely satisfactory. The first is clear but does not scale well, as the interface could grow too much. The second may pay a penalty at runtime. The last is not intuitive and should be documented.

Instead, TMP makes heavy use of *language-neutral idioms*, which are language constructs that have no impact on code generation.

A basic technique for this issue is overload resolution via *tag objects*. Each member of the overload set has a formal unnamed parameter of a different static type.

```
class X
{
public:
    void swap(T& that);
    void swap(T& that, unsafe);
};

    Here's a different example:

struct naive_algorithm_tag {};
struct precise_algorithm_tag {};

template <typename T>
inline T log1p(T x, naive_algorithm_tag)
{
    return log(x+1);
}

template <typename T>
inline T log1p(T x, precise algorithm tag)
```

struct unsafe {};

```
{
   const T xp1 = x+1;
   return xp1==1 ? x : x*log(xp1)/(xp1-1);
}

// later...

double t1 = log1p(3.14, naive_algorithm_tag());
double t2 = log1p(0.0000000314, precise_algorithm_tag());
```

Building a temporary tag is inexpensive (most optimizing compilers will do nothing and behave as if you had two functions named log1p_naive and log1p_precise, with one parameter each).

So, let's dig a bit further into the mechanisms of overload selection.

Recall that you are facing the problem of picking the right function at compile time, supplying an extra parameter that's human-readable.

The extra parameter is usually an unnamed instance of an empty class:

```
template <typename T>
inline T log1p(T x, selector<true>);

template <typename T>
inline T log1p(T x, selector<false>);

// code #1
return log1p(x, selector<PRECISE_ALGORITHM>());
```

You might wonder why a type is necessary, when the same effect can be achieved with simpler syntax:

```
// code #2
if (USE_PRECISE_ALGORITHM)
  return log1p_precise(x);
else
  return log1p_standard(x);
```

The key principle in tag dispatching is that the program compiles only the functions that are strictly necessary. In code #1, the compiler sees one function call, but in the second fragment, there are two. The if decision is fixed, but is irrelevant (as is the fact that the optimizer may simplify the redundant code later).

In fact, tag dispatching allows the code to select between a function that works and one that would not even compile (see the following paragraph about iterators).

This does not imply that *every* if with a static decision variable must be turned into a function call. Typically, in the middle of a complex algorithm, an explicit statement is cleaner:

```
do_it();
do_it_again();
```

```
if (my_options<T>::need_to_clean_up)
{
   std::fill(begin, end, T());
}
```

2.3.1. Type Tags

The simplest tags are just empty structures:

```
struct naive_algorithm_tag {};
struct precise_algorithm_tag {};

template <typename T>
inline T log1p(T x, naive_algorithm_tag);

template <typename T>
inline T log1p(T x, precise_algorithm_tag);
```

You can use template tags to transport extra parameters to the function:

```
template <int N>
struct algorithm_precision_level {};

template <typename T, int N>
inline T log1p(T x, algorithm_precision_level<N>);

// ...

double x = log1p(3.14, algorithm precision level<4>());
```

You can use derivation to build a tag hierarchy.

This example sketches what actual STL implementations do (observe that inheritance is public by default):

```
void somefunc(iter_t begin, iter_t end,
bidirectional_iterator_tag)
{
    // do the work here
}
```

In this case, the bidirectional and random_access iterators will use the last overload of somefunc. Alternatively, if somefunc is invoked on any other iterator, the compiler will produce an error.

A generic implementation will process all the tags that do not have an exact match⁹:

```
template <typename iter_t, typename tag_t>
void somefunc(iter_t begin, iter_t end, tag_t)
{
    // generic implementation:
    // any tag for which there's no *exact* match, will fall here
}
```

This generic implementation can be made compatible with the tag hierarchy using pointers:

The overload resolution rules will try to select the match that loses less information. Thus, the cast derived*-to-base* is a better match than a cast to void*. So, whenever possible (whenever the iterator category is at least bidirectional), the second function will be taken.

Another valuable option is:

```
template <typename iter t>
```

```
void somefunc(iter_t begin, iter_t end, ...)
{
    // generic
}
```

The ellipsis operator is the worst match of all, but it cannot be used when the tag is a class (and this is exactly why you had to switch to pointers and tags).

2.3.2. Tagging with Functions

A slightly more sophisticated option is to use function pointers as tags:

```
enum algorithm_tag_t
{
    NAIVE,
    PRECISE
};
inline static_value<algorithm_tag_t, NAIVE> naive_algorithm_tag()
{
    return 0; // dummy function body: calls static_value<...>(int)
}
inline static_value<algorithm_tag_t, PRECISE>
precise_algorithm_tag()
{
    return 0; // dummy function body: calls static_value<...>(int)
}
```

The tag is not the return type, but the function itself. The idea comes somehow from STL stream manipulators (that have a common signature).

```
typedef
   static_value<algorithm_tag_t, NAIVE> (*naive_algorithm_tag_t)();

typedef
   static_value<algorithm_tag_t, PRECISE>
   (*precise_algorithm_tag_t)();

template <typename T>
inline T log1p(T x, naive_algorithm_tag_t);

// later
// line 4: pass a function as a tag

double y = log1p(3.14, naive_algorithm_tag);
```

Since each function has a different unique signature, you can use the function name (equivalent to

a function pointer) as a global constant. Inline functions are the only "constants" that can be written in header files without causing linker errors.

You can then omit brackets from the tags (compare line 4 above with its equivalent in the previous example). Function tags can be grouped in a namespace or be static members of a struct:

```
namespace algorithm_tag
{
  inline static_value<algorithm_tag_t, NAIVE> naive()
  { return 0; }
  inline static_value<algorithm_tag_t, PRECISE> precise()
  { return 0; }
}

or:
  struct algorithm_tag
{
   static static_value<algorithm_tag_t, NAIVE> naive()
   { return 0; }
   static static_value<algorithm_tag_t, PRECISE> precise()
   { return 0; }
};

double y = log1p(3.14, algorithm_tag::naive);
```

Another dramatic advantage of function pointers is that you can adopt a uniform syntax for the same runtime and compile-time algorithms:

```
enum binary_operation
          difference, product, division
  sum,
};
#define mxt SUM
                     x+y
#define mxt DIFF
                     X - A
#define mxt_DIFF x-y
#define mxt PROD x*y
#define mxt DIV
                     x/y
// define both the tag and the worker function with a single macro
#define mxt DEFINE(OPCODE,
FORMULA)
                                                \
inline static value <br/>
value <br/>
operation, OPCODE >
static tag ##OPCODE()
{
```

```
0;
template <typename
T>
T binary (T x, T y, static value < binary operation,
OPCODE>)
  return
(FORMULA);
mxt DEFINE (sum, mxt SUM);
mxt DEFINE(difference, mxt DIFF);
mxt DEFINE(product, mxt PROD);
mxt DEFINE (division, mxt DIV);
template <typename T, binary operation OP>
inline T binary (T x, T y, static value < binary operation, OP> (*)
())
{
  return binary(x, y, static value<binary operation, OP>());
```

return

This is the usual machinery needed for the static selection of the function. Due to the way you defined overloads, the following calls produce identical results (otherwise, it would be quite surprising for the user), even if they are not identical. The first is preferred:

```
double a1 = binary(8.0, 9.0, static_tag_product);
double a2 = binary(8.0, 9.0, static tag product());
```

However, with the same tools, you can further refine the function and add a similar runtime algorithm¹⁰:

```
}
```

The latter would be invoked as:

```
double a3 = binary(8.0, 9.0, product);
```

This may look similar, but it's a completely different function. It shares some implementation (in this case, the four kernel macros), but it selects the right one *at runtime*.

- Manipulators (see Section 1.4.7 are similar to functions used as compile-time constants. However, they differ in a few ways too:
- Manipulators are more generic. All operations have a similar signature (which must be supported by the stream object) and any user can supply more of them, but they involve some runtime dispatch.
- Function constants are a fixed set, but since there's a one-to-one match between signatures and overloaded operators, there is no runtime work.

2.3.3. Tag Iteration

A useful feature of functions tagged with static values is that, by playing with bits and compile-time computations, it's possible to write functions that automatically unroll some "iterative calls".

For example, the following function fills a C array with zeroes:

```
template <typename T, int N>
void zeroize_helper(T* const data, static_value<int, N>)
{
   zeroize_helper(data, static_value<int, N-1>());
   data[N-1] = T();
}

template <typename T>
void zeroize_helper(T* const data, static_value<int, 1>)
{
   data[0] = T();
}

template <typename T, int N>
void zeroize(T (&data)[N])
{
   zeroize_helper(data, static_value<int, N>());
}
```

You can swap two lines and iterate backward:

```
template <typename T, int N>
```

```
void zeroize_helper(T* const data, static_value<int, N>)
{
   data[N-1] = T();
   zeroize_helper(data, static_value<int, N-1>());
}
```

This unrolling is called *linear* and with two indices, you can have *exponential* unrolling. Assume for simplicity that N is a power of two:

```
template <int N, int M>
struct index
};
template <typename T, int N, int M>
void zeroize helper(T* const data, index<N, M>)
  zeroize helper(data, index<N/2, M>());
  zeroize helper(data, index<N/2, M+N/2>());
}
template <typename T, int M>
void zeroize helper(T* const data, index<1, M>)
  data[M] = T();
template <typename T, int N>
void zeroize(T (&data)[N])
  zeroize helper(data, index<N, 0>());
double test[8];
zeroize(test);
                <1,0>
           <2,0>
                <1,1>
      <4.0>
```

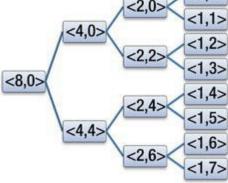


Figure 2-1. Exponential unrolling for N=8

As a more complex case, you can iterate over a set of bits.

Assume an enumeration describes some heuristic algorithms in increasing order of complexity:

```
enum
{
    ALGORITHM_1,
    ALGORITHM_2,
    ALGORITHM_3,
    ALGORITHM_4,
    // ...
};
```

For each value in the enumeration, you are given a function that performs a check. The function returns true when everything is okay or false if it detects a problem:

```
bool heuristic([[args]], static_value<size_t, ALGORITHM_1>);
bool heuristic([[args]], static_value<size_t, ALGORITHM_2>);
// ...
```

What if you wanted to run some or all of the checks, in increasing order, with a single function call?

First, you modify the enumeration using powers of two:

```
enum
{
    ALGORITHM_1 = 1,
    ALGORITHM_2 = 2,
    ALGORITHM_3 = 4,
    ALGORITHM_4 = 8,
    // ...
};
```

The user will use a static value as a tag, and algorithms will be combined with "bitwise or" (or +).

Here are the "private" implementation details:

```
template <size t K, size t J>
bool heuristic([[args]], VALUE(K), VALUE(J))
  static const size t JTH BIT = K & (size t(1) << J);
  // JTH BIT is either 0 or a power of 2.
  // try running the corresponding algorithm, first.
  // if it succeeds, the && will continue with new tags,
  // with the J-th bit turned off in K and J incremented by 1
  return
     heuristic([[args]], VALUE(JTH BIT)()) &&
     heuristic([[args]], VALUE(K-JTH BIT)(), VALUE(J+1)());
}
template <size t J>
bool heuristic([[args]], VALUE(0), VALUE(J))
  // finished: all bits have been removed from K
  return true;
}
template <size t K>
bool heuristic([[args]], VALUE(K))
  // this is invoked for all bits in K that do not have
  // a corresponding algorithm, and when K=0
  // i.e. when a bit in K is off
  return true;
}
```

#define VALUE(K) static value<size t, K>

2.3.4. Tags and Inheritance

Some classes inherit additional overloads from their bases. So an object that dispatches a tagged call might not know which of the bases will answer.

Suppose you are given a simple allocator class, which, given a fixed size, will allocate one block of memory of that length.

```
template <size_t SIZE>
struct fixed_size_allocator
{
   void* get_block();
};
```

You now wrap it up in a larger allocator. Assuming for simplicity that most memory requests have a size equal to a power of two, you can assemble a compound_pool<N> that will contain a fixed_size_allocator<J> for J=1,2,4,8. It will also resort to ::operator new when no suitable J exists (all at compile-time).

The syntax for this allocation is 11:

```
compound_pool<64> A;
double* p = A.allocate<double>();
```

The sketch of the idea is this. compound_pool<N> contains a fixed_size_allocator<N> and derives from compound_pool<N/2>. So, it can directly honor the allocation requests of N bytes and dispatch all other tags to base classes. If the last base, compound pool<0>, takes the call, no better match exists, so it will call operator new.

More precisely, every class has a *pick* function that returns either an allocator reference or a pointer.

The call tag is static_value<size_t, N>, where N is the size of the requested memory block.

```
template <size t SIZE>
class compound pool;
template < >
class compound pool<0>
protected:
  template <size t N>
  void* pick(static value<size t, N>)
     return :: operator new(N);
};
template <size t SIZE>
class compound pool : compound pool<SIZE/2>
  fixed_size_allocator<SIZE> p_;
protected:
  using compound pool<SIZE/2>::pick;
  fixed size allocator<SIZE>& pick(static value<SIZE>)
     return p ;
public:
  template <typename object t>
```

```
object_t* allocate()
{
    typedef static_value<size_t, sizeof(object_t)> selector_t;
    return static_cast<object_t*>(get_pointer(this-

>pick(selector_t()));
}

private:
    template <size_t N>
    void* get_pointer(fixed_size_allocator<N>& p)
{
        return p.get_block();
}

    void* get_pointer(void* p)
{
        return p;
    }
};
```

Note the using declaration, which makes all the overloaded pick functions in every class visible. Here, compound_pool<0>::pick has a lower priority because it's a function template, but it always succeeds. Furthermore, since it returns a different object, it ends up selecting a different get pointer.

¹Readers who are familiar with modern C++ will recognize that such a typedef already exists in namespace std. I will say more on this argument in Section 12.1.

²You could let selector derive from the other, but you can't assume explicitly that they are convertible. Under C++0x, you could also write a template typedef with the new using notation (see Section 12.6).

³See also Section 4.12.

⁴Only weak ordering is granted: 1=sizeof(char)≤sizeof(short)≤sizeof(int)≤sizeof(long).

⁵This remark will be clear in view of the material presented in Section 4.2.1.

⁶The level of a character is the number of open brackets minus the number of closed brackets in the string from the beginning of the line up to the character itself.

⁷However, a few compilers will generate a warning on MXT_INSTANCE_OF anyway, reporting that a null reference is not allowed.

⁸Actually, dereferencing the iterator returns std::iterator_traits<iterator_t>::reference, but value_type can be constructed from a reference.

⁹In particular, this will process random_access iterators as well. That is, it blindly ignores the base/derived tag hierarchy.

¹⁰This example anticipates ideas from Section 7.3.

¹¹Deallocation has been omitted on purpose.

PART 2

#include
#include <techniques>
#include <applications>

CHAPTER 3

Static Programming

Templates are exceptionally good at forcing the compiler and optimizer to perform some work only when the executable program is generated. By definition, this is called *static* work. This is as opposed to dynamic work, which refers to what is done when the program runs.

Some activities must be completed before runtime (computing integer constants) and some activities have an impact on runtime (generating machine code for a function template, which is later executed).

TMP can produce two types of code—metafunctions, which are entirely static (for example, a metafunction unsigned_integer<N>::type that returns an integer holding at least N bits) and mixed algorithms, which are part static and part runtime. (STL algorithms rely on iterator category or on the zeroize function explained in Section 4.1.2.

This section deals with techniques for writing efficient metafunctions.

3.1. Static Programming with the Preprocessor

The classic way to write a program that takes decisions about itself is through preprocessor directives. The C++ preprocessor can perform some integer computation tests and *cut off* portions of code that are not appropriate.

Consider the following example. You want to define fixed-length unsigned integer types, such as uint32_t, to be exactly 32-bits wide, and do the same for any bit length that's a power of two.

Define

```
template <size_t S>
struct uint_n;

#define mXT_UINT_N(T,N) \
   template <> struct uint_n<N> { typedef T type; }
```

and specialize uint_n for all sizes that are indeed supported on the current platform.

If the user tries uint_n<16>::type and there's no suitable type, she will get a proper and intelligible compiler error (about a missing template specialization).

So you have to ask the preprocessor to work out the sizes by trial and error¹:

```
#include <climits>
#define MXT I32BIT
                         0xffffffffU
#define MXT I16BIT
                         0xffffU
#define MXT I8BIT
                         0xffU
#if (UCHAR MAX == MXT I8BIT)
mXT UINT N (unsigned char, 8);
#endif
#if (USHRT MAX == MXT I16BIT)
mXT UINT N (unsigned short, 16);
\#elif UINT MAX == MXT I16BIT
mXT UINT N(unsigned int, 16);
#endif
#if (UINT MAX == MXT I32BIT)
mXT UINT N (unsigned int, 32);
\#elif (ULONG MAX == MXT I32BIT)
mXT UINT N(unsigned long, 32);
#endif
```

This code works, but it's rather fragile because interaction between the preprocessor and the compiler is limited.²

Note that this is not merely a generic style debate (macro versus templates), but a matter of correctness. If the preprocessor removes portions of the source file, the compiler does not have a chance to diagnose all errors until macro definitions change. On the other hand, if the TMP decisions rely on the fact that the compiler sees a whole set of templates, then it instantiates only some of them.

Note The preprocessor is not "evil".

Preprocessor-based "metaprogramming," like the previous example, usually compiles much faster and—if it's simple—it's highly portable. Many high-end servers still ship with old or custom compilers that do not support language-based (template) metaprogramming. On the other hand, I should mention that, while compilers tend to conform 100% to the standard, this is not true for preprocessors. Therefore, obscure preprocessor tricks may fail to produce the desired results, and bugs caused by misusing the preprocessor are quite hard to detect.³

An implementation of uint_n that does not rely on the preprocessor is shown and explained in Section 3.6.10.

3.2. Compilation Complexity

When a class template is instantiated, the compiler generates:

- Every member signature at class level
- All static constants and typedefs
- Only strictly necessary function bodies

If the same instance is needed again in the same compilation unit, it's found via lookup (which need not be particularly efficient, but it's still faster than instantiation).

For example, given the following code:

The initialization of n9 has a cost of 10 template instantiations, but the subsequent initialization of n8 has a cost of *one* lookup (not 9). Both instructions have zero runtime impact, as the assembly code shows.

As a rule, most metafunctions are implemented using recursion. The compilation complexity is the number of template instances recursively required by the metafunction itself.

This example has linear complexity, because the instantiation of X<N> needs X<N-1>... X<0>. While you'll usually want to look for the implementation with the lowest complexity (to reduce compilation times, not execution times), you can skip this optimization if there's a large amount of code reuse. Because of lookups, the first instantiation of X<N> will be costly, but it allows instantiation of X<M> for free in the same translation unit if M<N.

Consider this example of an optimized low-complexity implementation:

```
template <size_t N, size_t K>
struct static_raise
{
   static const size_t value = /* N raised to K */;
};
```

The trivial implementation has linear complexity:

```
template <size t N, size t K>
```

```
struct static_raise
{
   static const size_t value = N * static_raise<N, K-1>::value;
};

template <size_t N>
   struct static_raise<N, 0>
{
   static const size_t value = 1;
};
```

To obtain static_raise<N, K>::value, the compiler needs to produce K instances: static_raise<N, K-1>, static_raise<N, K-2>, . . .

Eventually static_raise<N, 1> needs static_raise<N, 0>, which is already known (because there's an explicit specialization). This stops the recursion.

However, there's a formula that needs only about log (K) intermediate types:

Note If the exponent is a power of two, you can save a lot of multiplications via repeated squaring. To compute X^8 , only three multiplications are needed if you can store only the intermediate results. Since $X^8 = ((X^2)^2)^2$, you need to execute

```
t = x*x; t = t*t; t = t*t; return t;
```

In general, you can use recursively the identity:

$$X^N = X^{N \mod 2} \cdot \left(X^{\lfloor N/2 \rfloor} \right)^2$$

```
public:
   static const size_t value = MXT_M_SQ(v0)*(K % 2 ? N : 1);
};
```

Note the use of MXT M SQ (see Section 1.3.2).

A final remark: Just because the natural implementation of metafunctions involves recursion, does not mean that *any* recursive implementation is equally optimal.⁴

Suppose N is an integer in base 10 and you want to extract the i-th digit (let's agree that digit 0 is the right-most) as digit<I, N>::value:

```
template <int I, int N>
struct digit;
```

template <int I>

Clearly, you have two choices. One is a "full" recursion on the main class itself

```
template <int I, int N>
struct digit
{
   static const int value = digit<i-1, N/10>::value;
};

template <int N>
struct digit<0, N>
{
   static const int value = (N % 10);
};
```

Or you can introduce an auxiliary class main class:

```
struct power_of_10
{
    static const int value = 10 * power_of_10<I-1>::value;
};

template <>
    struct power_of_10<0>
{
    static const int value = 1;
};

template <int I, int N>
    struct digit
{
    static const int value = (N / power_of_10<I>::value) % 10;
};
```

While the first implementation is clearly simpler, the second scales better. If you need to extract the 8th digit from 100 different random numbers, the former is going to produce 800 different

specializations because chances of reuse are very low. Starting with digit<8,12345678>, the compiler has to produce the sequence digit<7,1234567>, digit<6,123456>..., and each of these classes is likely to appear only once in the entire program.

On the other hand, the latter version produces eight different specialized powers of 10 that are reused every time, so the compiler workload is just 100+10 types.

3.3. Classic Metaprogramming Idioms

Metafunctions can be seen as functions that take one or more types and return types or constants. You'll see in this section how to implement some basic operations.

Binary operators are replaced by metafunctions of two variables. The concept T1==T2 becomes typeequal<T1, T2>::value:

```
template <typename T1, typename T2>
struct typeequal
{
   static const bool value = false;
};

template <typename T>
struct typeequal<T, T>
{
   static const bool value = true;
};
```

};

Whenever possible, you should derive from an elementary class that holds the result, rather than introduce a new type/constant. Remember that public inheritance is implied by struct

```
struct typeif<false, T1, T2>
   typedef T2 type;
};
  Or, according to the previous guideline:
template <bool STATEMENT, typename T1, typename T2>
struct typeif : instance of<T1>
};
template <typename T1, typename T2>
struct typeif<false, T1, T2>: instance of<T2>
};
   The strong motivation for derivation is an easier use of tagging techniques. Since you will often
"embed" the metafunction result in a selector, it will be easier to use the metafunction itself as a
selector. Suppose you have two functions that fill a range with random elements:
template <typename iterator t>
void random fill(iterator t begin, iterator t end, selector<false>)
   for (; begin != end; ++begin)
      *begin = rand();
}
template <typename iterator t>
void random fill(iterator t begin, iterator t end, selector<true>)
{
   for (; begin != end; ++begin)
      *begin = 'A' + (rand() % 26);
}
   Compare the invocation:
random fill(begin, end, selector<typeequal<T, char*>::value>());
with the simpler<sup>5</sup>:
random fill(begin, end, typeequal<T, char*>());
Note Note as a curiosity, that header files that store a version number in their guard macro can be
```

template <typename T1, typename T2>

used in a typeif. Compare the following snippets
#include "myheader.hpp"

```
typedef
typename typeif<MXT_MYHEADER_==0x1000, double,
float>::type float_t;

#if MXT_MYHEADER_ == 0x1000
typedef double float_t;
#else
typedef float float_t;
#endif
```

The first snippet will not compile if MXT_MYHEADER_ is undefined. The preprocessor instead would behave as if the variable were 0.

3.3.1. Static Short Circuit

As a case study of template recursion, let's compare the pseudo-code of a static and dynamic operator:

```
template <typename T>
struct F : typeif<[[CONDITION]], T, typename G<T>::type>
{
};
int F(int x)
{
   return [[CONDITION]] ? x : G(x);
}
```

These statements are *not* analogous:

- The runtime statement is short-circuited. It will not *execute* code unless necessary, so G(x) might never run.
- The static operator will always *compile* all the mentioned entities, as soon as one of their members is mentioned. So the first F will trigger the compilation of G<T>::type, regardless of the fact that the result is used (that is, even when the condition is true).

There is no automatic static short-circuit. If underestimated, this may increase the build times without extra benefits, and it may not be noticed, because results would be correct anyway.

The expression may be rewritten using an extra "indirection":

```
template <typename T>
```

```
struct F
{
   typedef
    typename typeif<[[CONDITION]], instance_of<T>, G<T> >::type
    aux_t;
   typedef typename aux_t::type type;
};
```

Here, only G<T> is mentioned, not G<T>::type. When the compiler is processing typeif, it needs only to know that the second and third parameters are valid types; that is, that they have been declared. If the condition is false, aux_t is set to G<T>. Otherwise, it is set to instance_of<T>. Since no member has been requested yet, nothing else has been compiled. Finally, the last line triggers compilation of either instance of<T> or G<T>.

So, if CONDITION is true, G<T>:: type is never used. G<T> may even lack a definition or it may not contain a member named type.

To summarize:

- Delay accessing members as long as possible
- Wrap items to leverage the interface

An identical optimization applies to constants:

```
static const size_t value = [[CONDITION]] ? 4
: alignment_of<T>::value;

typedef typename
  typeif<[[CONDITION]], static_value<size_t, 4>,
alignment_of<T>>::type
  aux_t;

static const size_t value = aux_t::value;
```

At first, it may look like there's no need for some special logic operator, since all default operators on integers are allowed inside of templates⁶:

The classic logical operators in C++ are short-circuited; that is, they don't *evaluate* the second operator if the first one is enough to return a result. Similarly, you can write a static OR that does not *compile* its second argument unnecessarily. If T1::value is true, T2::value is never accessed and it might not even exist (AND is obtained similarly).

```
// if (T1::value is true)
```

```
// return true;
// else
// return T2::value;

template <bool B, typename T2>
struct static_OR_helper;

template <typename T2>
struct static_OR_helper<false, T2> : selector<T2::value>
{
};

template <typename T2>
struct static_OR_helper<true, T2> : selector<true>
{
};

template <typename T2>
struct static_OR_helper<true, T2> : selector<true>
{
};

template <typename T1, typename T2>
struct static_OR : static_OR_helper<T1::value, T2>
{
};
```

3.4. Hidden Template Parameters

Some class templates may have undocumented template parameters, generally auto-deduced, that silently select the right specialization. This is a companion technique to tag dispatching, and an example follows:

The user of A will accept the default, as a rule:

```
A<char> c1;
```

```
A<char, true> c2; // exceptional case. do at own risk
   The following is a variation of an example that appeared in [3].
template <size t N>
struct fibonacci
   static const size t value =
      fibonacci<N-1>::value + fibonacci<N-2>::value;
};
template <>
struct fibonacci<0>
{
   static const size t value = 0;
};
template <>
struct fibonacci<1>
   static const size t value = 1;
};
   It can be rewritten using a hidden template parameter:
template \langle \text{size t N, bool TINY NUMBER} = (N < 2) \rangle
struct fibonacci
   static const size t value =
      fibonacci<N-1>::value + fibonacci<N-2>::value;
} ;
template <size t N>
struct fibonacci<N, true>
   static const size t value = N;
};
   To prevent the default from being changed, you can rename the original class by appending the
suffix helper and thus introducing a layer in the middle:
template <size t N, bool TINY NUMBER>
struct fibonacci helper
   // all as above
};
template <size t N>
class fibonacci : fibonacci helper<N, (N<2)>
```

```
{
};
```

3.4.1. Static Recursion on Hidden Parameters

Let's compute the highest bit of an unsigned integer x. Assume that x has type $size_t$ and, if x==0, it will conventionally return -1.

A non-recursive algorithm would be: set N = the number of bits of size_t; test bit N-1, then N-2..., and so on, until a non-zero bit is found.

First, as usual, a naive implementation:

As written, it works, but the compiler might need to generate a large number of different classes per static computation (that is, for any X, you pass to static highest bit).

First, you can rework the algorithm using bisection. Assume X has N bits, divide it in an upper and a lower half (U and L) having (N-N/2) and (N/2) bits, respectively. If U is 0, replace X with L; otherwise, replace X with U and remember to increment the result by $(N/2)^7$:

In pseudo-code:

```
size_t hibit(size_t x, size_t N = CHAR_BIT*sizeof(size_t))
{
    size_t u = (x>>(N/2));
    if (u>0)
        return hibit(u, N-N/2) + (N/2);
    else
        return hibit(x, N/2);
}
```

```
This means:
template <size t X, int N>
struct helper
   static const size t U = (X \gg (N/2));
   static const int value =
      U ? (N/2) +helper<U, N-N/2>::value : helper<X, N/2>::value;
};
   As written, each helper<X, N> induces the compiler to instantiate the template again twice—
namely helper<U, N-N/2> and helper<X, N/2>—even if only one will be used.
   Compilation time may be reduced either with the static short circuit, or even better, by moving all
the arithmetic inside the type.<sup>8</sup>
template <size t X, int N>
struct helper
   static const size t U = (X \gg (N/2));
   static const int value = (U ? N/2 : 0) +
            helper<(U ? U : X), (U ? N-N/2 : N/2)>::value;
};
   This is definitely less clear, but more convenient for the compiler.
   Since N is the number of bits of X, N>0 initially.
   You can terminate the static recursion when N==1:
template <size t X>
struct helper<X, 1>
   static const int value = X ? 0 : -1;
};
   Finally, you can use derivation from static value to store the result:
template <size t X>
```

```
template <size_t X>
struct static_highest_bit
: static_value<int, helper<X, CHAR_BIT*sizeof(size_t)>::value>
{
};
```

The recursion depth is fixed and logarithmic. static_highest_bit<X> instantiates at most five or six classes for every value of X.

3.4.2. Accessing the Primary Template

A dummy parameter can allow specializations to call back the primary template.

Suppose you have two algorithms, one for computing cos(x) and another for sin(x), where x is any floating-point type. Initially, the code is organized as follows:

```
template <typename float t>
struct trigonometry
  static float t cos(const float t x)
     // ...
  static float_t sin(const float t x)
    // ...
};
template <typename float t>
inline float t fast cos(const float t x)
  return trigonometry<float_t>::cos(x);
template <typename float t>
inline float t fast sin(const float t x)
  return trigonometry<float t>::sin(x);
  Later, someone writes another algorithm for cos<float>, but not for sin<float>.
  You can either specialize/overload fast cos for float or use a hidden template parameter,
as shown:
```

template <typename float t, bool = false>

static float t cos(const float t x)

static float t sin(const float t x)

struct trigonometry

// ...

// ...

};

```
template <>
struct trigonometry<float, false>
{
    static float_t cos(const float_t x)
    {
        // specialized algorithm here
    }
    static float_t sin(const float_t x)
    {
        // calls the general template
        return trigonometry<float, true>::sin(x);
    }
};
```

Note that in specializing the class, it's not required that you write <float, false>. You can simply enter:

```
template <>
struct trigonometry<float>
{
```

because the default value for the second parameter is known from the declaration.

Any specialization can access the corresponding general function by setting the Boolean to true explicitly.

This technique will appear again in Section 7.1.

A similar trick comes in handy to make partial specializations unambiguous.

C++ does not allow specializing a template twice, even if the specializations are identical. In particular, if you mix cases for standard typedefs and integers, the code becomes subtly non-portable:

```
template <typename T>
struct is_integer
{
    static const bool value = false;
};

template < > struct is_integer < short>
{    static const bool value = true; };

template < > struct is_integer < int>
{    static const bool value = true; };

template < > struct is_integer < long>
{    static const bool value = true; };

template < > struct is_integer < long>
{    static const bool value = true; };

template < > struct is_integer < ptrdiff_t> // problem:
{    static const bool value = true; }; // may or may not
```

```
compile
```

struct is integer<time t>

{ static const bool value = true; };

If ptrdiff_t is a fourth type, say long long, then all the specializations are different. Alternatively, if ptrdiff_t is simply a typedef for long, the code is incorrect. Instead, this works:

```
works:
template <typename T, int = 0>
struct is integer
  static const bool value = false;
};
template <int N> struct is integer<short, N>
{ static const bool value = true; };
template <int N> struct is integer<int</pre>
{ static const bool value = true; };
template <int N> struct is integer<long , N>
{ static const bool value = true; };
template <>
struct is integer<ptrdiff t>
  static const bool value = true;
};
  Since is integer<ptrdiff t, 0> is more specialized than is integer<long, N>,
it will be used unambiguously.<sup>9</sup>
  This technique does not scale well, <sup>10</sup> but it might be extended to a small number of typedefs,
by adding more unnamed parameters. This example uses int, but anything would do, such as bool
= false or typename = void.
template <typename T, int = 0, int = 0>
struct is integer
   static const bool value = false;
};
template <int N1, int N2>
struct is integer<long, N1, N2>
{ static const bool value = true; };
template <int N1>
struct is integer<ptrdiff t, N1>
{ static const bool value = true; };
template < >
```

3.4.3. Disambiguation

In TMP it's common to generate classes that derive several times from the same base (indirectly). It's not yet time to list a full example, so here's a simple one:

```
template <int N>
struct A {};

template <int N>
struct B : A<N % 2>, B<N / 2> {};

template <>
struct B<0> {};
```

For example, the inheritance chain for B < 9 > is illustrated in Figure 3-1.

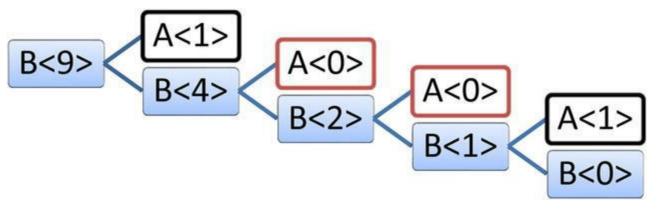


Figure 3-1. The inheritance chain for B<9>

Note that A<0> and A<1> occur several times. This is allowed, except that you cannot cast, explicitly or implicitly, B<9> to A<0> or A<1>:

```
template <int N>
struct A
{
   int getN() { return N; }
};

template <int N>
struct B : A<N % 2>, B<N / 2>
{
   int doIt() { return A<N % 2>::getN(); } // error: ambiguous
};
```

What you can do is add a hidden template parameter so that different levels of inheritance correspond to physically different types.

The most popular disambiguation parameters are counters:

```
template <int N, int FAKE = 0>
```

```
struct A {};

template <int N, int FAKE = 0>
struct B : A<N % 2, FAKE<sup>11</sup>>, B<N / 2, FAKE+1> {};

template <int FAKE>
struct B<0, FAKE> {};

B<9,0>
A<1,0>
B<4,1>
B<2,2>
B<1,3>
B<0,4>
```

Figure 3-2. The modified inheritance chain for B<9> using a counter

Another commonly used disambiguator tag is the type this:

Figure 3-3. The modified inheritance chain for B<9> using a tag-type

This idea is used extensively in Section 5.2

3.5. Traits

Traits classes (or simply, traits) are a collection of static functions, types, and constants that abstract the public interface of a type T. More precisely, for all T representing the same concept, traits<T> is a class template that allows you to operate on T uniformly. In particular, all traits<T> have the same public interface. 12

Using traits, it's possible to deal with type \mathbb{T} by ignoring partially or completely its public interface. This makes traits an optimal building layer for algorithms.

Why ignore the public interface of T? The main reasons are because it could have none or it could be inappropriate.

Suppose T represents a "string" and you want to get the length of an instance of T. T may be const char* or std::string, but you want the same call to be valid for both. Otherwise, it will be impossible to write template string functions. Furthermore, 0 may have a special meaning as a "character" for some T, but not for all.

The first rigorous definition of traits is an article by Nathan Myers, ¹³ dated 1995.

The motivation for the technique is that, when writing a class template or a function, you'll realize that some types, constants, or atomic actions are parameters of the "main" template argument.

So you could put in additional template parameters, but that's usually impractical. You could also group the parameters in a traits class. Both the next example and the following sentences are quotes from Myers' article¹⁴:

Because the user never mentions it, the [traits class] name can be long and descriptive.

```
template <typename char_t>
struct ios_char_traits
{
};

template <>
struct ios_char_traits<char>
{
    typedef char char_type;
    typedef int int_type;
    static inline int_type eof() { return EOF; }
};

template <>
struct ios_char_traits<wchar_t>
{
    typedef wchar_t char_type;
    typedef wint_t int_type;
    static inline int_type eof() { return WEOF; }
};
```

The default traits class template is empty. What can anyone say about an unknown character type? However, for real character types, you can specialize the template and provide useful semantics.

To put a new character type on a stream, you need only specialize ios_char_traits

for the new type.

Notice that ios_char_traits has no data members; it only provides public definitions. Now you can define the streambuf template:

```
template <typename char_t>
class basic streambuf
```

Notice that it has only one template parameter, the one that interests users.

In fact, Myers concludes his article with a formal definition and an interesting observation:

Traits class:

A class used in place of template parameters. As a class, it aggregates useful types and constants. As a template, it provides an avenue for that "extra level of indirection" that solves all software problems.

This technique turns out to be useful anywhere that a template must be applied to native types, or to any type for which you cannot add members as required for the template's operations.

Traits classes may be "global" or "local". Global traits are simply available in the system and they can be freely used anywhere. In particular, all specializations of a global traits class have system-wide scope (so specializations are automatically used everywhere). This approach is in fact preferred when traits express properties of the platform.

```
template <typename char_t>
class basic_streambuf
{
   typedef typename ios_char_traits<char_t>::int_type int_type;
   ...
};
```

■ **Note** For example, you could access the largest unsigned integer, of float, available. Consider the following pseudo-code:

```
template <typename T>
struct largest;

template <>
struct largest<int>
{
   typedef long long type;
};

template <>
```

```
struct largest<float>
{
  typedef long double type;
};

template <>
struct largest<unsigned>
{
  typedef unsigned long long type;
};
```

Evidently, a call such as largest<unsigned>::type is expected to return a result that's constant in the platform, so all customizations—if any—should be global to keep the client code coherent.

A more flexible approach is to use local traits, passing the appropriate type to each template instance as an additional parameter (which defaults to the global value).

```
template <typename char_t, typename traits_t
= ios_char_traits<char_t> >
class basic_streambuf
{
   typedef typename traits_t::int_type int_type;
   ...
};
```

The following sections focus on a special kind of traits—pure static traits, which do not contain functions but only types and constants. You will come back to this argument in Section 4.2.

3.5.1. Type Traits

Some traits classes provide typedefs only, so they are indeed multi-value metafunctions. As an example, consider again std::iterator traits.

Type traits¹⁵ are a collection of metafunctions that provide information about qualifiers of a given type and/or alter such qualifiers. Information can be deduced by a static mechanism inside traits, can be explicitly supplied with a full/partial specialization of the traits class, or can be supplied by the compiler itself.¹⁶

```
template <typename T>
struct is_const : selector<false>
{
};
```

```
template <typename T>
struct is_const<const T> : selector<true>
{
};
```

Note Today, type traits are split to reduce compile times, but historically they were large monolithic classes with many static constants.

```
template <typename T>
struct all_info_together
{
   static const bool is_class = true;
   static const bool is_pointer = false;
   static const bool is_integer = false;
   static const bool is_floating = false;
   static const bool is_unsigned = false;
   static const bool is_const = false;
   static const bool is_reference = false;
   static const bool is_volatile = false;
};
```

template <typename T>

As a rule, traits have a general implementation with conservative defaults, including partial specializations with meaningful values for classes of types and full specializations customized on individual types.

```
struct add_reference
{
   typedef T& type;
};

template <typename T>
struct add_reference<T&>
{
   typedef T& type;
};

template < >
struct add_reference<void>
{
   // reference to void is illegal. don't put anything here<sup>17</sup>
```

```
Traits are often recursive:
template <typename T>
struct is unsigned integer : selector<false>
};
template <typename T>
struct is unsigned integer < const T > : is unsigned integer < T >
};
template <typename T>
struct is unsigned integer<volatile T>: is unsigned integer<T>
};
template < >
struct is unsigned integer<unsigned int> : selector<true>
};
template < >
struct is unsigned integer < unsigned long> : selector < true>
};
// add more specializations...
  Traits can use inheritance and then selectively hide some members:
template <typename T>
struct integer traits;
template <>
struct integer traits<int>
  typedef long long largest type;
  typedef unsigned int unsigned type;
};
template <>
struct integer_traits<long> : integer_traits<int>
  // keeps integer traits<int>::largest type
  typedef unsigned long unsigned type;
};
```

};

```
template <typename T>
struct BASE
  typedef T type;
};
template <typename T>
struct DER : public BASE<T>
  type t; // error: 'type' is not in scope
};
However, from a static point of view, DER does contain a type member:
template <typename T>
struct typeof
  typedef typename T::type type;
};
typeof< DER<int> >::type i = 0;  // ok: int i = 0
  Type traits, if not carefully designed, are vulnerable to hard conceptual problems, as the C++ type
system is a lot more complex than it seems:
template <typename T>
struct is const : selector<false>
};
template <typename T>
struct is const<const T> : selector<true>
} ;
template <typename T>
struct add const : instance of<const T>
};
template <typename T>
```

Note In C++, a template base class is not in scope of name resolution:

```
struct add_const<const T> : instance_of<const T>
{
};
```

Here are some oddities:

- If N is a compile-time constant and T is a type, you can form two distinct array types: T [N] and T []. 18
- Qualifiers such as const applied to array types behave a bit oddly. If T is an array, for example, double [4], const T is an "array of four const double," not "const array of four double". In particular, const T is not const:

So, you should add more specializations:

```
template <typename T, size_t N>
struct is_const<const T [N]>
{
   static const bool value = true;
};

template <typename T >
struct is_const<const T []>
{
   static const bool value = true;
};
```

There are two possible criteria you can verify on types:

- A match is satisfied; for example, const int matches const T with T==int.
- A logical test is satisfied; for example, you could say that T is const if const T and T are the same type.

The C++ type system is complex enough that criteria may look equivalent in the majority of cases, but still not be identical. As a rule, whenever such a logical problem arises, the solution will come from more precise reasoning about your requirements. For any T, is_const<T&>::value is false because T& does not satisfy a match with a const type. However, add_const<T&>::type is again T& (any qualifiers applied to a reference are ignored). Does this mean that references are const?

Should you add a specialization of is_const<T&> that returns true? Or do you really want add const<T&>::type to be const T&?

In C++, objects can have different degrees of const-ness. More specifically, they can be

- Assignable
- Immutable
- const

Being assignable is a syntactic property. An assignable object can live on the left side of operator=. A const reference is not assignable. In fact, however, T& is assignable whenever T is. (Incidentally, an assignment would change the referenced object, not the reference, but this is irrelevant.)

Being *immutable* is a logical property. An immutable object cannot be changed after construction, either because it is not assignable or because its assignment does not alter the state of the instance. Since you cannot make a reference "point" to another object, a reference is immutable.

Being const is a pure language property. An object is const if its type matches const T for some T. A const object may have a reduced interface and operator= is *likely* one of the restricted member functions.

References are not the only entities that are both immutable and assignable. Such a situation can be reproduced with a custom operator=.

This also shows that const objects may be assignable, ¹⁹ but it does not imply that references are const, only that they can be simulated with const objects.

So the standard approach is to provide type traits that operate atomically, with minimal logic and just a match. is const<T&>::value should be false.

However, type traits are also easy to extend in user code. If an application requires it, you can introduce more concepts, such as "intrusive const-ness"

```
template <typename T>
struct is_const_intrusive : selector<false>
{
};

template <typename T>
struct is_const_intrusive<const T> : selector<true>
{
};

template <typename T>
struct is_const_intrusive<const volatile T> : selector<true>
{
};

template <typename T>
struct is_const_intrusive<const volatile T> : selector<true>
{
};

template <typename T>
struct is_const_intrusive<T&> : is_const_intrusive<T>
{
};
```

Type traits have infinite applications; this example uses the simplest. Assume that C<T> is a class template that holds a member of type T, initialized by the constructor. However, T has no restriction, and in particular it may be a reference.

```
template <typename T>
class C
{
    T member_;

public:
    explicit C(argument_type x)
    : member_(x)
    {
    }
};
```

You need to define argument_type. If T is a value type, it's best to pass it by reference-to-const. But if T is a reference, writing const T& is illegal. So you'd write:

```
typedef typename add_reference<const T>::type argument_type;
```

Here, add_reference<T> returns const T&, as desired.

If T is a reference or reference-to-const, const T is T and add_reference returns T. That

3.5.2. Type Dismantling

A type in C++ can generate infinitely many "variations" by adding qualifiers, considering references, pointers, and arrays, and so on. But it can happen that you have to recursively remove all the additional attributes, one at a time. This recursive process is usually called *dismantling*.²⁰

This section shows a metafunction, named <code>copy_q</code>, that shifts all the "qualifiers" from type T1 to type T2 so <code>copy_q<const double&, int>::type will be const int&.</code>

Type deduction is entirely recursive. You dismantle one attribute at a time and move the same attribute to the result. To continue with the previous example, const double a matches Ta where T is const double, so the result is "reference to the result of copy_q<const double, int>," which in turn is "const result of copy_q<double, int>". Since this does not match any specialization, it gives int.

```
template <typename T1, typename T2>
struct copy q
{
  typedef T2 type;
};
template <typename T1, typename T2>
struct copy q<T1&, T2>
  typedef typename copy q<T1, T2>::type& type;
};
template <typename T1, typename T2>
struct copy q<const T1, T2>
{
  typedef const typename copy q<T1, T2>::type type;
};
template <typename T1, typename T2>
struct copy q<volatile T1, T2>
  typedef volatile typename copy q<T1, T2>::type type;
};
template <typename T1, typename T2>
struct copy q<T1*, T2>
  typedef typename copy q<T1, T2>::type* type;
};
template <typename T1, typename T2, int N>
```

```
struct copy_q<T1 [N], T2>
{
   typedef typename copy_q<T1, T2>::type type[N];
};
```

A more complete implementation could address the problems caused by T2 being a reference:

```
copy_q<double&, int&>::type err1;  // error: reference to
reference
copy q<double [3], int&>::type err2; // error: array of 'int&'
```

However, it's questionable if such classes should silently resolve the error or stop compilation. Let's just note that declaring a std::vector<int&> is illegal, but the compiler error is not "trapped":

```
/usr/include/gcc/darwin/4.0/c++/ext/new_allocator.h: In
instantiation of '__gnu_cxx::new_allocator<int&>':
/usr/include/gcc/darwin/4.0/c++/bits/allocator.h:83:
instantiated from 'std::allocator<int&>'
/usr/include/gcc/darwin/4.0/c++/bits/stl_vector.h:80:
instantiated from 'std::_Vector_base<int&, std::allocator<int&>
>::_Vector_impl'
/usr/include/gcc/darwin/4.0/c++/bits/stl_vector.h:113:
instantiated from 'std::_Vector_base<int&, std::allocator<int&> >'
/usr/include/gcc/darwin/4.0/c++/bits/stl_vector.h:149:
instantiated from 'std::vector<int&, std::allocator<int&> >'
main.cpp:94: instantiated from here
/usr/include/gcc/darwin/4.0/c++/ext/new_allocator.h:55: error:
forming pointer to reference type 'int&'
```

3.6. Type Containers

So what is a typelist? It's got to be one of those weird template beasts, right?

—Andrei Alexandrescu

The maximum number of template parameters is implementation-defined, but it's usually large enough to use a class template as *a container of types*.²¹

This section shows how some elementary *static algorithms* work, because you'll reuse the same techniques many times in the future. Actually, it's possible to implement most STL concepts in TMP, including containers, algorithms, iterators, and functors, where complexity requirements are translated at compilation time.²²

This section shows the ideas of the elementary techniques; you'll see some applications later.

The simplest type containers are *pairs* (the static equivalent of linked lists) and *arrays* (resemble C-style arrays of a fixed length).

```
template <typename T1, typename T2>
struct typepair
{
   typedef T1 head_t;
   typedef T2 tail_t;
};
struct empty
{
};
```

In fact, you can easily store a list of arbitrary (subject to reasonable limitations) length using pairs of pairs. In principle, you could form a complete binary tree, but for simplicity's sake, a list of types (T1, T2... Tn) is represented as typepair<T1, typepair<T2, ...> >. In other words, you'll allow the second component to be a pair. Actually, it forces the second component to be a typepair or an empty, which is the list terminator. In pseudo-code:

```
P0 = empty
P1 = typepair<T1, empty >
P2 = typepair<T2, typepair<T1, empty> >
// ...
Pn = typepair<Tn, P<sub>n-1</sub>>
```

This incidentally shows that the easiest operation with typepair-sequences is push_front. Following Alexandrescu's notation (see [1]), I call such an encoding a *typelist*. You say that the first accessible type Tn is the *head* of the list and Pn-1 is the *tail*.

Alternatively, if you fix the maximum length to a reasonable number, you can store all the types in a row. Due to the default value (which can be empty or void), you can declare any number of parameters on the same line:

The properties of these containers are different. A typelist with J elements requires the compiler

to produce J different types. On the other hand, arrays are direct-access, so writing algorithms for type arrays involves writing many (say 32) specializations. Typelists are shorter and recursive but take more time to compile.

Note Before the theoretical establishment made by Abrahams in [3], there was some naming confusion. The original idea of type pairs was fully developed by Alexandrescu (in [1] and subsequently in CUJ), and he introduced the name *typelist*.

Apparently, Alexandrescu was also the first to use type arrays as wrappers for declaring long typelists in an easy way:

```
template <typename T1, typename T2, ..., typename Tn>
struct cons
{
   typedef typepair<T1, typepair<T2, ...> > type;
};
```

However, the name *typelist* is still widely used as a synonym of a more generic type container.

3.6.1. typeat

typeat is a metafunction that extracts the Nth type from a container.

If the Nth type does not exist, the result is ERR.

The same metafunction can process type arrays and typelists. As anticipated, arrays require all the possible specializations. The generic template simply returns an error, then the metafunction is specialized first on type arrays, and then on typelists.

```
template <size_t N, typename CONTAINER, typename ERR
= Error_UNDEFINED_TYPE>
struct typeat
{
   typedef ERR type;
};

template <typename T1, ... typename T32, typename ERR>
struct typeat<0, typearray<T1, ..., T32>, ERR>
{
```

```
typedef T1 type;
};

template <typename T1, ... typename T32, typename ERR>
struct typeat<1, typearray<T1, ..., T32>, ERR>
{
   typedef T2 type;
};

// write all 32 specializations
```

The same code for typelists is more concise. The Nth type of the list is declared equal to the (N-1)th type in the tail of the list. If N is 0, the result is the head type. However, if you meet an empty list, the result is ERR.

```
template <size_t N, typename T1, typename T2, typename ERR>
struct typeat<N, typepair<T1, T2>, ERR>
{
   typedef typename typeat<N-1, T2, ERR>::type type;
};

template <typename T1, typename T2, typename ERR>
struct typeat<0, typepair<T1, T2>, ERR>
{
   typedef T1 type;
};

template <size_t N, typename ERR>
struct typeat<N, empty, ERR>
{
   typedef ERR type;
};
```

Observe that, whatever index you use, typeat<N, typearray<...>> requires just one template instantiation. typeat<N, typepair<...>> may require N different instantiations.

Note also the shorter implementation:

```
template <size_t N, typename T1, typename T2, typename ERR>
struct typeat<N, typepair<T1, T2>, ERR> : typeat<N-1, T2, ERR>
{
};
```

3.6.2. Returning an Error

When a metafunction F<T> is undefined, such as with typeat<N, empty, ERR>, common options for returning an error include:

- Removing the body of F<T> entirely.
- Giving F<T> an empty body, with no result (type or value).
- Defining F<T>::type so that it will cause compilation errors, if used (void or a class that has no definition).
- Defining F<T>::type using an user-supplied error type (as shown previously).

Remember that forcing a compiler error is quite drastic; it's analogous to throwing exceptions. It's hard to ignore, but a bogus type is more like a return false. A false can be easily converted to a throw and a bogus type can be converted to a compiler error (a static assertion would suffice).

3.6.3. Depth

Dealing with type arrays can be easier with the help of some simple macros²³:

Surprisingly, you can write class declarations that look extremely simple and concise. Here is an example (before and after preprocessing).

```
template <MXT_LIST_32(typename T) >
struct depth< typelist<MXT_LIST_32(T) > >
template <typename T1, ..., typename T32>
struct depth< typelist<T1, ... T32> >
```

The metafunction called depth returns the length of the typelists:

```
template <typename CONTAINER>
struct depth;

template <>
struct depth< empty > : static_value<size_t, 0>
{
};

template <typename T1, typename T2>
struct depth< typepair<T1, T2> > : static_value<size_t, depth<T2>::value+1>
{
```

};

- The primary template is undefined, so depth<int> is unusable.
- If the depth of a typelist is K, the compiler must generate K different intermediate types (namely depth<P1>... depth<Pn> where Pj is the jth tail of the list).

For type arrays, you use macros again. The depth of typearray<> is 0; the depth of typearray<T1> is 1; and in fact the depth of typearray<MXT_LIST_N (T)> is N.

```
template <MXT_LIST_0(typename T) >
struct depth< typearray<MXT_LIST_0(T) > >
: static_value<size_t, 0> {};

template <MXT_LIST_1(typename T) >
struct depth< typearray<MXT_LIST_1(T) > >
: static_value<size_t, 1> {};

// ...

template <MXT_LIST_32(typename T) >
struct depth< typearray<MXT_LIST_32(T) > >
: static value<size t, 32> {};
```

Note that even if a malicious user inserts a fake empty delimiter in the middle, depth returns the position of the last non-empty type:

```
typedef typearray<int, double, empty, char> t4;
depth<t4>::value; // returns 4
```

In fact, this call will match depth<T1, T2, T3, T4>, where it happens that T3 = empty.

In any case, empty should be confined to an inaccessible namespace.

3.6.4. Front and Back

This section shows you how to extract the first and the last type from both type containers.

```
template <typename CONTAINER>
struct front;

template <typename CONTAINER>
struct back;
```

First, when the container is empty, you cause an error:

```
template <>
```

```
template <>
struct front<empty>
};
   While front is trivial, back iterates all over the list:
template <typename T1, typename T2>
struct front< typepair<T1, T2> >
   typedef T1 type;
};
template <typename T1>
struct back< typepair<T1, empty> >
{
   typedef T1 type;
};
template <typename T1, typename T2>
struct back< typepair<T1, T2> >
   typedef typename back<T2>::type type;
};
or simply:
template <typename T1, typename T2>
struct back< typepair<T1, T2> > : back<T2>
};
   For type arrays, you exploit the fact that depth and typeat are very fast and you simply do
what is natural with, say, a vector. The back element is the one at size-1. In principle, this would
work for typelists too, but it would "iterate" several times over the whole list (where each "iteration"
causes the instantiation of a new type).
template <MXT LIST 32(typename T)>
struct back< typearray<MXT LIST 32(T)> >
   typedef typelist<MXT LIST 32(T) > aux t;
   typedef typename typeat<depth<aux t>::value - 1, aux t>::type
type;
};
```

struct back<empty>;

template <>

struct back< typearray<> >

```
template <MXT_LIST_32(typename T)>
struct front< typearray<MXT_LIST_32(T)>>
{
   typedef T1 type;
};

template <>
struct front< typearray<>>
{
};
```

3.6.5. Find

You can perform a sequential search and return the index of the (first) type that matches a given T. If T does not appear in CONTAINER, you return a conventional number (say -1), as opposed to causing a compiler error.

The code for the recursive version basically reads:

- Nothing belongs to an empty container.
- The first element of a pair has index 0.
- The index is one plus the index of T in the tail, unless this latter index is undefined.

```
template <typename T, typename CONTAINER>
struct typeindex;

template <typename T>
struct typeindex<T, empty>
{
    static const int value = (-1);
};

template <typename T1, typename T2>
struct typeindex< T1, typepair<T1, T2> >
{
    static const int value = 0;
};

template <typename T, typename T1, typename T2>
struct typeindex< T, typename T1, typename T2>
struct typeindex< T, typepair<T1, T2> >
{
    static const int aux_v = typeindex<T, T2>::value;
```

```
The first implementation for type arrays is:

/* tentative version */

template <MXT_LIST_32(typename T) >
    struct typeindex< T1, typearray<MXT_LIST_32(T) > >
    {
       static const int value = 0;
    };

template <MXT_LIST_32(typename T) >
    struct typeindex< T2, typearray<MXT_LIST_32(T) > >
    {
       static const int value = 1;
    };

// ...
```

static const int value = (aux v==-1 ? -1 : aux v+1);

If the type you are looking for is identical to the first type in the array, the value is 0; if it is equal to the second type in the array, the value is 1, and so on. Unfortunately the following is *incorrect*:

```
typedef typearray<int, int, double> t3;
int i = typeindex<int, t3>::value;
```

There's more than one match (namely, the first two), and this gives a compilation error. I defer the solution of this problem until after the next section.

3.6.6. Push and Pop

It was already mentioned that the easiest operation with type pairs is push_front. It is simply a matter of wrapping the new head type in a pair with the old container:

```
template <typename CONTAINER, typename T>
struct push_front;

template <typename T>
struct push_front<empty, T>
{
   typedef typepair<T, empty> type;
};

template <typename T1, typename T2, typename T>
struct push_front<typepair<T1, T2>, T>
{
```

```
Quite naturally, pop_front is also straightforward:

template <typename CONTAINER>
struct pop_front;

template <>
struct pop_front<empty>;

template <typename T1, typename T2>
struct pop_front< typepair<T1, T2> >
{
   typedef T2 type;
```

};

};

typedef typepair< T, typepair<T1, T2> > type;

To implement the same algorithm for type arrays, you must adopt a very important technique named *template rotation*. This rotation shifts all template parameters by one position to the left (or to the right).

```
template <P1, P2 = some_default, ..., P_N = some_default> struct container { typedef container<P2, P3, ..., P_N, some_default> tail_t; 24};
```

The type resulting from a pop_front is called the *tail* of the container (that's why the source code repeatedly refers to tail t).

Parameters need not be types. The following class computes the maximum in a list of positive integers.

```
#define MXT_M_MAX(a,b) ((a)<(b) ? (b) : (a))

template <size_t S1, size_t S2=0, ..., size_t S32=0>
struct typemax : typemax<MXT_M_MAX(S1, S2), S3, ..., S32>
{
};

template <size_t S1>
struct typemax<S1,0,0,...,0> : static_value<size_t, S1>
{
};
```

As a side note, whenever it's feasible, it's convenient to *accelerate* the rotation. In the previous example, you would write

```
template <size_t S1, size_t S2=0, ..., size_t S32=0>
struct typemax
```

```
: typemax<MXT_M_MAX(S1, S2), MXT_M_MAX(S3, S4), ...,
MXT_M_MAX(S31, S32)>
{
};
```

To compute the maximum of N constants, you need only log2(N) instances of typemax, instead of N.

It's easy to combine rotations and macros with elegance²⁵:

```
template <typename T0, MXT_LIST_31(typename T) >
struct pop_front< typearray<T0, MXT_LIST_31(T) > >
{
   typedef typearray<MXT_LIST_31(T) > type;
};

template <MXT_LIST_32(typename T), typename T>
struct push_front<typearray<MXT_LIST_32(T) >, T>
{
   typedef typearray<T, MXT_LIST_31(T) > type;
};
```

Using pop_front, you can implement a generic sequential find. Note that for clarity, you want to add some intermediate typedefs. As in metaprogramming, types are the equivalent of variables in classic C++. You can consider typedefs as equivalent to (named) temporary variables. Additionally, private and public sections help separate "temporary" variables from the results:

The procedure you'll follow here is:

- The index of T in an empty container is -1.
- The index of T1 in array<T1, ...> is 0 (this unambiguously holds, even if T1 appears more than once).
- To obtain the index of T in array<T1, T2, T3, ...>, you compute its index in a rotated array and add 1 to the result.

```
template <typename T>
struct typeindex<T, typearray<> >
{
   static const int value = (-1);
};

template <MXT_LIST_32(typename T)>
struct typeindex< T1, typearray<MXT_LIST_32(T)> >
{
   static const int value = 0;
};

template <typename T, MXT_LIST_32(typename T)>
struct typeindex< T, typearray<MXT_LIST_32(T)> >
```

```
{
private:
   typedef typearray<MXT_LIST_32(T)> argument_t;
   typedef typename pop_front<argument_t>::type
tail_t;

   static const int aux_v = typeindex<T,
tail_t>::value;

public:
   static const int value = (aux_v<0) ? aux_v
: aux_v+1;
};</pre>
```

3.6.7. More on Template Rotation

Template arguments can be easily rotated; however, it's usually simpler to consume them left to right. Suppose you want to compose an integer by entering all its digits in base 10. Here's some pseudocode.

```
template <int D1, int D2 = 0, ..., int D_N = 0> struct join_digits { static const int value = join_digits<D2, ..., D_N>::value * 10 + D1; }; template <int D1> struct join_digits<D1> { static const int value = D1; }; join_digits<3,2,1>::value; // compiles, but yields 123, not 321
```

Observe instead that it's not so easy to consume DN in the rotation. This will not compile, because whenever DN is equal to its default (zero), value is defined in terms of itself:

```
template <int D1, int D2 = 0, ..., int D_{N-1} = 0, int D_{N} = 0> struct join_digits { static const int value = join_digits < D1, D2, ..., D_{N-1} >::value * 10 + D_{N}; };
```

Rotation to the right won't produce the correct result:

```
template <int D1, int D2 = 0, ..., int D_{N-1} = 0, int D_N = 0
struct join digits
 static const int value = join digits<0,D1,D2, ...,D_{N-1}>::value * 10
+ D<sub>N</sub>;
};
```

The solution is simply to store auxiliary constants and borrow them from the tail:

```
template <int D1 = 0, int D2 = 0, ..., int D_N = 0
struct join digits
  typedef join digits<D2, ..., D_N > next t;
  static const int pow10 = 10 * next t::pow10;
  static const int value = next t::value + D1*pow10;
} ;
template <int D1>
struct join digits<D1>
  static const int value = D1;
  static const int pow10 = 1;
};
join digits<3,2,1>::value;  // now really gives 321
```

Template rotation can be used in two ways:

• Direct rotation of the main template (as shown previously):

```
template <int D1 = 0, int D2 = 0, ..., int D_N = 0>
struct join digits
{ ... };
template <int D1>
struct join digits<D1>
{ ... };
```

• Rotation on a parameter. This adds an extra "indirection":

```
template <int D1 = 0, int D2 = 0, ..., int D_N = 0>
struct digit group
{
  // empty
};
```

The first solution is usually simpler to code. However, the second has two serious advantages:

- Type T, which "carries" the tuple of template parameters, can be reused. T is usually a type container of some kind.
- Suppose for the moment that <code>join_digits<...></code> is a true class (not a metafunction), and it is actually instantiated. It will be easy to write generic templates accepting any instance of <code>join_digits</code>. They just need to take <code>join_digits<X></code>. But, if <code>join_digits</code> has a long and unspecified number of parameters, clients will have to manipulate it as <code>X.26</code>

3.6.8. Agglomerates

The rotation technique encapsulated in pop_front can be used to create tuples as agglomerate objects.

In synthesis, an agglomerate A is a class that has a type container C in its template parameters. The class uses front<C> and recursively inherits from A< pop_front<C> >. The simplest way to "use" the front type is to declare a member of that type. In pseudo-code:

```
template <typename C>
class A : public A<typename pop_front<C>::type>
{
   typename front<C>::type member_;

public:
   // ...
};

template < >
class A<empty>
```

```
template < >
class A< typearray<> >
{
};
```

- Inheritance can be public, private, or even protected.
- There are two possible recursion stoppers: A<empty_typelist> and A<empty_typearray>.

So, an agglomerate is a package of objects whose type is listed in the container. If C is typearray<int, double, std::string>, the layout of A would be as shown in Figure 3-4.

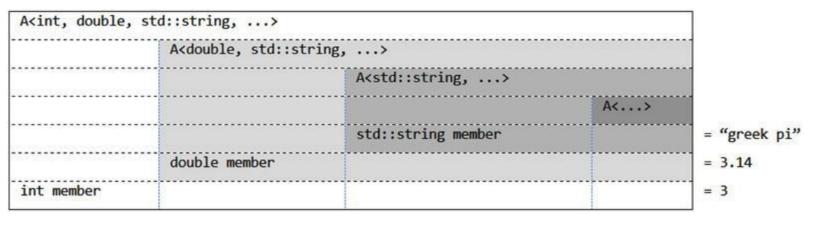


Figure 3-4. Layout of the agglomerate A

Note that in the implementation under review, the memory layout of the objects is reversed with respect to the type container.

To access the elements of the package, you use rotation again. Assume for the moment that all members are public. You'll get a reference to the Nth member of the agglomerate via a global function and the collaboration of a suitable traits class.

There are two equally good development strategies: *intrusive traits* and *non-intrusive traits*. *Intrusive traits* require the agglomerate to expose some auxiliary information:

```
template <typename C>
struct A : public A<typename pop_front<C>::type>
{
   typedef typename front<C>::type value_type;
   value_type member;

   typedef typename pop_front<C>::type tail_t;
};

template <typename agglom_t, size_t N>
struct reference_traits
```

```
{
  typedef reference traits<typename agglom t::tail t, N-1> next t;
  typedef typename next t::value type value type;
  static value type& ref(agglom t& a)
  {
     return next t::ref(a);
};
template <typename agglom t>
struct reference traits<agglom t, 0>
  typedef typename agglom t::value type value type;
  static value type& ref(agglom t& a)
     return a.member;
};
template <size t N, typename agglom t>
inline typename reference traits<agglom t, N>::value type&
ref(agglom t& a)
  return reference traits<agglom t, N>::ref(a);
  A quick example:
typedef typearray<int, double, std::string> C;
A < C > a;
ref<0>(a) = 3;
ref<1>(a) = 3.14;
ref<2>(a) = "3.14";
  Non-intrusive traits instead determine the information with partial specializations:
template <typename agglom t, size t N>
struct reference traits;
template <typename C, size t N>
struct reference traits < A<C>, N >
  typedef reference traits<typename pop front<C>::type, N-1> next t;
  typedef typename front<C>::type value type;
};
```

When feasible, non-intrusive traits are preferred. It's not obvious that the author of reference_traits can modify the definition of A. However it's common for traits to require reasonable "cooperation" from objects. Furthermore, auto-deduction code is a duplication of class A internals and auto-deduced values tend to be "rigid," so intrusiveness is not a clear loser.

A special case is an agglomerate modeled on a typelist *containing no duplicates*. The implementation is much simpler, because instead of rotation, a pseudo-cast suffices:

The cast works because the syntax ref<T> (a) fixes the first type of the pair and lets the compiler match the tail that follows. This is indeed possible, due to the uniqueness hypothesis.

In fact, the C++ Standard allows one derived-to-base cast before argument deduction, if it's a necessary and sufficient condition for an exact match.

Here, the only way to bind an argument of type A<C> to a reference to A< typepair<std::string, tail_t> > is to cast it to typepair<std::string, empty> and then deduce tail_t = empty.

To store a value extracted from an agglomerate, declare an object of type reference_traits<agglom_t, N>::value_type.

Finally, with a little more intrusiveness, you just add a member function to A:

```
template <typename C>
struct A : public A< typename pop_front<C>::type >

{
   typedef typename front<C>::type value_type;
   value_type member;

   typedef typename pop_front<C>::type tail_t;

   tail_t& tail() { return *this; }
};

template <typename agglom_t, size_t N>
struct reference_traits
{
   // ...
```

```
static value_type& get_ref(agglom_t& a)
{
    return next_t::get_ref(a.tail());
}
```

Invoking a member function instead of an implicit cast allows you to switch to private inheritance or even to a has-a relationship:

```
template <typename C>
class A
{
public:
    typedef typename pop_front<C>::type tail_t;
    typedef typename front<C>::type value_type;

private:
    A<tail_t> tail_;
    value_type member;

public:
    tail_t& tail() { return tail_; }

    // ...
};
```

The memory layout of the object is now in the same order as the type container.

3.6.9. Conversions

Many algorithms in fact require a linear number of recursion steps, both for typelists and for type arrays. In practice, the typepair representation suffices for most practical purposes except one—: the declaration of a typelist is indeed unfeasible.

As anticipated, it's very easy to convert from a type array to typelist and vice versa.

It is an interesting exercise to provide a unified implementation²⁷:

```
template <typename T>
struct convert
{
   typedef typename pop_front<T>::type tail_t;
   typedef typename front<T>::type head_t;

   typedef
     typename push_front<typename convert<tail_t>::type,
head_t>::type
     type;
};
```

```
template <>
struct convert< typearray<> >
{
   typedef empty type;
};

template <>
struct convert< empty >
{
   typedef typearray<> type;
};
```

Note that T in this code is a generic type container, not a generic type.

Before, you used partial template specialization as a protection against *bad static argument types*. For example, if you try front<int>::type, the compiler will output that front cannot be instantiated on int (if you did not define the main template) or that it does not contain a member type (if it's empty).

However, such a protection is not necessary here. convert is built on top of front and pop_front, and they will perform the required argument validation. In this case, the compiler will diagnose that front<int>, instantiated inside convert<int>, is illegal.

The problem is just a less clear debug message. Among the options you have to correct the problem, you can write type traits to identify type containers and then place assertions:

```
template <typename T>
struct type container
  static const bool value = false;
} ;
template <typename T1, typename T2>
struct type container< typepair<T1, T2> >
  static const bool value = true;
};
template <>
struct type container<empty>
  static const bool value = true;
};
template <MXT LIST 32(typename T)>
struct type container< typearray<MXT LIST 32(T)>>
  static const bool value = true;
};
```

```
template <typename T>
struct convert
  : static_assert< type_container<T>::value >
{
    //...
```

Very likely, the compiler will emit the first error pointing to the assertion line.

Note Section 5.2 is fully devoted to bad static argument types. You will meet function templates that statically restrict their template parameters to those having a particular interface.

It can be useful to extend type container traits by inserting a type representing the empty container of that kind (the primary template is unchanged).

```
template <typename T1, typename T2>
struct type_container< typepair<T1, T2> >

{
    static const bool value = true;
    typedef empty type;
};

template <>
struct type_container<empty>
{
    static const bool value = true;
    typedef empty type;
};

template <MXT_LIST_32(typename T)>
struct type_container< typearray<MXT_LIST_32(T)> >
{
    static const bool value = true;
    typedef typearray<> type;
};
```

When enough "low-level" metafunctions—such as front, back, push_front, and so on—are available, most meta-algorithms will work on arrays and lists. You just need two different recursion terminations, as well as a specialization for typearray<> and one for empty.

Another option is *the empty-empty idiom*: Let a helper class take the original type container as T and a second type, which is the empty container of the same kind (obtained from traits). When these are equal, you stop.

```
template <typename T>
struct some_metafunction
: static_assert<type_container<T>::value>
, helper<T, typename type_container<T>::type>
{
```

```
template <typename T, typename E>
struct helper
{
    // general case:
    // T is a non-empty type container of any kind
    // E is the empty container of the same kind
};

template <typename E>
struct helper<E, E>
{
    // recursion terminator
};
```

3.6.10. Metafunctors

template <typename T>

struct size of

User functors, predicates, and binary operations can be replaced by template-template parameters. Here is a simple metafunctor:

```
{
    static const size_t value = CHAR_BIT*sizeof(T);
};

template <>
    struct size_of<void>
{
        static const size_t value = 0;
};

    Here is a simple binary metarelation:

template <typename X1, typename X2>
    struct less_by_size : selector<(sizeof(X1) < sizeof(X2))>
{
};

template <typename X>
    struct less_by_size<void, X> : selector<true>
{
};

template <typename X>
    struct less_by_size<X, void> : selector<false>
```

```
template <>
struct less_by_size<void, void> : selector<false>
{
};
```

And here's the skeleton of a metafunction that might use it:

```
template <typename T, template <typename X1, typename X2> class LESS>
struct static_stable_sort
: static_assert< type_container<T>::value >
{
    // write LESS<T1, T2>::value instead of "T1<T2"

    typedef [[RESULT]] type;
};</pre>
```

Instead of describing an implementation, this section sketches a possible application of static_stable sort. Suppose our source code includes a collection of random generators that return unsigned integers:

```
class linear_generator
{
   typedef unsigned short random_type;
   ...
};

class mersenne_twister
{
   typedef unsigned int random_type;
   ...
};

class mersenne_twister_64bit
{
   typedef /* ... */ random_type;
   ...
};
```

The user will list all the generators in a type container, in order from the best (the preferred algorithm) to the worst. This container can be sorted by sizeof (typename T::random_type). Finally, when the user asks for a random number of type X, you scan the sorted container and stop on the first element whose random_type has at least the same size as X. You then use that generator to return a value. Since sorting is stable, the first suitable type is also the best in the user preferences.

As promised earlier, I turn now to the problem of selecting unsigned integers by size (in bit).

First, you put all candidates in a type container:

You have to scan the list from left to right and use the first type that has a specified size (it's also possible to append to the list a compiler-specific type).

Note A little algebra is necessary here. By definition of the sign function, for any integer, you have the identity $\delta \cdot \text{sign}(\delta) = |\delta|$. On the other hand, if S is a prescribed constant in $\{-1, 0, 1\}$, the equality $\delta \cdot S = |\delta|$ implies respectively $\delta \le 0$, $\delta = 0$, $\delta \ge 0$. This elementary relationship allows you to represent three predicates (less-or-equal-to-zero, equal-to-zero, and greater-or-equal-to-zero) with an integer parameter.

In the following code, T is any type container:

```
#define MXT M ABS(a) ((a) < 0 ? -(a) : (a))
enum
  LESS OR EQUAL = -1,
  EQUAL = 0,
  GREATER OR EQUAL = +1
};
template
  typename T,
  template <typename X> class SIZE_OF,
  int SIGN,
  size t SIZE BIT N
struct static find if
: static assertion< type container<T>::value >
  typedef typename front<T>::type head t;
  static const int delta = (int) SIZE OF<head t>::value -
(int) SIZE BIT N;
  typedef typename typeif
  <
     SIGN*delta == MXT M ABS(delta),
     front<T>,
     static_find_if<typename pop_front<T>::type,
                  SIZE OF, SIGN, SIZE BIT N>
  >::type aux t;
```

```
typedef typename aux t::type type;
};
// define an unsigned integer type which has exactly 'size' bits
template <size t N>
struct uint n
: static find if <all unsigned, size of, EQUAL, N>
};
// defines an unsigned integer type which has at least 'size' bits
template <size t N>
struct uint nx
: static find if <all_unsigned, size_of, GREATER_OR_EQUAL, N>
};
typedef uint n<8>::type uint8;
typedef uint n<16>::type uint16;
typedef uint n<32>::type uint32;
typedef uint n<64>::type uint64;
typedef uint nx<32>::type uint32x;
  Note that the order of template parameters was chosen to make clear the line that uses
static_find_if, not static find if itself. 28
  What happens if a suitable type is not found? Any invalid use will unwind a long error cascade
(the code has been edited to suppress most of the noise):
uint n<25>::type i0 = 8;
uint nx<128>::type i1 = 8;
error C2039: 'type' : is not a member of 'front<typearray<>>'
                             : see declaration of 'front<typearray<>>'
                             : see reference to class template
instantiation
'static find if<T,SIZE OF,SIZE BIT N,SIGN>' being compiled
       with
T=pop front<pop front<pop front<pop front<pop front<all unsigned>::
       1
                          : see reference to class template
instantiation 'static find if<T,SIZE OF,SIZE BIT N,SIGN>' being
compiled
       with
```

```
T=pop front<pop front<pop front<all unsigned>::type>::ty
                       : see reference to class template
instantiation
'static find if<T,SIZE OF,SIZE BIT N,SIGN>' being compiled
      with
T=pop front<pop front<all unsigned>::type>::type>::type,
[...]
                   : see reference to class template
instantiation
'static find if<T,SIZE OF,SIZE BIT N,SIGN>' being compiled
      with
      T=all unsigned,
                   : see reference to class template instantiation
'uint n<SIZE BIT N>' being compiled
      with
         SIZE BIT N=25
```

Basically, the compiler is saying that, during deduction of uint_n<25>::type, after applying pop_front to the type array five times, it ended up with an empty container, which has no front type.

However it's easy to get a more manageable report. You just add an undefined type as a result of the recursion terminator:

```
template
<
  template <typename X> class SIZE_OF,
  int SIGN,
  size_t SIZE_BIT_N
>
struct static_find_if<typearray<>, SIZE_OF, SIGN, SIZE_BIT_N>
{
  typedef error_UNDEFINED_TYPE type;
};
```

Now the error message is more concise:

```
error C2079: 'i0' uses undefined class 'error_UNDEFINED_TYPE' error C2079: 'i1' uses undefined class 'error_UNDEFINED_TYPE'
```

3.7. A Summary of Styles

When programming metafunctions, identify:

- A suggestive name and syntax.
- Which template parameters are needed to express the concept.
- Which atomic actions the algorithm depends on.
- A recursive efficient implementation.
- Special cases that must be isolated.

If the metafunction name is similar to a classic algorithm (say, find_if), then you can adopt a similar name (static_find_if) or even an identical one if it resides in a specific namespace (say, typelist::find if).

Some authors append an underscore to pure static algorithms, because this allows mimicking real keywords (typeif would be called if).

If several template parameters are necessary, write code so that the users will be able to remember their meaning and order. It's a good idea to give a syntax hint through the name:

```
: static_find_if<all_unsigned, size_of, GREATER_OR_EQUAL, N>
```

Many unrelated parameters should be grouped in a traits class, which should have a default implementation that is easy to copy.

Finally, the following table may help you translate a classic algorithm to a static one.

	Classic C++ Function	Static Metaprogramming
What they manipulate	Instances of objects	Types
Argument handling	Via argument public interface	Via metafunctions
Dealing with different arguments	Function overload	Partial template specializations
Return result	Zero or one return statement	Zero or more static data (type or constant), usually inherited
Error trapping	$ ext{try}$ catch $ ext{block}$	Extra template parameter ERR
User-supplied callbacks	Functors	Template-template parameters
Temporary objects	Local variables	Private typedef/static const
Function calls	Yes, as subroutines	Yes, also via derivation
Algorithm structure	Iteration or recursion	Static recursion, stopped with suitable full/partial template specializations
Conditional decisions	Language constructs (if, switch)	Partial specializations
	Throw an exception	Abort compilation

Set result to an incomplete type

¹Remember that the preprocessor runs *before* the compiler so it cannot rely on sizeof.

²Read the previous note again ⁽²⁾.

³See also http://www.boost.org/doc/libs/1 46 0/libs/wave/doc/preface.html.

⁴This example was taken from a private conversation with Marco Marcello.

⁵I do not always use the derivation notation in the book, mainly for sake of clarity. However, I strongly encourage adopting it in production code, as it boosts code reuse.

⁶Except casts to non-integer types. For example, N*1.2 is illegal, but N+N/5 is fine.

⁷In practice, N is always even, so N-N/2 = N/2.

⁸See also the double-check stop in Section 7.2.

⁹I insist that the problem is solvable because the implementations of is_integer<long> and is_integer<ptrdiff_t> are identical; otherwise, it is ill-formed. For a counterexample, consider the problem of converting a time_t and long to a string; even if time_t is long, the strings need to be different. Therefore, this issue cannot be solved by TMP techniques.

¹⁰This is a good thing, because a well-built template class shouldn't need it.

¹¹Here, FAKE and FAKE+1 both work.

¹²Same does not imply that all functions must be identical, as some differences may have a limited impact on "uniform use". As a trivial example, arguments may be passed by value or by const reference.

¹³Available at: cantrip.org/trails.html. The article cites as previous bibliography [10], [11] and [12].

¹⁴The sentences have been slightly rearranged.

¹⁵The term *type traits*, introduced by John Maddock and Steve Cleary, is used here as a common name, but it is also popular as a proper name, denoting a particular library implementation. See http://cppreference.com/header/type_traits or http://www.boost.org/doc/libs/1_57_0/libs/type_traits/doc/html/index.html.

¹⁶In modern C++, there's a dedicated <type_traits> header that contains most of the metafunctions described here, and many more that cannot be replicated in classic C++. For example, has_trivial_destructor<T> is indeducible without the cooperation of the compiler, and current implementations always return false, except for built-in types.

¹⁷ It's possible to define add_reference<void>::type to be void.

 $^{^{18}}$ This is actually used. Some smart pointers, including std::unique_ptr, use operator delete [] when the type matches T[] and single deletion in any other case.

¹⁹Alternatively, std::pair<const int, double> is neither const nor assignable.

²⁰The expression "type dismantling" was introduced by Stephen C. Dewhurst.

²¹The C++ Standard contains an informative section, called "Implementation Quantities," where a recommended minimum is suggested for the number of template arguments (1024) and for nested template instantiations (1024), but compilers do not need to respect these numbers.

²²The reference on the argument is [3].

²³The boost preprocessor library would be more suitable, anyway, but its description would require another chapter. Here, the focus is on the word *simple*: a strategic hand-written macro can improve the esthetics of code noticeably.

²⁴In principle, some_default should not be explicitly specified. All forms of code duplication can lead to maintenance errors. Here, I

show it to emphasize the rotation.

- ²⁵See Section 3.6.3.
- ²⁶This need not be a problem' if join_digits were a functor, clients would likely take it as X anyway.
- ²⁷It's another exercise of type dismantling; note also that using push_back instead of push_front would reverse the container.
- ²⁸I adopted the name find_if with some abuse of notation; a genuine static_find_if would be static_find_if<typename T, template <typename X> class F>, which returns the first type in T where F<X>::value is true.

CHAPTER 4

Overload Resolution

This chapter presents TMP techniques based on overload resolution.

The common underlying schema is as follows:

- You want to test if type T satisfies a condition.
- You write several static functions with the same name, say test, and pass them a dummy argument that "carries" type T (in other words, an argument that allows deduction of T, such as T^*).
- The compiler selects the best candidate, according to C++ language rules.
- You deduce which function was used, either using the return type or indirectly from a property of this type, and eventually make a decision.

The first section introduces some definitions.

4.1. Groups

A *group* is a class that provides optimized variants of a single routine. From the outside, a group acts as a monolithic function that automatically picks the best implementation for every call.

A group is composed of two entities:

- A template struct containing variants of a (single) static member function.
- A companion global function template that just forwards the execution to the correct member of the group, performing a static decision based on the auto-deduced template parameter and on some framework-supplied information.

The group itself is usually a template, even if formally unnecessary (it may be possible to write the group as a normal class with template member functions).

Finally, observe that groups and traits are somehow orthogonal. Traits contain all the actions of a specific type, while groups contain a single action for many types.

4.1.1. From Overload to Groups

A group is the evolution of a set of overloaded functions.

Step 1: You realize that a default template implementation can handle most cases, so you just add overloaded variants:

```
template <typename T>
bool is_product_negative(T x, T y)
{
   return x<0 ^ y<0;
}
bool is_product_negative(short x, short y)
{
   return int(x)*int(y) < 0;
}
bool is_product_negative(unsigned int x, unsigned int y)
{
   return false;
}
bool is_product_negative(unsigned long x, unsigned long y)
{
   return false;
}</pre>
```

Step 2: Implementation is clustered in several templates that are picked using tags.

```
template <typename T>
bool is_product_negative(T x, T y, selector<false>)
{
   return x<0 ^ y<0;
}
template <typename T>
bool is_product_negative(T x, T y, selector<true>)
```

```
{
  return int(x) *int(y) < 0;
template <typename T>
bool is product negative (T x, T y)
{
  typedef selector<(sizeof(T)<sizeof(int))> small int t;
  return is product negative(x, y, small int t());
  Step 3: Group all the auxiliary functions in a class and leave a single function outside that
dispatches the work:
// companion function
template <typename T>
bool is product negative (T x, T y)
  return is product negative t<T>::doIt(x, y);
}
template <typename T>
struct is product negative t
  static bool doIt(T x, T y)
  { . . . }
  static bool doIt(unsigned, unsigned)
  { return false; }
};
  Here is another very simple group:
struct maths
  template <typename T>
  inline static T abs(const T x)
     return x<0 ? -x : x;
  inline static unsigned int abs(unsigned int x)
     return x;
};
template <typename T>
```

```
inline T absolute_value(const T x)
{
   return maths::abs(x);
}
```

■ **Note** Remember that the group class, being a non-template, is always fully instantiated. Furthermore, a non-template function in a header file must be declared inline.

Suppose further that you have a metafunction named has_abs_method, such that has_abs_method<T>::value is true if the absolute value of an object x of type T is given by x.abs().

This allows your group to grow a bit more complex. In the next example, you'll specialize the whole group for double, and the specialization will ignore the actual result of has abs method<double>.2

```
template <typename scalar t>
struct maths
  static scalar t abs(const scalar t& x, selector<false>)
     return x<0 ? -x : x;
  static scalar t abs(const scalar t& x, selector<true>)
     return x.abs();
};
template <>
struct maths<double>
  template <bool UNUSED>
  static double abs(const double x, selector< UNUSED >)
     return std::fabs(x);
};
template <typename scalar t>
inline scalar t absolute value(const scalar t& x)
  typedef selector< has abs method<scalar t>::value > select t;
  return maths<scalar t>::abs(x, select t());
}
```

Too many overloads will likely conflict. Remember that a non-template function is preferred to a matching template, but this does not hold for a member function that uses the template parameter of the class:

```
template <typename scalar_t>
struct maths
{
    static scalar_t abs(const scalar_t& x, selector<false>)
    {
        return x<0 ? -x : x;
    }
    static int abs(const int x, selector<false>)
    {
        return std::abs(x);
    }
}
error: ambiguous call to overloaded function, during instantiation of absolute_value<int>
```

This is precisely the advantage of a "double-layer" template selection. "Layer one" is the automatic deduction of scalar_t in the companion function and "layer two" is the overload selection, performed inside a class template (the group) whose parameter has already been fixed:

```
template <typename scalar_t>
inline scalar_t absolute_value(const scalar_t& x)
{
    // collect auxiliary information, if needed
    return math<scalar_t>::abs(x, ...);
}
```

Combining them, you have fewer global function templates (too many overloads are likely to cause "ambiguous calls"). In addition, the group can have subroutines (private static member functions).

The user has several expansion choices:

- Specialize the whole group (if it's a template)
- Specialize the global companion function
- Model types to take advantage of the existing framework (for example, specialize has abs method)

The selection part can be even subtler, with additional layers in the middle. As the following example shows, the right member of the group is chosen via an implicit argument promotion:

```
#include <cmath>
```

```
struct tag floating
  tag floating() {}
  tag floating(instance of<float>) {}
  tag floating(instance of<double>) {}
  tag floating(instance of<long double>) {}
};
struct tag signed int
  tag signed int() {}
  tag signed int(instance of<short>) {}
  tag signed int(instance of<int>) {}
  tag signed int(instance of<long>) {}
};
struct tag unsigned int
  tag unsigned int() {}
  tag unsigned int(instance of<unsigned short>) {}
  tag unsigned int(instance of<unsigned int>) {}
  tag unsigned int(instance of<unsigned long>) {}
};
template <typename scalar t>
struct maths
  inline static scalar t abs(const scalar t x, tag signed int)
     return x<0 ? -x : x;
  inline static scalar t abs(const scalar t x, tag unsigned int)
     return x;
  inline static scalar t abs(const scalar t x, tag floating)
     return fabs(x);
};
template <typename scalar t>
inline scalar t absv(const scalar t& x)
  return maths<scalar t>::abs(x, instance of<scalar t>());
```

The same effect could be obtained with a reversed selector hierarchy (for example, letting instance_of<double> derive from scalar_floating), but instance_of is a general-purpose template and I treat it as non-modifiable.

You could also introduce intermediate selectors (unfortunately, you have to write the constructors by hand):

```
struct tag int
  tag int() {}
  tag int(instance of<short>) {}
  tag int(instance of<int>) {}
  tag int(instance of<long>) {}
  tag int(instance of<unsigned short>) {}
  tag int(instance of<unsigned int>) {}
  tag int(instance of<unsigned long>) {}
};
template <typename scalar t>
struct maths
  static scalar t mod(const scalar t x, const scalar t y, tag int)
     return x % y;
  static scalar t mod(const scalar t& x, const scalar t& y,
tag floating)
     return fmod(x, y);
};
template <typename scalar t>
inline scalar t mod(const scalar t& x, const scalar t& y)
  return maths<scalar t>::mod(x, y, instance of<scalar t>());
}
```

Note in this code that maths < double > contains a method that must not be called (there's no operator% for double). Had operation been a non-template class, it would have been instantiated anyway, thus yielding a compiler error.

However, when parsing an expression depending on a template parameter, the compiler, not knowing the actual type involved, will accept any formally legal C++ statement.³ So if at least one of the two arguments x and y has generic type T, $x \, y$ is considered valid until instantiation time.

The former example works unambiguously because the companion function restricts the call to members of maths <double > named mod, and for any type T, instance_of <T > can be promoted to at most one of either tag int or tag floating.

Sometimes groups are associated with a special header file that detects platform information using macro blocks and translates it in C++ using typedefs:

In different platforms, the same function could have a different "best" implementation, so you can select the most suitable one using compiler_type as a tag (but *all* functions must be legal C++ code):

```
template <typename scalar_t, typename compiler_t>
struct maths
{
   static scalar_t multiply_by_two(const scalar_t x)
     { return 2*x; }
};

template < >
struct maths<unsigned int, msvc>
{
   static unsigned int multiply_by_two(const unsigned int x)
     { return x << 1; }
};

template <typename scalar_t>
inline scalar_t multiply_by_two(const scalar_t& x)
{
   return maths<scalar_t, compiler_type>::multiply_by_two(x);
}
```

Note that you can branch the selection of member functions as you wish—either simultaneously on multiple tags or hierarchically.

As a rule, you might want to use the "compiler tag" whenever you need to manipulate the result of a standard function that is defined as compiler-specific to some extent, for example, to pretty-print a string given by typeid (...) .name ().

Consider a real-world example. According to the standard, if A and B are both signed integers,

not both positive, the sign of A % B is undefined (if instead A>0 and B>0, the standard guarantees that A % B > 0).

For example, -10 % 3 can yield either -1 or +2, because -10 can be written as 3*(-3)+(-1) or 3*(-4)+(+2) and both |-1| < 3 and |2| < 3. In any case, both solutions will differ by 3.

However, operator% is often implemented so that A and (A % B) both have the same sign (which, in fact, is the same rule used for fmod). It therefore makes sense to write a reminder function that grants this condition.

Since (-A) % B == -(A % B) and A % (-B) == A % B, you can deduce that you can return sign (A) * (|A| % |B|) when the native implementation of A % B yields a different result.

A simple implementation can rely on (-3) % 2 being equal to +1 or -1. (Note that the following code is not 100% bulletproof, but it's a good compromise.)

```
template <typename T, int X = (-3)\%2, int Y = (-3)\%(-2), int Z =
3% (-2) >
struct modgroup;
// if X=+1, Y=-1, Z=+1 then operator% already does what we want
// (strictly speaking, we tested only int)
template <typename T>
struct modgroup<T, 1, -1, 1>
  static scalar t mod(const T x, const T y)
     return x % y;
};
// in any other case, fall back to the safe formula
template <typename T, int X, int Y, int Z>
struct modgroup
  static scalar t mod(const T x, const T y)
  {
     const T result = abs(x) % abs(y);
     return x<0 ? -result : result;
  }
};
template <typename scalar t>
struct maths
{
  static scalar t mod(const scalar t x, const scalar t y,
                   tag int)
  {
```

4.1.2. Runtime Decay

A type tag may implement a special cast operator so that if no overload in the group matches the tag exactly, the execution continues in a default function, which usually performs some work at runtime. The prototype is a static integer that decays into a normal integer if there's no better match.

Suppose you want to fill a C array with zeroes:

```
template <typename T, T VALUE>
struct static_value
{
    // ...
    operator T() const
    {
        return VALUE;
    }
};

template <typename T>
struct zeroize_helper
{
    static void apply(T* const data, static_value<int, 1>)
        {
          *data = T();
    }

    static void apply(T (&data)[2], static_value<int, 2>)
        {
             data[0] = data[1] = T();
        }
}
```

```
static void apply(T* const data, const int N)
{
    std::fill_n(data, N, T());
};

template <typename T, int N>
void zeroize(T (&data)[N])
{
    zeroize_helper<T>::apply(data, static_value<int, N>());
}
```

- Instead of 0, you write T (), which works for a broader range of types.
- If N is larger than 2, the best match is the third member.
- Each function in the group can decide freely to cast, or even to ignore, the static value.
- The default case may accept every static_value not necessarily performing all the work at runtime, but with another template function:

```
template <>
struct zeroize helper<char>
  template <int N>
  struct chunk
     char data[N];
  };
  template <int N>
  static void apply(char* const data, static value<int, N>,
selector<true>)
     *reinterpret cast<chunk<N>*>(data) = chunk<N>();
  template <int N>
  static void apply(char* const data, static value<int, N>,
selector<false>)
    memset(data, N, 0);
  template <int N>
  static void apply(char* const data, static value<int, N> S)
     apply(data, S, selector<sizeof(chunk<N>) == N>());
```

```
};
```

4.2. More Traits

This section completes the review of traits.

This time you are going to use traits restricted for static programming, but also as function groups. Let's start with a concrete case.

4.2.1. A Function Set for Strings

Suppose you are going to write some generic algorithms for strings. Surely you can use iterators, in particular random-access iterators, right? Most STL implementations have char-optimized algorithms, such as std::find, std::copy, and so on.

The only burden on the user is a large number of calls to strlen to find the end of range. strlen is a *very* fast function, but this is a violation of STL assumptions, as "end" is assumed to be obtained in constant time, not linear time.

```
const char* c_string = "this is an example";
// can we avoid this?
std::copy(c_string, c_string+strlen(c_string), destination);
```

You can squeeze in even more optimization using traits:

```
template <typename string_t>
struct string_traits
{
   typedef /* dependent on string_t */ const_iterator;
   typedef const string_t& argument_type;

   const_iterator begin(argument_type s);
   const_iterator end (argument_type s);

   static bool is_end_of_string(const_iterator i, argument_type s);
};
```

Assuming that for every meaningful string, string_traits has the same interface, you can write an algorithm as follows:

```
template <typename string_t>
void loop_on_all_chars(const string_t& s)
{
   typedef string_traits<string_t> traits_t;
```

```
typename traits_t::const_iterator i = traits_t::begin(s);
while (!traits_t::is_end_of_string(i, s))
{
    std::cout << *(i++);
}</pre>
```

The code is verbose but clear. Yet at this point it may not be evident what you accomplished. The semi-opaque interface of string traits gives more freedom in doing comparisons:

```
template <typename char t>
struct string traits< std::basic string<char t> >
  typedef char t char type;
  typedef
    typename std::basic string<char type>::const iterator
    const iterator;
  typedef const std::basic string<char type>& argument type;
  static const iterator begin (argument type text)
  {
     return text.begin();
  static const iterator end(argument type text)
  {
     return text.end();
  static bool is end of string(const iterator i, argument type s);
     return i == s.end();
};
template <>
struct string traits<const char*>
  typedef char char type;
  typedef const char* const iterator;
  typedef const char* argument type;
  static const iterator begin (argument type text)
     return text;
```

Since end is now constant-time, you save a linear-time pass (you'll meet this very same problem again and solve it with a different technique in Section 6.2.2.

You can easily extend string_traits to a full interface (some words have been renamed for ease of reading):

```
template <typename string t>
struct string traits
typedef /* ... */ char_type;
typedef /* ... */ const iterator;
typedef /* ... */ argument type; // either string t or const
string t&
static size t npos();
static size t find1st(arg t txt, const char t c, size t offset=0);
static size t find1st(arg t txt, const arg t s, size t offset=0);
static size t findlast(arg t txt, const char t s, size t offset);
static size t findlast(arg t txt, const arg t s, size t offset);
static size t find1st in(arg t txt, const char t* charset, size t
offs=0);
static size t find1st out(arg t txt, const char t* charset, size t
offs=0);
static size t size (arg t txt);
static const iterator begin(arg t txt);
static const iterator end(arg t txt);
static const char t* c str(arg t txt);
static bool empty(const iterator begin, const iterator end);
static bool less (const iterator begin, const iterator end);
```

```
static size_t distance(const_iterator begin, const_iterator end);
};
```

To leverage the interface and take advantage of std::string member functions, consider the following convention:

- All iterators are random-access.
- The find functions return either the index of the character (which is portable in all kind of strings) or npos(), which means "not found".

```
static size_t find1st(arg_t text, const char_type c, size_t
offset=0)
{
  const char_t* pos = strchr(text+offset, c);
  return pos ? (pos-text) : npos();
}
```

In the specialization for const char*, you carry on the ambiguity on the end iterator, which can be a null pointer to mean "until char 0 is found". Thus, you could implement distance as follows:

```
static size_t distance(const_iterator begin, const_iterator end)
{
   return end ? end-begin : (begin ? strlen(begin) : 0);
}
```

Finally, you can inherit function sets via public derivation, as usual with traits, because they are stateless (so the protected empty destructor can be omitted):

```
template <>
struct string_traits<char*> : string_traits<const char*>
{
};
```

4.2.2. Concept Traits

As you repeatedly saw in the first chapters, traits classes prescribe syntax, not precise entities. Code may borrow from traits in such a way that several different implementations are possible.

Suppose you have some kind of smart pointer class whose traits class is also responsible for freeing memory:

```
template <typename T, typename traits_t = smart_ptr_traits<T> >
class smart_ptr
{
   typedef typename traits_t::pointer pointer;
   pointer p_;
```

```
public:
   ~smart ptr()
           traits_t::release(p_);
   // ...
};
   traits::release can be:
        • A public static function (or functor); the relevant code is in the function body.
template <typename T>
struct smart ptr traits
   typedef T* pointer;
   static void release (pointer p)
      delete p;
   }
        • A public static function that triggers a conversion operator, which in fact runs the
          code.
template <typename T>
struct smart ptr traits
   static void release (bool)
   };
   class pointer
      // ...
      public:
         operator bool()
         { . . . }
   };
   // ....
   Using a slightly different syntax, you can rewrite this as follows:
template <typename T, typename traits t = smart ptr traits<T> >
class smart ptr
   typedef typename traits t::pointer pointer;
```

```
pointer p_;
   static void traits release(typename traits t::release)
                 // note: empty body
   };
public:
  ~smart ptr()
          traits release(p );
  Release can now be a type, and the relevant code is in the (non-explicit) constructor body.
template <typename T>
struct smart ptr traits
  typedef T* pointer;
   struct release
     release(pointer p)
        delete p;
  };
  The code can, again, trigger a conversion operator:
template <typename T>
struct smart ptr traits
   struct release
   {
   };
  class pointer
      // ...
     public:
        operator release()
           delete p ;
           return release();
         }
   };
};
```

All these implementations are valid and you can choose the best positioning of the code that is actually executed.⁴

If traits::release is provided as a type, it may have static data that is easily shared with the rest of the program (you could, for example, log all the released pointers).

4.2.3. Platform-Specific Traits

Recall that traits classes can be "global" or "local". Global traits classes are visible everywhere and local traits should be passed as parameters.

Global traits are preferred to make some platform properties easily accessible to clients:

```
template <typename char_t>
struct textfile_traits
{
    static char_t get_eol() { return '\n'; }
    // ...
};
```

The following full example represents a timer object with a class template and borrows additional information from a "timer traits" class:

- How to get current time (in an unspecified unit)
- How to convert time into seconds (using a frequency)

```
template <typename traits t>
class basic timer
{
  typedef typename traits t::time type tm t;
  typedef typename traits t::difference type diff t;
  tm t start ;
  tm t stop ;
  inline static tm t now()
  {
     return traits t::get time();
  }
  double elapsed (const tm t end) const
  {
     static const tm t frequency = traits t::get freq();
     return double(diff t(end-start ))/frequency;
  }
public:
  typedef tm t time type;
```

```
typedef diff t difference type;
  basic timer()
  : start ()
  { }
  difference type lap() const
  { return now()-start ; }
  time type start()
  { return start = now(); }
  difference type stop()
  { return (stop = now())-start; }
  difference type interval() const
  { return stop -start ; }
  double as seconds() const
  { return elapsed(stop ); }
  double elapsed() const
  { return elapsed(now()); }
};
  Here is a sample traits class that measures clock time (in seconds):
#include <ctime>
struct clock time traits
  typedef size t time type;
  typedef ptrdiff t difference type;
  static time type get time()
     time t t;
     return std::time(&t);
  }
  static time type get freq()
     return 1;
};
  Here's a different traits class that accounts for CPU time:
struct cpu time traits
  typedef size t time type;
```

{

```
typedef ptrdiff_t difference_type;

static time_type get_time()
{
    return std::clock();
}

static time_type get_freq()
{
    return CLOCKS_PER_SEC;
}
};

And a short use case:

basic_timer<clock_time_traits> t;
t.start();
// ...
t.stop();
std::cout << "I ran for " << t.as_seconds() << " seconds.";</pre>
```

The fundamental restriction of traits is that all member functions must contain valid C++ code, even if unused. You cannot use compiler-specific code in one of the functions.

Since different operating systems can expose more precise APIs for time measurement, you might be tempted to write specialized traits:

```
#include <windows.h>
struct windows clock time traits
  typedef ULONGLONG time type;
  typedef LONGLONG difference type;
  static time type get time()
  {
     LARGE INTEGER i;
     QueryPerformanceCounter(&i);
     return i.QuadPart;
  }
  static time type get freq()
  {
     LARGE INTEGER value;
     QueryPerformanceFrequency(&value);
     return value. QuadPart;
};
#include <sys/time.h>
```

```
struct macosx_clock_time_traits
{
  typedef uint64_t time_type;
  typedef int64_t difference_type;

  static time_type get_time()
  {
    timeval now;
    gettimeofday(&now, 0);
    return time_type(now.tv_sec) * get_freq() + now.tv_usec;
  }

  static time_type get_freq()
  {
    return 10000000;
  }
};
```

Apart from the typedefs for large integers, this traits interface is standard C++, so you might are tempted to isolate the preprocessor in a "factory header" and rely entirely on template properties later:

```
// platform detect.hpp
struct windows { };
struct macosx {};
struct other os {};
#if defined(WIN32)
typedef windows platform type;
#elif defined( APPLE )
typedef macosx platform type;
#else
typedef other os platform type;
#endif
// timer traits.hpp
template <typename platform t>
struct clock time traits;
template < >
struct clock time traits<windows>
  // implementation with QPC/QPF
} ;
template < >
```

```
struct clock_time_traits<macosx>
{
    // implementation with gettimeofday
};

template < >
    struct clock_time_traits<other_os>
{
        // implementation with std::time
};

typedef basic_timer< clock_time_traits<platform_type> >
native_timer_type;
```

Unfortunately, the code is *non-portable* (if it compiles, however, it runs correctly).

According to the standard, a compiler is not required to diagnose errors in unused template member functions, but if it does, it requires that all mentioned entities be well-defined. In particular, GCC will report an error in clock_time_traits<windows>::get_time, because no function named QueryPerformanceCounter has been declared.

As the approach is attractive, some workarounds are possible:

• Define a macro with the same name and as many arguments as the function:

```
// define as nothing because the return type is void
// otherwise define as an appropriate constant, e.g. 0
#define QueryPerformanceCounter(X)
#if defined(WIN32)
#undef QueryPerformanceCounter // remove the fake...
#include <windows.h> // ...and include the true
function
#endif
```

• Declare—but do not define—the function. This is the preferred solution, because Windows traits should not link in other operating systems.

```
#if !defined(WIN32)
    void QueryPerformanceCounter(void*);
#endif
```

Note A common trick, if the function returns void, is to define the name of the function itself to <nothing>. The comma-separated argument list will be parsed as a comma operator.

This also allows ellipsis functions to be used:

```
#define printf
```

```
printf("Hello world, %f", cos(3.14));
```

However, there are a couple of potential issues. First, the macro changes the return type of the expression to double (the last argument). Furthermore, the program is still evaluating cos (3.14). An alternative that also minimizes the runtime effort—although it's not totally bulletproof—is:

```
inline bool discard_everything(...) { return false };

#define printf false && discard_everything
```

4.2.4. Merging Traits

Especially when you're dealing with large traits, it's good practice to enable the users to customize smaller parts of the traits class. Typically, the problem is solved by splitting the traits class into parts and recombining them using public inheritance to form a traits default value.

Suppose you are grouping some comparison operators in traits:

```
template <typename T>
struct binary_relation_traits
{
   static bool gt(const T& x, const T& y) { return x>y; }
   static bool lt(const T& x, const T& y) { return x<y; }

   static bool gteq(const T& x, const T& y) { return x>=y; }
   static bool lteq(const T& x, const T& y) { return x<=y; }

   static bool eq(const T& x, const T& y) { return x==y; }
   static bool ineq(const T& x, const T& y) { return x!=y; }
};</pre>
```

The general implementation of $binary_relation_traits$ assumes that T defines all six comparison operators, but this example supports two important special cases, namely:

- T defines operator< only
- T defines operator< and operator== only

Without your support, the users will have to implement all the traits structure from scratch. So you must rearrange the code as follows:

```
template <typename T>
struct b_r_ordering_traits
{
   static bool gt(const T& x, const T& y) { return x>y; }
   static bool lt(const T& x, const T& y) { return x<y; }</pre>
```

```
static bool gteq(const T& x, const T& y) { return x>=y; }
  static bool lteq(const T& x, const T& y) { return x<=y; }
};
template <typename T>
struct b r equivalence traits
  static bool eq(const T& x, const T& y) { return x==y; }
  static bool ineq(const T& x, const T& y) { return x!=y; }
};
template <typename T>
struct binary relation traits
: public b r ordering traits<T>
, public b r equivalence traits<T>
};
  Then you have to write the alternative blocks, which can be combined:
template <typename T>
struct b r ordering less traits
  static bool gt(const T& x, const T& y) { return y<x; }
  static bool lt(const T& x, const T& y) { return x<y; }
  static bool gteq(const T& x, const T& y) { return !(x<y); }
  static bool lteq(const T& x, const T& y) { return !(y<x); }
};
template <typename T>
struct b r equivalence equal traits
  static bool eq(const T& x, const T& y) { return x==y; }
  static bool ineq(const T& x, const T& y) { return !(x==y); }
};
template <typename T>
struct b r equivalence less traits
  static bool eq(const T& x, const T& y) { return !(x<y) && !
(y < x);
  static bool ineq(const T& x, const T& y) { return x<y || y<x; }
};
```

Finally, you combine the pieces via derivation and a hidden template parameter.

```
{
  HAS JUST OPERATOR LESS,
  HAS OPERATOR LESS AND EQ,
  HAS ALL 6 OPERATORS
};
template <typename T, int = HAS ALL 6 OPERATORS>
struct binary relation traits
: b r ordering traits<T>
, b r equivalence traits<T>
};
template <typename T>
struct binary relation traits<T, HAS JUST OPERATOR LESS>
: b r ordering less traits<T>
, b r equivalence less traits<T>
};
template <typename T>
struct binary relation traits<T, OPERATOR LESS AND EQ>
: b_r_ordering_less_traits<T>
, b r equivalence equal traits<T>
};
  Further, traits can be chained using appropriate enumerations and "bitwise-or" syntax.<sup>5</sup>
  What if you wanted to provide an enumeration set, containing powers of two that will be
combined using the standard flags idiom, but at compile time:
fstream fs("main.txt", ios::in | ios:out);
typedef binary_relation_traits<MyType, native::less | native::eq>
MyTraits;
  First, you let the flags start at 1, since you need powers of two.
namespace native
{
  enum
     = 1,
     lt_eq = 2,
     gt
              = 4,
     gt_eq = 8,
              = 16,
     eq
     ineq = 32
```

```
};
```

Second, you split the traits class into atoms, using partial specialization:

```
template <typename T, int FLAG>
struct binary_relation_traits; // no body!

template <typename T>
struct binary_relation_traits<T, native::lt>
{
   static bool lt(const T& x, const T& y) { return x<y; }
};

template <typename T>
struct binary_relation_traits<T, native::lt_eq>
{
   static bool lteq(const T& x, const T& y) { return x<=y; }
};

// and so on...</pre>
```

If the user-supplied bitmask FLAG is set to (native::ineq | ...), traits<T, FLAGS> should derive from both traits<T, native::ineq> and traits <T, FLAGS - native::ineq>.

You need an auxiliary metafunction called static_highest_bit<N>::value, which returns the index of the highest bit set in a (positive) integer N, such as the exponent of the largest power of two less or equal to N.⁶

Having this tool at your disposal, you come up with an implementation:

```
template <typename T, unsigned FLAG>
struct binary_relation_traits;

template <typename T>
struct binary_relation_traits<T, 0>
{
    // empty!
};

template <typename T>
struct binary_relation_traits<T, native::lt>
{
    static bool lt(const T& x, const T& y) { return x<y; }
};

template <typename T>
struct binary_relation_traits<T, native::gt>
{
    struct binary_relation_traits<T, native::gt>
{
}
```

```
static bool gt(const T& x, const T& y) { return x>y; }
};
// write all remaining specializations
// then finally...
template <typename T, unsigned FLAG>
struct binary relation traits
: binary relation traits<T, FLAG & (1 <<
static highest bit<FLAG>::value)>
, binary relation traits<T, FLAG - (1 <<
static highest bit<FLAG>::value)>
  // empty!
};
  Now the user can select binary relation traits members at compile time:
typedef binary relation traits<MyType, native::less
| native::eq> MyTraits;
MyType a, b;
MyTraits::lt(a,b); // ok.
MyTraits::lteq(a,b); // error: undefined
  This technique is interesting in itself, but it does not meet the original requirements, since you can
only pick "native" operators. But you can add more flags:
namespace native
{
  enum
         = 1,
     lt
     lt_eq = 2,
     qt
              = 4,
     gt_eq = 8,
     eq
             = 16,
     ineq
              = 32
  };
namespace deduce
{
  enum
     ordering = 64,
     equivalence = 128,
          = 256
     ineq
  };
```

```
}
template <typename T>
struct binary relation traits<T, deduce::ordering>
  static bool gt(const T& x, const T& y) { return y<x; }
  static bool gteq(const T& x, const T& y) { return !(x<y); }
  static bool lteq(const T& x, const T& y) { return !(y<x); }
};
template <typename T>
struct binary relation traits<T, deduce ::ineq>
{
  static bool ineq(const T& x, const T& y) { return !(x==y); }
};
template <typename T>
struct binary relation traits<T, deduce::equivalence>
  static bool eq(const T& x, const T& y) { return !(x<y) &&!
(y < x);
  static bool ineq(const T& x, const T& y) { return x<y || y<x; }
};
typedef
 binary relation traits
    MyType,
    native::less | deduce::ordering | deduce::equivalence
 >
 MyTraits;
```

Note that any unnecessary duplication (such that native::ineq | deduce::ineq) will trigger a compiler error at the first use. If traits<T, N> and traits<T, M> both have a member x, traits<T, N+M>::x is an ambiguous call.

4.3. SFINAE

The "substitution failure is not an error" (or SFINAE) principle is a guarantee that the C++ standard offers. You will see precisely what it means and how to remove function templates from an overload set when they do not satisfy a compile-time condition.

Remember that when a class template is instantiated, the compiler generates:

• Every member signature at class level

• Only strictly necessary function bodies

As a consequence, this code does not compile:

```
template <typename T>
struct A
{
   typename T::pointer f() const
   {
     return 0;
   }
};
```

As soon as A<int> is met, the compiler will try to generate a signature for *every* member function, and it will give an error because int::pointer is not a valid type. Instead, this would work:

```
template <typename T>
struct A
{
   int f() const
   {
     typename T::type a = 0;
     return a;
   }
};
```

As long as A<int>::f() is unused, the compiler will ignore its body (and that is good news, because it contains an error).

Furthermore, when the compiler meets f(x) and x has type X, it should decide which particular f(x) is being invoked, so it sorts all possible candidates from the best to the worst and tries to substitute X in any template parameter. If this replacement produces a function with an invalid signature (signature, not body!), the candidate is silently discarded. This is the SFINAE principle.

```
template <typename T>
typename T::pointer f(T*);
int f(void*);
int* x = 0;
f(x);
```

The first f would be preferred because T* is a better match than void*; however, int has no member type called pointer, so the second f is used. SFINAE applies only when the substitution produces an expression that is formally invalid (like int::pointer). Instead, it does not apply

when the result is a type that does not compile:

```
template <typename T, int N>
struct B
{
   static const int value = 100/N;
};

template <typename T>
B<T, 0> f(T*);
int f(void*);
```

B<T, 0> is a valid type, but its compilation gives an error. The first f will be picked anyway, and the compiler will stop.

To take advantage of SFINAE, when you want to "enable" or "disable" a particular overload of a function template, you artificially insert in its signature a dependent name that may resolve to an invalid expression (a non-existent type like int::pointer).

If all candidates have been discarded, you get a compiler error (trivial uses of SFINAE look in fact like static assertions).

There are two main applications of SFINAE: when f runs after being selected and when f is not executed at all.

4.3.1. SFINAE Metafunctions

Using SFINAE and sizeof, you can write metafunctions that take a decision based on the interface of a type T. This is very close to what is called *reflection* in different programming languages.

The basic ingredients are:

- Two (or more) types with different sizes; let's call them YES and NO.
- A set of overloaded functions f, where at least one must be a template, returning either YES or NO.
- A static constant defined in terms of sizeof (f (something)).

The following paradigm helps clarify this:

```
static YES<[[condition on x]]> test(X);
static NO test(...);
static T this_type();

public:
    static const bool value = sizeof(test(this_type())) !=
sizeof(NO);
};
```

The compiler has to decide which test is being called when the argument has type T. It will try to evaluate YES<[[condition on T]]> first (because void* and the ellipsis . . . have very low priority). If this generates an invalid type, the first overload of test is discarded and it will select the other.

Note some important facts:

- The static functions *need not have a body*; only their signature is used in sizeof.
- YES<T> need not have size 2. It would be an error to write sizeof(test(this_type())) == 2. However, char *must* have size 1, so you could verify if sizeof(test(this type()))>1.
- At least one of the test functions should be a template that depends on a *new* parameter X. It would be wrong to define test in terms of T (the parameter of MF), since SFINAE would not apply.
- You use a dummy function that returns T instead of, say, invoking test (T()) because T might not have a default constructor.

Some compilers will emit a warning because it's illegal to pass an object to an ellipsis function. Actually, the code does not run, since sizeof wraps the whole expression, but warnings may be long and annoying. A good workaround is to pass pointers to functions:

```
template <typename X>
static YES<[[condition on X]]> test(X*);
static NO test(...);
static T* this_type();
```

If you switch to pointers:

- void becomes an admissible type (since T* exists).
- References become illegal (a pointer to a reference is an error).

So either way, you'll have to write some explicit specialization of MF to deal with corner cases. SFINAE applies if *any* substitution of the template parameter produces an invalid type, not

necessarily in the return type. Sometimes, in fact, it's more convenient to use arguments:

```
template <typename T>
class MF
{
  template <typename X>
  static YES<void> test([[type that depends on X]]*);

  template <typename X>
  static NO test(...);

public:
  static const bool value = sizeof(test<T>(0)) != sizeof(NO);
};
```

If the substitution of X in the first expression produces a valid type, thus a valid pointer, test<T>(0) takes it as the preferred call. (It casts 0 to a typed pointer and returns YES<void> or whatever yes-type.) Otherwise, 0 is passed without any cast (as integer) to test(...), which returns NO.

The explicit call test<T> works because the ellipsis test function has a dummy template parameter; otherwise, it would never match.⁷

As a simple example, you can test if type T has a member type named pointer:

```
template <typename T>
class has pointer type
  template <typename X>
  static YES<typename X::pointer> test(X*);
  static NO test(...);
  static T* this type();
public:
  static const bool value = sizeof(test(this type())) !=
sizeof(NO);
};
or (almost) equivalently:<sup>8</sup>
template <typename T>
class has pointer type
  template <typename X>
  static YES<void> test(typename X::pointer*);
  template <typename X>
  static NO test(...);
```

```
public:
    static const bool value = sizeof(test<T>(0)) == sizeof(YES);
};
```

By modifying the template parameter of YES, you can check if T has a static constant named value. Once again, it's convenient to derive from a common yes-type:

```
// copied from Section 2.1.4
typedef char no_type;
typedef larger_than<no_type> yes_type;

template <int VALUE>
struct YES2 : yes_type
{
};

template <typename T>
class has_value
{
   template <typename X>
   static YES2<X::value> test(X*);

   // ...
};
```

Or you can check for the presence of a member function with a fixed name and signature⁹:

```
template <typename T, void (T::*F)(T&)>
struct YES3 : yes_type
{
};

template <typename T>
class has_swap_member
{
   template <typename X>
   static YES3<x, &x::swap> test(X*);

   // ...
};
```

Finally, a popular idiom checks if T is a class or a fundamental type using a fake pointer-to-member. (Literal zero can be cast to int T::* if T is a class, even if it has no member of type int.)

```
template <typename T>
class is_class
{
  template <typename X>
```

```
static yes_type test(int X::*);

template <typename X>
    static no_type test(...);

public:
    static const bool value = (sizeof(test<T>(0))!=sizeof(no_type));
};
```

4.3.2. Multiple Decisions

The examples shown so far take a single yes/no decision path, but some criteria can be more complex. Let's write a metafunction that identifies all signed integers ¹⁰:

```
if (T is a class)
       return false
if (T is a pointer)
       return false
if (T is a reference)
       return false
if (we can have a non-type template parameter of type T)
       if (the expression "T(0) > T(-1)" is well-formed and true)
             return true
       else
             return false
else
{
       return false
template <typename X, bool IS CLASS = is class<X>::value>
class is_signed_integer;
template <typename X>
class is signed integer<X*, false> : public selector<false>
};
template <typename X>
class is signed integer<X&, false> : public selector<false>
};
```

```
template <typename X>
class is signed integer<X, true> : public selector<false>
};
template <typename X>
class is signed integer < X, false >
  template <typename T>
  static static parameter<T, 0>* decide int(T*);
  static void* decide int(...);
  template <typename T>
   static selector<(T(0) > T(-1))>
decide signed(static parameter<T, 0>*);
  static selector<false> decide signed(...);
  static yes type cast(selector<true>);
  static no type cast(selector<false>);
  static X* getX();
public:
  static const bool value =
sizeof(cast(decide signed(decide int(getX())))) == sizeof(yes type);
};
  cast maps all possible intermediate return types to yes type or no type, for the final
sizeof test.
  In general, it's possible to stretch this idea and return an enumeration (more precisely, a
size t), instead of bool. Suppose you had more intermediate decision cases:
static T1 decide(int*);
static T2 decide (double*);
static Tn decide (void*);
  Then you can map T1, T2, ... In to an enumeration using fixed_size:
  static fixed size<1>::type& cast(T1);
  static fixed size<2>::type& cast(T2);
  // ...
public:
  static const size t value = sizeof(cast(decide(...)));
};
```

4.3.3. Only_If

Another interesting use of SFINAE is in excluding elements from a set of overloaded (member) functions that are not compliant with some condition:

```
template <bool CONDITION>
struct static_assert_SFINAE
{
   typedef void type;
};

template <>
struct static_assert_SFINAE<false>
{
};
```

If a function has an argument of type pointer-to-X, where X is defined as static_assert_SFINAE<...>::type, substitution of any CONDITION that evaluates to false generates an invalid expression. So that particular function is removed from the set of overloads.

The fake pointer argument has a default value of 0, which means the user can safely ignore its existence.¹¹

```
#define ONLY_IF(COND) typename static_assert_SFINAE<COND>::type*
= 0

template <typename T>
void f(T x, ONLY_IF(is_integer<T>::value))
{
}

void f(float x)
{
}

// later...

double x = 3.14;
```

This technique is often useful in universal-copy constructors of class templates:

```
template <typename T1>
class MyVector
{
public:
   // not used if T2 is T1
```

f(x); // calls f(float)

```
template <typename T2>
  MyVector(const MyVector<T2>& that)
};
  Restrictions on T2 may be easily introduced using ONLY IF (has conversion is fully
documented in Section 4.4.
template <typename T2>
MyVector(const MyVector<T2>& that,
        ONLY IF((has conversion<T2,T1>::L2R)))
}
  Another application is the "static cast" of static value. You might need to convert, say,
static value<int, 3> to static value<long, 3>:
template <typename T, T VALUE>
struct static value
   static const T value = VALUE;
   static value(const int = 0)
   }
   template <typename S, S OTHER>
      static value(const static value<S, OTHER>,
                  typename only_if<VALUE==OTHER, int>::type = 0)
};
  Sometimes it can be useful to apply the idiom, not to arguments, but to the return value:
template <bool CONDITION, typename T = void>
struct only if
   typedef T type;
};
template <typename T>
struct only if<false, T>
};
template <typename T>
typename only if < is integer < T>:: value, T>:: type multiply by 2 (const
```

```
T x)
{
   return x << 1;
}
```

This function is either ill-formed or takes a const T and returns T.

4.3.4. SFINAE and Returned Functors

template <typename T, typename S>

The various test functions you've seen so far have no use for their return type, whose size is all that matters. Sometimes they will instead return a functor that is immediately invoked. Consider a simple example, where the function number_of_elem returns x.size() if x has a type member called size type and otherwise returns 1.

```
struct get size
  S operator()(const T& x) const { return x.size(); }
  get size(int) {}
};
struct get one
  template <typename T>
  size t operator()(const T&) const { return 1; }
  get one(int) {}
};
template <typename T>
get size<T, typename T::size_type> test(const T* x) // SFINAE
  return 0;
get one test(const void*)
  return 0;
template <typename T>
size t number of elem(const T& x)
  return test(&x)(x);
```

You can use some techniques from the previous paragraph to describe an implementation of a logging callback, with a variable log level, based on metaprogramming.

In scientific computing, you can meet functions that run for a long time. So it's necessary to maintain some interaction with the function even while it's running, for example, to get feedback on the progress or to send an abort signal. Since there is no hypothesis on the environment (computational routines are usually portable), you cannot pass a pointer to a progress bar, and you have to design an equally portable interface.

A possible solution follows. The function internally updates a structure (whose type is known to its caller) with all the meaningful information about the state of the program, and it invokes a user functor regularly on the structure:

```
struct algorithm info
  int iteration current;
  int iteration max;
  double best tentative solution;
  size t time elapsed;
  size t memory used;
};
template <..., typename logger_t>
void algorithm(..., logger t LOG)
  algorithm info I;
  for (...)
     // do the work...
     I.iteration current = ...;
     I.best tentative solution = ...;
     LOG(I);
  }
}
```

You can try to design some static interaction between the logger and the algorithm so that only some relevant portion of the information is updated. If LOG does nothing, no time is wasted updating $^{\text{T}}$

First, all recordable information is partitioned in levels. logger_t will declare a static constant named log_level and the algorithm loop will not update the objects corresponding to information in ignored levels.

By convention, having no member log_level or having log_level=0 corresponds to skipping the log.

```
template <int LEVEL = 3>
struct algorithm info;
template <>
struct algorithm info<0>
};
template <>
struct algorithm info<1> : algorithm info<0>
  int iteration current;
  int iteration max;
} ;
template <>
struct algorithm info<2>: algorithm info<1>
  double best value;
};
template <>
struct algorithm info<3> : algorithm info<2>
  size t time elapsed;
  size t memory used;
};
  Second, you use SFINAE to query logger t for a constant named log level:
template <int N>
struct log level t
  operator int () const
     return N;
};
template <typename T>
log level t<T::log level> log level(const T*)
```

```
return log_level_t<T::log_level>();
}
inline int log_level(...)
{
  return 0;
}
```

Finally, a simple switch will do the work. If logger_t does contain log_level, SFINAE will pick the first overload of log_level, returning an object that's immediately cast to integer. Otherwise, the weaker overload will immediately return 0.

```
switch (log_level(&LOG))
{
   case 3:
        I.time_elapsed = ...;
        I.memory_used = ...;

   case 2: // fall through
        I.best_value = ...;

   case 1: // fall through
        I.iteration_current = ...;
        I.iteration_max = ...;

   case 0: // fall through
   default:
        break;
}
LOG(I);
```

This implementation is the simplest to code, but LOG still has access to the whole object I, even the part that is not initialized.

The static information about the level is already contained in log_level_t, so it's appropriate to transform this object into a functor that performs a cast.

```
template <int N>
struct log_level_t
{
  operator int () const
  {
    return N;
  }

  typedef const algorithm_info<N>& ref_n;
  typedef const algorithm_info< >& ref;

  ref_n operator()(ref i) const
```

```
return i;
}
};

template <typename T>
log_level_t<T::log_level> log_level(const T*)
{
    return log_level_t<T::log_level>();
}

inline log_level_t<0> log_level(...)
{
    return log_level_t<0>();
}

    switch (log_level(&LOG))
{
        // as above...
}

LOG(log_level(&LOG)(I));
```

This enforces LOG to implement an operator () that accepts exactly the right "slice" of information.

4.3.5. SFINAE and Software Updates

One of the many uses of SFINAE-based metafunctions is conditional requirement detection.

TMP libraries often interact with user types and user functors, which must usually satisfy some (minimal) interface constraint. New releases of these libraries could in principle impose additional requirements for extra optimizations, but this often conflicts with backward compatibility.

Suppose you sort a range by passing a custom binary relation to an external library function, called nonstd::sort:

```
struct MyLess
{
   bool operator()(const Person& x, const Person & y) const
   {
       // ...
   }
};
std::vector<Person> v;
nonstd::sort(v.begin(), v.end(), MyLess());
```

Version 2.0 of the sorting library requires MyLess to contain an additional function called

static void CompareAndSwap (Person& a, Person& b), so this code will not compile.

Instead, the library could easily detect if such a function is provided, and, if so, automatically invoke a faster parallel CAS-based algorithm.

This "self-detection" of features allows *independent upgrades* of the underlying libraries. This applies also to traits:

```
struct MyTraits
{
    static const bool ENABLE_FAST_ALLOCATOR = true;
    static const bool ENABLE_UTF8 = true;
    static const bool ENABLE_SERIALIZATION = false;
};

typedef nonstd::basic_string<char, MyTraits> MyString;
    Version 2.0 of the string library has a use for an extra member:

struct MyTraits
{
    static const bool ENABLE_FAST_ALLOCATOR = true;
    static const bool ENABLE_UTF8 = true;
    static const bool ENABLE_SERIALIZATION = false;

    static const size_t NUMBER_OF_THREADS = 4;
};
```

But the author of the library should not assume that this new constant is present in the traits class he receives. However, he can use SFINAE to indirectly extract this value, if it exists, or use a default:

```
template <typename T, size_t DEFAULT>
class read_NUMBER_OF_THREADS
{
  template <typename X>
    static static_value<size_t, X::NUMBER_OF_THREADS> test(X*);
  static static_value<size_t, DEFAULT> test(void*);
  template <size_t N>
    static typename fixed_size<N+1>::type&
  cast(static_value<size_t,N>);
  static T* getT();

public:
   static const size_t value = sizeof(cast(test(getT()))) - 1;
};
```

The $\pm 1/-1$ trick is necessary to avoid arrays of length zero.

The author of nonstd::basic_string will write:

```
template <typename char_t, typename traits_t>
class basic_string
{
    // ...
int n = read_NUMBER_OF_THREADS<traits_t, 4>::value;
```

So this class compiles even with older traits.

As a rule, you don't need to check that NUMBER_OF_THREADS has indeed type (static const) size_t. Any integer will do. It's possible to be more rigorous, but it is generally not worth the machinery. I am going to show all the details, but you should consider the rest of this section an exercise. You need *three* additional metafunctions:

- Detect if T has any constant named NUMBER_OF_THREADS, with the usual techniques.
- If this is false, the result is immediately false (line #2).
- Otherwise, use a different specialization, where it's legal to write
 T::NUMBER_OF_THREADS. You pass this "item" to a test function (line #1).
 The best choice is a non-template function with an argument of type
 REQUIRED_T; the other option is a template that will match everything else, so no cast can occur.

```
template <typename T>
struct has_any_NUMBER_OF_THREADS
{
  template <typename X>
    static static_value<size_t, X::NUMBER_OF_THREADS> test(X*);
  static no_type test(void*);
  template <size_t N>
    static yes_type cast(static_value<size_t, N>);
  static no_type cast(no_type);
  static T* getT();
  static const bool value = (sizeof(cast(test(getT()))) > 1);
};

template <typename REQUIRED_T, typename T, bool>
struct check_NUMBER_OF_THREADS_type;

template <typename REQUIRED_T, typename T>
struct check_NUMBER_OF_THREADS_type
template <typename REQUIRED_T, typename T>
struct check_NUMBER_OF_THREADS_type
```

4.3.6. Limitations and Workarounds

SFINAE techniques ultimately rely on the compiler handling an error gracefully, so they are especially vulnerable to compiler bugs.

If the correct code does not compile, here's a checklist of workarounds:

- Give all functions a body.
- Move static functions outside of the class, in a private namespace.
- Remove private and use struct.
- Think of a simpler algorithm.

Table 4-1. side-by-side comparison of the code, before and after the workarounds

```
namespace priv {

template <typename T>
static_value<T, 0>* decide_int(T*);

template <typename X>
class is_signed_integer
{
template <typename T>
static static_value<T, 0>* decide_int(...);

template <typename T>
static static_value<T, 0>* decide_int(T*);

selector<(T(0)>T(-1))>
decide_signed(static_value<T, 0>*);

selector<false> decide_signed(...);
```

```
template <typename T>
 static selector<(T(0)>T(-1))>
                                                             yes type cast(selector<true>);
  decide_signed(static_value<T,0>*);
                                                             no_type cast(selector<false>);
 static selector<false> decide_signed(...);
                                                             template <typename X>
                                                             struct is signed integer helper
 static yes_type cast(selector<true>);
 static no type cast(selector<false>);
                                                                X* getX();
 static X* getX();
                                                                static const bool value =
                                                                   sizeof(cast(decide signed(decide int(getX()))))
 public:
                                                                    ==sizeof(yes_type);
  static const bool value =
                                                             };
     sizeof(cast(decide_signed(decide_int(getX()))))
     = size of (yes type);
                                                             } // end of namespace
};
                                                             template <typename T>
                                                             struct is_signed_integer
                                                             : public selector<priv::is signed integer helper<T>::value>
```

A corner case in the standard is a substitution failure inside a sizeof that should bind to a template parameter. The following example usually does not compile:

```
template <typename T>
class is_dereferenceable
{
  template <size_t N>
    class YES { char dummy[2]; };

  template <typename X>
       static YES<sizeof(*x())> test(X*);

  static NO test(...);

  static T* this_type();

public:
  static const bool value = sizeof(test(this_type()))>1;
};
```

Detection of member functions is extremely problematic. Let's rewrite the metafunction here.

```
template <typename S>
class has_swap_member
{
  template <typename T, void (T::*)(T&) >
  class YES { char dummy[2]; };
  typedef char NO;
```

```
template <typename T>
  static YES<T, &T::swap> test( T* );

static NO test(...);

static S* ptr();

public:
  static const bool value = sizeof(test(ptr()))>1;
};
```

Suppose that classes D1 and D2 have a public template base called B<T1> and B<T2>, and they have no data members of their own. swap will likely be implemented only once in B, with signature void B<T>::swap (B<T>&), but the users will see it as D1::swap and D2::swap (an argument of type D1 will be cast to B<T1>&). 12

However, has_swap_member<D1>::value is false because YES<D1, &D1::swap>
does not match YES<T, void (T::*F) (T&)>.

Hypothetically, it would match either YES<T1, void (T2::*F) (T2&) > or even YES<T1, void (T1::*F) (T2&) >, but this pointer cast is out of scope, because T2 is unknown.

Furthermore, the standard explicitly says that you cannot take a pointer to a member function of a library object, because the implementation is allowed to modify the prototype, as long as the syntax works as expected. For example, you could have a perfectly valid void T::swap(T&, int = 0).

So the fact that has _swap _member <T>::value is false does not mean that the syntax a.swap(b) is illegal.

The best you can do is integrate the detection phase with the swap itself and create a function that swaps two references with the best-known method. When swap detection fails, ADL will usually find an equivalent routine in the right namespace (at least for all STL containers; see Section 1.4.2.

```
using std::swap;
struct swap_traits
{
  template <typename T>
  inline static void apply(T& a, T& b)
  {
    apply1(a, b, test(&a));
  }
private:
  template <typename T, void (T::*F)(T&)>
  struct yes : public yes_type
  {
    yes(int = 0)
    {}
};
```

```
template <typename T>
   static yes<T, &T::swap> test(T*)
   { return 0; }
  static no type test(void*)
   { return 0; }
  template <typename T>
   inline static void apply1(T& a, T& b, no type)
     swap(a, b);
  template <typename T>
  inline static void apply1(T& a, T& b, yes type)
     a.swap(b);
};
template <typename T>
inline void smart swap(T& x, T& y)
  swap traits::apply(x, y);
  Note that all functions have a body, as they are truly invoked.
  The workflow is as follows. smart swap (x, y) invokes apply, which in turn is
apply (x, y, [[condition on T]]). apply 1 is an ADL swap when the condition is no
and a member swap invocation otherwise.
#include <map>
struct swappable
  void swap(swappable&)
};
int main()
  std::map<int, int> a, b;
                               // if it fails detection of map::swap
  smart swap(a, b);
                          // then it uses ADL swap, which is the same
   swappable c, d;
```

// correctly detects and uses

smart swap(c, d);

swappable::swap

```
int i = 3, j = 4;
smart_swap(i, j);  // correctly uses std::swap
}
```

■ **Note** The true solution requires the C++0x keyword decltype. See Section 12.2.

One final caveat is to avoid mixing SFINAE with private members.

```
The C++ 2003 Standard says that access control occurs after template deduction. So, if
T:: type exists but it's private, SFINAE will select an action based on the information that
T:: type actually exists, but a compiler error will generally occur immediately after (since
T:: type is inaccessible). ^{13}
template <typename T>
typename T::type F(int);
template <typename T>
char F(...);
class X
        typedef double type; // note: private, by default
};
// A condensed version of the usual SFINAE machinery...
// We would expect the code to compile and N==1.
// This occurs only in C++0x
int N = sizeof(F < X > (0));
error: type "X::type" is inaccessible
  typename T::type F(int);
          detected during instantiation of "F" based on template
```

4.3.7. SFINAE with Partial Specializations

argument <X>

SFINAE applies also to partial specialization of class templates. When a condition that should be used to select the partial specialization is ill-formed, that specialization is silently removed from the set of candidates. This section shows a practical application with an example.¹⁴

Suppose you have a template class called A<T>, which you want to specialize when type T contains a typedef called iterator.

You start by adding a second template parameter to A and a partial specialization on the second (you will define DEFAULT TYPE and METAFUNC later):

```
template <typename T, typename X = DEFAULT_TYPE>
struct A
{ ... };

template <typename T>
struct A<T, typename METAFUNC<typename T::iterator>::type >
{ ... };
```

According to SFINAE, when T::iterator does not exist, the specialization is ignored and the general template is used. However, when T::iterator indeed exists (and METAFUNC is well defined), both definitions are valid. But according to the C++ language rules, if DEFAULT_TYPE happens to be the same as METAFUNCTION<T::iterator>::type, the specialization of A is used. Let's rewrite the example more expliticly:

```
template <typename T>
struct METAFUNC
{
   typedef int type;
};

template <typename T, typename X = int>
struct A
{ ... };

template <typename T>
struct A<T, typename METAFUNC<typename T::iterator>::type >
{ ... };

A<int> a1; // uses the general template
A<std::vector<int>> a2; // uses the specialization
```

4.4. Other Classic Metafunctions with Sizeof

```
An overload may be selected because the argument can be cast successfully.

This section shows a metafunction that returns three Boolean constants

—has_conversion<L,R>::L2R is true when L (left) is convertible to R (right) and has_conversion<L,R>::identity is true when L and R are the same type. 15

template <typename L, typename R> class has_conversion
{
    static yes_type test(R);
    static no_type test(...);
    static L left();
```

public:

```
static const bool L2R = (sizeof(test(left())) ==
sizeof(yes_type));
   static const bool identity = false;
};

template <typename T>
class has_conversion<T, T>
{
   public:
    static const bool L2R = true;
   static const bool identity = true;
};
```

This code passes a fake instance of L to test. If L is convertible to R, the first overload is preferred, and the result is yes_type.

Following Alexandrescu, ¹⁶ you can deduce whether a type publicly derives from another:

```
template <typename B, typename D>
struct is_base_of
{
   static const bool value =
   (
     has_conversion<const D*, const B*>::L2R &&
    !has_conversion<const B*, const void*>::identity
   );
};
```

Implicit promotion techniques have been extensively used by David Abrahams. ¹⁷ The key point is to overload an operator at namespace level, not as a member.

```
sizeof(no_type);
};
```

The ++getT() statement can either resolve to x's own operator++ or (with lower priority) resolve to a conversion to fake_incrementable, followed by fake_incrementable increment. This latter function is visible, because, as anticipated, it is declared as a global entity in the namespace, not as a member function.

To test post-increment, replace line #1 with:

```
fake_incrementable operator++(fake_incrementable, int);
```

Note that the computation of sizeof(test(++x)) must be done in the namespace where fake incrementable lives. Otherwise, it will fail:

You can also move the computation inside the namespace and recall the result outside:

```
namespace aux {
// ... (all as above)

template <typename T>
struct has_preincrement_helper
{
   static T& getT();
   static const bool value = sizeof(test(++getT())) ==
```

```
sizeof(no_type);
};

template <typename T>
struct has_preincrement
: selector<aux::has_preincrement_helper<T>::value>
{
};
```

4.5. Overload on Function Pointers

One of the most convenient tag objects used to select an overloaded function is a function pointer, which is then discarded.

A pointer is cheap to build yet can convey a lot of static information, which makes it suitable for template argument deduction.

4.5.1. Erase

The following is the primary example. It iterates over an STL container, so you need to erase the element pointed to by iterator i. Erasure should advance (not invalidate) the iterator itself. Unfortunately, the syntax differs. For some containers, the right syntax is i = c.erase(i), but for associative containers it is c.erase(i++).

Taking advantage of the fact that C::erase must exist (otherwise you wouldn't know what to do and the call to erase_gap would be ill formed), you just pick the right one with a dummy pointer:

```
template <typename C, typename iterator_t, typename base_t>
void erase_gap2(C& c, iterator_t& i, iterator_t (base_t::*)
(iterator_t))
{
    i = c.erase(i);
}

template <typename C, typename iterator_t, typename base_t>
void erase_gap2(C& c, iterator_t& i, void (base_t::*)(iterator_t))
{
    c.erase(i++);
}

template <typename C>
void erase_gap(C& c, typename C::iterator& i)
{
    erase gap2(c, i,&C::erase);
```

```
int main()
{
   for (i = c.begin(); i != c.end(); )
   {
      if (need_to_erase(i))
          erase_gap(c, i);
      else
          ++i;
   }
}
```

Observe that erasure is *not* invoked via the pointer. It's just the type of the pointer that matters. Also, the type of erase may not be . . . (C::*) (...), because a container could have a "hidden base". The exact type is therefore left open to compiler deduction.

4.5.2. Swap

}

The previous technique can be extended via SFINAE to cases where it's unknown if the member function exists. To demonstrate, you need to extend swap_traits (introduced in Section 4.3.6) to perform the following 18:

- If T has void T::swap(T&), use a.swap(b).
- If T has static void swap (T&, T&), use T:: swap (a, b).
- If T has both swaps, the call is ambiguous.
- In any other case, use ADL swap.

The first part simply reuses the techniques from the previous sections. In particular, observe that all yes-types derive from a common "yes-base," because the first test is meant only to ensure that the possible swap member functions exist.

```
struct swap_traits
{
  template <typename T, void (T::*F)(T&)>
  class yes1 : public yes_type {};

  template <typename T, void (*F)(T&, T&)>
  class yes2 : public yes_type {};

  template <typename T>
  inline static void apply(T& a, T& b)
  {
    apply1(a, b, test(&a));
  }
}
```

```
private:
    // first test: return a yes_type* if any allowed T::swap exists

template <typename T>
    static yes1<T, &T::swap>* test(T*)
{ return 0; }

template <typename T>
    static yes2<T, &T::swap>* test(T*)
{ return 0; }

static no_type* test(void*)
{ return 0; }
```

When the test is false, call ADL swap. Otherwise, perform a *function-pointer based* test. Call apply2 by taking the address of swap, which is known to be possible because at least one swap exists.

```
private:
  template <typename T>
  inline static void apply1(T& a, T& b, no type*)
     swap(a, b);
  }
  template <typename T>
  inline static void apply1(T& a, T& b, yes type*)
  {
     apply2(a, b,&T::swap);
  }
  template <typename T>
  inline static void apply2(T& a, T& b, void (*)(T&, T&))
     T::swap(a, b);
  }
  template <typename T, typename BASE>
  inline static void apply2(T& a, T& b, void (BASE::*)(BASE&))
  {
     a.swap(b);
  }
  template <typename T>
  inline static void apply2 (T& a, T& b, ...)
     swap(a, b);
```

```
};
```

4.5.2. Argument Dominance

When a function template has several arguments whose type must be deduced, you may incur ambiguities:

```
template <typename T>
T max(T a1, T a2) { ... }

max(3, 4.0); // error: ambiguous, T may be int or double
```

It's often the case that one argument is more important, so you can explicitly instruct the compiler to ignore everything else during type deduction:

```
// here T must be the type of arg1
template <typename T>
void add_to(T& a1, T a2) { ... }
double x = 0;
add_to(x, 3); // we would like this code to compile
```

The solution to this is to replace T with an indirect metafunction that yields the same result. Type deduction is performed only on non-dependent names, and the compiler then ensures that the result is compatible with any other dependent name:

```
template <typename T>
void add_to(T& a1, typename instance_of<T>::type a2)
{ ... }
```

In this example, T& is viable for type-detection. T=double is the only match. instance_of<double> does indeed contain a type called type (which is double), so the match is feasible. So the function automatically casts a 2 to double.

This idiom is very popular when a1 is a function pointer and a2 is the argument of a1:

```
template <typename A, typename R>
R function_call(R (*f)(A), R x)
{ return f(x); }
```

The function pointer is a dominating argument, because you can call f on everything that is convertible. You should therefore consider disabling the detection on x:

```
template <typename A, typename R>
R function_call(R (*f)(A), typename instance_of<R>::type x)
{ return f(x); }
```

- ¹Sections 5.3 and 5.3.1 show how to detect if T has a member function T T::abs() const.
- ²Of course, you could have written a method that takes selector<false>, but using a template as a replacement for C ellipsis can be of some interest.
- ³An illegal statement would be, for example, a call to an undeclared function. Recall that compilers are not required to diagnose errors in templates that are not instantiated. MSVC skips even some basic syntax checks, while GCC does forbid usage of undeclared functions and types. See also Section 5.2.3 about platform specific traits.
- ⁴Mostly, the choice will depend on release and pointer being independent or provided by the same traits.
- ⁵See Section 2.3.3.
- ⁶The details of static highest bit are in Section 3.4.1.
- ⁷See Section 1.2.1.
- ⁸This would fail if X::pointer were a reference; at the moment, you don't need to worry about this.
- ⁹The swap-detection problem is actually much more difficult; it's discussed later in this section.
- 10The "main algorithm" alone would not suffice. It will work when T is a fundamental type. Some compilers evaluate the expression T (0) < T (-1) as true when T is a pointer; other compilers will give errors if T is a type with no constructor. That's why you add explicit specializations for pointers, references, and class types. Note, however, that this approach is superior to an explicit list of specializations, because it's completely compiler/preprocessor independent.
- Sometimes it's desirable to document C++ code, not literally, but just as the user is supposed to use it. This kind of functional documentation is also a part of C++ style. The example illustrated here documents that f(T) is a single argument function, even if it's not. All the implementation details should be hidden.
- 12This may look like a corner case, but it's quite common. In popular STL implementation, let D1=std::map, D2=std::set and B<T> be an undocumented class that represents a balanced tree.
- ¹³This was changed in the C++11 Standard. See http://www.open-std.org/jtc1/sc22/wg21/docs/cwg_defects.html#1170.
- ¹⁴Walter Brown recently made this technique popular. See http://www.openstd.org/jtc1/sc22/wg21/docs/papers/2014/n3911.
- ¹⁵The left-right notation may not be the most elegant, but it's indeed excellent for remembering how the class works.
- ¹⁶See the bibliography.
- ¹⁷boost::is_incrementable correctly strips qualifiers from T, but it allows operator++ to return void, which in general is not desirable. In this case, the simpler version presented here gives a compile-time error.
- ¹⁸This extension is to be considered an exercise, but not necessarily a good idea.

CHAPTER 5

Interfaces

Templates are used as interfaces in two different ways: to provide sets of atomic functions and to obtain *compile-time polymorphism*.

If several functions use the same portion of the interface of an object, you can factor them out in a single template:

```
void do something(std::vector<double>& v)
  if (v.empty())
     // ...
   ... v.size();
  for each(v.begin(), v.end(), my functor());
}
void do_something(std::list<double>& L)
  if (L.empty())
     // ...
  ... L.size();
  for each(L.begin(), L.end(), my functor());
becomes:
template <typename T>
void do_something(T& L)
  if (L.empty())
   ... L.size();
```

```
for_each(L.begin(), L.end(), my_functor());
...
}
```

This code unification is simpler when you follow common guidelines for containers (as listed in Section 1.4).

If necessary, as described in Section 3.4.3, you can replace calls to *member* functions with calls to small *global* functions. Assume you have a third do_something that executes a slightly different test:

```
void do_something(MyContainer<double>& M)
{
  if (M.size() == 0)
  ...
```

It's better to isolate the test for "emptiness" in a different function:

```
template <typename T>
bool is_empty(const T& c)
{
   return c.empty();
}

template <typename T>
bool is_empty(const MyContainer<T>& c)
{
   return c.size() == 0;
}

template <typename T>
void do_something(T& L)
{
   if (is_empty(L))
   ...
```

5.1. Wrapping References

A class template and its specializations can be used to make interfaces uniform:

```
class Dog
{
   public:
     void bark();
     void go_to_sleep();
};
```

```
class Professor
  public:
     void begin lesson();
     void end lesson();
} ;
template <typename T>
class Reference
  T& obj ;
public:
  Reference(T& obj) : obj (obj) {}
  void start talking() { obj .talk(); }
  void end talking() { obj .shut(); }
};
template <>
class Reference<Dog>
  Dog& obj ;
public:
  Reference(Dog& obj) : obj (obj) {}
  void start talking() { for (int i=0; i<3; ++i) obj .bark(); }
  void end talking() { obj .go to sleep(); }
};
template <>
class Reference < Professor >
  Professor& obj ;
public:
  Reference(Professor& obj) : obj (obj) {}
  void start talking() { obj .begin lesson(); }
  void end talking() { obj .end lesson(); }
};
  Note that the wrapper may indeed contain some logic. Finally:
template <typename T>
void DoIt(T& any)
  Reference<T> r(any);
  r.start talking();
```

```
// ...
r.end_talking();
}
```

5.2. Static Interfaces

When a function template manipulates an object of unspecified type T, it actually forces the object to implement an interface. For example, this very simple function contains a lot of hidden assumptions about the (unknown) types involved:

```
template <typename iter1_t, typename iter2_t>
iter2_t copy(iter1_t begin, const iter1_t end, iter2_t output)
{
  while (begin != end)
    *(output++) = *(begin++),

  return output;
}
```

Here, iter1_t and iter2_t must have a copy constructor, called operator++ (int). iter1_t also needs operator!=. Furthermore, every operator++ returns a dereferenceable entity, and in the case of iter2_t, the final result is an l-value whose assignment blindly accepts whatever * (begin++) returns.

In short, template code pretends that all instructions compile, until the compiler can prove they don't.

In general, it's too verbose and/or generally not useful to list the assumptions on a type interface. In the previous example, iter1_t::operator++ will likely return iter1_t, which also implements operator*, but it need not be *exactly* the case (for instance, copy would work if, say, iter1 t::operator++ returned int*).

So you must try to list explicitly a minimal set of *concepts* that the template parameter must satisfy. Informally, a concept is a requirement on the type that implies that a C++ statement is legal, whatever its implementation.¹

For example, this object will happily play the role of iter2_t:

```
struct black_hole_iterator
{
   const black_hole_iterator& operator++ () const
   {
     return *this;
   }
   const black_hole_iterator& operator++ (int) const
   {
     return *this;
   }
}
```

```
const black_hole_iterator& operator* () const
{
    return *this;
}

template <typename T>
    const black_hole_iterator& operator= (const T&) const
{
    return *this;
}
};
```

Here, the concept of "the object returned by operator* must be an l-value" is satisfied, even if in an unusual way (the assignment does not modify the black hole).

Generally, you won't list the exact concepts for any generic function. However, some sets of concepts have a standard name, so whenever possible, you'll adopt it, even if it's a superset of what is actually needed.

In the previous copy template, it's best to use an *input iterator* and an *output iterator*, because these are the smallest universally known labels that identify a (super-)set of the concepts. As you will read in Chapter 6, a true output iterator satisfies a few more properties (for example, it must provide some typedefs, which are irrelevant here); however, this is a fair price for reusability.²

Authors of template code often need to make concepts explicit. If they have a simple name, they can be used as template parameters:

```
template <typename FwdIter, typename RandIter>
FwdIter special_copy(RandIter beg, RandIter end, FwdIter output);
```

Note that in this function, nothing constrains beg to be an iterator except names (which are hints for humans, not for the compiler). The template argument FwdIter will match *anything*, say double or void*, and if you are lucky, the body of the function will report errors. It may happen that you pass a type that works, but it does not behave as expected.³

On the other hand, classic C++ does offer a tool to constrain types: inheritance. You write pieces of code that accept a BASE* and at runtime they invoke the right virtual functions.

Static interfaces are their equivalent in TMP. They offer less generality than a "flat" type \mathbb{T} , but have the same level of static optimizations.

A *static interface* is a skeleton class that limits the scope of validity of a template to types derived from the interface, and at the same time it provides a default (static) implementation of the "virtual" callback mechanism.

The details follow.

5.2.1. Static Interfaces

The original language idiom was called the "curiously recurring template" pattern (*CRTP*) and it is based on the following observation: a static_cast can traverse a class hierarchy using only

compile-time information. Put simply, static_cast can convert BASE* to DERIVED*. If the inheritance relationship between DERIVED and BASE is incorrect or ambiguous, the cast will not compile. However, the result will be valid only if at runtime BASE* is pointing to a true DERIVED object.

As a special case, there's an easy way to be sure that the cast will succeed; that is, when each derived class inherits from a "personal base":

```
template <typename DERIVED_T>
class BASE
{
   protected:
     ~BASE() {};

class DERIVED1 : public BASE<DERIVED1>
{
};

class DERIVED2 : public BASE<DERIVED2>
{
};
```

An object of type BASE<T> is guaranteed to be the base of a T, because thanks to the protected destructor, nobody except a derived class can build a BASE<T>, and only T itself derives from BASE<T>.

So BASE<T> can cast itself to T and invoke functions:

```
template <typename DERIVED_T>
struct BASE
{
    DERIVED_T& true_this()
    {
       return static_cast<DERIVED_T&>(*this);
    }

    const DERIVED_T& true_this () const
    {
       return static_cast<const DERIVED_T&>(*this);
    }

    double getSomeNumber() const
    {
       return true_this().getSomeNumber();
    }
};

struct DERIVED_rand : public BASE<DERIVED_rand>
{
```

```
double getSomeNumber() const
     return std::rand();
};
struct DERIVED circle: public BASE<DERIVED circle>
   double radius ;
  double getSomeNumber() const
     return 3.14159265359 * sq(radius);
};
  Exactly as for virtual functions, normal calls via the derived class interface are inexpensive:
DERIVED rand d;
d.getSomeNumber();
                              // normal call; BASE is completely
ignored
  However, you can write a function template that takes a reference-to-base and makes an
inexpensive call to the derived member function. true this will produce no overhead.
template <typename T>
void PrintSomeNumber(BASE<T>& b) // crucial: pass argument by
reference
   // here BASE methods will dispatch to the correct T equivalent
   std::cout << b.getSomeNumber();</pre>
}
DERIVED circle C;
DERIVED rand R;
PrintSomeNumber(C); // prints the area of the circle
PrintSomeNumber(R); // prints a random number
  Conceptually, the previous function is identical to the simpler (but vaguer) function here:
template <typename T>
void PrintSomeNumber(T& b)
  std::cout << b.getSomeNumber();</pre>
```

However, the replacement looks acceptable because PrintSomeNumber is a named function, not an operator (think about writing a global operator+ with two arguments of type T). The

following example demonstrates the use of static interfaces with operators.⁴ It will implement only operator+= and have operator+ for free, simply deriving from the summable<...> interface.

```
template <typename T>
struct summable
  T& true_this()
     return static cast<T&>(*this);
  const T& true this () const
     return static cast<const T&>(*this);
  T operator+ (const T& that) const
      T result(true this());
     result += that; // call dispatch to native
T::operator+=
     return result;
  }
};
struct complex number : public summable < complex number >
{
  complex number& operator+= (const complex number& that)
};
complex number a;
complex number b;
complex number s = a+b;
```

The (apparently simple) last line performs the following compile-time steps:

- a does not have an operator+ of its own, so cast a to its base that has it, namely const summable<complex number>&.
- const summable<complex_number>& can be summed to a complex_number, so b is fine as is.
- summable<complex_number>::operator+builds a complex number named result, which is a copy of true this, because

```
true this is a complex number.
```

• Dispatching execution to complex_number::operator+=, the result is computed and returned.

Note that you could rewrite the base class as:

```
template <typename T>
struct summable
{
    // ...

    T operator+ (const summable<T>& that) const
    {
        T result(true_this());
        result += that.true_this();
        return result;
    }
};
```

Let's call *interface* the base class and *specializations* the derived classes.

5.2.2. Common Errors

You just met a situation where the interface class makes a specialized copy of itself:

```
T result(true this());
```

This is not a problem, since the interface, which is static, knows its "true type" by definition. However, the correct behavior of true this can be destroyed by *slicing*:

Usually, it's necessary to declare BASE destructor non-virtual and protected, and sometimes it's a good idea to extend protection to the copy constructor. Algorithms should not need to make a copy of the static interface. If they need to clone the object, the correct idiom is to call the DERIVED_T constructor and pass true this(), as shown previously.

```
template <typename DERIVED_T>
struct BASE
{
    DERIVED_T& true_this()
    {
```

```
return static cast<DERIVED T&>(*this);
  }
  const DERIVED T& true this() const
    return static cast<const DERIVED T&>(*this);
protected:
  ~BASE()
  BASE (const BASE &)
};
  The interface of DERIVED is visible only inside the body of BASE member functions:
template <typename DERIVED T>
struct BASE
  // ...
  void f()
    typedef DERIVED T::someType someType; // ok here
};
class DERIVED : public BASE < DERIVED >
```

Typedefs and enums from DERIVED are not available at class level in BASE. This is obvious, because DERIVED is compiled after its base, which is BASE<DERIVED>. When BASE<DERIVED> is processed, DERIVED is known, but still incomplete.

It's a good idea (not an error) to make BASE expose a typedef for DERIVED_T. This allows external functions to make a specialized copy of BASE.

```
template <typename DERIVED_T>
struct BASE
{
   typedef DERIVED_T static_type;
```

However, DERIVED cannot access BASE members without full qualification, because a template

```
base class is out of scope for the derived objects.<sup>5</sup>
template <typename DERIVED T>
struct BASE
  typedef double value type;
  value type f() const
     return true this().f();
  // ...
struct DERIVED1 : public BASE<DERIVED1>
  value type f() const // error: value type is undefined
                               // error: true this is undefined
     true this();
     return 0;
};
struct DERIVED2 : public BASE<DERIVED2>
  BASE<DERIVED2>::value type f() const // ok
     this->true this();
                                             // ok
     return 0;
};
  Note once again that scope restriction holds only "inside" the class. External users will correctly
see DERIVED1::value type:
template <typename T>
struct value type of
  typedef typename T::value type type;
};
double
  Finally, the developer must ensure that all derived classes correctly announce their names to the
```

class DERIVED1 : public BASE<DERIVED1>

base in order to avoid a classic copy and paste error:

```
{
};
class DERIVED2 : public BASE<derived1>
{
};
```

Benefits	Problems
Write algorithms that take "not too generic objects" and use them with a statically known interface.	The developer must ensure that all algorithms take arguments by reference and avoid other common errors.
Implement only some part of the code in the derived (specialized) class and move all common code in the base.	Experimental measurements suggest that the presence of non-virtual protected destructors and multiple inheritance may inhibit or degrade code optimizations.

5.2.3. A Static_Interface Implementation

Many of the previous ideas can be grouped in a class:

template <typename T>

```
struct clone of
  typedef const T& type;
};
template <typename static type, typename aux t = void>
class static interface
public:
  typedef static type type;
  typename clone of < static type >::type clone() const
     return true this();
protected:
  static interface() {}
  ~static_interface() {}
  static type& true this()
     return static cast<static type&>(*this);
  const static type& true this() const
```

```
return static cast<const static type&>(*this);
};
```

You'll come back to the extra template parameter later in this chapter.

The helper metafunction clone of can be customized and returning const reference is a

```
reasonable default choice. For small objects, it may be faster to return a copy:
template <typename T, bool SMALL OBJECT = (sizeof(T)
<sizeof(void*))>
struct clone of;
template <typename T>
struct clone of<T, true>
   typedef T type;
};
template <typename T>
struct clone of<T, false>
   typedef const T& type;
};
   First, you make some macros available to ease interface declaration.
   An interface is defined by
#define MXT INTERFACE (NAME)
template <typename static type>
class NAME : public static interface < static type >
#define MXT SPECIALIZED this->true this()
   Here's a practical example. The interface macro is similar to a normal class declaration.<sup>6</sup>
MXT INTERFACE(random)
protected:
   ~random()
```

public:

typedef double random type;

random type max() const

```
return MXT_SPECIALIZED.max();
}

random_type operator()() const
{
   return MXT_SPECIALIZED(); // note operator call
}
};
```

- random can access true_this() only with explicit qualification (as MXT SPECIALIZED does).
- random needs to declare a protected destructor.
- static_type is a valid type name inside random, even if static_interface is out of scope, because it's the template parameter name.

Now let's implement some random algorithms:

```
#define MXT SPECIALIZATION(S, I)
                                                   class S : public I<</pre>
S >
MXT SPECIALIZATION (gaussian, random)
  public:
     double max() const
     {
        return std::numeric limits<double>::max();
     }
     double operator()() const
        // ...
};
MXT SPECIALIZATION (uniform, random)
  public:
     double max() const
     {
        return 1.0;
     }
     // ...
```

What if you need a template static interface, such as:

```
template <typename RANDOM T, typename SCALAR T>
class random
  public:
     typedef SCALAR T random type;
     // ...
};
template <typename T>
class gaussian : public random < gaussian < T>, T>
  // ...
};
  It's easy to provide more macros for template static interfaces (with a small number of
parameters). A naïve idea is:
#define MXT TEMPLATE INTERFACE (NAME, T)
template <typename static type, typename T>
class NAME : public static interface < static type >
#define MXT TEMPLATE SPECIALIZATION(S,I,T)
template <typename T>
class S : public I< S<T> >
  Which is used like this:
MXT TEMPLATE INTERFACE (pseudo array, value t)
protected:
  ~pseudo array()
public:
  typedef value t value type;
  value type operator[](const size t i) const
     return MXT SPECIALIZED.read(i, instance of < value type > ());
   }
  size t size() const
```

```
{
    return MXT_SPECIALIZED.size(instance_of<value_type>());
};
```

A non-template class can use a template static interface. For example, you could have a bitstring class that behaves like an array of bits, an array of nibbles, or an array of bytes:

```
typedef bool bit;
typedef char nibble;
typedef unsigned char byte;
class bitstring
: public pseudo_array<bitstring, bit>
, public pseudo_array<bitstring, nibble>
, public pseudo_array<bitstring, byte>
{
```

An interface need not respect the same member names as the true specialization. In this case, operator[] dispatches execution to a function template read. This makes sense, because the underlying bitstring can read the element at position i in many ways (there are three distinct i-th elements). But inside pseudo_array, the type to retrieve is statically known, so using a bitstring as a pseudo_array is equivalent to "slicing" the bitstring interface. This makes code much simpler.

The first problem you need to solve is that when the macro expands, the compiler reads:

```
template <typename static_type, typename value_t>
class pseudo_array : public static_interface<static_type>
```

Thus bitstring inherits multiple times from static_interface
bitstring>, which will make the static cast in true this ambiguous.

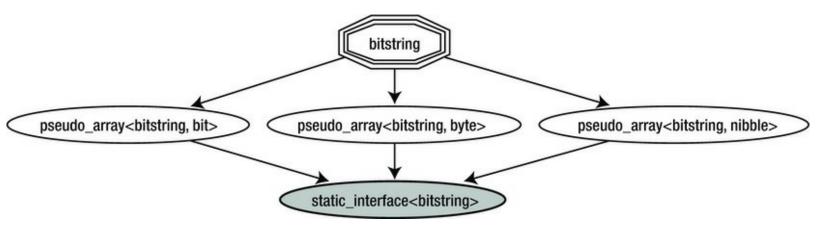


Figure 5-1. Ambiguous inheritance diagram

To avoid this issue, use an extra parameter in the static interface for disambiguation. The most unambiguous type names are either T or the whole interface (pseudo_array
bitstring, T>). The macro becomes:

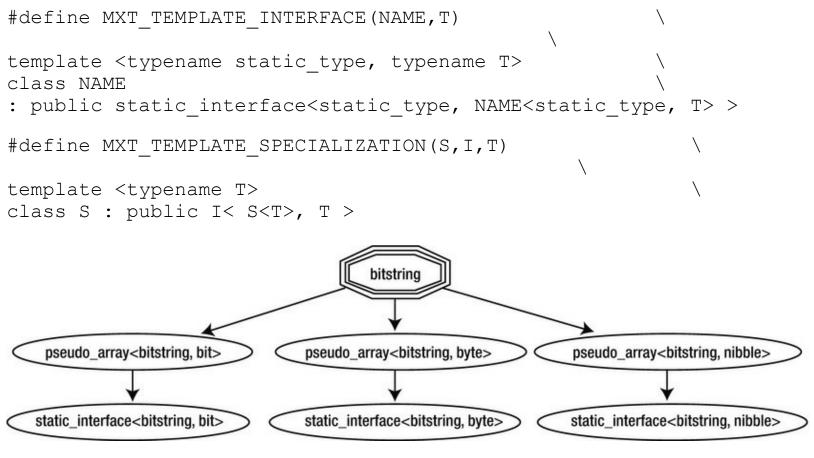


Figure 5-2. Improved inheritance diagram

5.2.4. The Memberspace Problem

Up to now, static interfaces have been described as techniques that limit the scope of some template parameters. So instead of F(T), you write F(random < T) where T is a special implementation of a random generator. This is especially useful if F is indeed a (global) operator.

A second application of static interfaces is the *memberspace* problem.⁷ The name memberspace is the equivalent of a namespace, relative to the member functions of a class. In other words, it's sort of a subspace where a class can put member functions with duplicate names.

Assume that C is a container that follows the STL conventions, so the first element of C is *begin() and the last is *rbegin().

This is the classic solution to partition an interface where function names have a unique prefix/suffix, such as push+front, push+back, r+begin, and so on.

It's better to have a real partition, where front and back are both containers with their own interfaces:⁸

```
C MyList;
// ...

first = MyList.front.begin();
last = MyList.back.begin();

MyList.front.push(3.14);
MyList.back.push(6.28);
```

```
MyList.back.pop();
  Indeed, you can use static interfaces to write code such as:<sup>9</sup>
class bitstring
: public pseudo array<br/>bitstring, bit>
, public pseudo array<br/>bitstring, nibble>
, public pseudo array<br/>bitstring, byte>
  char* data ;
  size t nbits ;
public:
  pseudo array<bitstring, bit>& as bit() { return *this; }
  pseudo array<bitstring, nibble>& as nibble() { return *this; }
  pseudo array<bitstring, byte>& as byte() { return *this; }
  size t size(instance of<byte>) const { return nbits / CHAR BIT;
  size t size(instance of<bit>) const { return nbits ; }
  size t size(instance of<nibble>) const { return nbits
/ (CHAR BIT / 2); }
  bit read(size t n, instance of<byte>) const { return ...; }
  nibble read(size t n, instance of<bit>) const { return ...; }
  byte read(size_t n, instance of<nibble>) const { return ...; }
};
bitstring b;
int n1 = b.as bit().size();
int n2 = b.as byte().size();
  Compare that with:
bitstring b;
int n1 = b.size(instance of<bit tag>());
  b.as bit () is also sort of a container of its own, and it can be passed by reference to
algorithms:
template <typename T, typename X>
X parity(pseudo array<T, X>& data)
  X \text{ result} = 0;
  for (size t i=0; i<data.size(); ++i)</pre>
     result ^= data[i];
  return result;
```

}

This technique is excellent, but it suffers from a limitation. As mentioned, typedefs provided in the specialization are not available in the static interface, thus you have no way of declaring a member function returning an iterator. This is because the static interface has to borrow the iterator type from the specialization.

```
MXT INTERFACE (front)
  return MXT SPECIALIZED.begin();
                                         // <-- error again
  typename static type::iterator end()
     return MXT SPECIALIZED.end();
} ;
MXT INTERFACE (back)
  typename static type::reverse iterator begin() // <-- another error
     return MXT SPECIALIZED.rbegin();
  typename static_type::reverse_iterator end() // <-- lots of errors</pre>
     return MXT SPECIALIZED.rend();
};
class C : public front<C>, public back<C>
  // ...
public:
  front<C>& front()
  { return *this; }
  back<C>& back()
  { return *this; }
};
C MyList;
MyList.front().begin(); // error
```

```
MyList.back().begin();  // error
// ...
```

Note that it's not a matter of syntax. Since C is still incomplete, C::iterator does not yet exist. However, there are some design fixes:

• Define iterator before C:

```
class C_iterator
{
    // ...
};
class C
{
    // container implementation
    typedef C_iterator iterator;
};
```

class C

• Insert an additional layer between C and the interfaces, so that the static interface compiles after C (and before the wrapper class):

```
// container implementation
  class iterator { ... };
} ;
MXT TEMPLATE INTERFACE (front, impl t)
  typename impl t::iterator begin()
     return MXT SPECIALIZED.begin();
  typename impl t::iterator end()
     return MXT SPECIALIZED.end();
};
// ...
class C WRAPPER: public front<C WRAPPER, C>, public
back<C WRAPPER, C>
  C c ;
```

```
public:
    // reproduce C's interface
    // dispatch all execution to c_
    typename C::iterator begin()
    {
      return c_.begin();
    }
    // ....
};
```

5.2.5. Member Selection

The same technique used in merging traits (see Section 4.2.4 can be successfully applied to value objects. The next listing, which is intentionally incomplete, suggests a possible motivation:

```
enum
{
  empty = 0, year = 1,
  month = 2,
  day = 4,
  // ...
};
template <unsigned CODE> struct time val;
template <> struct time val<empty> { }; // empty, I really mean it
template <> struct time val<year> { int year; };
template <> struct time val<month> { short month; };
// ...
template <unsigned CODE>
struct time val
: public time val<CODE & static highest bit<CODE>::value>
, public time val<CODE - static highest bit<CODE>::value>
};
// an algorithm
template <unsigned CODE>
time val<(year | month | day)> easter(const time val<CODE>& t)
```

```
time_val<(year | month | day) > result;

result.year = t.year;
result.month = compute_easter_month(t.year);
result.day = compute_easter_day(t.year);

return result;
}

time_val<year | month > tv1;
time_val<month | day > tv2;
easter(tv1);  // ok.
easter(tv2);  // error: tv2.year is undefined.
```

Note that the algorithm acts unconditionally as if any time_val<CODE> had a member year. When necessary, you can isolate this assumption using a wrapper:

```
template <unsigned CODE>
time_val<year | month | day> easter(const time_val<CODE>& t,
selector<true>)
{
    // implementation
}

template <int CODE>
time_val<year | month | day> easter(const time_val<CODE>& t,
selector<false>)
{
    // put whatever here: throw exception, static assert...
}

template <int CODE>
time_val<year | month | day> easter(const time_val<CODE>& t)
{
    return easter(t, selector<CODE & year>());
}
```

5.3. Type Hiding

Classic C++ programs transform instances of objects into other instances that have possibly different types (via function calls).

```
int i = 314; double x = f(i); // transform an instance of int into an instance of double
```

Using templates, C++ can manipulate instances, compile-time constants, and types (constants are in the middle because they share some properties with both). You can transform types and constants into instances (trivially), types into types (via traits and metafunctions), types into constants (via metafunctions and other operators, such as sizeof), instances into constants (via sizeof), and types into some special system objects (using typeid). However, classic C++ has very limited language tools to transform an instance into a type. 10

The most common example comes from iterator handling:

```
T t = *begin; // store a copy of the first element // who is T?
```

At the moment, a suitable type is provided by metafunctions:

```
typename std::iterator_traits<iterator_t>::value_type t = *begin;
```

There are tricks, which essentially avoid direct knowledge of T. The simplest option is to pass *begin as a dummy unused parameter to a template function that will deduce its type:

```
template <typename iterator_t>
void f(iterator_t beg, iterator_t end)
{
   if (beg == end)
        return;

   f_helper(beg, end, *beg);
}

template <typename iterator_t, typename value_t>
void f_helper(iterator_t beg, iterator_t end, const value_t&)
{
   // for most iterators,
   // value_t ~ iterator_traits<iterator_t>::value_type

   // however if *beg returns a proxy, value_t is the type of the proxy
   // so this may not work with std::vector<bool> and in general,
   // where value_t just stores a reference to the value.
}
```

In classic C++, there are two ways to store an object without knowing its type:

- Pass it to a template function, as shown previously. However, the lifetime of the object is limited.
- Cancel its interface, possibly via a combination of templates and virtual functions. In the simplest case, the object can be *merely stored* and nothing else: 11

```
class wrapper_base
{
```

```
virtual ~wrapper base() {}
     virtual wrapper base* clone() const = 0;
};
template <typename T>
class wrapper : public wrapper_base
     T obj ;
  public:
     wrapper(const T& x)
     : obj (x) {}
     wrapper<T>* clone() const
       return new wrapper<T>(obj );
     }
};
template <typename T>
wrapper base* make clone(const T& x)
  return new wrapper<T>(x);
```

public:

Sometimes it's desirable to provide a common interface for several types. The most famous example is given by *variant objects (also known as discriminated unions)*, which are classes whose static types are fixed, but whose internal storage can transport different types.

The rest of this section discusses in detail the problem of command-line parsing. Assume you are coding a tool that gets options from the command line. Each option has a name and an associated value of some fixed type. Options come first, and everything else is an argument:

```
tool.exe -i=7 -f=3.14 -d=6.28 -b=true ARGUMENT1 ARGUMENT2 ... ARGUMENTn
```

where i is an int, f is a float, and so on.

Ideally, you need a sort of map<string, T>, where T can vary for each pair. Also, you should be able to query such a map for values having the right type, so that you can accept -f=3.14 but reject -f="hello world".

Assume, for extra simplicity, that you start with an array of strings, where each string is either [prefix] [name] or [prefix] [name] = [value], 12 and that each parameter value will be obtained via stream extraction (operator>>).

You can produce two containers. The first, named option_map, stores name-value pairs, like std::map, but each value has an arbitrary type. The second container, named option_parser, is another map that knows the desired pairing name-type (for example, "f" is a float) before parsing

```
int main(int argc, char* argv[])
  option parser PARSER;
  PARSER.declare as<float>("f"); // we tell the parser what it
should
  PARSER.declare as<int>("i"); // expect, i.e. that "d" is
a double,
  PARSER.declare as<double>("d");// etc. etc.
  option map<std::string> CL; // only key type is a template
parameter
  try
  {
     const char* prefix = "-";
     char** opt begin = argv+1;
     char** opt end = argv+argc;
     // finally we ask the parser to fill a map with the actual
values
     // this may throw an exception...
     char** arg begin = PARSER.parse(CL, opt begin, opt end,
prefix);
     double d;
     if (!CL.get(d, "d"))
       // the user did not specify a value for "d"
       d = SOME DEFAULT_VALUE;
     }
  catch (std::invalid argument& ex)
     // ...
```

5.3.1. Trampolines

the command line. The target is writing code like:

The core technique for this kind of "polymorphism" is the use of *trampolines*.

Formally, a trampoline is a local class inside a function template, but the meaning of "local"

should not be taken literally.

}

The class has only static member functions. Its public interface accepts parameters of a fixed type (say, void*), but being nested in a template, the body of the trampoline is aware of the "outer" template parameter and uses it to perform safe static casts.

Here is a bare bones example—a naked struct that holds untyped pointers and a function template that knows the static type of the object and apparently loses information.

```
struct generic t
  void* obj;
  void (*del)(void*);
};
template <typename T>
                                        // outer template parameter
generic t copy to generic (const T& value)
                                       // local class
  struct local cast
     static void destroy(void* p)
                                      // void*-based interface
       delete static cast<T*>(p);  // static type knowledge
  };
  generic t p;
  p.obj = new T(value);
                                 // information loss: copy T* to
void*
  p.del = &local cast::destroy;
  return p;
```

Actually, p.obj alone does not know how to destroy its attached object, but p.del points to (in pseudo-code) copy_to_generic<T>::local_cast::destroy and this function will do the right thing, namely cast the void* back to T* just before deleting it.

```
p.del(p.obj);  // it works!
```

del is the functional equivalent of a virtual destructor. The analogy between trampolines and virtual function tables is correct, but:

- Trampoline techniques allow you to work with *objects*, not with *pointers* (a classic factory would have returned pointer-to-base, while copy_to_generic produces an object).
- Trampoline pointers can be tested and modified at runtime. For example, del can be replaced anytime with a do-nothing function if ownership of the pointer is transferred.

• Trampolines are much less clear (that is, more difficult to maintain) than abstract class hierarchies.

The advantage of structures like <code>generic_t</code> is that their type is statically known, so they can be used in standard containers, and they are classes, so they can manage their own resources and invariants.

Unfortunately, while type T is known *internally*, it cannot be exposed. Function pointers like del cannot have T anywhere in their signature. The interface of the trampoline class must be independent of T and it cannot have template member functions (thus, for example, you cannot have a trampoline member that takes a functor and applies it to the pointee).

Next, you'll need another tool—a wrapper for std::type info.

5.3.2. Typeinfo Wrapper

The typeid operator is a less-known C++ operator that determines the type of an expression at runtime and returns a constant reference to a system object of type std::type info.

type_info::before is a member function that can be used to simulate a total (but unspecified) ordering on types.

Several wrappers have been proposed to give std::type_info value semantics. This code is similar to the elegant implementation found in [1] but the comparison operator ensures that a default-constructed (null) typeinfo is less than any other instance.¹³

```
class typeinfo
  const std::type info* p ;
public:
  typeinfo()
     : p_(0)
  { }
  typeinfo(const std::type info& t)
     : p_(&t)
  { }
  inline const char* name() const
  {
     return p_ ? p_->name() : "";
  inline bool operator<(const typeinfo& that) const
     return (p != that.p ) &&
        (!p_ || (that.p_ && static_cast<bool>(p_-
>before(*that.p ))));
  }
```

```
inline bool operator==(const typeinfo& that) const
{
   return (p_ == that.p_) ||
        (p_ && that.p_ && static_cast<bool>(*p_ == *that.p_));
};
```

5.3.3. Option_Map

Recall that option_map was introduced in Section 5.3 as a container to store values parsed from the command line, together with their type. The interface for option map is indeed very simple.

```
template <typename userkey t>
class option map
public:
// typed find:
// MAP.find<T>("name") returns true
// if "name" corresponds to an object of type T
  template <typename T>
  bool find(const userkey t& name) const;
// typeless find:
// MAP.scan("name") returns true if "name" corresponds to any
object
  bool scan(const userkey t& name) const;
// checked extraction:
// MAP.get(x, "name") returns true
// if "name" corresponds to an object of type T;
// in this case, x is assigned a copy of such object;
// otherwise, x is not changed
  template <typename T>
  bool get (T& dest, const userkey t& name) const;
// unchecked extraction:
// MAP.get<T>("name") returns either the object of type T
// corresponding to "name", or T().
  template <typename T>
  T get(const userkey t& name) const;
// insertion
// MAP.put("name", x) inserts a copy of x into the map
```

```
template <typename T>
bool put(const userkey_t& name, const T& value);
size_t size() const;

~option_map();
};
```

Now for the implementation details—the idea of generic_t is developed a bit further, giving it the ability to copy and destroy:

```
template <typename userkey_t>
class option_map
{
   struct generic_t
   {
      void* obj;
      void (*copy) (void* , const void*);
      void (*del) (void*);
   };
```

Since you'll want to search the container both by name and by pair (name, type), you should pick the latter structure as key, using the typeinfo wrapper class.

```
typedef std::pair<userkey_t, typeinfo> key_t;
typedef std::map<key_t, generic_t> map_t;
typedef typename map_t::iterator iterator_t;
map t map;
```

The insertion routine is almost identical to the prototype example:

```
template <typename T>
bool put(const userkey_t& name, const T& value)
{
   struct local_cast
   {
      static void copy(void* dest, const void* src)
      {
            *static_cast<T*>(dest) = *static_cast<const T*>(src);
      }
      static void destroy(void* p)
      {
            delete static_cast<T*>(p);
      }
    };
```

```
generic_t& p = map_[key_t(name, typeid(T))];
  p.obj = new T(value);
  p.copy = &local cast::copy;
  p.del = &local cast::destroy;
  return true;
}
  Some functions come for free on the top of std::map:
size t size() const
  return map .size();
  Here is the typed find:
template <typename T>
bool find(const userkey t& name) const
  return map_.find(key_t(name, typeid(T))) != map .end();
  To retrieve data from the option map, you use the copy function. First, you do a typed find.
If it succeeds and the object is non-null, you perform the copy over the user-supplied reference:
template <typename T>
bool get(T& dest, const userkey t& name) const
  const typename map t::const iterator i = map .find(key t(name,
typeid(T)));
  const bool test = (i != map .end());
  if (test && i->second.obj)
     i->second.copy(&dest, i->second.obj);
  return test;
}
  The unchecked retrieval is a shortcut implemented for convenience:
template <typename T>
T get(const userkey t& name) const
   initialized value<T> v;
  get(v.result, name);
  return v.result;
}
```

At this moment, you simply let the destructor wipe out all the objects. 14

```
~option_map()
{
   iterator_t i = map_.begin();
   while (i != map_.end())
   {
      generic_t& p = (i++)->second;
      if (p.del)
        p.del(p.obj);
   }
}
```

Finally, you can take advantage of the ordering properties of typeinfo for the typeless find. Due to the way pairs are ordered, the map is sorted by name and entries with the same names are sorted by typeinfo. First, you search for the upper bound of (name, typeinfo()). Any other pair with the same name will be larger, because typeinfo() is the least possible value. So, if the upper bound exists and has the same name you are looking for, it returns true.

Note that the container may hold more objects of different types having the same name.

5.3.4. Option_Parser

option_parser is not described in full, since it does not add anything to the concepts used in building option_map. However, note that a trampoline may have parameters whose type is not void*. We leave some details for exercise.

```
class option_parser
{
  typedef option_map<std::string> option_map_t;
  typedef bool (*store_t)(option_map_t&, const char*, const char*);
  typedef std::map<std::string, store_t> map_t;
  map_t map_;

public:
  template <typename T>
    void declare_as(const char* const name)
```

```
{
   struct local store
     static bool store (option map t& m,
                      const char* name, const char* value)
     {
       std::istringstream is(value);
       T temp;
       return (is >> temp) && m.put(name, temp);
   };
   map [name] = &local store::store;
  Note that local store::store does not take void* arguments. The only requirement for a
trampoline is to publish an interface independent of T.
template <typename iterator t>
iterator t parse (option map t& m, iterator t begin, iterator t
end)
  for every iterator i=begin...end
    get the string S = *i;
    if S has no prefix
        stop and return i;
    else
        remove the prefix
    if S has the form "N=V"
        split S in N and V
    else
        set N = S
        set V = <empty string>
    if N is not contained in map
        throw exception "unknown option"
    else
        set F := local store::store
```

if it fails, throw exception "illegal value"

5.3.5. Final Additions

execute F(m, N, V)

Due to the way declare_as works, every type that can be extracted from a string stream is acceptable in the command-line parser.

To include parameterless options, simply add an empty class:

```
struct option
{
};

inline std::istream& operator>>(std::istream& is, option&)
{
   return is;
}
```

This will enable a command-line switch, such as:

```
tool.exe -verbose
```

If the name is unique, the simplest way to retrieve the value of the switch is using a typeless find. This will yield false if the switch is omitted.

```
PARSER.declare_as<option>("verbose");
char** arg_begin = PARSER.parse(CL, opt_begin, opt_end, prefix);
if (CL.scan("verbose"))
{
    // ...
}
```

Trampoline techniques can be easily optimized for space. Instead of creating one pointer for each "virtual function," you can group functions for type \mathbb{T} in a static instance of a structure and therefore have a single pointer, exactly as in the traditional implementation of virtual function tables.

This approach is also scalable. Should you need to add an extra "capability" to the interface, it requires fewer modifications and almost no extra memory (since you have a single pointer table, as opposed to many pointers per instance).

```
struct virtual_function_table
{
    void (*copy)(void* , void*);
    void (*del)(void*);
    void* (*clone)(const void*);
};

    struct generic_t
    {
        void* obj;
        const virtual_function_table* table; // single pointer-to-const
     };
```

```
// identical implementation, but not a local class any more...
  template <typename T>
  struct local cast
     static void copy(void* dest, void* src)
        *static cast<T*>(dest) = *static cast<T*>(src);
     }
     static void destroy(void* p)
        delete static cast<T*>(p);
     }
     static void* clone(const void* p)
        return new T(*static cast<const T*>(p));
  };
  template <typename T>
     bool put(const userkey t& name, const T& value)
   {
     static const virtual function table pt =
        &local cast<T>::copy,
        &local cast<T>::destroy,
        &local cast<T>::clone
     };
     generic_t& p = map [key t(name, typeid(T))];
     p.obj = new T(value);
     p.table = &pt;
     return true;
  }
  Of course, instead of p.del, you should write p.table->del and pay an extra indirection.
  Finally, you make generic t a true value by the rule of three: implementing copy constructor,
assignment, and destructor.
struct generic t
  void* obj;
  const virtual function table* table;
```

generic t()

```
: obj(0), table(0)
  }
  generic t(const generic t& that)
     : table(that.table)
     if (table)
       obj = table.clone(that.obj);
  }
  generic t& operator=(const generic t& that)
     generic t temp(that);
     swap(obj, temp.obj);
     swap(table, temp.table);
     return *this;
  }
  ~generic t()
     if (table && obj)
        (table->del)(obj);
};
```

5.3.6. Boundary Crossing with Trampolines

This section briefly summarizes the last paragraphs. A trampoline function is used as a companion to a void pointer when it contains enough information to recover the original type:

```
void* myptr_;
void (*del_)(void*);

template <typename T>
struct secret_class
{
    static void destroy(void* p)
    {
        delete static_cast<T*>(p);
    }
};

myptr_ = [[a pointer to T]];
del_ = &secret_class<T>::destroy;
```

The information about T cannot be returned to the caller, because T cannot be present in the

trampoline interface.

So you will generally tackle the issue requiring the caller to specify a type \mathbb{T} , and the trampoline just ensures it's the same as the original type (calling typeid, for example, see the "typed find"). This is informally called an *exact cast*.

In short, an exact cast will fail if the type is not precisely what the program expects:

```
template <typename T>
T* exact cast() const
  return &secret class<T>::destroy == del ?
     static cast<T*>(myptr ) : 0;
}
  A second possibility is to throw an exception:
  template <typename T>
  struct secret class
     static void throw T star(void* p)
     {
       throw static cast<T*>(p);
  };
struct myobj
  void* myptr ;
  void (*throw )(void*);
  template <typename T>
  myoby(T*p)
  {
     myptr_ = p;
     throw = &secret class<T>::throw T star;
  }
  template <typename T>
  T* cast via exception() const
  {
     try
     {
        (*throw ) (myptr );
     catch (T^* p) // yes, it was indeed a T^*
     {
        return p;
     catch (...) // no, it was something else
```

```
{
    return 0;
}
};
```

This approach is several orders of magnitude slower (a try...catch block may not be cheap), but it adds an interesting new feature. You can cast not only to the original type T, but also to any *base class* of T. When the trampoline function throws DERIVED*, the exception handler will succeed in catching BASE*.

Remember that it's not possible to dynamic_cast a void* directly, so this is actually the best you can do. If efficiency is an issue, in practice you might want to adopt a scheme where you perform an exact cast to BASE* using trampolines and execute a dynamic cast on the result later (after the trampoline code).

Observe also that, depending on the precise application semantics, you can sometimes limit the number of "destination" types to a small set and hardcode them in the trampoline:

```
struct virtual function table
  bool (*safe to double) (void*, double&);
  std::string (*to string)(void*);
} ;
template <typename T1, typename T2>
struct multi cast
  static T2 cast(void* src)
     return has conversion<T1,T2>::L2R ?
       T2(*static cast<T1*>(src)) : T2();
  }
  static bool safe cast(void* src, T2& dest)
     if (has conversion<T1,T2>::L2R)
       dest = *static cast<T1*>(src);
     return has conversion<T1,T2>::L2R;
  }
};
  to double = &multi cast<T, double>::safe cast;
  to string = &multi cast<T, std::string>::cast;
```

5.4. Variant

The key point in type-hiding techniques is deciding who remembers the correct type of the objects. In this example, the client of option_map is responsible for declaring and querying the right types, by calling option map::get<T>("name").

In some cases, the client needs or prefers to ignore the type and blindly delegate the "opaque" object. This way, it performs the right action, whatever the stored object is.

5.4.1. Parameter Deletion with Virtual Calls

If you simply need to transport a copy of an object of arbitrary type, you can wrap it in a custom class template, thereby "hiding" the template parameter behind a non-template abstract base class.

The following rough code snippet will help clarify this idea:

```
struct wrapper base
  virtual ~wrapper base()
  }
  virtual wrapper base* clone() const = 0;
  // add more virtual functions if needed
  virtual size t size() const = 0;
};
template <typename T>
struct wrapper : wrapper_base
{
  T obj ;
  wrapper(const T& that)
     : obj (that)
  {
  }
  virtual wrapper base* clone() const
  {
     return new wrapper<T>(obj );
  }
     implement virtual functions delegating to obj
  virtual size t size() const
     return obj .size();
};
```

```
class transporter
  wrapper base* myptr ;
public:
  ~transporter()
     delete myptr ;
  transporter(const transporter& that)
    myptr (that.myptr ? that.myptr ->clone() : 0)
  transporter()
  : myptr (0)
  template <typename T>
  transporter (const T& that)
     : myptr (new wrapper<T>(that))
  }
     implement member functions delegating to wrapper base
  size t size() const
     return myptr ? myptr ->size() : 0;
};
  You can also add a custom (friend) dynamic cast:
template <typename T>
static T* transporter cast(transporter& t)
  if (wrapper<T>* p = dynamic cast<wrapper<T>*>(t.myptr ))
     return & (p->obj );
  else
     return 0;
```

5.4.2. Variant with Visitors

}

Opaque interfaces often make use of the visitor pattern. The *visitor* is a functor of unspecified type that is accepted by the interface and is allowed to communicate with the real objects, whose type is otherwise hidden.

In other words, you need a way to pass a generic functor through the non-template trampoline interface.

As a prototype problem, you will code a concept class that can store any object of size not greater than a fixed limit.¹⁵

```
template <size_t N>
class variant;
```

First, you define the required trampolines. variant will have some fixed-size storage where you place the objects:

```
template <size_t N>
class variant
{
   char storage_[N];
   const vtable* vt;
};
```

Again from the rule of three, the tentative interface has three functions:

```
struct vtable
  void (*construct)(void*, const void*);
  void (*destroy) (void*);
  void (*assign) (void*, const void*);
};
template <typename T>
struct vtable impl
  static void construct (void* dest, const void* src)
     new(dest) T(*static cast<const T*>(src));
  }
  static void destroy(void* dest)
  {
     static cast<T^*>(dest)->\sim T();
  }
  static void assign(void* dest, const void* src)
     *static cast<T*>(dest) = *static cast<const T*>(src);
};
```

```
template <>
struct vtable impl<void>
  static void construct(void* dest, const void* src)
  static void destroy(void* dest)
  static void assign(void* dest, const void* src)
};
template <typename T>
struct vtable singleton
  static const vtable* get()
     static const vtable v =
        &vtable impl<T>::construct,
       &vtable impl<T>::destroy,
       &vtable impl<T>::assign
     };
     return &v;
};
template <size_t N>
class variant
  char storage [N];
  const vtable* vt;
public:
  ~variant()
     (vt->destroy) (storage );
  variant()
     : vt(vtable singleton<void>::get())
```

```
variant(const variant& that)
   : vt(that.vt)
{
    (vt->construct)(storage_, that.storage_);
}

template <typename T>
   variant(const T& that)
    : vt(vtable_singleton<T>::get())
{
    MXT_ASSERT(sizeof(T)<=N);
    (vt->construct)(storage_, &that);
};
```

The constructors initialize the "virtual function table pointer" and invoke the construction over raw memory. 16

The assignment operator depends on a subtle issue: exceptions. If a constructor throws an exception, since the object was never fully constructed it won't be destroyed either, and that's exactly what you need. However, if you need to overwrite an instance of T1 with an instance of T2, you destroy T1 first, but construction of T2 may fail.

Thus, you need to reset the virtual table pointer to a no-op version, destroy T1, construct T2, and then eventually store the right pointer.

```
void rebuild(const void* src, const vtable* newvt)
{
  const vtable* oldvt = vt;
  vt = vtable_singleton<void>::get();
  (oldvt->destroy) (storage_);

  // if construct throws,
  // then variant will be in a consistent (null) state
  (newvt->construct) (storage_, src);
  vt = newvt;
}

Thanks to rebuild, you can copy another variant and any other object of type T:
  variant& operator=(const variant& that)
  {
    if (vt == that.vt)
        (vt->assign) (storage_, that.storage_);
    else
        rebuild(that.storage , that.vt);
```

```
return *this;
 template <typename T>
 variant& operator=(const T& that)
  {
   MXT ASSERT(sizeof(T) <= N);</pre>
   if (vt == vtable singleton<T>::get())
     (vt->assign) (storage , &that);
   else
     rebuild(&that, vtable singleton<T>::get());
   return *this;
};
  This variant is only pure storage, but consider this addition:
class variant
 // ...
 template <typename visitor t>
 void accept visitor(visitor t& v)
    // ???
};
```

Since trampolines need to have a fixed non-template signature, here the solution is *virtual inheritance*. You define an interface for any unspecified visitor and another interface for a visitor who visits type \mathbb{T} . Since the trampoline knows \mathbb{T} , it will try one dynamic cast.

Virtual inheritance is necessary because visitors may want to visit more than one type.

```
class variant_visitor_base
{
  public:
    virtual ~variant_visitor_base()
    {
      }
};

template <typename T>
  class variant_visitor : public virtual variant_visitor_base
{
  public:
    virtual void visit(T&) = 0;
```

```
virtual ~variant visitor()
};
struct bad visitor
struct vtable
 // ...
 void (*visit) (void*, variant visitor base*);
} ;
template <typename T>
struct vtable impl
  // ...
 static void visit(void* dest, variant visitor base* vb)
   if (variant visitor<T>* v = dynamic cast<variant visitor<T>*>
(vb))
     v->visit(*static cast<T*>(dest));
   else
     throw bad visitor();
};
template <>
struct vtable impl<void>
  // ...
 static void visit(void* dest, variant visitor base* vb)
};
template <size t N>
class variant
public:
 variant& accept visitor(variant visitor base& v)
    (vt->visit) (storage , &v);
```

```
return *this;
}
```

Finally, here's a concrete visitor (which will visit three types, hence the importance of the virtual base class):

```
struct MyVisitor
: public variant visitor<int>
, public variant visitor<double>
, public variant visitor<std::string>
  virtual void visit(std::string& s)
     std::cout << "visit: {s}" << s << std::endl;
  }
  virtual void visit(int& i)
  {
     std::cout << "visit: {i}" << i << std::endl;
  }
  virtual void visit(double& x)
  {
     std::cout << "visit: {d}" << x << std::endl;
};
     variant<64> v1, v2, v3;
     std::string s = "hello world!";
     double x = 3.14;
     int j = 628;
     v1 = s;
     v2 = x;
     v3 = j;
     MyVisitor mv;
     v1.accept visitor(mv);
     v2.accept visitor(mv);
     v3.accept visitor(mv);
visit: {s}hello world!
visit: {d}3.14
visit: {i}628
```

Note for the sake of completeness that *bounded* discriminated unions, such as boost::variant, adopt a different approach. variant is a class template with N type parameters T1...Tn. At any time, variant holds exactly one instance of Tj. The constructor can

take any object of type T having an unambiguous conversion to exactly one Tj, or it fails.

5.5. Wrapping Containers

New containers are often built on top of classical STL objects:

```
template <typename T, typename less_t = std::less<T> >
class sorted_vector
{
   typedef std::vector<T> vector_t;
   vector_t data_;
```

The sorted vector basically is equivalent to a set of functions that manipulates a vector, enforcing some invariant (namely, preserving the ordering). So, a sorted_vector is a kind of a *container* adapter, because it delegates the actual storage and it only alters the way data is stored.

Suppose you already have a vector and want to treat it as a sorted vector. Remember that it's a bad idea to expose the internals of the class (recall Section 1.4.4).

You can instead have an additional parameter that defines the storage, analogous to the container type of std::stack and similar to the allocator parameter for std::vector.

```
template <typename T, typename less_t = ..., typename vector_t
= std::vector<T> >
class sorted_vector
{
   vector_t data_;
public:
   sorted_vector(vector_t data)
```

```
: data_(data)
{
  }
};

void treatVectorAsSorted(vector<double>& v)
{
  sorted_vector<double, less<double>, vector<double>&>
sorted_v(v);
  // ...
}
```

When you write the code of sorted_vector, you should behave as if vector_t is std::vector, whose interface is defined unambiguously. Any replacement type will have to satisfy the same contract.

Anyway, this solution is the most complex to code, since you need reference-aware functions. You should explicitly support the case of vector_t being a reference to some vector, and this is likely to cause problems when deciding to take arguments by value/by reference. This would be a good case for traits.

```
template <typename T>
struct s v storage traits
  typedef const T& argument type;
  typedef T value type;
};
template <typename T>
struct s v storage traits<T&>
  typedef T& argument type;
  typedef T& value type;
};
template <typename T, typename less t = ..., typename vector t
= vector<T> >
class sorted vector
  typename s v storage traits<vector t>::value type data ;
public:
  sorted vector (typename
              s v storage traits<vector t>::argument type data)
    data (data)
};
```

A strong need to isolate the storage parameter comes from serialization. Modern operating systems can easily map memory from arbitrary storage devices into the program address space at no cost.

In other words, you can get a pointer (or a pointer-to-const) from the OS that *looks* like ordinary memory, but that's actually pointing somewhere else, for example, to the hard drive.

Now you can create a sorted_vector that points directly to the mapped memory, plugging a suitable class as $vector_t$:

```
template <typename T>
class read_only_memory_block
{
   const T* data_;
   size_t size_;

public:
   // constructor, etc....

   // now implement the same interface as a const vector
   typedef const T* const_iterator;

   const_iterator begin() const { return data_; }
   const_iterator end() const { return data_ + size_; }

   const T& operator[](size_t i) const { return data_[i]; }
   // ...
};
```

Observe that you don't need a true drop-in replacement for vector. If you just call const member functions, a subset of the interface will suffice.

¹The notion of *concept* was introduced in Section 2.2.4.

²The black hole iterator is a hack, not a perfect output iterator.

³This is why, for example, the standard describes carefully what happens to functors passed to STL algorithms, such as how many times they are copied, and so on.

⁴The boost library contains some more general code. See http://www.boost.org/doc/libs/1_57_0/libs/utility/operators.htm.

⁵See [2] page 135.

⁶The downside of this technique is that the macro may confuse some IDEs that parse headers to build a graphical representation of the project.

⁷Apparently, the term "memberspace" was introduced by Joaquín M López Muñoz in "An STL-Like Bidirectional Map" (see www.codeproject.com/vcpp/stl/bimap.asp). Also, the double-end queue example is from the same author.

⁸In the pseudo-code that follows, you should pretend that C is a class; of course a non-template container would be an unusual beast.

⁹This code does not compile, because for conciseness, we removed all const versions of the member functions. However, the fix

should be obvious.

- Modern C++ offers two new keywords: decltype and auto. The former returns the exact type of any expression, similarly to sizeof. The latter allows an instance to "copy" the type of its initializer, so auto i = f() would declare a variable i having the best possible type to store the result of f() locally. See Chapter 12 for more details.
- ¹¹This example is important and it will be analyzed again in Section 5.4.1.
- ¹²The prefix is a fixed character sequence, usually "-", "--", or "/".
- ¹³The implementation uses short-circuit to prevent null pointer dereferencing, and it's extremely concise. See also an exercise in Appendix B.
- ¹⁴The implementation is obviously faulty; option map cannot be safely copied/assigned. To keep the code as simple as possible, and even simpler, the discussion of this topic is deferred to Section 5.35.
- 15 This is also known as an *unbounded discriminated union*. The code should be taken as a proof-of-concept, not as production ready. Two big issues are not considered: const-ness and aligned storage. I suggest as a quick-and-dirty fix that you put variant::storage_in a union with a dummy structure having a single member double. See. A. Alexandrescu's "An Implementation of Discriminated Unions in C++".
- 16 We got rid of the " if pointer is null " tests initializing members with a dummy trampoline.
- 17The mapped memory area is usually not resizable. Merely for simplicity, we assume it's const, but that need not be the case. A vector needs to store three independent pieces of data (for example, "begin," "size," and "capacity"; everything else can be deduced). A read_write_memory_block would also need these members, but the capacity would be a constant and equal to "max size" from the beginning.

CHAPTER 6

Algorithms

The implementation of an algorithm needs a generic I/O interface. You need to decide how and where functions get data and write results, and how and what intermediate results are retained. Iterators are an existing abstraction that helps solve this problem.

An *iterator* is a small data type that offers a sequential view over a dataset. Put simply, it's a class that implements a subset of the operations that pointers can perform.

The importance of iterators is that they decouple functions from the actual data storage. An algorithm reads its input via a couple of iterators [begin...end) of unspecified type and often writes its output to another range:

```
template <typename iterator_t>
... sort(iterator_t begin, iterator_t end);
template <typename iter1_t, typename iter2_t >
... copy(iter1_t input_begin, iter1_t input_end, iter2_t output_begin);
```

It's possible to write a rough classification of algorithms based on their I/O interface. *Non-mutating algorithms* iterate on one or more read-only ranges. There are two sub-families:

- "find" algorithms return an iterator that points to the result (such as std::min element) or to end if no result exists.
- "accumulate" algorithms return an arbitrary value, which need not correspond to any element in the range.

Selective copying algorithms take an input read-only range and an output range, where the results are written. The output range is assumed to be writable. If the output range can store an arbitrary number of elements or simply as many elements as the input range, only the left-most position is given (such as std::copy).

• Usually each algorithm describes what happens if input and output ranges overlap. "transform" algorithms accept an output range that's either disjoint or entirely coincident with begin...end.

Reordering algorithms shuffle the elements of the input range and ensure that the result will be in

some special position, or equivalently, that the result will be a particular sub-range (for example, std::nth element and std::partition).

• "shrinking" algorithms (std::remove_if) rearrange the data in begin...end and, if the result is shorter than the input sequence, they return a new end1, leaving unspecified elements in range end1..end.

Writing algorithms in terms of iterators can offer significant advantages:

- All containers, standard and nonstandard, can easily provide iterators, so it's a way to decouple algorithms and data storage.
- Iterators have a default and a convenient way to signal "failure," namely by returning end.
- It may be feasible, depending on the algorithm's details, to ignore the actual "pointed type" under the iterators.

On the other hand, there are two difficulties:

- Iterators are a view over a dataset. You'll often need to adapt a given view to match what another algorithm expects. For example, you write a function that gets as input a sequence of pair<X, Y> but internally you may need to invoke a routine that needs a sequence of X. In some cases, changing the view requires a lot of effort.
- The exact type of the iterator should be avoided whenever possible. Assume that v is a const-reference to a container and compare the following two ways to iterate over all elements (let's informally say two "loops").

```
for (vector<string>::const_iterator i = v.begin();
i != v.end(); ++i)
{ ... }
std::for each(v.begin(), v.end(), ...);
```

The first "loop" is less generic and more verbose. It is strongly tied to the exact container that you are using (namely, vector<string>). If you replace the data structure, this code won't compile any more. Additionally, it recomputes v.end() once per iteration.¹

The second loop has its disadvantages as well. You have to pass a function object as the "loop body," which may be inconvenient.

6.1. Algorithm I/O

Algorithms are usually functions that perform their input/output operations via generic ranges. In this case, a range is represented by a pair of iterators of generic type iterator t, and the function

assumes that iterator_t supports all the required operations. You'll see, however, that this assumption is not only a convenient simplification, it's often the best you have, as it's extremely hard to *detect* if a generic type T is an iterator.

The hypotheses are:

- *i returns std::iterator_traits<T>::reference, which behaves like a reference to the underlying object.²
- Whatever *i returns, a copy of the pointed value can be stored as an std::iterator_traits<T>::value_type; often, you'll impose further that this type is assignable or swappable.³
- Any elementary manipulation of i (copy, dereference, increment, and so on) is inexpensive.
- You can dispatch specialized algorithms for iterators of different types using std::iterator traits<T>::iterator category as a type tag.
- All increment/decrement operators that are valid on i return a dereferenceable object (usually, another instance of T). This allows you to write safely *(i++).

Sometimes you'll implicitly assume that two copies of the same iterator are independent. This is usually violated by I/O-related iterators, such as objects that read/write files or memory, like std::back_insert_iterator, because *i conceptually allocates space for a new object; it does not retrieve an existing element of the range.

6.1.1. Swap-Based or Copy-Based

As a consequence of the basic assumptions, most (if not all) I/O in the algorithm should be written without explicitly declaring types. Using reference and value_type should be minimized, if possible, usually via swaps and direct dereference-and-assign.

For example, copy tackles the problem of output. It simply asks a valid iterator where the result is written:

Not knowing what the elements are, you can assume that a swap operation is less heavy than an ordinary assignment. A POD swap performs three assignments, so it is slightly worse, but if objects contain resource handles (such as pointers to heap allocated memory), swaps are usually optimized to avoid construction of temporary objects (which may fail or throw). If s1 is a short string and s2 is a

very long string, the assignment s1=s2 will require a large amount of memory, while swap (s1, s2) will cost nothing.

For example, an implementation of std::remove_if could overwrite out-of-place elements with a smart swap.

move is a destructive copy process, where the original value is left in a state that's consistent but unknown to the caller.

```
template <typename iterator t>
void move iter(iterator t dest, iterator t source)
  if (dest == source)
     return;
  if
(is class<std::iterator traits<iterator t>::value type>::value)
     smart swap(*dest, *source);
  else
     *dest = *source;
}
template <typename iterator t, typename func_t>
iterator t remove if (iterator t begin, iterator t end, func t F)
  iterator t i = begin;
  while (true)
     while (i != end && F(*i))
       ++i;
     if (i == end)
       break;
     move iter(begin++, i++);
  }
  return begin;
```

This algorithm returns the new end of range. It will use an assignment to "move" a primitive type and a swap to "move" a class. Since the decision rule is hidden,⁴ it follows that the algorithm will leave unpredictable objects between the new and old ends of range:

```
struct less_than_3_digits
{
  bool operator()(const std::string& x) const
  {
    return x.size()<3;</pre>
```

}

After executing this code, the arrays A1 and A2 will be different. The trailing range is filled with unspecified objects, and they do vary.

```
[0]
             111
                                   "111"
[1]
             4444
                                   "4444"
             555555
                                   "555555"
[2]
[3]
             4444
                                   "2"
                                   "3"
             555555
[4]
             66
                                   "66"
```

Note C++0x has a language construct for move semantics: R-value references.

A function argument declared as an R-value reference-to-T (written T&&) will bind to a non-constant temporary object. Being temporary, the function can freely steal resources from it. In particular, you can write a special "move constructor" that initializes a new instance from a temporary object.

Furthermore, casting a reference to an R-value reference has the effect of marking an existing object as "moveable" (this cast is encapsulated in the STL function std::move).

Combining these features, the three-copy swap can be rewritten as:

```
void swap(T& a, T& b)
{
    T x(std::move(a));
    a = std::move(b);
    b = std::move(x);
}
```

So if T implements a move constructor, this function has the same complexity as a native swap.

Other implementations of move iter could:

- Test if (!has_trivial_destructor<...>::value). It's worth swapping a class that owns resources, and such a class should have a non-trivial destructor. Observe, however, that if the type is not swappable, this approach may be slower, because it will end up calling the three-copy swap, instead of one assignment.
- Test the presence of a swap member function and use assignment in any other case.

```
template <typename iterator_t>
void move_iter(iterator_t dest, iterator_t source, selector<true>)
{
    dest->swap(*source);
}

template <typename iterator_t>
void move_iter(iterator_t dest, iterator_t source,
selector<false>)
{
    *dest = *source;
}

template <typename iterator_t>
void move_iter(iterator_t dest, iterator_t source)
{
    typedef typename std::iterator_traits<iterator_t>::value_type
val_t;
    if (dest != source)
        move_iter(dest, source, has_swap<val_t>());
}
```

6.1.2. Classification of Algorithms

Recall the distinction between non-mutating, selective copy and reordering algorithms. This section shows how sometimes, even when the mathematical details of the algorithm are clear, several implementations are possible, and it discusses the side effects of each.

Let's say you want to find the minimum and the maximum value of a range simultaneously. If the range has N elements, a naïve algorithm uses ~2N comparisons, but it's possible to do better. While iterating, you can examine two consecutive elements at a time and then compare the larger with the max and the smaller with the min, thus using three comparisons per two elements, or about 1.5*N comparisons total.

First, consider a non-mutating function (the macro is only for conciseness)⁵:

```
#define VALUE T typename
```

```
std::iterator_traits<iterator_t>::value_type

template <typename iterator_t, typename less_t>
std::pair<VALUE_T, VALUE_T> minmax(iterator_t b, iterator_t e, less_t less)
```

minmax (begin, end) scans the range once from begin to end, without changing any element, and it returns a pair (min, max). If the range is empty, you can either return a default-constructed pair or break the assumption that result.first < result.second, using std::numeric limits.

Here's a reasonable implementation, which needs only forward iterators:

```
template <typename scalar_t, typename less_t>
inline scalar t& mmax(scalar t& a, const scalar t& b, less t less)
{
  return (less(a, b) ? a=b : a);
}
template <typename scalar t, typename less t>
inline scalar t& mmin(scalar t& a, const scalar t& b, less t less)
  return (less(b, a) ? a=b : a);
}
template <typename iterator t, typename less t>
std::pair<...> minmax(iterator t begin, const iterator t end,
                 less t less)
{
  typedef
    typename std::iterator traits<iterator t>::value type
value type;
  std::pair<value type, value type> p;
  if (begin != end)
    p.first = p.second = *(begin++);
  while (begin != end)
    const value type& x0 = *(begin++);
    const value type x1 = (begin != end) ? *(begin++) : x0;
    if (less(x0, x1))
      mmax(p.second, x1, less);
      mmin(p.first, x0, less);
```

```
}
else
{
    mmax(p.second, x0, less);
    mmin(p.first , x1, less);
}
return p;
}
```

As a rule, it's more valuable to return iterators, for two reasons. First, the objects may be expensive to copy, and second, if no answer exists, you return end.

So, given that dereferencing an iterator is inexpensive, a possible refinement can be:

```
template <typen.ator t, typename less t>
std::pair<iterator t, iterator t> minmax(...)
     std::pair<iterator t, iterator t> p(end, end);
     if (begin != end)
     {
       p.first = p.second = begin++;
     }
     while (begin != end)
       iterator t i0 = (begin++);
       iterator t i1 = (begin != end) ? (begin++) : i0;
        if (less(*i1, *i0))
          swap(i0, i1);
        // here *i0 is less than *i1
       if (less(*i0, *p.first))
          p.first = i0;
       if (less(*p.second, *i1))
          p.second = i1;
     }
     return p;
}
```

Note that you never mention value_type any more. Finally, you can outline the reordering variant:

```
template <typename iterator_t>
void minmax(iterator_t begin, iterator_t end);
```

The function reorders the range so that, after execution, *begin is the minimum and * (end-1) is the maximum. All the other elements will be moved to an unspecified position. Iterators are bidirectional, so end-1 is just a formal notation.

Suppose F takes a range [begin...end). It compares the first and the last element, swaps them if they are not in order, and then it proceeds to the second and the second to last. When the iterators cross, it stops and it returns an iterator H, which points to the middle of the range. F executes about N/2 "compare and swap" operations, where N is the length of the range.

Obviously, the maximum cannot belong to the left half and the minimum cannot belong to the right half. You must invoke again F on both half-intervals and let HL=F (begin, HL) and HR=F (HR, end).

When there's a single element in one of the intervals, it has to be the extreme.

If a unit of complexity is a single "compare and swap," the algorithm performs N/2 at iteration 0 to find H, $2 \cdot (N/4)$ for the second partition, $2 \cdot (N/8)$ for the third, and so on, so the total number of operations is again about $3/2 \cdot N$.

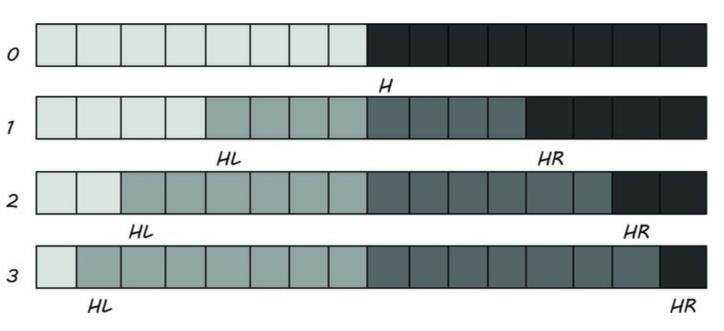


Figure 6-1. A graphical representation of the reordering minmax algorithm

6.1.3. Iterator Requirements

Algorithms have requirements about the kind of operation the iterator must provide. As a rule of thumb, the "average" iterator is bidirectional. It supports single increment and decrement (++ and - -), equality/inequality (== and !=). However, it does not offer additions of arbitrary integers, difference, and operator<. Random-access iterators are used wherever maximum speed is needed, and for sorting, so they usually deserve a special treatment with specialized algorithms.

As previously mentioned, you can ensure that a requirement on the iterator is met by dispatching to another function that accepts an additional formal argument of type "iterator category":

```
template <typename iter_t>
void do_something(iter_t begin, iter_t end)
{
  return do_something(begin, end,
```

```
typename std::iterator_traits<iter_t>::iterator_category());

template <typename iter_t>
void do_something(iter_t begin, iter_t end,
std::bidirectional_iterator_tag)
{
    // do the work here
}
```

This technique was invented for invoking optimized versions of the algorithm for any iterator type, but it can be used to restrict the invocation as well. Standard iterator tags form a class hierarchy, so a "strong" tag can be cast nicely to a "weaker" requirement.

Here are some guidelines:

- Sometimes you'll write an algorithm first and *then* deduce which iterator is required for the algorithm to work. While deduction a posteriori is perfectly acceptable, it is easy to underestimate the requirements imposed by subroutines.
- It's usually good design to separate algorithms that have different requirements. For example, instead of sorting *and then* iterating, just prescribe that the range should be already sorted. This may bring down the requirements to bidirectional iterators.

```
template <typename iterator_t>
void do_something(iterator_t begin, iterator_t end)
{
    // the following line has stronger requirements than all the rest

    std::sort(begin, end, std::greater<...>());
    std::for_each(begin, end, ...);
}

template <typename iterator_t>
void do_something_on_sorted_range(iterator_t begin, iterator_t end)
{
    // much better: all lines have the same complexity
    std::reverse(begin, end);
    std::for_each(begin, end, ...);
}
```

6.1.4. An Example: Set Partitioning

Suppose you are given a set of integers X and you need to partition it into two subsets so that the sum

in each has roughly the same value.⁷

For this problem, heuristic algorithms are known that quickly find an acceptable partition (possibly suboptimal). The simplest is the greedy algorithm, which states that:

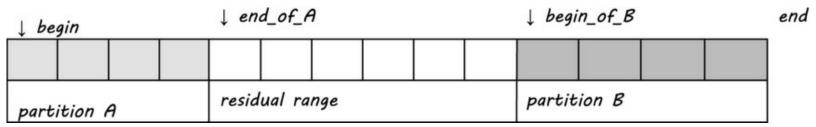
```
Let P1={} and P2={} be empty sets;
While X is not empty, repeat:
{
    Assign the largest remaining integer in X to the set Pi which currently has the lower sum
    (break ties arbitrarily);
}
```

This prescription sounds like reordering, so you can consider a mutating algorithm. You reorder the input range and return an iterator h, so that [begin, h) and [h, end) are the required partitions. Also, as an additional bonus, you can compute the difference of the sums of both partitions

$$\left| \sum_{i \in [begin,h)} *i - \sum_{i \in [h,end)} *i \right|, \text{ which is the objective to be minimized. Thus, the result will be std::pair.}$$

The implementation behaves like this:

- The range is divided in three logical blocks: partition A [begin, end_of_A) on the left, partition B on the right [begin_of_B, end), and a residual middle block M.
- A and B are initially empty and M = [begin, end).
- While M is non-empty, repeat:
 - Elements of M are sorted in decreasing order.
 - Iterate over the elements of M. Objects assigned to A are swapped to the right of A (in position "end of A") and objects assigned to B are swapped to the left of B (in position "begin of B minus 1").8



This code is a concise example of a mutating algorithm:

- It does not allocate temporary memory
- Its runtime complexity is documented

```
#define mxt_value_type(T) typename
std::iterator traits<T>::value type
```

```
template <typename iterator t>
std::pair<iterator t, mxt value type(iterator t)>
          equal partition (iterator t begin, iterator t end)
  typedef mxt value type(iterator t) > scalar t;
  scalar t sum a = 0;
  scalar t sum b = 0;
  iterator t end of A = begin;
  iterator t beg of B = end;
  while (end of A != beg of B)
  {
     std::sort(end of A, beg of B, std::greater<scalar t>());
     iterator t i = end of A;
     do
     {
       if (sum b < sum a)
          sum a = sum a - sum b;
          sum b = *i;
          smart swap(*i, *(--beg of B));
        }
       else
        {
          sum b = sum b - sum a;
          sum a = *i;
          smart swap(*i, *(end of A++));
        }
     while ((i != beg of B) && (++i != beg of B));
  return std::make pair(end of A,
                      sum a < sum b ? sum b - sum a : sum a - sum b);</pre>
```

Let's examine the implementation to determine the requirements on iterators and types.

- At a first glance, it may look like a bidirectional iterator t suffices, because the code only uses copy construction, inequality, ++, and --. However, std::sort requires random access iterators.
- The underlying scalar t needs to implement operator< and binary operator -. Note that there's a small difference between these lines:

```
sum b = sum b - sum a;
```

{

}

```
sum b -= sum a;
```

The second option would introduce a new requirement (namely operator ==).

6.1.5. Identifying Iterators

The metafunction std::iterator_traits<T> returns several types if T is an iterator or a pointer (in this case, types are trivially deduced). It also is the most reliable way to ensure that T is an iterator, because for most other types it will not compile:

```
template
<
   typename T,
   typename IS_ITERATOR = std::iterator_traits<T>::value_type
>
class require_iterator
{
   // similar to a static assertion,
   // this will compile only if T is a compliant iterator
};
```

You can take an educated guess as to whether a type is a conforming iterator by using the *five* basic typedefs that iterators are required to supply. 10

Using the SFINAE techniques¹¹ again, you would write:

```
template <typename T>
struct is iterator
  static const bool value =
     static AND
     <
       has type value type<T>,
       static_AND
          has type reference<T>,
          static AND
          <
            has type pointer<T>,
             static AND
               has type iterator category<T>,
               has type difference type<T>
             >
          >
     >::value;
```

```
template <typename T>
struct is_iterator<T*>
{
   static const bool value = true;
};
```

};

The rational for the heuristic is as follows:

- std::map is not an iterator, but it defines all types except iterator_category. Therefore, you really need to test that all five types are present.
- You cannot test if std::iterator_traits is well defined, because it will not compile if T is invalid.
- There exist types where is_iterator is true, but they are not even dereferenceable (trivially, let T be std::iterator_traits<int*>).

Here's a test that, with good precision, will identify non const-iterators.¹² The key motivation is the following:

- An iterator will define a value type T and a reference type, usually T& or const T&.
- T& is convertible to T, but not vice versa
- const T& and T are mutually convertible. 13

There are several possible cases:

- If \mathbb{T} is not an iterator, it's not even a mutable iterator (that's handled by the last partial specialization).
- If reference is value_type& then the answer is true (this case is handled by the helper class).
- If reference is convertible to value_type, but not vice versa, the answer is again true.

```
template <typename T1, typename T2>
struct is_mutable_iterator_helper
{
   static const bool value = false;
};
template <typename T>
struct is_mutable_iterator_helper<T&, T>
{
```

```
static const bool value = true;
};
template <typename T, bool IS ITERATOR = is iterator<T>::value>
class is mutable iterator
  typedef typename std::iterator traits<T>::value type val t;
  typedef typename std::iterator traits<T>::reference ref t;
public:
  static const bool value =
     static OR
     <
       is mutable iterator helper<ref t, val t>,
       selector
          has conversion<ref t, val t>::L2R &&
          !has conversion<val t, ref t>::L2R
     >::value;
};
template <typename T>
class is mutable iterator<T, false>
public:
  static const bool value = false;
};
```

- Has_conversion<ref_t, val_t>::L2R should be true by definition of value type.
- You wrap a static bool in a selector, since static_OR needs two types, not constants.

Some iterators are known to be views on sorted sets, for example, set<T>::iterator. Can you detect them?

As is, the question is ill-formed: set<T>::iterator is a dependent type, and in C++ there's no "reverse matching" to deduce T, given iterator. 14

However, you can make the problem easier if you limit the options to some special candidates. In

fact, a set is declared as set<T, L, A>, but in some contexts you might take a guess on L and A. A practical example is given in the following code:

```
template <typename T, typename less_t, typename alloc_t
= std::allocator<T> >
class sorted_vector
{
    std::vector<T, alloc_t> data_;
    less_t less_;

public:
    template <typename iterator_t>
    sorted_vector(iterator_t begin, iterator_t end, less_t less
= less_t())
    : data_(begin, end), less_(less)
    {
        // this is unnecessary if begin...end is already sorted
        std::sort(data_.begin(), data_.end(), less_);
    }
};
```

Since the underlying sort algorithm will consume CPU even when the range is already sorted, try to guess if this step can be avoided.¹⁵

There are two distinct tests. First, some iterators guarantee that the range they point to is sorted (this is a "static test," as it depends only on the iterator type); second, any iterator pair can happen to point at a sorted range (this is a "runtime test"). You can combine the static and runtime tests in this order:

```
if (!is_sorted_iterator<iterator_t, less_t>::value)
{
   if (!is_sorted(begin, end, less_)
      std::sort(begin, end, less_);
}
```

A very important observation is that is_sorted_iterator<iterator_t, less_t> is allowed to return false negatives but not false positives. You can tolerate unnecessary sorts, but you must not let an unsorted range pass.

Note Testing if a range is already sorted takes linear time.

In C++0x, there's a dedicated algorithm:

```
template <typename FwdIter>
bool is_sorted(FwdIter begin, FwdIter end);
```

```
template <typename FwdIter, typename less_t>
bool is_sorted(FwdIter begin, FwdIter end, less_t LESS);

In classic C++, the implementation of the latter function is extremely concise:
```

```
using std::adjacent find;
using std::reverse iterator;
return
 adjacent find(reverse iterator<FwdIter>(end),
reverse iterator<FwdIter>(begin), LESS)
 == reverse_iterator<FwdIter>(begin);
  is sorted iterator<iterator t, less t> could simply try to match
iterator t against some special standard iterators:
#define ITER(C,T1)
                          typename std::C<T1,less t>::iterator
#define CONST ITER(C,T1) typename
std::C<T1,less t>::const iterator
template
  typename iter t,
  typename less t,
  typename value t = typename
std::iterator traits<iter t>::value type
>
struct is sorted iterator
static const bool value =
  static OR
  <
     static OR
     typeequal<iter t, ITER(set, value t)>,
     typeequal<iter t, CONST ITER(set, value t)>
     >,
     static OR
     typeequal<iter t, ITER(multiset, value t)>,
     typeequal<iter t, CONST ITER(multiset, value t)>
  >::value;
```

```
};
```

There's a partial specialization for maps:

```
#define ITER(C,T1,T2)
                               typename
std::C<T1,T2,less t>::iterator
#define CONST ITER(C,T1,T2)
                             typename
std::C<T1,T2,less t>::const iterator
template
<
  typename iter t,
  typename less t,
  typename T1,
  typename T2
>
struct is sorted iterator< iter t, less t, std::pair<const T1, T2>
>
{
  static const bool value =
  static OR
     static OR
     <
       static OR
        typeequal<iter t, ITER(map, T1, T2)>,
        typeequal<iter t, CONST ITER(map,T1,T2)>
       >,
       static OR
        typeequal<iter t, ITER(multimap, T1, T2)>,
        typeequal<iter t, CONST ITER(multimap,T1,T2)>
     >,
     static OR
     <
       static OR
        typeequal<iter t, ITER(map,const T1,T2)>,
        typeequal<iter t, CONST ITER(map,const T1,T2)>
       >,
       static OR
        typeequal<iter t, ITER(multimap,const T1,T2)>,
        typeequal<iter t, CONST ITER(multimap,const T1,T2)>
     >
```

```
>::value;
};
```

6.1.6. Selection by Iterator Value Type

A function that takes iterators may want to invoke another template, tagging the call with the iterator value type. In particular, this allows some mutating algorithms to deal with anomalies, such as mutable iterators that have constant references (for example, std::map).

```
template <typename iterator_t>
iterator_t F(iterator_t b, iterator_t e)
{
   typedef typename std::iterator_traits<iterator_t>::value_type
value_type;
   return F(b, e, instance_of<value_type>());
}

template <typename iterator_t, typename T1, typename T2>
iterator_t F(iterator_t b, iterator_t e, instance_of<
std::pair<const T1, T2> >)
{
   // modify only i->second
}

template <typename iterator_t, typename T>
iterator_t F(iterator_t b, iterator_t e, instance_of<T>)
{
   // modify *i
}
```

Selective-copy algorithms may use the *output* iterator value type to decide what to return. Suppose a computation produces a series of values and the corresponding weights; if the output type is a pair, the dump writes both; otherwise, it writes only the value:

```
template <[...], typename iterator_t>
void do_it([...], iterator_t out_begin, iterator_t out_end)
{
   typedef typename
     std::iterator_traits<iterator_t>::value_type value_type;

   // ...
   dump([...], out_begin, out_end, instance_of<value_type>());
}
private:
template <[...], typename iterator_t, typename T1, typename T2>
```

Note that the implementations may be unified using accessors. See the next section for details.

6.2. Generalizations

This section discusses alternative ways of coding functions, with different I/O interfaces. Since iterators offer a view over data, they may not be flexible enough, especially for algorithms that have special semantics.

Some computations may be described in terms of properties, such as "find the object whose *price* is the minimum". Surely, you will need iterators to scan the objects, but how is a price read?

6.2.1. Properties and Accessors

An algorithm that accepts iterators may not use the actual interface of the pointed type. Usually, they have two versions, one where the required operations are handled directly by the pointed type, and one that takes an extra functor that completely supersedes the object interface.

For example, std::sort(b, e) assumes that the pointed type is less-than comparable and it uses the pointer's operator<, while std::sort(b, e, LESS) uses an external binary predicate for all the comparisons, and operator< may not exist at all.

As a generalization of this concept, algorithms may be defined in terms of *properties*.

Properties generalize data members: a (read-only) *property* is simply a non-void function of a single (const) argument, which by default invokes a const member function of the argument or returns a copy of a data member of the argument. Here's a trivial example.

```
template <typename T>
struct property_size
{
  typedef size_t value_type;
  value_type operator()(const T& x) const
  {
```

```
return x.size();
};
```

An instance of property_size passed to an algorithm is called the *accessor* of the property. Many computational algorithms can be defined in terms of properties. They ignore the pointed type, but they need to read "its size"; thus, they require a suitable accessor.

By hypothesis, applying an accessor is inexpensive.

Note A property is a functor, thus the right typedef should be result_type, but the user needs to store a copy of the property value, which conceptually lies in the object and is only "accessed" by the functor. Therefore, value type is preferred.

A read-write property has an additional member that writes back a value:

```
template <typename T>
struct property_size
{
  typedef size_t value_type;

  value_type operator()(const T& x) const
  {
    return x.size();
  }

  value_type operator()(T& x, const value_type v) const
  {
    x.resize(v);
    return x.size();
  }
};
```

Accessors are useful in different contexts. In the simplest case, they save calls to std::transform or to custom binary operators.

Suppose you have a range of std::string, and need to find the total size and the maximum size. Using the classical STL, you can write a custom "sum" and a custom "less". The transformation from string to integer (the size) is performed inside these functors.

```
struct sum_size
{
    size_t operator()(size_t n, const std::string& s) const
    {
       return n + s.size();
    }
};
```

```
struct less by size
  bool operator()(const std::string& s1, const std::string& s2)
const
     return s1.size() < s2.size();
};
// assume beg!=end
size t tot = std::accumulate(beg, end, OU, sum size());
size t max = std::max element(beg, end, less by size())->size();
  Using accessors, you would have more code reuse:
#define VALUE
              typename accessor t::value type
template <typename iterator t, typename accessor t>
VALUE accumulate(iterator t b, iterator t e, accessor t A, VALUE
init = 0
  while (b != e)
     init = init + A(*b++);
  return init;
template <typename iterator t, typename accessor t>
iterator t max element (iterator t b, iterator t e, accessor t A)
  if (b == e)
     return e;
  iterator t result = b;
  while ((++b) != e)
     if (A(*result) < A(*b))
       result = b;
  }
  return result;
}
size t tot = accumulate(beg, end, property size<std::string>());
size t max = max element(beg, end, property size<std::string>());
```

The default accessor returns the object itself:

```
template <typename T>
struct default_accessor
{
   typedef T value_type;

   T& operator()(T& x) const
   {
     return x;
   }
};
```

Accessors offer a good degree of abstraction in complex computational algorithms that need several "named properties" at a time. We cite the knapsack problem as an example.

Each object has two properties: a (nonnegative) price and a (nonnegative) quality. You are given an initial amount of money and your objective is to buy the subset of objects having maximal total quality. At the end of the computation, (part of) the result is a subset of the original range, so you choose a reordering algorithm. You return an iterator that partitions the range, paired with an additional by-product of the algorithm—in this case, the total quality.

The function prototype in terms of accessors is long, but extremely clear:

price_t and quality_t are accessors for the required properties.

So the price of the element *i is simply price (*i), and it can be stored in a variable having type typename price_t::value_type.

To illustrate the usage, here's a non-mutating function that simply evaluates the total quality of the solution, assuming that you buy all elements possible starting from begin:

```
break;
     money -= p;
     total q += quality(*begin++);
  return total q;
}
  For algorithm testing, you will usually have the accessors fixed. It can be convenient to generate a
placeholder structure that fits:
struct property price
  typedef unsigned value type;
  template <typename T>
  value type operator()(const T& x) const
   {
     return x.price();
};
struct price tag t {};
struct quality_tag_t {};
struct knapsack object
  property<unsigned, price tag t> price;
  property<unsigned, quality tag t> quality;
};
  The property class is described next.
  The extra tag forbids assignment between different properties having the same underlying type
(for example, unsigned int).
template <typename object t, typename tag t = void>
class property
{
  object t data ;
public:
  property()
   : data () // default-constructs fundamental types to zero
   { }
  property(const object t& x)
   : data (x)
```

{ }

```
const object_t& operator()() const
{
    return data_;
}

const object_t& operator()(const object_t& x)
{
    return data_ = x;
}

const char* name() const
{
    return typeid(tag_t).name();
}
```

6.2.2. Mimesis

Some general-purpose algorithms accept a range—that is, two iterators [begin, end) and an additional value—or a unary predicate. These algorithms are implemented twice. The latter version uses the predicate to test the elements and the former tests for "equality with the given value".

A classic example is std::find versus std::find if.

```
template <typename iter t, typename
                                            template <typename iter t, typename
object t>
                                            functor t>
iter t find(iter t begin, iter t end,
                                            iter t find if (iter t begin, iter t end,
object tx)
                                            functor t f)
    for (; begin != end; ++begin)
                                                for (; begin != end; ++begin)
        if (*begin == x)
                                                     if (f(*begin))
            break;
                                                         break;
    return begin;
                                                return begin;
```

In principle, find could be rewritten in terms of find if:

```
template <typename iter_t, typename object_t>
iter_t find(iter_t begin, const iter_t end, object_t x)
{
    std::equal_to<object_t> EQ;
    return std::find_if(begin, end, std::bind2nd(EQ, x));
}
```

But the converse is also possible:

```
template <typename functor t>
class wrapper
  functor t f ;
public:
  wrapper(functor t f = functor t())
  : f (f)
    // verify with a static assertion that
    // functor t::result type is bool
  }
  bool operator == (const typename functor t:: argument type& that)
const
  {
     return f (that);
};
template <typename iter t, typename functor t>
iter t find if (iter t begin, const iter t end, functor t F)
  return std::find(begin, end, wrapper<functor t>(F));
}
```

A mimesis object for type T informally behaves like an instance of T, but internally it's a unary predicate. A mimesis implements operator== (const T&), operator!=(const T&), and operator "cast to T" (all operators being const).

To invoke a predicate, you write:

```
if (f(x))
```

To invoke a mimesis, the equivalent syntax would be:

```
if (f == x)
```

These requirements are slightly incomplete:

- The equality and inequality should take the mimesis itself and a T in any order, to prevent the undesired usage of "cast to T" for comparisons (you'll read more about this later).
- The cast operator should return a prototype value that satisfies the same criteria.

In other words, if M is a mimesis for type T, then the fundamental property for M is:

```
assert(m == static cast<T>(m));
   The simplest mimes is for type T is T itself.
   As a very simple case, let's implement a mimesis that identifies positive numbers:
template <typename scalar t >
struct positive
   bool operator == (const scalar t& x) const
   {
      return 0<x;
   }
   bool operator!=(const scalar t& x) const
   {
      return !(*this == x);
   operator scalar t() const
      return 1; // an arbitrary positive number
};
   Here's the first application where you don't need find if any more.
double a[] = \{ -3.1, 2.5, -1.0 \};
std::find(a, a+3, positive<double>()); // fine, returns pointer to
2.5
   The key is that the value parameter in find has an independent template type, so
positive < double > is passed over as is, without casting.
   A deduced template type, such as:
template <typename I>
iter t find(I, I, typename std::iterator traits<I>::value type x)
would have caused the mimesis to decay into its default value (and consequently, would return a
wrong find result).
   The mimes is interface can in fact be richer:
template <typename scalar t, bool SIGN = true>
struct positive
   bool operator == (const scalar t& x) const
      return (0 < x) ^ (!SIGN);
```

```
bool operator!=(const scalar t& x) const
     return !(*this == x);
  operator scalar t() const
     // arbitrary positive and non-positive numbers
     return SIGN ? +1: -1;
  positive < scalar t, !SIGN > operator!() const
     return positive < scalar t, !SIGN>();
};
template <typename scalar t, bool SIGN>
inline bool operator == (const scalar t& x, const positive < scalar t,
SIGN> p)
  return p == x;
template <typename scalar t, bool SIGN>
inline bool operator!=(const scalar t& x,
                    const positive<scalar t, SIGN> p)
{
  return p != x;
}
```

Thus positive double, true will compare equal to any strictly positive double and it will convert to 1.0 when needed. On the other hand, positive double, false will compare equal to non-positive numbers, and it will return -1.0.

Note that the user will simply write positive<T>() or !positive<T>().

```
std::find(a, a+3, !positive<double>());
```

You have seen that writing a mimesis takes more effort than a functor, but it's worth it, especially for generalizing functions that take a special value as an argument. The next section offers another application.

6.2.3. End of Range

Iterator-based algorithms cannot compute the end of a range dynamically. For example, you cannot express the concept "find 5.0 but stop on the first negative number," because the range is precomputed. You need two function calls.

```
using namespace std;
find(begin, find if(begin, end, bind2nd(less<double>(), 0.0)),
5.0)
```

The canonical example of range inefficiency is given by C strings. Suppose you are copying a C string and you get an output iterator to the destination:

```
const char* c string = "this is an example";
// can we avoid strlen?
std::copy(c string, c string+strlen(c string), destination);
```

strlen has to traverse the string looking for the terminator, then copy traverses it again. The process in practice is extremely fast, but it does an unnecessary pass.

Suppose for a moment that you rewrite copy. You don't change the function body, but just allow

```
the endpoints of the range to have different types.
template <typename iter1 t, typename iter2 t, typename end t>
iter2 t copy 2(iter1 t begin, end t end, iter2 t output)
  while (begin != end)
     *(output++) = *(begin++),
  return output;
}
  This is equivalent to asking that end be a mimesis for the type iter1 t.
  Compare with the following code:
template <typename char_t, char t STOP = 0>
struct c string end
  typedef char t* iterator t;
  operator iterator t() const { return 0; }
  bool operator!=(const iterator t i) const
   {
     return ! (*this == i);
   }
  bool operator == (const iterator t i) const
   {
     return i==0 || *i==STOP;
};
```

// implement operator== and != with arguments in different order

// ...

```
const char* begin = "hello world!";
copy_2(begin, c_string_end<const char>(), output); // ok and
efficient!
copy 2(begin, begin+5, output); // also ok!
```

The latter invocation of copy_2 is equivalent to std::copy.

To sum up, a mimesis has two uses:

- Algorithms that accept a "test," which can be either a value or a predicate.
- Algorithms that process a range begin...end, where end is just a termination criteria (that is, it's not decremented).

Note the difference between the two. The "test" mentioned in the first point is a skip-and-continue condition *on elements*; end is a terminate-and-exit criterion *on iterators*.

The cast operator in the interface of a mimesis turns out to be useful when the object acts as a skip-and-continue filter. Assume you are computing the average of all the elements that satisfy some criteria. First, you write a tentative "classical" version.

```
template <class iter_t, class predicate_t>
typename std::iterator_traits<iter_t>::value_type
   average_if(iter_t begin, iter_t end, predicate_t f)
{
   size_t count = 0;
   typename std::iterator_traits<iter_t>::value_type result = 0;
   for (; begin != end; ++begin)
   {
      if (f(*begin))
      {
        result += *begin;
        ++count;
      }
   }
   return count>0 ? result/count : [[???]];
}
```

If the predicate rejects all elements, you don't know what to return, except possibly std::numeric_limits<...>::quiet_NaN() (hoping that has quiet NaN is true).

However, the best choice is to ask the functional object what to return. If F is seen as the rejection logic (not the acceptance), it should be also responsible for providing a prototype of the rejected element, and that's exactly the fundamental property for a mimesis.

That's why you rewrite the algorithm using a mimesis standing for a quiet NaN: 16

```
template <typename iter_t, typename end_t, typename nan_t>
typename std::iterator_traits<iter_t>::value_type
    average(iter_t begin, const end_t end, nan_t NaN)
```

```
{
  size t count = 0;
  typename std::iterator traits<iter t>::value type result = 0;
   for (; begin != end; ++begin)
   {
     if (NaN != *begin)
        result += *begin;
        ++count;
      }
  return count>0 ? result/count : NaN;
}
  A typical role of a mimesis is to represent an "exclusion filter":
template <typename scalar t>
struct ieee nan
  operator scalar t() const
     return std::numeric limits<scalar t>::quiet NaN();
   }
  bool operator!=(const scalar t& x) const
   {
     return x == x;
   }
  bool operator == (const scalar t& x) const
   {
     return x != x;
};
  The dangerous downside of a cast operator is that it can be called unexpectedly. Consider again
the example four pages ago:
template <typename iterator t, char STOP = 0>
struct c string end
   // ...
// ooops. forgot to implement operator == and !=
// with arguments in different order
// later...
```

```
while (begin != end)
{
    // ...
}
```

begin!=end will actually call bool operator!=(const char*, const char*) passing begin, which is already a pointer, and applying a cast to end (which produces a null pointer). Therefore, the loop will never exit.

Note also that it's possible to wrap a mimesis and turn it into a predicate and vice versa.

6.3. Iterator Wrapping

Writing STL-compliant iterators is a complex activity and it involves a lot of code duplication. Luckily, writing const iterators is far easier.

A wrapped iterator, const or non-const, is a class that contains another iterator as a member. The wrapper forwards every "positioning operation" (for example, increments and decrements) to the member, but it intercepts dereferencing, changing the result so as to express a logical view on the underlying dataset.

Since the end user may not see actual data, but a custom-forged value, it's often impossible to modify the original objects through the view, so that wrapped iterators are mostly const_iterators.

Suppose you have a vector of integers and an iterator wrapper that returns the actual value multiplied by 5.

}

Even if multiplier_iterator could physically write an integer at position data[0], what should it do? Were it smart enough, it would write 25/5=5, so that *i returns 25 from that point on.

However, the instruction *i = 24 is even more problematic. Should it throw an exception? Or do nothing? Or set data [0] = (24 + (5-1)) / 5 anyway?

A correct implementation of operator-> is indeed the hardest issue. A lucky wrapper will simply dispatch the execution to the wrapped iterator recursively, but since this usually reveals the "real" data underneath, it may not be compatible with the wrapping logic.

Consider instead *omitting* operators that are less likely to be used. operator-> is the first candidate, unless their implementation is both trivial and correct.¹⁷

The arrow operator is used to access members of the pointed type, but in portions of code where this type is generic (for example, a template parameter may be deduced), these members are usually *not known*, so the arrow should not be used.¹⁸

For example, std::vector::assign will generally work even on iterators having no operator->.

6.3.1. Iterator Expander

Iterator wrappers will delegate most operations to the wrapped object, such as operator++.

The dispatching part is extremely easy to automate using a static interface, ¹⁹ which is named iterator expander:

```
class wrapper
: public iterator_expander<wrapper>
, public std::iterator_traits<wrapped>
{
    wrapped w_;

public:
    wrapped& base()
    {
       return w_;
    }

    const wrapped& base() const
    {
       return w_;
    }

    wrapper(wrapped w)
      : w_(w)
    {
    }
}
```

```
[...] operator* () const
{
    // write code here
}
[...] operator-> () const
{
    // write code here
}
};
```

The iterator_expander interface (listed next) is responsible for all possible positioning (++, ++, +=, -=, +, and -) and comparison operators. They are all implemented, and as usual, they will be compiled only if used. If wrapped does not support any of them, an error will be emitted (no static assertion is necessary, as the cause of the error will be evident).

Note also that every operator in the interface returns true_this(), not *this, because otherwise a combined expression such as *(i++) would not work. iterator_expander does not implement operator*, but true this() returns the actual wrapper.

```
template <typename iterator t, typename diff t = ptrdiff t>
class iterator expander
protected:
// the static interface part, see Section 6.2
~iterator expander() {}
 iterator expander() {}
iterator t& true this()
{ return static cast<iterator t&>(*this); }
const iterator t& true this() const
 { return static cast<const iterator t&>(*this); }
public:
iterator t& operator++() { ++true this().base(); return
true this(); }
iterator t& operator--() { --true this().base(); return
true this(); }
iterator t& operator+=(diff t i)
{ true this().base() += i; return true this(); }
iterator t& operator-=(diff t i)
{ true this().base() -= i; return true this(); }
iterator t operator++(int)
 { iterator t t(true this()); ++(*this); return t; }
iterator t operator--(int)
```

```
{ iterator_t t(true_this()); --(*this); return t; }
 iterator t operator+(diff t i) const
 { iterator t t(true this()); t+=i; return t; }
iterator t operator-(diff t i) const
 { iterator t t(true this()); t-=i; return t; }
diff t operator-(const iterator expander& x) const
 { return true this().base() - x.true this().base(); }
bool operator<(const iterator expander& x) const
{ return true this().base() < x.true this().base(); }
bool operator == (const iterator expander & x) const
{ return true this().base() == x.true this().base(); }
bool operator!=(const iterator expander& x) const
{ return ! (*this == x); }
bool operator> (const iterator expander& x) const
{ return x < *this; }
bool operator<=(const iterator expander& x) const</pre>
{ return !(x < *this); }
bool operator >= (const iterator expander & x) const
{ return ! (*this < x); }
};
  You also need an external operator:
template <typename iterator t, typename diff t>
iterator t operator+(diff t n, iterator expander<iterator t,
diff t> i)
  return i+n;
}
  Note that difference type is taken, not deduced. iterator expander <T > cannot read
types defined in T, because it's compiled before T since it's a base of T.
  So the wrapper will be declared as follows:
template <typename iterator t>
class wrapper
: public iterator expander
          wrapper<iterator t>,
          typename
std::iterator traits<iterator t>::difference type
  // ...
```

```
};
```

Here's a trivial practical example that also shows that the iterator base can be a simple integer.

```
class random iterator
: public iterator expander<random iterator>
 public std::iterator traits<const int*>
  int i ;
public:
  int& base() { return i ; }
  const int& base() const { return i ; }
  explicit random iterator(const int i=0)
     : i (i)
  }
  int operator*() const
     return std::rand();
};
int main()
  std::vector<int> v;
  v.assign(random iterator(0), random iterator(25));
  // now v contains 25 random numbers
  //...
```

Note that this example skips the arrow operator and dereferencing returns a value, not a reference (but since the class inherits const int* traits, it's still possible to bind a reference to *iterator, as reference is const int&).

Note Don't store copies of values in iterators. While this actually allows returning genuine references and pointers, the referenced entity has a lifetime that is bound to the iterator, not to the "container" (in other words, destroying the iterator, the reference becomes invalid), and this will lead to subtle bugs. Here's some *bad* code:

```
class random_iterator
: public iterator_expander<random_iterator>
, public std::iterator_traits<const int*>
{
```

```
int i_;
int val_; // bad

public:
   const int& operator*() const
   {
     return val_ = std::rand(); // bad
   }

   const int* operator->() const
   {
     return &*(*this); // even worse
   }
};
```

Iterator wrappers solve the problem of iterating over values in a map (or equivalently, the problem of const-iterating over keys).

This time, the example is going to be a true non-const iterator implementation, because you iterate over existing elements, so you can return pointers and references.

```
template <typename T, int N>
struct component;
template <typename T1, typename T2>
struct component<std::pair<T1, T2>, 1>
  typedef T1 value type;
  typedef T1& reference;
  typedef const T1& const reference;
  typedef T1* pointer;
  typedef const T1* const pointer;
};
template <typename T1, typename T2>
struct component<std::pair<const T1, T2>, 1>
  typedef T1 value type;
  typedef const T1& reference;
  typedef const T1& const reference;
  typedef const T1* pointer;
  typedef const T1* const_pointer;
};
template <typename T1, typename T2>
```

```
struct component<std::pair<T1, T2>, 2> : component<std::pair<T2,
T1>, 1>
{
};
```

Assume that iterator_t (the wrapped type) points to a std::pair-like class. If that's not the case, the compiler will give an error when compiling one of the ref overloads.

```
template <typename iterator t, int N>
class pair iterator
: public iterator expander< pair iterator<iterator t, N> >
  static const bool IS MUTABLE =
     is mutable iterator<iterator t>::value;
  iterator t i ;
  typedef std::iterator traits<iterator t> traits t;
  typedef component<typename traits t::value type, N> component t;
  typedef typename component t::reference ref t;
  typedef typename component t::const reference cref t;
  typedef typename component t::pointer ptr t;
  typedef typename component_t::const_pointer cptr_t;
  template <typename pair t>
  static ref t ref(pair t& p, static value<int, 1>)
  { return p.first; }
  template <typename pair t>
  static ref t ref(pair t& p, static value<int, 2>)
  { return p.second; }
  template <typename pair t>
  static cref t ref(const pair t& p, static value<int, 1>)
  { return p.first; }
  template <typename pair t>
  static cref t ref(const pair t& p, static value<int, 2>)
  { return p.second; }
public:
  explicit pair iterator(iterator t i)
  : i_(i)
  { }
  iterator t& base() { return i ; }
```

```
const iterator t& base() const { return i ; }
  typedef typename typeif<IS MUTABLE, ref t, cref t>::type
reference;
  typedef typename typeif<IS MUTABLE, ptr t, cptr t>::type
pointer;
  typedef typename component t::value type value type;
  typedef typename traits t::iterator category iterator category;
  typedef typename traits t::difference type difference type;
  reference operator* () const
     return ref(*i , static value<int, N>());
  pointer operator->() const
     return &*(*this);
};
  Here's a driver function:
template <int N, typename iterator t>
inline pair iterator<iterator t, N> select(iterator t i)
  return pair iterator<iterator t, N>(i);
  And finally some example code. The syntax for the driver is select<N>(i) where N is 1 or 2
and i is an iterator whose value type is a pair:
template <typename T>
struct Doubler
  void operator()(T& x) const
     x *= 2;
};
template <typename T>
struct User
  void operator()(const T& x) const
     std::cout << x << ';';
};
```

```
typedef std::map<int, double> map t;
MXT ASSERT(!is mutable iterator<map t::const iterator>::value);
MXT ASSERT(is mutable iterator<map t::iterator>::value);
map t m;
const map t \& c = m;
m[3] = 1.4;
m[6] = 2.8;
m[9] = 0.1;
// print 3;6;9; via iterator
std::for each(select<1>(m.begin()), select<1>(m.end()), User<int>
());
// print 3;6;9; via const iterator
std::for each(select<1>(c.begin()), select<1>(c.end()), User<int>
());
// multiplies by 2 each value in the map
std::for each(select<2>(m.begin()), select<2>(m.end()),
Doubler<double>());
std::vector<double> v1;
v1.assign(select<1>(c.begin()), select<1>(c.end()));
std::vector< std::pair<int, double> > v2(m.begin(), m.end());
// multiplies by 2 each key in the vector (the key is not
constant)
std::for each(select<1>(v2.begin()), select<1>(v2.end()),
Doubler<int>());
// these two lines should give an error:
// std::for each(select<1>(m.begin()), select<1>(m.end()),
Doubler<int>());
// std::for each(select<1>(c.begin()), select<1>(c.end()),
Doubler<int>());
```

6.3.2. Fake Pairs

The inverse problem is "merging" two logical views and obtaining a single iterator that makes them look like pairs. With pair_iterator, you can build a vector of keys and a vector of values reading a map, but not the other way around.

```
std::vector<int> key;
```

```
std::vector<double> value;
std::map<int, double> m = /* ??? */;
```

struct plusplus

Actually, you can extend the interface of iterator expander to allow the possibility that the derived class has more than one base. Simply let base have N overloads, taking a static_value<size_t, N>, and each can possibly return a reference to an iterator of a different kind.

You can isolate the elementary modifiers to be applied to the bases and code a very simple statically-recursive method.²⁰

Since you do not know in advance what base (static_value<size_t, K>) is, you must introduce some auxiliary "modifier" objects with template member functions, as follows:

```
template <typename any t>
     void operator()(any t& x) const { ++x; }
};
class pluseq
  const diff t i ;
public:
  pluseq(const diff t i) : i (i) {}
  template <typename any t>
     void operator()(any_t& x) const { x += i ; }
} ;
template <typename iterator t, size t N, typename diff t>
class iterator pack
protected:
  typedef static value<size t, N> n times;
  ~iterator pack() {}
  iterator pack() {}
  iterator t& true this()
     return static cast<iterator t&>(*this);
  const iterator t& true this() const
     return static cast<const iterator t&>(*this);
```

```
/* static recursion */
  template <typename modifier t, size t K>
  void apply(const modifier t modifier, const static value<size t,
K > )
  {
     modifier(true this().base(static value<size t, K-1>()));
     apply(modifier, static value<size t, K-1>());
  }
  template <typename modifier t>
  void apply(const modifier t modifier, const static value<size t,
0 > )
public:
  typedef diff t difference type;
  iterator t& operator++()
     apply(plusplus(), n times());
     return true this();
  }
  iterator t& operator+=(const diff t i)
     apply(pluseq(i), n times());
     return true this();
   }
  You need to add a few more member functions. For simplicity, some operators, such as
comparisons, make use of the first element only: <sup>21</sup>
typedef static value<size t,0> default t;
diff t operator-(const iterator pack& x) const
  const default t d;
  return true this().base(d) - x.true this().base(d);
}
bool operator<(const iterator pack& x) const
  const default t d;
  return true this().base(d) < x.true this().base(d);
}
```

```
bool operator==(const iterator_pack& x) const
{
   const default_t d;
   return true_this().base(d) == x.true_this().base(d);
}
```

All other operators derive from the basic ones in the usual way—postfix ++ and operator+ from prefix ++ and += and other comparisons from < and ==.

With the new tool at your disposal, here's a not-fully-standard iterator that pretends to iterate over std::pair.

First, some highlights:

- Here, pointer is void, because you don't want to support operator->, but
 to compile std::iterator_traits< iterator_couple<...>>
 pointer needs to be defined; however this definition will prevent any other
 use.
- iterator_category is the weaker of the two categories; however, you can statically-assert that both categories should be comparable so as to avoid unusual pairs (such as input/output iterators). Of course, the restriction could be removed.
- The main problem is how to define reference. Obviously, you have to rely on r1_t and r2_t but cannot use std::pair<r1_t, r2_t>. (Mainly because, in classic C++, std::pair does not support it and it will not compile.)²²

```
typedef typename
     typeif
     <
        is base of < cat1 t, cat2 t>::value,
        cat1 t,
        cat2 t
     >::type iterator category;
  typedef std::pair<v1 t, v2 t> value type;
  typedef void pointer;
  struct reference
     /* see below... */
  };
  iterator1 t& base(static value<size t, 0>) { return i1 ; }
  iterator2 t& base(static value<size t, 1>) { return i2 ; }
  const iterator1 t& base(static value<size t, 0>) const
  { return i1 ; }
  const iterator2 t& base(static value<size t, 1>) const
  { return i2 ; }
  reference operator* () const
  {
     MXT ASSERT
        (is base of < cat1 t, cat2 t>::value
             || is base of<cat2 t, cat1 t>::value)
     );
     return reference(*i1_, *i2_);
  }
private:
  iterator1 t i1 ;
  iterator2 t i2;
} ;
  You have to emulate a pair of references, which std::pair does not allow:
struct reference
  rl t first;
  r2 t second;
```

```
}
  operator std::pair<v1 t, v2 t>() const
     return std::pair<v1 t, v2 t>(first, second);
  template <typename any1 t, typename any2 t>
     operator std::pair<any1 t, any2 t>() const
  {
     return std::pair<any1 t, any2 t>(first, second);
  reference& operator= (const std::pair<v1 t, v2 t>& p)
     first = p.first;
     second = p.second;
     return *this;
  }
  void swap(reference& r)
     swap(first, r.first);
     swap(second, r.second);
  }
  void swap(std::pair<v1 t, v2 t>& p)
     swap(first, p.first);
     swap(second, p.second);
};
  The template cast-to-pair operator is needed since std::map will likely cast the reference not
to pair<V1, V2>, but to pair<const V1, V2>.
  This implementation may suffice to write code like this:
template <typename iter1 t, typename iter2 t>
iterator couple<iter1 t, iter2 t> make couple(iter1 t i1, iter2 t
i2)
  return iterator couple<iter1 t, iter2 t>(i1, i2);
std::vector<int> k;
```

reference(r1_t r1, r2_t r2)
 : first(r1), second(r2)

Note that the first insert gets a bidirectional iterator, whereas the last assign gets a random-access iterator.²³

6.4. Receipts

Receipts are empty classes that can be created only by "legally authorized entities" and unlock the execution of functions, as required parameters. Some receipts can be stored for later use, but some instead must be passed on immediately.

In a very simple case, when you need to enforce that function F is called before G, you modify F and let it return a receipt R, which cannot be constructed otherwise. Finally, G takes R as an additional—formal—parameter.

Receipts are mostly useful in connection with hierarchy of classes, when a virtual member function foo in every DERIVED should invoke BASE::foo at some point.

Assume for the moment that foo returns void.

There are two similar solutions:

• In the public non-virtual/protected virtual technique, the base class implements a public non-virtual foo, which calls a protected virtual function when appropriate.

```
class BASE
{
protected:
    virtual void custom_foo()
    {
    }

public:
    void foo()
    {
        /* ... */
```

```
};
        • Use receipts. BASE:: foo returns a secret receipt, private to BASE.
class BASE
protected:
  class RECEIPT_TYPE
      friend class BASE;
     RECEIPT TYPE() {} // constructor is private
   };
public:
  virtual RECEIPT TYPE foo()
      /* ... */
     return RECEIPT TYPE();
};
class DERIVED : public BASE
public:
  virtual RECEIPT TYPE foo()
      /* ... */
      // the only way to return is...
      return BASE::foo();
};
  If RECEIPT TYPE has a public copy constructor, DERIVED can store the result of
BASE:: foo at any time. Otherwise, it's forced to invoke it on the return line.
```

Note that a non-void return type T can be changed into std::pair<T, RECEIPT TYPE>, or a custom class, which needs a receipt, but ignores it.

Receipts are particularly useful in objects, where you want to control the execution order of member functions (algors are described in Section 8.6):

```
class an algor
public:
  bool initialize();
```

custom foo();

```
void iterate();
  bool stop() const;
  double get result() const;
} ;
double execute correctly algor(an_algor& a)
  if (!a.initialize())
     throw std::logic error("something bad happened");
  do
     a.iterate();
  } while (!a.stop());
  return a.get result();
}
double totally crazy execution (an algor& a)
  if (a.stop())
     a.iterate();
  if (a.initialize())
     return a.get result();
  else
     return 0;
}
  In general, you want initialize to be called before iterate, and get result after at
least one iteration. So you need to modify the interface as follows:
template <int STEP, typename T>
class receipt t : receipt t<STEP-1, T>
  friend class T;
                                             // note: private
  receipt t() {}
};
template <typename T>
class receipt t<0, T>
  friend class T;
                                             // note: private
  receipt t() {}
};
```

```
class a_better_algor
{
public:
    typedef receipt_t<0, a_better_algor> init_ok_t;
    typedef receipt_t<1, a_better_algor> iterate_ok_t;
    init_ok_t initialize();
    iterate_ok_t iterate(init_ok_t);
    bool stop(iterate_ok_t) const;
    double get_result(iterate_ok_t) const;
};
```

With the necessary evil of a template friendship declaration (which is non-standard yet), the idea should be clear: since the user cannot forge receipts, she must store the return value of initialize and pass it to iterate. Finally, to get the result, it's necessary to prove that at least one iteration was performed:²⁴

```
a_better_algor A;
a_better_algor::init_ok_t RECEIPT1 = A.initialize();
while (true)
{
    a_better_algor::iterate_ok_t RECEIPT2 = a.iterate(RECEIPT1);
    if (a.stop(RECEIPT2))
        return a.get_result(RECEIPT2);
}
```

Note The code:

```
template <typename T>
class ...
{
    friend class T;
```

is not standard in classic C++ because the statement could be nonsensical when T is a native type |(say, int). However, it's accepted by some compilers as an extension. In C++0x, it's legal, but the syntax is:

```
friend T:
```

As a workaround, some (but not all) classic C++ compilers accept this. The rationale for this workaround is that it introduces an additional indirection, which

allows the compiler to treat T as an "indirect" type, not as a template parameter.

```
template <typename T>
class ...
{
   struct nested_t { typedef T type; };
   friend class nested t::type;
```

6.5. Algebraic Requirements

6.5.1. Less and NaN

Objects of generic type T are often assumed LessThanComparable.

This means that either T: : operator< is defined, or an instance of a binary predicate "less" is given as an extra argument.²⁵

An algorithm should avoid mixing different comparison operators, such as operator<= and operator>, because they could be inconsistent. The best solution is to replace them with operator< (or with the binary predicate "less").

It's questionable if operator== should be assumed valid or replaced with the *equivalence* test. In fact, two calls to operator< may be significantly slower (as in std::string). However, in some cases, with additional hypotheses, one of the tests may be omitted. For example, if a range is sorted, a test with iterators *i == *(i+k) can be replaced by !less(*i, *(i+k)).

A NaN (Not-a-Number) is an instance of T that causes any comparison operator to "fail". In other words, if at least one of x and y is NaN, then $x \in P$ y returns false if OP is <,>,<=,>=,== and returns true if OP is !=. In fact, a NaN can be detected by this simple test:

```
template <typename T>
bool is_nan(const T& x)
{
   return x != x;
}
```

Types double and float have a native NaN.

If T has a NaN, it can create problems with sorting algorithms. Two elements are equivalent if neither is less than the other, ²⁶ so a NaN is equivalent to any other element. If you write, for example:

```
std::map<double, int> m;
// insert elements...
m[std::numeric_limits<double>::quiet_NaN()] = 7;
```

you are effectively overwriting a random (that is, an implementation dependent) value with 7.

The right way to deal with ranges that may contain NaN is to partition them out before sorting, or modify the comparison operator, so for example they fall at the beginning of range:

```
template <typename T>
struct LessWithNAN
{
  bool operator() (const T& x, const T& y) const
  {
    if (is_nan(x))
      return !is_nan(y);
    else
      return x<y;
  }
};</pre>
```

6.6. The Barton-Nackman Trick

Knuth wrote that a *trick* is a clever idea that is used once and a *technique* is a trick that is used at least twice. The Barton-Nackman *technique*, also known as *restricted template expansion*, is a way to declare non-member functions and operators inside a class, marking them as friends:

};

The *non-member* function and operator shown here are *non-template* functions that are injected in the scope of X<T>, when the class is instantiated. In other words, they are found with ADL, so at least one of the arguments must have type X<T>.

The main use of this technique is to declare global functions that take an inner class of a template class.

```
template <typename T>
struct outer
{
  template <int N>
  struct inner {};
};
```

You cannot write a template that takes outer<T>::inner<N> for arbitrary T, because T is non-deducible. However the Barton-Nackman trick will do:

```
template <typename T>
struct outer
{
  template <int N>
    struct inner
  {
     friend int f(inner<N>)
      { return N; }
};
```

Regardless of the fact that f is not a template, you can manipulate the template parameters at will:

```
template <typename T>
struct outer
{
   template <int N>
   struct inner
   {
      friend inner<N+1> operator++(inner<N>)
      { return inner<N+1>(); }
};

outer<double>::inner<0> I1;
++I1; // returns outer<double>::inner<1>
```

You can also write template functions in the same way, but *all parameters should be deducible* because ADL does not find functions with explicit template parameters. The following code is correct:

```
template <typename T>
struct outer
{
  template <int N>
    struct inner
  {
    inner(void*) {}

    template <typename X>
      friend inner<N+1> combine(inner<N>, X x)
      { return inner<N+1>(&x); }
  };
};

outer<double>::inner<0> I;
combine(I, 0);
```

Instead, the following example works only when outer is in the global scope, but not if it's enclosed in a namespace:

```
template <typename T>
struct outer
{
   template <typename S>
   struct inner
   {
      template <typename X>
        friend inner<X> my_cast(inner<S>)
        { return inner<X>(); }
   };
};
outer<double>::inner<int> I;
outer<double>::inner<float> F = my cast<float>(I);
```

The only workaround for having a functional my_cast as shown here would be a static interface, with the base class at namespace level, but the required machinery is non-negligible:

```
// note: global scope
template <typename T, typename S>
struct inner_interface
{
};
```

```
namespace XYZ
  template <typename T>
  struct outer
     template <typename S>
     struct inner : inner interface<T, S>
        inner(int = 0)  {}
     };
  };
// note: global scope
template <typename X, typename T, typename S>
typename XYZ::outer<T>::template inner<X>
  my cast(const inner interface<T,S>& x)
{
  // cast x to outer<T>::inner<S> if necessary
  return 0;
int main()
  XYZ::outer<double>::inner<int> I;
  my cast<float>(I);
  Obviously, my cast could be simply a template member function of inner, but this may force
clients to introduce a template keyword between the dot and the function name:
template <typename S>
struct inner
  template <typename X>
  inner<X> my_cast() const
  { return inner<X>(); }
};
outer<double>::inner<int> I;
outer<double>::inner<float> F = I.my cast<float>(); // Ok.
template <typename T>
void f(outer<T>& o)
  o.get inner().my cast<float>();
  // error: should be
  // o.get inner().template my cast<float>()
```

//...

In modern C++, there are member functions called cbegin () and cend that always return const-iterators, and so the loop would be for (auto i = v.cbegin(); i != v.cend(); ++i).

The C++ Standard guarantees that *i is a real reference, but it may be useful not to assume this unless necessary. Not all containers are standard-compliant, and in fact not even std::vector<bool>is.

³value_type is granted by the standard to be non-const qualified, but this is not enough. For example, std::map's value type is pair<const key type, mapped type> which is not assignable; more about this problem is in Section 6.3.

⁴And possibly suboptimal, but this is not really relevant here.

 $^{^5}$ A similar function is part of the modern C++ standard. See http://en.cppreference.com/w/cpp/algorithm/minmax.

⁶This is of course fair, but arbitrary. About half the STL containers have bidirectional iterators and half random-access. However, if you weight them by usage and include plain pointers, the average iterator would be more random-access.

⁷Ideally, you would like these sums to differ by 0 or 1. This is a NP-hard problem, anyway.

⁸The swapping process will leave some elements behind. That's the reason for the loop on the size of M. In the worst case, about half of the members of M will be skipped, so the algorithm complexity is still superlinear; that is, it takes time proportional to $\sim n \cdot \log(n)$ when processing n elements.

⁹The fastest classical algorithms, *quicksort* and *heapsort*, can sort a random-access container in place in superlinear time (usually std::sort is a combination of both). *mergesort* is a third superlinear method that works with forward iterators, but it requires extra memory to hold a copy of the whole input. In practice, however, if such extra memory is available, it may be convenient to copy/swap the input in a vector, sort the vector, and put the result back.

¹⁰Refer to the excellent description in Chapter 3 of [6].

¹¹See Section 4.3.1 on SFINAE.

¹²Of course, if a metafunction is known to fail for some specific type, it's always possible for the user to specialize it explicitly. Note also that the boost library takes another approach: if x is an instance of T, it checks if *x can be converted to T's $value_type$. See boost::is readable iterator.

¹³Remember that X is convertible to Y if given two functions void F(X) and Y(G), the call F(G) is legal. If X=T and Y=T& or Y=const T&, the call is fine. Alternatively, if X=T& and Y=T, the call is invalid. That's precisely the way has conversion works.

¹⁴As remarked in Chapter 1, there is actually a way to do this, but it does not scale well and it's intrusive. It requires cooperation from the author of std::set; such a technique is the topic of Section 6.6.

¹⁵Among all superlinear algorithms, only some mergesort variant may take advantage of a sorted input; quicksort and heapsort, on the other hand, do not depend significantly on the "entropy" of the initial data.

 $^{^{16}{}m The~algorithm~should~now~be~named~average_if_not.}$

¹⁷As a rule, iterator wrappers need not be 100% standard-conforming, as there are some common issues that won't compromise their functionality. The most common are lack of operator-> and operator* returning a value and not a reference (in other words: iterator::reference and iterator::value_type are the same). On the other hand, the implementation with these simplified features may be much easier. See the random_iterator example later in this chapter.

¹⁸An exception is std::map, which could legitimately call i->first and i->second.

¹⁹As usual, the Boost library offers a more complete solution, but this is simpler and fully functional.

²⁰For brevity, all "subtractive" functions have been omitted.

- ²¹In synthesis, an iterator pack is an iterator-like class that maintains synchronization between N different iterators. If P is such a pack, you can call P += 2 only if all iterators are random-access. Otherwise, the code will not compile. However, if the first component is a random-access iterator, the pack will have a constant-time difference.
- ²²std::pair takes arguments in the constructor by const reference, but if either type is a reference, it creates a reference to a reference, and that's forbidden.
- ²³Interestingly, the use of a global helper function avoids all the nasty ambiguities between a constructor and a function declaration. The problem is described and solved in "Item 6" of [7].
- ²⁴A receipt system based on *types* does not deal with *instances*. For example, you could have two algors and unlock the second with receipts from the first. This can be mitigated (at runtime!) by adding a state to the receipt. For example, you may want to store a pointer to the algor in the receipt itself and add assertions in the algor to enforce that the pointer in the receipt is indeed equal to "this".
- As a rule of thumb, if T is such that there's a single way to decide if A<B, because the comparison is trivial or fixed, then you should supply T::operator<(e.g., T = RomanNumber). Conversely, if there's more than one feasible comparison, you should not implement operator< and pass the right functional every time, to make your intentions explicit (e.g., T = Employee). These functionals may be defined inside T.
- ²⁶Alternatively, if both x<y and y<x are true, then the comparison operator is invalid.

CHAPTER 7

Code Generators

This chapter deals with templates that generate code—partly static, partly executed at runtime. Suppose you have to perform a simple comparison of powers:

```
int x = ...;
if (3^4 < x^5 < 4^7)
```

Clearly, you would like to have static constants for 3^4 and 4^7 and a corresponding runtime powering algorithm to obtain x^5 . However, a call to std:pow(x, 5) may be suboptimal, since 5 is a compile-time constant that might possibly be "embedded" in the call.

One of the goals of TMP is in fact to make the maximum information available to the compiler, so that it can take advantage of it.

7.1. Static Code Generators

Iteration can be used in a purely static context; recall the repeated squaring algorithm from Chapter 3:

```
static const size_t v0 = static_raise<X, Y/2>::value;
static const size_t value = ((Y % 2) ? X : 1U) * MXT_M_SQ(v0);
};
double data[static_raise<3, 4>::value]; // an array with 81
numbers
```

static_raise does not generate any code, only a compile-time result (namely, a numeric constant).

The same algorithm is now used to implement *static code generation*. Static recursion generates a *function* for any specified value of the exponent.

Assume that 1 is a valid scalar.

```
template <typename scalar t, size t N>
struct static pow
 static inline scalar t apply(const scalar t& x)
   return ((N % 2) ? x : 1) *
static pow<scalar t,2>::apply(static pow<scalar t,N/2>::apply(x));
};
template <typename scalar t>
struct static pow<scalar t, 2>
 static inline scalar t apply(const scalar t& x)
 { return x*x; }
};
template <typename scalar t>
struct static pow<scalar t, 1>
 static inline scalar t apply(const scalar t& x)
 { return x; }
};
template <typename scalar t>
struct static pow<scalar t, 0>
 static inline scalar t apply(const scalar t& x)
 { return 1; }
};
size t x = 3;
size t n = static pow<size t, 4>::apply(x); // yields 81
```

Here, template recursion does not produce a compile-time result, but a compile-time algorithm; in fact, static pow is a *code generator template*.

Note also that you can avoid multiplication by 1, which is implied by the ternary operator:

```
template <typename scalar_t, size_t N>
struct static_pow
{
   static inline scalar_t apply(const scalar_t& x, selector<false>)
   {
     return static_pow<2>::apply(static_pow<N/2>::apply(x));
   }
   static inline scalar_t apply(const scalar_t& x, selector<true>)
   {
     return x*apply(x, selector<false>());
   }
   static inline scalar_t apply(const scalar_t& x)
   {
     return apply(x, selector<(N % 2)>());
   }
};
```

In particular, this code generator is *strongly typed*. The user must specify the argument type in advance. This is not necessary for the algorithm to work properly. In fact, a weaker version that deduces its arguments is fine too:

```
template <size_t N>
struct static_pow
{
   template <typename scalar_t>
        static inline scalar_t apply(const scalar_t& x)
        { ... }
};

template <>
struct static_pow<2>
{
   template <typename scalar_t>
   static inline scalar_t apply(const scalar_t& x) { return x*x; }
};

// ...
```

The invocation of strongly typed templates is more verbose, since the user explicitly writes a type that could be deduced:

```
size_t x = 3;
size_t n1 = static_pow<size t, 4>::apply(x);  // verbose
```

```
size t n2 = static pow<4>::apply(x); // nicer
```

However, it sometimes pays to be explicit. A cast on the argument is quite different from a cast of the result, because the code generator will produce an entirely new function:

```
double x1 = static_pow<double, 4>::apply(10000000); // correct
double x2 = static_pow<4>::apply(10000000); // wrong (it
overflows)

double x3 = static_pow<4>::apply(10000000.0); // correct again
```

Usually it's possible to code both a strong and a weak code generator at the same time by borrowing a trick from groups. You move the weak generator into a partial specialization, which is recalled by the general template.

```
struct deduce
};
template <size t N, typename scalar t = deduce>
struct static pow;
template <>
struct static pow<2, deduce>
  template <typename scalar t>
  static inline scalar t apply(const scalar t& x)
  { . . . }
};
template <size t N>
struct static_pow<N, deduce>
  template <typename scalar t>
  static inline scalar t apply(const scalar t& x)
  { . . . }
};
// primary template comes last
template <size_t N, typename scalar t>
struct static pow
  static inline scalar t apply(const scalar t& x)
     return static pow<n>::apply(x);
};
```

A strict argument check is actually performed only by the primary template, which immediately calls the deduce specialization. The order of declarations matters: static_pow<N, deduce> will likely use static_pow<2, deduce>, so the latter must precede the former in the source file.

7.2. Double checked Stop

Compile-time recursion is usually obtained by having a template call "itself" with a different set of template parameters. Actually, there's no recursion at all, since a change in template parameters generates a different entity. What you get is static "loop unrolling".

The advantage of static recursion is that explicitly unrolled code is easier to optimize.

The following snippets perform a vector-sum of two arrays of known length:

```
template <size_t N, typename T>
void vector_sum_LOOP(T* a, const T* b, const T* c)
{
   for (int i=0; i<N; ++i)
      a[i] = b[i] + c[i];
}

template <size_t N, typename T>
void vector_sum_EXPLICIT(T* a, const T* b, const T* c)
{
   a[0] = b[0] + c[0];
   a[1] = b[1] + c[1];
   // ...
   // assume that it's possible to generate exactly N of these lines
   // ...
   a[N-1] = b[N-1] + c[N-1];
}
```

The explicitly unrolled version will be faster for small N, because modern processors can execute a few arithmetic/floating point operations in parallel. Even without specific optimizations from the compiler, the processor will perform the sums, say, four at a time.¹

However, for large N, the code would exceed the size of the processor cache, so the first version will be faster from some point on.

The ideal solution in fact is a mixture of both:

```
static const int THRESHOLD = /* platform-dependent */;
template <size_t N, typename T>
void vector_sum(T* a, const T* b, const T* c)
{
  if (N>THRESHOLD)
```

```
{
     int i=0;
     for (; (i+4) < N; i+=4)
                                         // the constant 4 and...
       a[i+0] = b[i+0] + c[i+0];
                                        //
       a[i+1] = b[i+1] + c[i+1];
                                        // ...the number of lines in
this block
       a[i+2] = b[i+2] + c[i+2];
                                       // are platform-dependent
       a[i+3] = b[i+3] + c[i+3];
                                        //
     }
     for (; i<N; ++i)
                                         // residual loop
       a[i] = b[i] + c[i];
     }
  else
  {
     vector sum EXPLICIT<N>(a, b, c);
  }
}
```

This implementation has a problem anyway. Suppose THRESHOLD is 1000. When the compiler instantiates, say, vector_sum<1000, double>, it wastes time generating 1,000 lines that will never be called:

```
if (true)
{
    // ...
}
else
{
    a[0] = b[0] + c[0];
    a[1] = b[1] + c[1];
    // ...
    a[999] = b[999] + c[999];
}
```

To fix this issue, you add a *double check*:

```
else
{
   vector_sum_EXPLICIT<(N>THRESHOLD ? 1 : N)>(a, b, c);
}
```

The double check is not simply an optimization. Static recursion can yield an *unlimited* number of lines, Assume again you have an array of length N and need to fill it with consecutive integers. You hope to be able to write a function template integrize whose call produces native machine code

```
that is logically equivalent to:
{
   data[0] = 0;
   data[1] = 1;
   // ...
   data[N-1] = N-1;
}
   But you guess that when N is very large, due to the effect of processor caches, the unrolled loop
will generate a huge amount of bytes, whose mass will eventually slow down the execution.<sup>2</sup>
   So you use integrize to select a compile-time strategy or a runtime strategy:
template<typename T, int N>
void integrize(T (&data)[N])
   if (N<STATIC LOWER BOUND)
      integrize helper<N>(data);
   else
      for (size t i=0; i<N; ++i)
         data[i] = i;
}
   First, start with an incorrect function:
template <int N, typename T>
void integrize helper(T* const data)
{
   data[N-1] = N-1;
   integrize helper<N-1>(data);
}
   The recursion has no limit, so it will never compile successfully.
   You might be tempted to make the following improvement:
template <int N, typename T>
void integrize helper (T* const data)
   data[N-1] = N-1;
   if (N>1)
```

This version still doesn't work, since the compiler will produce a sequence of calls with unlimited depth. From some point on, the condition if (N>1) is always false, but it doesn't matter—such code would be pruned by the optimizer, but the compiler will complain and stop much earlier!

```
data[2-1] = 2-1; // here N=2 if (true) // 2>1?
```

integrize helper<N-1>(data);

}

In other words, the compiler sees that integrize_helper<1> depends on integrize helper<0>, hence the unlimited recursion (at compile time).

The *double checked stop* idiom is again the solution:

```
template <int N, typename T>
void integrize_helper(T* const data)
{
   data[N-1] = N-1;
   if (N>1)
     integrize_helper<(N>1) ? N-1 : 1>(data);
}
```

Note the extra parentheses around N>1 (otherwise, the > between N and 1 will be parsed as the angle bracket that closes the template).

Thanks to the double check, the compiler will expand code like this:

The expansion is finite, since integrize_helper<1> mentions only itself (which is a well-defined entity, not a new one) and the recursion stops. Of course, integrize_helper<1> will never call itself at runtime. The optimizer will streamline the if (true) branches and remove the last if (false).

In general, the double checked stop idiom prescribes to stop a recursion, mentioning a template that has been already instantiated (instead of a new one) and preventing its execution at the same time.

Finally, you again apply the idiom as an optimization against code bloat:

```
template<typename T, int N>
void integrize(T (&data)[N])
{
   if (N<STATIC_LOWER_BOUND)</pre>
```

```
integrize_helper<(n<static_lower_bound) ? N : 1>(data);
else
  for (size_t i=0; i<N; ++i)
    data[i] = i;
}</pre>
```

7.3. Static and Dynamic Hashing

Sometimes it's possible to share an algorithm between a static and runtime implementation via kernel macros. The following example shows how to hash a string statically.

Assume as usual that a hash is an integer stored in a size_t and that you have a macro. Taking x, the old hash and a new character called c, here are some possibilities:

```
#define MXT_HASH(x, c) ((x) << 1) ^ (c) #define MXT_HASH(x, c) (x) + ((x) << 5) + (c) #define MXT_HASH(x, c) ((x) << 6) ^ ((x) & ((~size_t(0)) << 26)) ^ (c)
```

Note The hashing macros require that c be a positive number. You could replace c with (c−CHAR_MIN), but this would make the hash platform-dependent. Where char is signed, 'a'-CHAR_MIN equals 97-(-128) = 225 and where char is unsigned, the same expression yields 97-0 = 97.

Furthermore, the same text in a std::string and in a std::wstring should not return two different hash codes.

Given that you disregard what happens for non-ASCII characters, an elegant workaround is to cast charc to unsigned char.

Constants should not be hard-coded, but rather generated at compile time.

You could replace the classic code

```
const char* text = ...;
if (strcmp(text, "FIRST") == 0)
{
    // ...
} else if (strcmp(text, "SECOND") == 0)
{
    // ...
} else if (strcmp(text, "THIRD") == 0)
{
```

// ...

- Hashing will save a lot of string comparisons, even if it could produce false positives.³
- If static_hash produces duplicate values, the switch won't compile, so it will never produce false negatives (that is, the words "FIRST", "SECOND", and so on will always be matched without ambiguities).

The static algorithm uses template rotation and a very neat implementation:

```
template
<
  char C0=0, char C1=0, char C2=0, char C3=0, ..., char C23=0,
  size_t HASH = 0
>
struct static_hash
: static_hash<C1,C2...,C23,0, MXT_HASH(HASH, static_cast<unsigned char>(C0))>
{
};

template <size_t HASH>
struct static_hash<0,0,0,0,...,0, HASH>
: static_value<size_t, HASH>
{
};
```

The only degree of freedom in dynamic_hash is the function signature. Here's a fairly general one, with some plain old vanilla C tricks:

```
{
 size t h = 0;
  const char* const end1 = separ ? text+strcspn(text, separ)
  const char* const end2 = (end && end<end1) ? end : end1;</pre>
  while (end2 ? text<end2 : (*text != 0))</pre>
     const size t c = static cast<unsigned char>(*(text++));
    h = MXT HASH(h, c);
  return std::make pair(h, text);
}
int main()
  const char* text = "hello, dynamic hash";
                                      // hash all string, up to
  dynamic hash(text);
char(0)
  dynamic hash(text, ";,");
                                // hash up to any of the
separators
  dynamic hash(text, ";,", text+10); // up to separator, at most
10 chars
}
```

I chose to return a composite result, the hash value and the updated "iterator".

7.3.1. A Function Set for Characters

The selection of the correct function set can be done either by a deduced template parameter (as seen in string traits in Section 4.2.1) or by an environment template parameter.

A natural example is the problem of a character set: some string-conversion functions can be accelerated, given that some set of characters, say $\{ '0', '1' \dots '9' \}$, is contiguous. If c belongs to the set, you can convert c to integer via a simple subtraction c - '0', but if the digit character set is arbitrarily scattered, a more complex implementation is needed.

You scan sets of characters with template rotation:

```
namespace charset {
template
<
   typename char_t,
   char_t C0,
   char_t C1 = 0,
   char_t C2 = 0,
   // ...
   char_t C9 = 0</pre>
```

```
struct is contiquous
  static const bool value = (C0+1==C1) &&
     is contiguous<char t, C1, C2, C3, C4, C5, C6, C7, C8, C9>::value;
};
template <char C0>
struct is contiguous<char, C0>
  static const bool value = true;
};
template <wchar t C0>
struct is contiguous < wchar t, C0>
  static const bool value = true;
};
  Next, the result of a static test can be saved in a global traits structure:
struct ascii
  static const bool value lowerc =
   charset::is contiguous<char,</pre>
     'a','b','c','d','e','f','g','h','i','j'>::value
   & &
   charset::is contiquous<char,
     'j','k','l','m','n','o','p','q','r','s'>::value
   & &
   charset::is contiguous<char,
     's','t','u','v','w','x','y','z'>::value;
  static const bool value upperc =
   charset::is contiguous<char,
     'A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J'>::value
   & &
   charset::is contiquous<char,
     'J','K','L','M','N','O','P','Q','R','S'>::value
   & &
   charset::is contiquous<char,
     'S', 'T', 'U', 'V', 'W', 'X', 'Y', 'Z'>::value;
  static const bool value 09 =
   charset::is contiquous<char,
     '0','1','2','3','4','5','6','7','8','9'>::value;
```

>

```
static const bool value = value_09 && value_lowerc &&
value_upperc;
};
```

Suppose for the moment that ascii::value is true. You can write a function set to deal with the special case:

```
template <typename T, T lower, T upper>
inline bool is between(const T c)
  return !(c<lower) && !(upper<c);
struct ascii traits
  typedef char char type;
  static inline bool isupper(const char type c)
     return is between<char, 'A', 'Z'>(c);
  static inline bool islower (const char type c)
     return is between<char, 'a', 'z'>(c);
  static inline bool isalpha(const char type c)
     return islower(c) || isupper(c);
  static inline bool isdigit (const char type c)
     return is between<char,'0','9'>(c);
  }
  //...
  static inline char tolower (const char c)
     return isupper(c) ? c-'A'+'a' : c;
  static inline char toupper(const char c)
     return islower(c) ? c-'a'+'A' : c;
};
```

```
In a different implementation, you use std::locale:
template <typename char t>
struct stdchar traits
  typedef char t char type;
  static inline bool isupper(const char t c)
     return std::isupper(c, locale());
  }
  static inline bool islower(const char t c)
  {
     return std::islower(c, locale());
  static inline bool isalpha(const char t c)
     return std::isalpha(c, locale());
  }
  static inline bool isdigit(const char t c)
     return std::isdigit(c, locale());
  }
  static inline char t tolower(const char t c)
     return std::tolower(c, std::locale());
  static inline char t toupper(const char t c)
     return std::toupper(c, std::locale());
};
  And eventually combine these types:
struct standard { };
struct fast {};
template <typename char t, typename charset t = fast>
struct char traits : stdchar traits<char t>
};
```

```
template <>
struct char_traits<char, fast>
: typeif<ascii::value, ascii_traits, stdchar_traits<char> >::type
{
};
```

The environment parameter charset_t is by default set to fast. If it's possible in the current platform, the fast set is preferred; otherwise, the standard set is used.⁴

7.3.2. Changing Case

This section lists some utilities used to change the case of characters. First, it introduces some tags. Note that "case_sensitive" is treated as a "no conversion" label.⁵

```
struct case_sensitive {};
struct upper_case {};
struct lower case {};
```

This example exploits the fact that char_traits offers a leveraged interface to mutate characters at runtime (the example is limited to char). The classic part of the work is a collection of functors.

```
template <typename mutation_t, typename traits_t
= char_traits<char> >
struct change_case;

template <typename traits_t>
struct change_case<case_sensitive, traits_t>
{
   typedef typename traits_t::char_type char_type;
   char_type operator()(const char_type c) const
   {
      return c;
   }
};

template <typename traits_t>
struct change_case<lower_case, traits_t>
{
   typedef typename traits_t::char_type char_type;
   char_type operator()(const char_type char_type;
   char_type operator()(const char_type c) const
   {
      return traits_t::tolower(c);
   }
}
```

```
template <typename traits t>
struct change case<upper case, traits t>
  typedef typename traits t::char type char type;
  char type operator()(const char type c) const
     return traits t::toupper(c);
};
int main()
   std::string s = "this is a lower case string";
   std::transform(s.begin(), s.end(), s.begin(),
change case<upper case>());
  Now you move to the analogous conversion at compile time.
template <typename case t, char C, bool FAST = ascii::value>
struct static change case;
  FAST is a hidden parameter; regardless of its value, a case-sensitive conversion should do
nothing:
template <char C, bool FAST>
struct static change case < case sensitive, C, FAST>
   static const char value = C;
};
  If FAST is true, the transformation is trivial. If FAST is false, unfortunately, every character that
can change case needs its own specialization. Macros will save a lot of typing here.
template <char C>
struct static change case<lower case, C, true>
  static const char value = ((C>='A' \&\& C<='Z') ? C-'A'+'a' : C);
};
template <char C>
struct static change case < upper case, C, true>
  static const char value = ((C)='a' \&\& C<='z') ? C-'a'+'A' : C);
};
```

};

```
template <char C>
struct static change case<lower case, C, false>
  static const char value = C; // a generic char has no case
};
template <char C>
struct static change case < upper case, C, false >
  static const char value = C; // a generic char has no case
};
#define mxt STATIC CASE GENERIC(C LO,
C UP)
template <> struct static change case<lower case, C UP, false>
{ static const char value = C LO; };
template <> struct static change case<upper case, C LO, false>
{ static const char value = C UP; }
mxt STATIC CASE GENERIC('a', 'A');
mxt STATIC CASE GENERIC('b', 'B');
mxt STATIC CASE GENERIC('z', 'Z');
#undef mxt STATIC CASE GENERIC
  This has an immediate application to both static hash and dynamic hash.
  As usual, the macro is merely for convenience. Note that a non-deduced template parameter is
introduced in the dynamic hash.
#define mxt FIRST CHAR(c)
     static cast<unsigned char>(static_change case<case t, C>::value)
template
 typename case t,
 char C0=0, char C1=0, char C2=0, char C3=0, ..., char C23=0,
 size t HASH = 0
struct static hash
: static hash<case t,C1,C2,...,C23,0, MXT HASH(HASH,
```

mxt FIRST CHAR(C0))>

```
} ;
template <typename case t, size t HASH>
struct static hash<case t,0,0,0,0,...,0, HASH>
: static value<size t, HASH>
};
template <typename case t>
inline ... dynamic hash(const char* text, ...)
 const change case<case t> CHANGE;
 size t h = 0;
 const char* const end1 = (separ ? text+strcspn(text, separ)
: end);
 const char* const end2 = (end && end<end1) ? end : end1;</pre>
 while (end2 ? text<end2 : (*text != 0))
   const size t c = static cast<unsigned char>(CHANGE(*(text++)));
   h = MXT HASH(h, c);
 return std::make pair(h, text);
```

Such a modified algorithm will alter the case of a string *inside the computation* of the hash value, so an "upper case hash" is effectively a case-insensitive value:

```
switch (dynamic_hash<upper_case>(text).first)
{
  case static_hash<'F','I','R','S','T'>::value:
    // will match "First", "FIRST", "first", "fiRST"...
    break;
}
```

7.3.3. Mimesis Techniques

This section rewrites the dynamic_hash using mimes techniques. In the new prototype, end is not optional, so you have to provide more overloads to get a flexible syntax. As for the original C version:

```
template <typename case_t, typename iterator_t, typename end_t>
std::pair<size_t, iterator_t>
   dynamic_hash(iterator_t begin, const end_t end, size_t h = 0)
{
   typedef typename std::iterator_traits<iterator_t>::value_type
   char_t;
```

```
const change case< case t, char traits<char t> > CHANGE;
 while (end != begin)
   const size t c = static cast<unsigned char>(CHANGE(*
(begin++)));
   h = MXT HASH(h, c);
 return std::make_pair(h, begin);
template <typename case t, typename iterator t>
inline std::pair<size t, iterator t>
 dynamic hash(iterator t begin, size t h = 0)
 return dynamic hash (begin, c string end<iterator t>(), h);
  You can plug in some useful mimesis-like objects<sup>6</sup>:
template <typename char t, char t CLOSE TAG>
struct stop at
 template <typename iterator t>
 inline bool operator!=(const iterator t i) const
   return (*i != 0) && (*i != CLOSE TAG);
};
size t h = dynamic hash<case insensitive>(text, stop at<char, ';'>
()).first;
template <bool (*funct)(const char), bool NEGATE>
struct apply f
 template <typename iterator t>
 inline bool operator!=(const iterator t i) const
   return funct(*i) ^ NEGATE;
};
typedef apply f<char traits<char>::isspace, true> end of word;
typedef apply f<char traits<char>::isalpha, false> all alpha;
  end_of_word stops at the first space, and all_alpha stops at the first non-alphabetical
```

character.

7.3.4. Ambiguous Overloads

The evolution of the dynamic_hash has led to adding more template parameters and more overloads. You need to be careful not to cause compilation problems because of *ambiguous overload* resolution.

The exact overload resolution rules are described in Appendix B of [2], but a rough summary is described here.

When the compiler meets a function call, it must pick, from the set of all functions with the same name, the most specialized set that matches the given arguments. It must emit an error if no such function exists or if the best match is ambiguous.

If you have several function templates named F, you denote them as F[1], F[2], and so on.⁷ You say that F[1] is more specialized than F[2] if F[2] can be used wherever F[1] is used, with an exact argument match, but not vice versa.

For example:

```
template <typename T1, typename T2>
void F(T1 a, T2 b); // this is F[1]

template <typename T>
void F(T a, T b); // this is F[2]

template <typename T>
void F(T a, int b); // this is F[3]
```

The second template, F[2], is more specialized than F[1], because the call F(X, X) can refer to either one, but only F[2] matches F(X, Y) exactly. Similarly, F[3] is more specialized than F[1].

However, this is a partial ordering criterion. If no function is more specialized than the other(s), the compiler will abort, reporting an ambiguous overload. In fact, in the previous example, F[2] and F[3] are not comparable. F[3] will not match exactly F(X, X) and F[2] will not match exactly F(X, int).

```
int z = 2;

F(z, z); // error: could be F[2] with T=int or F[3] with T=int
```

Informally, an easy unambiguous special case is total replacement. If a template parameter is completely replaced by fixed types or previous template parameters, the resulting function is more specialized than the original. Take F[1], replace every occurrence of T2 with T1, and obtain F[2]; replace T2 with int and obtain F[3].

A library writer usually provides a set of overloads, where one or more elements are function templates. One of the problems, often underestimated or ignored, is to decide in advance if the set is *well-ordered*. A well-ordered set will never generate ambiguity errors.

The combination of default arguments and templates often makes deduction very hard.

```
template <typename case t, typename iterator t, typename end t>
```

```
[...] dynamic hash(iterator t begin, const end t end,
                size t crc = \overline{0}); // dynamic has\overline{h}[1]
template <typename case t, typename iterator t>
[...] dynamic hash(iterator t begin, size t crc = 0); //
dynamic hash[2]
  To determine if this set is well-ordered, you need only to consider the case of a call with two
arguments, and it's evident that the total replacement condition holds (replace end t with
size t).
  However, note that dynamic hash (T, int) will invoke dynamic hash [1]:
dynamic hash(text, 123); // invokes (1) [with end t = int]
  A user-friendly library will try to avoid ambiguities, first by using additional types:
struct hash type
  size t value;
  hash type() : value(0) {}
  explicit hash type(const size t c) : value(c) {}
};
template <typename case t, typename iterator t, typename end t>
[...] dynamic hash(iterator t begin, end t end, hash type h =
hash type());
template <typename case t, typename iterator t>
[...] dynamic hash(iterator t begin, hash type h = hash type());
  While this does not change the way the compiler picks functions, it will make the error more
evident to the user, because now dynamic hash (text, 123) will not even compile.
dynamic hash(text, hash type(123)); // this instead is
correct
  A radical change instead is obtained by wrapping the original return type in a typename
only if < [[condition]], ...>::type clause (See Section 4.3.3).
template <typename T1, typename T2>
struct different : selector<true>
{ };
template <typename T>
struct different<T, T> : selector<false>
{ };
template <typename case t, typename iterator t, typename end t>
typename only if<different<end t, hash type>::value, [...]>::type
```

```
dynamic_hash(iterator_t begin, const end_t end, hash_type h =
hash_type());

Suppose that you add the C version back in (denoted as dynamic_hash[3]):

template <typename case_t>
[...] dynamic_hash(const char* text, const char* const separator
= 0, const char* const end = 0, size_t h = 0)
```

This function, as is, can generate an ambiguous call. dynamic_hash (const char*) matches either dynamic_hash[2] (with iterator_t = const char*) or dynamic_hash[3]. The error depends on both functions being templates. Because case_t: had dynamic_hash[3] was a classic function, it would have been picked with higher priority.

To avoid the problem, remove the default arguments to separator and end.

7.3.5. Algorithm I/O

You can let dynamic_hash return a pair that contains the updated iterator position and the hash value.

Often the user will need to store the result just to split it:

```
std:pair<size_t, const char*> p = dynamic_hash(text);
text = p.second;
switch (p.first)
{
    //...
}
```

This can be verbose, especially if the iterator has a long type. 9
C++11 gave a new meaning to the keyword auto exactly for this purpose:

```
auto p = dynamic hash(text);
```

But observe that auto cannot refer to a part of an object. The following line is illegal:

```
std::pair<auto, const char*> p = dynamic_hash(text);
```

You could take an iterator by reference and update it, but this is not a fair solution, as it forces the caller to duplicate the iterator if you want to save the original value.

Instead, you modify the return type. It will be an object conceptually similar to a pair, with the option to overwrite a reference with the result:

```
template <typename iterator_t>
struct dynamic_hash_result
{
   size_t value;
   iterator_t end;
```

```
dynamic_hash_result(const size_t v, const iterator_t i)
    : value(v), end(i)
{
    dynamic_hash_result& operator>>(iterator_t& i)
    {
        i = end;
        return *this;
    }
};
```

You change the return statement in the dynamic_hash functions accordingly (namely, replace std::make pair(...) with dynamic hash result(...)).

The final function call is indeed compact. It updates text and returns the hash at the same time. Additionally, the .value suffix reminds you of static_hash<>::value. Of course, more variations are possible. 10

```
switch ((dynamic_hash(text) >> text).value)
{
  case static_hash<'a','b','c'>::value:
    //...
}
```

7.3.6. Mimesis Interface

template <typename iterator t>

Mimesis objects are lightweight and conceptually similar to functors, but their expressivity is close to a scalar. Since they are indeed instantiated, let's investigate the possibility of combining them with operators:

inline bool operator!=(const iterator t i) const

```
{
     return true this() != i;
};
// note the CRTP
template <bool (*funct)(const char), bool NEGATE>
struct apply f : public hash end type< apply f<funct, NEGATE> >
  template <typename iterator t>
  inline bool operator!=(const iterator t i) const
     return funct(*i) ^ NEGATE;
};
// note again the CRTP
template <typename char t, char t CLOSE TAG>
struct stop at : public hash end type< stop at<char t, CLOSE TAG>
>
  template <typename iterator t>
  inline bool operator!=(const iterator t i) const
     return (*i != CLOSE TAG);
};
  Having all objects inherit the same interface, you can define "combo type" and logic operators:
struct logic AND {};
struct logic OR {};
template <typename T1, typename T2, typename LOGICAL OP>
class hash end type combo
: public hash end type< hash end type combo<T1, T2, LOGICAL OP> >
  T1 t1 ;
  T2 t2 ;
public:
  hash end type combo(const T1& t1, const T2& t2)
  : t1_(t1), t2 (t2)
  template <typename iterator t>
  inline bool operator!=(const iterator t i) const
```

```
{
     return combine(i, LOGICAL OP());
private:
  template <typename iterator t>
     bool combine (const iterator t i, logic AND) const
     return (t1 != i) && (t2 != i);
  template <typename iterator t>
  bool combine (const iterator t i, logic OR) const
     return (t1 != i) || (t2 != i);
};
template <typename K1, typename K2>
inline hash end type combo<K1, K2, logic AND>
 operator&& (const hash end type<K1>& k1, const hash end type<K2>&
k2)
{
 return hash end type combo<K1, K2, logic AND>(k1.true this(),
k2.true this());
template <typename K1, typename K2>
inline hash_end_type combo<K1, K2, logic OR>
 operator | | (const hash end type<K1>& k1, const hash end type<K2>&
k2)
{
 return hash end type combo<K1, K2, logic OR>(k1.true this(),
k2.true this());
  Note the counterintuitive use of the operation tag. You may be tempted to replace logic AND
```

with an "active tag," such as std::logical and<bool>, drop combine entirely, and just use the tag as a function call to produce the result:

```
template <typename iterator t>
inline bool operator!=(const iterator t i) const
  return LOGICAL_OP()(t1_ != i, t2 != i);
}
```

This is *incorrect*, as it would blow short-circuit (when you express, say, A && B as F (A, B), all the arguments must be evaluated before calling F).

Note also that the check for null char is removed in stop_at. It now has to be added explicitly, but it's performed only once.

This syntax is an example of a *lambda expression*, which is the main topic of Section 9.2.

7.4. Nth Minimum

{

template <typename scalar t, size t N>

This section gives a step-by-step example of a simple recursive compile-time function that involves a data structure.

You write a container called $nth_min<T$, N>. An instance of this container receives values of type T, one at a time, 12 via an insert member function, and it can be asked for the smallest N elements met so far.

For a reason to be discussed later, let's impose the extra requirement that the container should not allocate its workspace from dynamic memory.

```
class nth_min
{
    scalar_t data_[N];

public:
    void insert(const scalar_t& x)
    {
        update(data_, x);
    }

    const scalar_t& operator[](const size_t i) const
    {
        return data_[i];
    }
};

The following paragraphs produce a suitable update function. 13

template <typename scalar_t, int N>
inline void update(scalar t (&data)[N], const scalar t& x)
```

// now N is known, start iterations here

First, you need to visualize the algorithm in recursive form. Assume as the induction hypothesis that data_contains the N smallest values met so far, in ascending order.

```
if (x ≥ data_[N-1])
    // x is not in the N minima
    discard x and return;

Else

// here x < data_[N-1], so

// data_[N-1] will be replaced either by x or by data_[N-2]

if (x ≥ data_[N-2])
    data_[N-1] = x and return;

Else

data_[N-1] = data_[N-2];

if (x ≥ data_[N-3])
    data_[N-2] = x and return;

Else

data_[N-2] = data_[N-3];

...</pre>
```

```
15
                       17
                                       24
                                               31
                                                      35
                                                                             29
data
                                                      N-1
               0
                       1
                                       N-3
                                              N-2
                                                                      overwrite element N-1
               15
                       17
                                       24
                                              31
                                                      31
data
                                                                      with N-2
                                       N-3
                                              N-2
                                                      N-1
                                       24
                                                      31
                                                                      overwrite element N-2
               15
                       17
                                               29
data
                                                                      with x
                                      N-3
                                              N-2
                                                      N-1
                       1
```

Now observe that "discard x" is equivalent to "write x in the non-existent position N". You factor out the write operation using a custom selector:

```
struct nth
{
};

template <typename scalar_t, int N, int SIZE>
void write(scalar_t (&data)[SIZE], const scalar_t& x, nth<N>)
```

template <int N>

```
{
  data[N] = x;
}

template <typename scalar_t, int size>
void write(scalar_t (&data)[size], const scalar_t& x, nth<size>)
{
}
```

The second overload uses the dimension of the array. So write (data, x, nth<I>()) actually means "write x in the Ith position of array data, if possible; otherwise, do nothing".

This small abstraction permits you to extend the same recursive pattern to the whole algorithm:

```
if (x \ge data [N-1])
       // x is not in the N minima
      data [N] = x and return;
else
       if (x \ge data [N-2])
             data [N-1] = x and return;
       else
template <typename scalar t, int N, int SIZE>
void iterate(scalar t (&data)[SIZE], const scalar t& x, nth<N>)
  if (x < data[N])
     data[N] = data[N-1];
     iterate(data, x, nth<N-1>());
  else
     write(data, x, nth<N+1>()); // write x at position N+1
  }
}
```

Next, you have to write an iteration terminator, and you can begin identifying values of template parameters that make the rest of the code meaningless. When N==0, data [N-1] is for sure not well-formed, so you specialize/overload the case where N is 0. In fact, if you have to track down only the smallest element of the sequence, there's no shift involved:

```
template <typename scalar_t, int SIZE>
void iterate(scalar_t (&data)[SIZE], const scalar_t& x, nth<0>)
{
    // here N=0, after this point, stop iterations
    // if x is less than minimum, keep x, else discard it
    if (x < data[0])</pre>
```

```
data[0] = x;
else
    write(data, x, nth<1>());
}
```

The else branch cannot be omitted, but if SIZE is 1, the optimizing compiler will wipe it out. Finally, the recursion starts backwards on the last element of the array, so you pass N-1:

```
template <typename scalar_t, int N>
void update(scalar_t (&data)[N], const scalar_t& x)
{
  iterate(data, x, nth<N-1>());
}
```

What's not elegant in this implementation is that iterate<0> contains duplicated code from iterate<N>. The most elegant solution would end with an empty function.

Another generalization is needed. All write operations involve either a shift data [K] = data[K-1] or the insertion data [K] = x, respecting array bounds. Can a single function template represent both?

Yes, if you are able to identify x with an element of data and specify only the index of the element to pick:

```
template <typename scalar_t, int N, int SIZE, int J>
void write(scalar_t (&data)[SIZE], const scalar_t& x, nth<N>,
nth<J>)
{
   data[N] = data[J];
}

template <typename scalar_t, int SIZE, int J>
void write(scalar_t (&data)[SIZE], const scalar_t& x, nth<SIZE>,
nth<J>)
{
}
```

If you compare the instructions data[K] = data[K-1] and data[0] = x from the implementation, you see that x is naturally identified with data[-1].

So you add two more specializations:

template <typename scalar t, int N, int SIZE>

```
void write(scalar_t (&data)[SIZE], const scalar_t& x, nth<N>,
nth<-1>)
{
   data[N] = x;
}
template <typename scalar_t, int SIZE>
void write(scalar_t (&data)[SIZE], const scalar_t& x, nth<SIZE>,
```

```
nth<-1>)
{
}
```

To sum up, write (data, x, N, J) is a complicated way to say data [N] = data[J]; N and J are selectors, not integers. As usual, the function deduces the length of the array, so out-of-bounds accesses become no-ops.

When N=0 in the code, write translates to data [0] = x, as required, and iteration -1 is empty.

Note that you pay the price of generality in line 1, which is rather unclear at first sight, since you have to explicitly use nth<-1> to access x.

If N is large, the fastest algorithm would possibly store objects in a large chunk of memory and sort them when necessary, doing all the work at runtime. In the worst case, if K is the number of items inserted, execution time is proportional to K.N for the static version, but for small values of N and simple POD types (that is, when operator< and assignment do not have significant overhead), the static version will usually perform faster, due to its compactness and absence of hidden constants.¹⁴

Finally, you can replace the write function call, whose hidden meaning is an assignment, with a real assignment. Just use a proxy:

```
struct null_reference
{
   template <typename scalar_t>
   null_reference& operator= (const scalar_t&)
   {
     return *this;
   }
};
```

```
template <int K>
struct nth
  template <typename scalar t, int SIZE>
  static scalar_t& element(scalar_t (&data)[SIZE], const scalar t&
X)
     return data[K];
  template <typename scalar t>
  static null reference element(scalar t (&data)[K], const
scalar t& x)
  {
     return null reference();
};
template <>
struct nth<0>
  template <typename scalar t, int SIZE>
  static scalar t& element(scalar t (&data)[SIZE], const scalar t&
\times)
     return data[0];
} ;
template <>
struct nth<-1>
  template <typename scalar t, int SIZE>
  static const scalar t& element(scalar t (&data)[SIZE], const
scalar t& x)
     return x;
};
struct nth min
  template <typename scalar t, int SIZE>
  static void update(scalar t (&data)[SIZE], const scalar t& x)
     iterate(data, x, nth<SIZE-1>());
```

```
private:
  template <typename scalar t, int N, int SIZE>
  static void iterate(scalar t (&data)[SIZE], const scalar t& x,
nth<N>)
  {
     if (x < data[N])
       nth<N>::element(data, x) = nth<N-1>::element(data, x);
       iterate(data, x, nth<N-1>());
     }
     else
       nth<N+1>::element(data, x) = nth<-1>::element(data, x);
     }
  }
  template <typename scalar t, int SIZE>
  static void iterate(scalar t (&data)[SIZE], const scalar t& x,
nth < -1 > )
};
```

7.5. The Template Factory Pattern

Templates are good at making compile-time decisions, but all programs need to take runtime decisions.

The *factory pattern* solves the runtime decision problem via polymorphism. An isolated function, called the *factory*, embeds all the logic and returns a pointer to a dynamically-created object, which drives the program flow with its virtual member function calls:

```
class abstract_task
{
   public:
      virtual void do_it() = 0;

      virtual ~abstract_task()
      {
       }
};
class first_task : public abstract_task
{
   public:
      first_task(/* parameters */)
```

```
{
        // ...
     virtual void do it()
      {
};
enum task type
  FIRST TASK, SECOND TASK, THIRD TASK
};
abstract task* factory(task type t)
   switch (t)
   {
     case FIRST_TASK: return new first_task(...);
case SECOND_TASK: return new second_task(...);
     case THIRD_TASK: return new third_task(...);
     default:
                           return 0;
int main()
{
  task_type t = ask_user();
  abstract task* a = factory(t);
  a->do it();
  delete a;
  return 0;
```

Note that the only switch...case construct, that is, the link between the user choice and the program flow, is hidden inside the factory.

As expected, templates have no exact equivalent, but the following pattern is definitely similar:

```
template <typename TASK_T>
void do_the_work(TASK_T task)
{
   task.loadParameters(...);
   task.run();
   task.writeResult(...);
}
```

```
enum task_type
{
   FIRST_TASK, SECOND_TASK, THIRD_TASK
};

void factory(task_type t)
{
   first_task t1;
   second_task t2;
   third_task t3;

   switch (t)
   {
      case FIRST_TASK: do_the_work(t1); break;
      case SECOND_TASK: do_the_work(t2); break;
      case THIRD_TASK: do_the_work(t3); break;
      default: throw some_exception();
   }
}
```

The function do_the_work is an example of *static polymorphism*. The usage of an object determines its interface and vice versa. Every static type for which the syntax is valid is automatically usable.

This approach offers the advantage of a unified workflow. There's a single function to debug and maintain, Obviously, having three overloads of do the work would minimize this benefit.

Here's another example—a function that takes an array and computes either the sum or the product of all elements.

```
enum compute_type { SUM, MULTIPLY };

double do_the_work(compute_type t, const double* data, size_t
length)
{
   switch (t)
   {
    case SUM:
      return std::accumulate(data,data+length,0.0);

   case MULTIPLY:
      return
std::accumulate(data,data+length,1.0,std::multiplies<double>());

   default:
      throw some_exception();
   }
}
```

You want to rework the code so that it takes numbers from a given text file and performs the

requested operation on all elements, and all computations should be performed with a user-supplied precision.

This requires a *multi-layer template factory*. Roughly speaking, you have N function templates. The Kth function has N-K arguments and K template parameters and it uses a switch block to branch execution to one of the possible (K+1) th functions.

```
enum result type { SUM, MULTIPLY };
enum data type { FLOAT, DOUBLE };
template <typename T>
T factory LAYER3 (result type t, const std::vector<T>& data)
 switch (t)
 case SUM:
   return std::accumulate(data.begin(),data.end(),T(0));
 case MULTIPLY:
   return
std::accumulate(data.begin(),data.end(),T(1),std::multiplies<T>
());
 default:
   throw some exception();
 }
}
template <typename T>
T factory LAYER2(result type t, std::istream& i)
  std::vector<T> data;
  std::copy(std::istream iterator<T>(i), std::istream iterator<T>
(),
           std::back inserter(data));
  return factory LAYER3(t, data);
}
double ML factory(result_type t, data_type d, const char*
filename)
{
  std::ifstream i(filename);
  switch (d)
     case FLOAT:
       return factory LAYER2<float>(t, i);
     case DOUBLE:
       return factory LAYER2<double>(t, i);
```

```
default:
     throw some_exception();
}
```

The hardest design problem in template factories is usually *the type of the result*. Here the code silently exploits the fact that all functions return a result convertible to double.

7.6. Automatic Enumeration of Types

It's possible to exploit the __LINE__ macro to create an easily extensible collection of types that can be accessed as an enumeration.

Consider the following prototype—you can trivially map an integer index into a selector:

```
template <int N>
struct single_value : selector<false>
};
template <>
struct single value<7>: selector<true> // user supplied
case #1
};
template <>
struct single value<13> : selector<true> // user supplied
case #2
};
// ...
template <>
struct single value<128>
                              // terminator, equivalent to max
size;
                               // it will be useful shortly
};
```

With greater generality we can write:

In fact, single_value is a metafunction that maps a range of integers, say [0...127] for simplicity, on types, which always returns selector<false>, except $7 \rightarrow \text{MyType1}$ and $13 \rightarrow \text{MyType2}$.

Assume again that MyType is just selector<true>.

Now you will see a template class, enum_hunter, that maps *consecutive indices* to the user-supplied cases, so that enum_hunter<1> is 15 single_value<7>, enum_hunter<2> is single_value<13>, and so on.

The key idea is as follows:

- Since a default implementation is given, any single value<N> exists.
- User-supplied specializations have their member :: value == true.
- enum_hunter<N> will inspect all single_value<J>, starting at J==0, until it finds the Nth user-supplied value.
- enum hunter<N> is actually enum hunter<N, 0>.
- enum_hunter<N, J> inspects single_value<J>::value. If it's false, it inherits from enum_hunter<N, J+1>. Otherwise, it inherits from enum_hunter<N-1, J+1> (except when N-1 would be zero, where you pick <0, J> because the final result is precisely single value<J>).
- When N reaches 0, you are done. You met exactly N user-supplied values. If the initial N is too large, J will reach the terminator before N drops to 0, and since the terminator is an empty class, the compiler will complain.

All this yields a surprisingly compact implementation (for the moment, ignore the fact that everything is hard-coded):

```
template <int N, int J=0>
struct enum_hunter
: enum_hunter<N-single_value<J>::value, J+1-(N ==
single_value<J>::value)>
{
};

template <int J>
struct enum_hunter<0, J> : single_value<J>
{
};

template <>
struct enum_hunter<0, 0> : single_value<0>
{
};
```

```
} ;
```

This skeleton technique can lead to a couple of different applications—the simplest is to build a sparse compile-time *array* between arbitrary (but small) integers and types:

```
#define MXT ADD ENUMERATION(N, TYPE) \
  template <> struct single value<N> : public TYPE, selector<true>
{ }
struct Mapped1
 static double do_it() { return 3.14; }
};
struct Mapped2
 static double do it() { return 6.28; }
};
MXT ADD ENUMERATION (7, Mapped1);
MXT ADD ENUMERATION (13, Mapped2);
double xx1 = enum hunter<1>::do it(); // == 3.14
double xx2 = enum hunter < 2 > :: do it(); // == 6.28
  Polishing up the macros, you parameterize the name of enum hunter as ENUM and rename
single value as ENUM## case.
#define
MXT BEGIN ENUMERATION (ENUM)
template <int N> struct ENUM## case : static value<int, 0>
{ } ;
template <int N, int J=0> struct
ENUM
: ENUM<N-ENUM##_case<J>::value, J+1-(N == ENUM##_case<J>::value)>
{ };
template <int N> struct ENUM<0, N> : ENUM## case<N>
{ };
template <> struct ENUM<0, 0> : ENUM## case<0> {}
struct empty class {};
```

```
#define MXT_END_ENUMERATION(ENUM, K) \
  template <> struct ENUM##_case<K> : {}

// we explicitly add a member "value" without using derivation.
// this allows TYPE itself to be selector<true>

#define MXT_ADD_ENUMERATION(ENUM, TYPE, K) \
  template <> struct ENUM##_case<K> : TYPE \
  { static const int value = 1; }
```

When using the macros, every directive in the sequence between begin/end will be added automatically using line numbers as a progressive index. Two directives on the same line won't compile, since you cannot specialize a class template twice.

```
MXT_BEGIN_ENUMERATION(MyTypeEnum);

MXT_ADD_ENUMERATION(MyTypeEnum, Mapped1, 7);  // this gets index 1

MXT_ADD_ENUMERATION(MyTypeEnum, Mapped2, 13);  // this gets index 2

MXT_END_ENUMERATION(MyTypeEnum, 128);
```

So MyTypeEnum<1> is Mapped1, MyTypeEnum<2> is Mapped2, but MyTypeEnum_case<...> is still available to the code. Observe that 7 and 13 in the example may not be needed, if you plan to use the enumeration via contiguous indices. However, you need to provide unique and ascending values. So you can just pass __LINE__ as parameter K.

Another application of type enumeration is that, unlike classic enums, several headers can add their own values. So you can "distribute" a function between different files.

Suppose you want to gather the list of files included in a cpp and you don't want each header to access a global variable:

```
#include "H1.hpp"
#include "H2.hpp"
#include "H3.hpp"

int main(int argc, const char* argv[])
{
   std::vector<std::string> global_list;
   // here initialize global_list
}

   Arough solution could be as follows:

// flag_init.hpp
#define MXT_INIT_LIST
```

// equivalent to BEGIN ENUMERATION

```
template <int N> struct flag init
  static void run(std::vector<std::string>& v)
};
template <int N>
void run flag init(std::vector<std::string>& v, static value<int,</pre>
N > )
{
  flag init<N>::run(v);
  run flag init(v, static value<int, N+1>());
}
// magic constant, terminator
inline void run flag init(std::vector<std::string>& v,
static value<int, 64>)
// H1.hpp
#ifdef MXT INIT LIST
// equivalent to ADD ENUMERATION
// pick a random number < 64
template < > struct flag init<7>
  static void run(std::vector<std::string>& v)
     v.push back("hello, I am " FILE );
};
#endif
// the rest of H1.hpp, then write similarly H2 and H3
// main.cpp
#include "flag init.hpp"
#include "H1.hpp"
#include "H2.hpp"
#include "H3.hpp"
int main(int argc, const char* argv[])
```

```
{
    std::vector<std::string> global_list_of_flags;
    run_flag_init(global_list_of_flags);
}
```

7.7. If-Less Code

Sometimes program logic can be embedded in "smart objects" that know what to do, thus eliminating the need for if/switch blocks.

7.7.1. Smart Constants

As an example, suppose you need to code a suitable print function for a date class:

```
class date
  public:
     int day() const;
     int month() const;
     int year() const;
};
enum dateformat t
  YYYYMMDD,
  YYMMDD,
  DDMMYYYY,
  // many more...
};
void print(date d, dateformat t f)
  switch (f)
     case YYYYMMDD:
        // Very verbose...
  }
```

Instead, you can write branch-free code. As usual, TMP techniques take advantage of storing information in places where it's not evident that meaningful data can be stored!

Suppose the format constants like YYYYMMDD are actually numbers with six decimal digits in the form [f1 e1 f2 e2 f3 e3], where fi is the index of the "date field to print" (say, 0=year, 1=month and 2=day) and ei is the width as a number of digits.

For example, 041222 would be "year with four digits (04), month with two digits (12), and day with two digits (22)," or simply YYYY-MM-DD. This would enable you to write:

```
const int pow10[] = \{ 1, 10, 100, 1000, 10000, ... \};
const int data[3] = { d.year(), d.month(), d.day() };
const char* sep[] = { "-", "-", "" );
for (int i=0; i<3; ++i)
 std::cout << std::setw(e[i]) << (data[f[i]] % pow10[e[i]]) <<</pre>
sep[i];
  Generating such constants is easy:
enum { Y, M, D };
template <unsigned F, unsigned W = 2
struct datefield : static value<unsigned, F*10 + (W % 10)>
};
template <typename T1, typename T2 = void, typename T3 = void>
struct dateformat
  static const unsigned pow10 = 100 * dateformat<T2,T3>::pow10;
  static const unsigned value = pow10 * T1::value
+ dateformat<T2,T3>::value;
};
template < >
struct dateformat<void, void, void>
  static const unsigned value = 0;
  static const unsigned pow10 = 1;
};
enum
  YYYYMMDD = dateformat<datefield<Y,4>, datefield<M>, datefield<D>
>::value,
    DDMMYY = dateformat<datefield<D>, datefield<M>, datefield<Y>
>::value,
    YYYYMM = dateformat<datefield<Y,4>, datefield<M> >::value,
  // ...
} ;
```

For simplicity, this implementation uses rotation on three parameters only. ¹⁶ The print function follows:

```
void print(date d, dateformat_t f)
{
  const unsigned pow10[] = { 1, 10, 100, 1000, 10000, ... };
  const int data[3] = { d.year(), d.month(), d.day() };

  for (unsigned int fc = f; fc != 0; fc /= 100)
  {
    unsigned w = fc % 10;
    unsigned j = (fc % 100) / 10;

    std::cout << std::setw(w) << (data[j] % pow10[w]);
  }
}</pre>
```

7.7.2. Converting Enum to String

template <char C> struct char2int;

Similarly to what you did in the previous paragraph, you can encode a short string inside the value of an enumeration. The C++ standard guarantees that an enum is represented by a large unsigned integer, if any of the values are large. In practice, you can assume the enum will be a 64-bit integer. Since $2^{64} > 40^{12}$, you can store a string of length 12 as an integer in base 40, where A=1, B=2, and so on.

First you define the "alphabet":

```
template <size t N> struct int2char;
#define C2I(C, I) \
 template <> struct char2int<C> { static const size t value = I; }
\#define I2C(C, I)
 template <> struct int2char<I> { static const char value = C; }
#define TRANSLATE1(C1, N) \
 C2I(C1, N); I2C(C1, N)
#define TRANSLATE2(C1, C2, N) \
 C2I(C1, N); C2I(C2, N); I2C(C1, N)
TRANSLATE2('a', 'A', 1); // convert both 'A' and 'a' to 1, and
1 to 'a'
TRANSLATE2('b', 'B', 2);
// ...
TRANSLATE2 ('z', 'Z', 26);
TRANSLATE1 ('0', 27);
TRANSLATE1 ('1', 28);
// ...
```

```
TRANSLATE1 ('9', 36);
TRANSLATE1(' ', 37);
static const size t SSTRING BASE = 40;
template <size t N, bool TEST = (N<SSTRING BASE)>
struct verify num
static const size t value = N;
};
template <size t N>
struct verify num<N, false>
// this will not compile if a number >= 40 is used by mistake
};
template <char C1, char C2 = 0, ..., char C12 = 0>
struct static string
 static const size t aux
             = verify num< char2int<C1>::value >::value;
 static const size t value
             = aux + static string<C2,...,C12>::value
* SSTRING BASE;
};
template <>
struct static string<0>
 static const size t value = 0;
};
template <size t VALUE>
std::string unpack(static value<size t, VALUE>)
 std::string result(1, char(int2char<VALUE</pre>
% SSTRING BASE>::value));
 return result + unpack(static value<size t, VALUE/SSTRING BASE>
());
std::string unpack(static value<size t, 0>)
 std::string result;
 return result;
```

```
#define MXT_ENUM_DECODER(TYPE) \
template <TYPE VALUE> \
std::string decode() \
{ return unpack(static_value<size_t, VALUE>()); }
```

Note that you separate the generic code from the "implementation". Now you define an enum:

```
enum MyEnum
{
   first = static_string<'f','i','r','s','t'>::value,
   verylong = static_string<'v','e','r','y','l','o','n','g'>::value
};

MXT_ENUM_DECODER(MyEnum); // Write this to get a "decode" function
std::cout << decode<first>(); // prints "first"
```

For simplicity, this example implements static decoding (that is, the decoded enum value is known at compile time). However, the same operation can be performed at runtime.¹⁷

In general this technique is effective when the actual value of the enum is not meaningful to the program.

7.7.3. Self-Modifying Function Tables

Consider a trivial example of a circular container, where elements are "pushed back" (at the moment, pretend anything is public):

```
template <typename T, size_t N>
struct circular_array
{
   T data_[N];
   size_t pos_;
   circular_array()
     : data_(), pos_(0)
   {
   }
   void push_back(const T& x)
   {
      data_[pos_] = x;
      if (++pos_ == N)
          pos_ = 0;
   }
};
```

You can convert push_back into a sort of self-modifying function, similar to trampolines (see

Section 5.3.1). You will use a function pointer initialized with a suitable function template.

```
template <typename T, size t N>
struct circular array
  T data [N];
  typedef void (*push back t) (circular array<T, N>& a, const T&
x);
  push back t pb ;
  template <size t K>
  struct update element at
  {
     static void apply(circular array<T, N>& a, const T& x)
       a.data [K] = x;
       a.pb = &update element at<(K+1) % N>::apply;
  };
  circular array()
     : data (), pb (&update element at<0>::apply)
  {
  }
  void push back(const T& x)
     pb (*this, x);
```

The key point of this pattern is that you have a collection of functions where all elements know the action that follows, and so they can update a pointer with this information.

Updating the function pointer is not mandatory. A function may select itself as the next candidate. Suppose you change the container policy so as to keep the first N-1 elements and then constantly overwrite the last:

```
if ((K+1)<N)
   a.pb = &update element at<K+1>::apply;
```

};

Self-modifying functions are usually elegant, but slightly less efficient than a classic switch, mostly because of technology factors such as caches or program flow predictors.

Applications include data structures whose behavior during initialization is different (a "warm-up" phase) until a minimum number of elements have been inserted.

[■] **Note** *ICF*, identical code folding, is a widespread optimization technique applied by compilers.

Put simply, the linker will try to find "duplicate functions" and generate binary code only once. For example, vector<int*> and vector<double*> will probably generate identical code, so they can be merged.

While this reduces the size of the binary, it has the side effect that equality of function pointers may not work as expected. If F and G are identical (suppose they have an empty body), it's possible that F != G in a debug build and F == G in an ICF-optimized build.

Be careful when writing a self-modifying function that relies on equality/inequality of function pointers (obviously, comparisons with NULL pointers work fine).

¹Usually, this requires the additional assumption that a, b, and c point to unrelated areas of memory, but modern compilers will try to understand if these optimizations can be safely applied.

²This is commonly called code bloat.

³There are 26^N sequences of N letters, and "only" say 2⁶⁴ different hash values, so for N>14, no hash can be injective; however a good hashing algorithm will "scatter" conflicts, so strings having the same hash will be *really* different.

⁴As an exercise, the reader might generalize the idea to wchar_t, which in this implementation always picks the locale-based function set

⁵The motivation will be evident when you see an application to string hashing, later in the paragraph.

⁶There's no need for a complete mimesis implementation: a cast operator is not needed.

⁷This syntax will be used only in this section, where there's no possibility of confusion.

⁸Another common error is the argument crossover. Suppose a class C has two template parameters T1 and T2. If you partially specialize C<T1, Y> and C<X, T2> for some fixed X and Y, C<X, Y> is ambiguous, so it must be explicitly specialized too.

⁹The problem actually falls under the opaque type principle. If the return type of a function is "complex," you should either publish a convenient typedef to the users or allow them to use the object by ignoring its type (refer to Chapter 9 for more details).

¹⁰As follows from the opaque type principle, it's not necessary to document what the exact return type is, just state that it works like a std::pair with an extra operator>>. In principle, it would be reasonable to add a conversion operator from dynamic hash result to std::pair<size t, iterator t>.

¹¹ Since there's a single function in the class, this example does not derive from static_interface but replicates the code.

¹² This is an *online* problem. In *offline* problems, all the input values are given at the same time. There's a data structure by David Eppstein (see http://www.ics.uci.edu/~eppstein/pubs/kbest.html) that solves the online problem using memory proportional to N and exhibits amortized constant-time operations. This example focuses on how to improve a naive implementation, not on creating an efficient algorithm.

¹³Here, update and its auxiliary subroutines are global functions. This just makes the illustration easier, because it allows you to focus on one feature at a time. You can safely declare all these functions as private static members of the container.

 $^{^{14}}$ It seems that this kind of "greedy compact style" for small values of N gets most benefit from an aggressive optimizing compiler. A rudimentary stress test with 10.000.000 insertions and N<32 showed a very large runtime difference (30–40%) between a "normal" and an "extreme" release build. Greedy algorithms and compact code take advantage of technological factors, such as processor caches.

¹⁵is means derives from.

¹⁶So you cannot generate patterns like YYYYMMDDYY.

¹⁷Hint: Use a const array of char of length SSTRING_BASE and initialize it with { 0, int2char<1>::value, int2char<2>::value, ... }.

CHAPTER 8

Functors

This chapter focuses on several techniques that help when writing (or when not writing) functors.

Most STL algorithms require compile-time function objects and this usually requires some manual coding:

```
struct Person
{
  unsigned int age;
  std::string home_address;

  double salary() const;
};

std::vector<Person> data;
std::sort(data.begin(), data.end(), /* by age */);
std::partition(data.begin(), data.end(), /* by salary */);
```

If you can modify Person, sometimes an elegant and quick solution is to write a public static member function and a member functor. This simultaneously attains the maximum efficiency and control, as your code has access to private members:

```
struct Person
{
private:
    unsigned int age;

public:
    static bool less_by_age(const Person& a, const Person& b)
    {
        return a.age < b.age;
    }

    struct BY_AGE
    {
        bool operator()(const Person& a, const Person& b) const
        {
            return Person::less_by_age(a, b);
    }
}</pre>
```

```
};
};

std::vector<Person> data;
std::sort(data.begin(), data.end(), Person::less_by_age); //
suboptimal
std::sort(data.begin(), data.end(), Person::BY_AGE()); // good
```

A static member function has access to private data. However it will be much harder for the compiler to inline the comparison, so a functor is usually better.

You can even factor out some code that converts the former to the latter:

```
template <typename T, bool (*LESS)(const T&, const T&)>
struct less compare t
  typedef T first argument type;
  typedef T second argument type;
  typedef bool result_type;
  bool operator()(const T& x, const T& y) const
     return LESS(x, y);
};
struct Person
private:
  unsigned int age;
public:
  static bool less by age(const Person& a, const Person& b)
     return a.age < b.age;
  }
  typedef less compare t<Person, Person::less by age> BY AGE;
};
```

The name of the function/functor is chosen to make the expression *clear at the point of instantiation*, not at the point of definition.

Note that non-generic functors (whose arguments have a fixed type) are usually members of the class.

It's generally fair to assume that a functor can be freely copied and passed by value. If a functor needs many data members, you had better collect them in a separate structure and store only a reference. The caller of the functor will be responsible for keeping the extra information alive:

```
struct information needed to sort elements
  // ...
class my less
  const information needed to sort elements& ref;
public:
  explicit functor (const information needed to sort elements& ref)
    ref (ref)
  }
  bool operator()(const Person& p1, const Person& P2) const
  { . . . }
};
int main()
  information needed to sort elements i;
  // build a suitable container data...
  std::sort(data.begin(), data.end(), my less(i));
```

STL algorithms do not provide any guarantee concerning the number of copies of function objects. Another interesting feature is that a functor static type is irrelevant, because it's always deduced. If the functor is returned from a function, it will be used immediately (see Section 4.3.4); if it's passed to a function template, it will bind to an argument that accepts anything.

This allows clients to generate anonymous instances of complex function objects at the call site:

```
i = std::find_if(begin, end, std::bind2nd(std::less<double>(), 3.14));

// the exact type of the functor is irrelevant
// since find_if has an argument that binds to anything:

// template <typename I, typename F>
// I find_if(I begin, I end, F func)
```

■ **Note** C++0x includes support for creation of lambda objects.

It is a new syntax that can pass anonymous "pieces of code" in curly brackets as if they were functors. This mitigates the problem of name pollution. In other words, it's not necessary to give a name to an entity that is not reused.

See Section 12.4 for more details.

8.1. Strong and Weak Functors

Some functors are strongly typed. This means that the user fixes the argument of the function call when determining the template arguments. All standard functionals are strongly typed.

```
template <typename T>
struct less
{
   bool operator() (const T& lhs, const T& rhs) const
   {
      return lhs < rhs;
   }
};

std::sort(data.begin(), data.end(), less<Person>());

Alternatively, you can have a weak functor that accepts arguments with more freedom!:
struct weak_less
{
   template <typename T>
   bool operator() (const T& lhs, const T& rhs) const
   {
      return lhs < rhs;
   }
};

std::sort(data.begin(), data.end(), weak less());</pre>
```

A strongly typed functor statically blocks all types that are incompatible with T, but since this is limited to the interface, it can actually share the implementation with a weak functor:

```
template <typename T>
struct less : private weak_less
{
   bool operator()(const T& lhs, const T& rhs) const
   {
     return static_cast<const weak_less&>(*this)(lhs, rhs);
   }
};
```

8.2. Functor Composition Tools

The STL offers facilities to compose functors and values. For example, std::bind2nd turns a

binary operation and an operand into a unary function. Often, you'll need tools that perform the reverse.

The prefix by in by age is actually the composition of a binary relation with an accessor. age extracts the age from a person and by compares two ages. Here's a minimal implementation that abstracts this composition concept.

```
template <typename functor t>
class by t
  functor t f ;
public:
  by t(functor t f)
  : f_(f)
  { }
  template <typename argument t>
  bool operator()(const argument t& a, const argument t& b) const
     return f(a) < f(b);
} ;
template <typename functor t>
inline by t<functor t> by(const functor t& f)
  return f;
// see Section 1.1.4
template <typename R, typename A>
inline by t < R (*)(A) > by(R (*f)(A))
  return f;
struct age t
  unsigned int operator()(const Person& p) const
  {
     return p.age;
  }
  age t(int = 0)
};
```

```
static const age t AGE = 0^2;
int main()
  std::vector<Person> data;
  std::sort(data.begin(), data.end(), by(AGE));
}
   by is a functor composition tool. Since it does not impose any requirement on functor t, it
will accept suitable static member functions, which are convenient if Person: : age is private:
struct Person
private:
   unsigned int age;
public:
   static int AGE (const Person& p)
      return p.age;
};
std::sort(data.begin(), data.end(), by(Person::AGE));  // ok!
   A functor/accessor may be given powerful lambda semantics.
   Here is another preview of Section 9.2. In pseudo-intuitive notation, comparator (A, S) is a
predicate that returns true on object 0 if A (0) is "less" than S. "less" is a generic binary
predicate.
template
   typename scalar t,
   typename accessor t,
   template <typename T> class less t
class comparator
   scalar t x ;
   accessor t a ;
public:
   comparator(scalar t x, accessor t a = accessor t())
   : x (x), a (a)
   template <typename argument t>
```

```
bool operator()(const argument_t& obj) const
{
    less_t<scalar_t> less_;
    return less_(a_(obj), x_);
}
```

Using a template-template parameter instead of a normal binary predicate saves you from typing scalar t twice and makes an anonymous instance quite clear to read:

```
comparator<double, SALARY, std::greater>(3.0)
```

Another minor point is the class layout: x_i is declared before a_i , because a_i will often be stateless and is therefore a small object. x_i might have stronger alignment constraints.

Now you can add operators to the functor and promote it to a lambda predicate³:

```
struct age_t
{
  int operator()(const Person& a) const
  {
    return a.age;
}

template <typename T>
  comparator<T,age_t,std::less> operator<(const T& x) const
  {
    return comparator<T,age_t,std::less>(x, *this);
}

template <typename T>
  comparator<T,age_t,std::equal_to> operator==(const T& x) const
  {
    return comparator<T,age_t,std::equal_to>(x, *this);
  }
};

std::partition(data.begin(), data.end(), Person::AGE < 35);
std::partition(data.begin(), data.end(), Person::AGE == 18);</pre>
```

With a little effort, you can add more syntactic tricks to the chaining operator:

```
const selector<true> INCREASING;
const selector<false> DECREASING;

template <typename T>
bool oriented_less(const T& x, const T& y, selector<true>)
{
    return x<y;
}</pre>
```

```
template <typename T>
bool oriented less(const T& x, const T& y, selector<false>)
  return y<x;
}
  oriented less can flip operator < and simulate operator >.
template <typename functor t, bool ASCENDING = true>
class by t
  functor t f ;
public:
  by t(functor t f) : f (f) {}
  template <typename argument t>
  bool operator()(const argument t& a, const argument t& b) const
     return oriented less(f (a), f (b), selector<ASCENDING>());
  // inversion operators:
  by t<functor t, true> operator+() const
     return f ;
  by t<functor t, false> operator-() const
     return f ;
};
  And finally, there's another by helper function:
template <bool DIRECTION, typename functor t>
by t<functor t, DIRECTION> by (selector<DIRECTION>, const
functor t& v)
  return by t<functor t, DIRECTION>(v);
  All this allows writing:
std::sort(data.begin(), data.end(), +by(Person::AGE));
std::sort(data.begin(), data.end(), -by(Person::AGE));
std::sort(data.begin(), data.end(), by(DECREASING, Person::AGE));
```

■ Note I chose operator+ and operator- because by deals with numeric properties; the logical inversion of a unary predicate is better expressed with operator!

Also, lines #2 and #3 are identical. It's only a matter of style to pick the clearest.

The last improvement to by t is to perform strict type checking in operator ().

The function call operator accepts almost anything, so more type checking will trap errors arising from code that compiles merely by chance:

```
std::vector<Animal> data;
std::sort(data.begin(), data.end(), by(Person::AGE));
```

A convenient approach is to exploit cooperation from the functor. If functor_t has a member argument_type, it also will be the argument of a strong operator(). Otherwise, you use the weak function call operator.

As usual, you hide the decision in a template parameter and provide two partial specializations. First, some traits:

```
template <typename T>
struct argument type of
  typedef typename T::argument type type;
};
template <typename A, typename R>
struct argument type of <R (*)(A)>
  typedef A type;
};
template <typename A, typename R>
struct argument type of <R (*) (const A&)>
  typedef A type;
};
template <typename T>
struct has argument type
: selector<[[ true if T::argument type exists4 ]]>
};
template <typename A, typename R>
struct has argument type<R (*)(A) >
: selector<true>
};
```

// ... The first specialization performs strict type checking. template typename functor t, bool ASCENDING = true, bool STRICT CHECK = has argument type<functor t>::value struct by t; template <typename functor t, bool ASCENDING> struct by t<functor t, ASCENDING, true> // ... typedef typename argument type of<functor t>::type argument type; // note: strong argument type bool operator()(const argument type& a, const argument type& b) const return oriented less(f (a), f (b), selector<ASCENDING>()); }; template <typename functor t, bool ASCENDING> struct by t<functor t, ASCENDING, false> // ... // note: weak argument type. This will accept anything template <typename argument t> bool operator()(const argument t& a, const argument t& b) const return oriented less(f (a), f (b), selector<ASCENDING>()); **}**; To minimize code duplication, you factor out the function call operator in a template base and use a static cast, as in CRTP: template <typename functor t, bool ASCENDING = true> struct by t; template <typename functor t, bool ASCENDING, bool STRICT CHECK>

struct by base t;

```
template <typename functor t, bool ASCENDING>
struct by base t<functor t, ASCENDING, true>
 const functor t& f() const
   typedef by t<functor t, ASCENDING> real type;
   return static cast<const real type&>(*this).f;
 typedef typename argument type of<functor t>::type argument type;
 bool operator()(const argument type& a, const argument type& b)
const
   return oriented less(f()(a), f()(b), selector<ASCENDING>());
};
template <typename functor t, bool ASCENDING>
struct by base t<functor t, ASCENDING, false>
 const functor t& f() const
   typedef by t<functor t, ASCENDING> real type;
   return static cast<const real type&>(*this).f;
 }
 template <typename argument t>
 bool operator()(const argument t& a, const argument t& b) const
   return oriented less(f()(a), f()(b), selector<ASCENDING>());
};
template <typename functor t, bool ASCENDING = true>
struct by t
by base t<functor t, ASCENDING, has argument type<functor t>::value>>
   // ...
```

8.3. Inner Template Functors

};

Functor wrappers may be used as interface-leveraging tools.

Syntactically, you take advantage of the fact that inner class templates know template parameters of the outer class.

8.3.1. Conversion of Functions to Functors

Assume for simplicity that you have a collection of functions with a similar signature \mathbb{T} $f(\mathbb{T}, \mathbb{T}, \dots, \mathbb{T})$, where the number of arguments varies. Suppose further that the list of functions to be executed will be known at runtime, so you need a base class with a virtual call whose unique signature could be (const \mathbb{T}^* , size t).

Let's look for an automatic way of performing the conversion:

```
template <typename T>
struct base
{
   virtual T eval(const T*, size_t) const = 0;
   virtual ~base() {}
};
```

Given a function, say double F (double, double), you could embed it in a functor, but you would have to deduce T and F simultaneously:

```
template <typename T, T (*F)(T,T)>
struct functor : public base<T>
{
    // ...
}:
```

Actually, you need T before F, so you can build a class template on T only, and after that an inner template class:

```
template <typename T>
struct outer
{
  template <T (*F)(T,T)>
  struct inner : public base<T>
  {
```

First you identify outer<T>, then you build inner:

```
template <typename T>
struct function_call_traits
{
   template <T (*F)()>
   struct eval_0 : public base<T>
   {
      virtual T eval(const T* , size_t) const { return F(); }
```

```
};
  template \langle T (*F)(T) \rangle
  struct eval 1 : public base<T>
     virtual T eval(const T* x, size t) const { return F(x[0]);
}
   };
  template \langle T (*F) (T, T) \rangle
  struct eval 2 : public base<T>
     virtual T eval(const T* x, size t) const { return F(x[0],
x[1]);
  };
  // ...
  template \langle T (*F) () \rangle
  eval 0<F>* get ptr() const
     return new eval 0<F>;
   }
  template \langle T (*F)(T) \rangle
  eval 1<F>* get ptr() const
   {
     return new eval 1<F>;
   }
  template \langle T (*F) (T, T) \rangle
  eval 2<F>* get ptr() const
     return new eval 2<F>;
   // ...
};
template <typename T>
inline function call traits<T> get function call(T (*F)())
  return function call traits<T>();
}
template <typename T>
inline function call traits<T> get function call(T (*F)(T))
  return function call traits<T>();
```

```
template <typename T>
inline function_call_traits<T> get_function_call(T (*F)(T, T))
{
   return function_call_traits<T>();
}
// ...
#define
MXT_FUNCTION_CALL_PTR(F) get_function_call(F).get_ptr<F>()
```

Note that:

double add0()

- F is used twice, first as a pointer, then as a template argument.
- The get_ptr functions are not static, bizarre as it may look, this is an example of a traits class that's actually meant to be instantiated (but used anonymously).

```
return 6.28;
}
double add1(double x)
  return x+3.14;
}
double add2 (double x, double y)
{
  return x+y;
}
int main()
{
  double x[5] = \{1, 2, 3, 4, 5\};
  base < double > * f[3] =
  {
     MXT FUNCTION CALL PTR (add0),
     MXT FUNCTION CALL PTR(add1),
     MXT FUNCTION CALL PTR (add2)
  };
  for (int i=0; i<3; ++i)
     std::cout << f[i]->eval(x, 5);
```

```
// normal destruction code has been omitted for brevity
```

The previous example executes add0(), add1(\times [0]), and add2(\times [0], \times [1]) via calls to the same interface.

8.3.2. Conversion of Members to Functors

The very same technique seen in the previous section can transform pointers into functors.⁶

In C++, simple structures with no access restrictions are often used to transport small pieces of data. Ideally, you'll want to maintain this simplicity and be able to write code with no overhead:

```
struct Person
{
   unsigned int age;
   double salary() const;
};

std::vector<Person> data;

// warning: pseudo-c++
std::sort(data.begin(), data.end(), by(Person::age));
std::sort(data.begin(), data.end(), by(Person::salary));
```

Because you can use a pointer-to-member as a template argument, it's not too hard to write an auxiliary wrapper that can help. Unfortunately, the instantiation is too verbose to be useful.

```
template <typename from_t, typename to_t, to_t from_t::* POINTER>
struct data_member
{
   const to_t& operator() (const from_t& x) const
   {
      return x.*POINTER;
   }
};

template <typename from_t, typename to_t, to_t (from_t::*POINTER)
() const>
struct property_member
{
   to_t operator() (const from_t& x) const
   {
      return (x.*POINTER) ();
   }
};
```

```
struct TEST
  int A;
  int B() const { return -A; }
};
TEST data[3] = \{2,1,3\};
// very verbose...
std::sort(data, data+3, by(data member<TEST, int, &TEST::A>()));
std::sort(data, data+3, by(property member<TEST, int, &TEST::B>
()));
  However, it's not possible to write a generic class pointer as the only template parameter:
template <typename A, typename B, B A::*POINTER>
                                                           // illegal:
struct wrapper<POINTER>
not c++
  You have to resort again to a nested class template:
template <typename from t, typename to t>
struct wrapper
  template <to t from t::*POINTER>
                                                           // legal!
  struct dataptr t
     const to t& operator()(const from t& x) const
        return x.*POINTER;
     }
  };
  template <to t from t::*POINTER>
  dataptr t<POINTER> get() const
     return dataptr t<POINTER>();
} ;
template <typename from t, typename to t>
wrapper<from t, to t> get wrapper(to t from t::* pointer)
{
  return wrapper<from t, to t>();
}
```

The example includes a function that takes the pointer to perform the first deduction, and again you have to supply the same pointer twice, once at runtime (whose *value* is basically ignored, but whose *type* is used for deduction) and once at compile-time:

```
#define MEMBER(PTR) get wrapper(PTR).get<PTR>()
```

- get_wrapper deduces arguments T1 and T2 automatically from PTR, so get_wrapper(PTR) will return wrapper<T1, T2>.
- Then you ask this wrapper to instantiate its member function get again on PTR, which returns the right object.

If PTR has type int TEST::*, the macro will produce a functor of type dataptr_t<PTR> (technically, wrapper<TEST, int>:: dataptr_t<PTR>).

However, any other overload will do. Here's an extended version:

```
template <typename from t, typename to t>
struct wrapper
{
  template <to t from t::* POINTER>
  struct dataptr t
     // optional:
     // typedef from t argument type;
     const to t& operator()(const from t& x) const
     {
       return x.*POINTER;
     }
  };
  template <to t (from t::*POINTER)() const>
  struct propptr t
     // optional:
     // typedef from t argument type;
     to t operator()(const from t& x) const
       return (x.*POINTER)();
     }
  };
  template <to t from t::* POINTER>
  dataptr t<POINTER> get() const
     return dataptr t<POINTER>();
  template <to t (from t::*POINTER)() const>
  propptr t<POINTER> get() const
```

```
return propptr t<POINTER>();
};
template <typename from t, typename to t>
wrapper<from t, to t> get wrapper(to t from t::* pointer)
  return wrapper<from t, to t>();
template <typename from t, typename to t>
wrapper<from t, to t> get wrapper(to t (from t::*pointer)() const)
  return wrapper<from t, to t>();
#define mxt create accessor(PTR) get wrapper(PTR).get<PTR>()
struct TEST
  int A;
  int B() const { return -A; }
};
TEST data[3] = \{2,1,3\};
std::sort(data, data+3, by(mxt create accessor(&TEST::A)));
std::sort(data, data+3, by(mxt create accessor(&TEST::B)));
```

As usual, if the name of the class contains a comma (such as std::map<int, float>), you need to typedef it before calling the macro.

The & is not strictly necessary. It's possible to redefine the macro as get wrapper (PTR) .get<&PTR>() in order to invoke it on the plain qualified name.

According to the Standard, the macro does not work inside templates as written. An additional template keyword is necessary for the compiler to deduce correctly what get is, so the best option is to define a second macro named (say)

```
mxt_create_accessor_template
get_wrapper(PTR).template get<&PTR>()
```

This version needs to be used whenever PTR depends on a template parameter that has impact on the line where the macro expands. On the other hand, it is forbidden whenever PTR does not depend on anything else. 7

8.3.3. More on the Double Wrapper Technique

In the previous paragraph, you saw a macro that looks like this one:

```
#define MEMBER(PTR) get wrapper(PTR).get<PTR>()
```

The argument PTR is used twice—the first time as an argument of a template function, which ignores its value but uses only its type and returns an "intermediate functor"; the second time as a template parameter of the functor itself, which produces the final object that you need.

Let's revise this technique, to face an apparently unrelated problem.⁸ In classic C++, enumeration values decay automatically to integers. This may cause bugs:

```
enum A { XA = 1 };
enum B { XB = 1 };

int main()
{
    A a = XA;
    B b = XB;
    a == b; // compiles and returns true, even if enums are unrelated
}
```

Let's introduce a simple helper functor: an object of type <code>enum_const</code> is a static value that compares exactly equal to one value from the same (non-anonymous) enumeration, but it cannot be compared to an integer or to a different type.

```
template <typename T, T VALUE>
struct enum_const
{
    bool operator==(T that) const
    {
       return VALUE == that;
    }

    // Barton-Nackman, see section 6.6
    friend inline bool operator==(T lhs, enum_const<T, VALUE> rhs)
    {
       return rhs == lhs;
    }
};

template <typename T>
struct enum_const_helper
{
    template <T VALUE>
    enum_const<T, VALUE> get() const
    {
}
```

```
return enum const<T, VALUE>();
};
template <typename T>
inline enum const helper<T> wrap(T)
   return enum const helper<T>();
  So you can write code like:
#define enum static const(X) wrap(X).get<X>()
int main()
  A a = XA;
  B b = XB;
                                       // ok
  a == b;
  enum static const(XB) == a;
                                     // error
}
error: invalid operands to binary expression ('int' and
'enum const<A, (A) 1U>')
   b == enum static const(XA); // fails
   note: candidate template ignored: deduced conflicting types for
parameter 'T' ('B' vs. 'A')
inline bool operator==(T lhs, enum const<T, VALUE> rhs)
error: invalid operands to binary expression ('enum const<B,
(B) 1U>' and 'int')
   enum static const(XB) == a; // fails
   ^~~~~~~~~~~~~~~~~~
note: candidate function not viable: no known conversion from 'A'
to 'B' for 1st argument;
   bool operator==(T that) const
  The macro as written works, but it needs X to be a compile-time constant:
#define enum static const(X) wrap(X).get<X>()
  Let's look for a workaround. The first question is, can wrap detect if X is a constant or a
variable? It can partially —a variable can bind to a reference.<sup>9</sup>
```

template <typename T>
inline enum_const_helper<T> wrap(T, ...)

```
{
    return enum_const_helper<T>();
}

template <typename T>
inline enum_var_helper<T> wrap(T& x, int)
{
    return enum_var_helper<T>(x);
}
```

Note the additional argument to wrap. Suppose X is a variable and you write wrap (X); both wrap (T&) and wrap (T) are valid matches, so the overload resolution is ambiguous. On the other hand, the expression wrap (X, 0) will prefer to match (T&, int) when possible, because 0 has exactly type int (which is better than ellipsis). So the macro becomes:

```
#define enum static const(X) wrap(X, 0).get<X>()
```

The second question is, if X is a variable, can you give a meaning to get < X > ()? Again, let's introduce a dummy argument of type int:

```
template <typename T>
struct enum_const_helper
{
   template <T VALUE>
   enum_const<T, VALUE> get(int) const
   {
     return enum_const<T, VALUE>();
   }
};
```

And here's the final version of the macro:

```
#define enum static const(X) wrap(X, 0).get<x>(0)
```

Now the syntax is different: get may be a member object and get<X>(0) is actually (get.operator<(X)). operator>(0). This is valid, since the object returned by wrap has no dependency on other template parameters.

Here's the missing piece of code:

```
template <typename T>
struct enum_var
{
   const T value_;
   explicit enum_var(T val)
   : value_(val) {}
   bool operator==(T that) const
   {
```

```
return value == that;
   }
       // Barton-Nackman again
   friend inline bool operator==(T lhs, enum var<T> rhs)
   {
       return rhs == lhs;
   enum var operator<(T) const // dummy operator<</pre>
       return *this; }
   enum var operator>(int) const // dummy operator>
      return *this;
};
template <typename T>
struct enum var helper
{
   enum var<T> get;
                                // surprise: data member called get
   enum var helper(T& x)
   : get(x) {}
};
                              // picks
enum static const(XB) == b;
enum const<B,1>::operator==(b)
enum static const(b) == XB;
                                // picks enum var<B>(b).operator==
(XB)
```

8.4. Accumulation

An *accumulator* is a functor that performs a logical "pass" over a sequence of elements and is updated via operator+= or operator+. This is implemented in the STL algorithm std::accumulate.

```
template <typename iterator_t, typename accumulator_t>
accumulator_t accumulate(iterator_t b, iterator_t e, accumulator_t
x)
{
  while (b != e)
    x = x + *(b++);
  return x;
}
```

If x is value type (0), this actually produces the sum over the range.

Accumulators can be classified as *online* or *offline*. Offline objects may accumulate only once over a range, and no more values can be added. On the other hand, online objects can accumulate disjoint ranges. (An ordinary sum is an online accumulation process, because the new total depends only on the previous total and the new values. An exact percentile would be an offline process, because the P-th percentile over two disjoint ranges depends on *all* the values at once.¹⁰)

The first step in a generalization is to accumulate F (*i), not necessarily *i.11

```
template <typename T>
struct identity
  T operator()(T x) const { return x; }
};
template <typename iter t, typename accumulator t, typename
accessor t>
accumulator t accumulate(iter t b, iter t e, accumulator t x,
accessor t F)
  while (b != e)
     x = x + F(*(b++));
  return x;
}
template <typename iter t, typename accumulator t>
accumulator_t accumulate(iterator_t b, iterator_t e, accumulator t
X)
{
  return accumulate(b, e, x,
           identity<typename
std::iterator traits<iter t>::reference>());
```

With TMP it's possible to build multi-layered accumulators on the fly:

- Recognize a set of similar operations that will get a performance boost from being performed simultaneously, rather than sequentially.¹²
- Define a reasonable syntax for instantiating an unnamed multiple accumulator.
- Define a reasonable syntax for extracting the results.

8.4.1. A Step-by-Step Implementation

The rest of the section will write a suitable function named collect that will make it possible to write the following:

```
// collect F(*i) for each i in the range
// and produce sum, gcd and max
std::accumulate(begin, end, collect(F)*SUM*GCD*MAX)
```

You'll take advantage of the fact that std::accumulate returns the accumulator to dump the desired results, either one or many at a time:

Let's restart from the beginning.

First, you identify the elementary operations and assign a code to each:

Again, you'll use template rotation. The main object contains the list of operations; it executes the first, then rotates the list and dispatches execution. T is the accessor.

Then you implement the binary operations (some code is omitted for brevity):

```
template <int N>
struct op_t;
template <>
struct op t<op void>
private:
  explicit op t(int = 0) {}
};
template <>
struct op t<op sum>
  explicit op t(int = 0) {}
  template <typename scalar t>
  scalar t operator()(const scalar t a, const scalar t b) const
     return a+b;
};
  You create some global constant objects; the explicit constructor has exactly this purpose.
const op t< op gcd > GCD(0);
const op t< op sum > SUM(0);
const op t< op max > MAX(0);
const op t< op min > MIN(0);
  Note that nobody can construct op t<op void>.
  Since you can perform exactly four different operations, you put four as the limit of template
parameters:
template
typename accessor t,
int O1 = op_void, int O2 = op_void, int O3 = op void, int O4
= op void
class accumulate t
  typedef typename accessor t::value type scalar t;
  typedef accumulate t<accessor t,02,03,04> next t;
  template <typename T, int I1, int I2, int I3, int I4>
     friend class accumulate t;
```

```
static const int OP COUNT = 1 + next t::OP COUNT;
   scalar t data [OP COUNT];
   size t count ;
   accessor t accessor;
   Every object is constructed via an instance of the accessor:
public:
   accumulate t(const accessor t \& v = accessor t())
   : accessor (v), count (0), data ()
   // more below...
};
  You have an array of results named data_. The i-th operation will store its result in
data [i].
   The recursive computation part is indeed simple. There's a public operator+= that calls a
private static member function:
template <typename object t>
accumulate t& operator+=(const object t& t)
   apply(data_, accessor_(t), count); // <-- static
   return *this;
and a global operator+:
template <typename accessor t, int N1, ..., int N4, typename
scalar t>
accumulate t<accessor t, N1, N2, N3, N4>
  operator+(accumulate t<accessor t,N1,N2,N3,N4> s, const scalar t
X)
   return s += x;
   accessor (t) yields the value to be accumulated over the memory cell *data. If count is
0, which means that the cell is "empty," just write the value. Otherwise, invoke the first binary
operation that merges the previous cell value and the new one. Then, advance the pointer to the next
cell and forward the call to next t:
static void apply(scalar t* const data, const scalar t x, size t&
```

{

}

count)

```
*data = (count>0) ? op t<01>()(*data, x) : x;
  next t::apply(data+1, x, count);
   The recursion is stopped when all operations are op void. At this point, you update the counter.
template <typename accessor t>
class accumulate t <accessor t, op void, op void, op void,
op void>
  /* ... */
  static const int OP COUNT = 0;
  static void apply(scalar t* const, const scalar t, size t& count)
    ++count;
   You need another static recursion to retrieve the result:
private:
   template <int N>
   static scalar t get(const scalar t* const data, op t<N>)
   {
     return O1==N ? data[0] : next t::get(data+1, op t<N>());
   }
public:
   template <int N>
   scalar t result(op t<N>) const
     return get(data , op t<N>());
   The recursion stopper is not expected to be invoked. However, it's necessary because
next t::get is mentioned (and thus, fully compiled anyway). It will be executed only if one asks
for result (op t<K>) for an object of type accumulate t<K1...Kn> and K is not in the
list.
   In this case, you can induce any suitable runtime error:
template <typename accessor t>
class accumulate t <accessor t, op void, op void, op void,
op void>
```

static scalar t get(const scalar t* const, op t<N>)

private:

template <int N>

```
// if you prefer,
    // throw std::runtime_error("invalid result request");
    return std::numeric_limits<scalar_t>::quiet_NaN();
}

public:
    /* nothing here */
};
```

Since SUM is a global constant of the right type, you are eventually going to call std::accumulate(begin,end,[...]).result(SUM).

At this point, you can write code that computes the result and code that retrieves the result, but you're still missing the accumulator factory. As frequently happens for all objects based on template rotation, you give the user a helper function that initially produces an "empty accumulator" (namely, accumulate_t<T>, or more precisely, accumulate_t<T, 0, 0, ..., 0>) and this empty object can be combined repeatedly with one or more op_t. In other words: there's an operator that combines accumulate_t<T> and an operation N1, performing a static "push-front" and returning accumulate t<T, N1>.

If you pick operator* (binary multiplication) for chaining, the function looks like this:

```
template <int N, int N1, ... int Nk> accumulate_t<T, N, N1, N2,...,N_{k-1}> operator*(accumulate_t<T, N1,...,N_{k-1}, Nk>, op_t<N>)
```

This chaining operator will contain a static assertion to ensure that the "dropped term" Nk is op void.

Here's the global helper function:

```
template <typename accessor_t>
inline accumulate_t<accessor_t> collect(const accessor_t& v)
{
   return v;
}
```

Finally, here is a listing of the whole class, side by side with the recursion stopping specialization:

```
template <typename T, int I1, int I2, int I3,
typedef typename accessor t::value type
scalar t;
template <typename T, int I1, int I2, int
                                           friend class accumulate t;
friend class accumulate t;
typedef
accumulate t<accessor t,02,03,04,op void>
next t;
static const int OP COUNT =
                                             static const int OP COUNT = 0;
1+next_t::OP_COUNT;
scalar_t data_[OP_COUNT];
size_t count_;
                                             accessor t accessor;
accessor t accessor;
static void apply(scalar_t* const data,
                                             static void apply(scalar t* const,
const scalar t x, size t& count)
                                             const scalar t, size t& count)
*data = (count>0) ? op t<01>() (*data, x)
                                             ++count;
next t::apply(data+1, x, count);
template <int N>
                                             template <int N>
static scalar t get(const scalar t* const
                                             static scalar t get(const scalar t* const,
data, op t<N>)
                                             op t < N > )
return O1==N ?
                                             assert (false);
data[0] : next t::get(data+1, op t<N>());
                                             return 0;
public:
                                             public:
accumulate t(const accessor t& v =
                                             accumulate t(const accessor t& v = accessor t)
accessor t())
                                             : accessor (v)
: accessor (v), count (0), data ()
template <int N>
                                             template <int N>
accumulate t<accessor t,N,01,02,03>
                                             accumulate t<accessor t, N>
operator* (op t<N>) const
                                             operator* (op t<N>) const
MXT ASSERT(O4 == op_void);
                                             return accessor;
return accessor_;
template <typename object t>
                                             template <typename object t>
accumulate t& operator+=(const object t&
t)
                                             accumulate t& operator+=(const object t& t)
```

The last feature provides the ability to retrieve more results at one time. This is extremely important, since it avoids storing the result of the accumulation.

You simply introduce an operator that binds a reference to each op_t (this example uses operator>> since it resembles an arrow). Another possible choice is operator<=, since <= can be seen as ←) and builds a reference wrapper of unique type. From this temporary, an overloaded accumulator::result will extract both operands and perform the assignment.

```
accumulator.result(SUM >> r1, MAX >> r2);
   The implementation is as follows:

template <typename scalar_t, int N>
struct op_result_t
{
    scalar_t& value;
    op_result_t(scalar_t& x)
    : value(x)
    {
    }
};

template <typename scalar_t, int N>
inline op_result_t<scalar_t, N> operator>> (const op_t<N>, scalar_t& x)
{
    return op_result_t<scalar_t, N>(x);
}
```

RESULT1 r1; RESULT2 r2;

```
Then you add these methods to the general template (the macro is for brevity only): #define ARG(J) const op result t<scalar t, N##J> o##J
```

```
// ARG(1) expands to "const op result t<scalar t, N1> o1"
  template <int N1>
  const accumulate t& result(ARG(1)) const
     o1.value = result(op t<N1>());
     return *this;
  }
  template <int N1, int N2>
  const accumulate t& result(ARG(1), ARG(2)) const
  {
     result (o2);
     return result(o1);
  }
  template <int N1, int N2, int N3>
  const accumulate t& result(ARG(1), ARG(2), ARG(3)) const
     result(o3);
     return result(o1, o2);
  }
  template <int N1, int N2, int N3, int N4>
  const accumulate t& result(ARG(1), ARG(2), ARG(3), ARG(4)) const
     result(04);
     return result(o1, o2, o3);
  }
```

#undef ARG

The expression MAX>>x silently returns op_result_t<[[type of x]], op_max>(x). If x does not have the same type as the accumulated results, it will not compile.

A couple of extra enhancements will save some typing. Instead of having many result, you just add the first one and chain the subsequent calls via operator (). 13

```
template <int N1>
const accumulate_t& result(const op_result_t<scalar_t,N1> o1)
const
{
   o1.value = result(op_t<N1>());
   return *this;
}
```

```
const accumulate t& operator()(const op_result_t<scalar_t,N1> o1)
const
  return result(o1);
  So instead of:
int q sum, q gcd, q max;
std::accumulate(...).result(SUM >> q sum, GCD >> q gcd, MAX >>
q max);
the new syntax is:
std::accumulate(...).result(SUM >> q sum)(GCD >> q gcd)(MAX >>
q max);
or even:
std::accumulate(...)(SUM >> q sum)(GCD >> q gcd)(MAX >> q max);
  Second, you add an overload that returns the first result for functions that accumulate a single
quantity:
   scalar t result() const
```

```
// MXT_ASSERT(02 == op_void);
   return result(op_t<01>());
}

// now .result(SUM) is equivalent to .result()
int S = std::accumulate(data, data+7, collect(...)*SUM).result();
```

8.5. Drivers

template <int N1>

A well-written algorithm avoids unnecessary multiplication of code. To rewrite an existing algorithm for greater generality, you have to remove some "fixed" logic from it and plug it in again through a template parameter, usually a functor:

```
template <typename iterator_t>
void sort(iterator_t begin, iterator_t end)
{
   for (...)
   {
      // ...
      if (a<b) // operator< is a good candidate for becoming</pre>
```

```
a functor
     { }
     }
}
```

So you rewrite this as:

```
template <typename iterator_t, typename less_t>
void sort(iterator_t begin, iterator_t end, less_t less)
{
   for (...)
   {
      // now we ask the functor to "plug" its code in the algorithm
      if (less(a,b))
      {}
   }
}
```

A *driver* is an object that can guide an algorithm along the way.

The main difference between a functor and a driver is that the former has a general-purpose function-like interface (at least, operator()), which is open to user customization. On the other hand, a driver is a low level object with a verbose interface, and it's not meant to be customized (except for its name, it might not even be documented, as if it were a tag type). The framework itself will provide a small fixed set of drivers.

Consider the following example. You need an sq function that optionally logs the result on std::cerr. Because you cannot enforce such a constraint if you receive a generic logger object, you switch to drivers and then provide some:

```
struct dont_log_at_all
{
   bool may_I_log() const { return false; }
};

struct log_everything
{
   bool may_I_log() const { return true; }
};

struct log_ask_once
{
   bool may_I_log() const
   {
     static bool RESULT = AskUsingMessageBox("Should I log?",
MSG_YN);
     return RESULT;
   }
};
```

```
template <typename scalar_t, typename driver_t>
inline scalar_t sq(const scalar_t& x, driver_t driver)
{
  const scalar_t result = (x*x);
  if (driver.may_I_log())
    std::cerr << result << std::endl;
  return result;
}

template <typename scalar_t>
inline scalar_t sq(const scalar_t& x)
{
  return sq(x, dont_log_at_all());
}
```

Note that driver_t::may_I_log() contains neither code about squaring, nor about logging. It just makes a decision, driving the flow of the algorithm.

The big advantage of drivers is to reduce debugging time, since the main algorithm is a single function. Usually drivers have minimal runtime impact. However nothing prevents a driver from performing long and complex computations.

As a rule, you always invoke drivers through instances. An interface such as

```
template <typename driver_t>
void explore(maze_t& maze, driver_t driver)
{
   while (!driver.may_I_stop())
   { ... }
}
```

is more general than its stateless counterpart¹⁴:

```
template <typename driver_t>
void explore(maze_t& maze)
{
   while (driver_t::may_I_stop())
   { ... }
}
```

A driver is somehow analogous to the "public non-virtual / protected virtual" classic C++ idiom (see Section 6.3). The key similarity is that the structure of the algorithm is fixed. The user is expected to customize only specific parts, which run only when the infrastructure needs them to.¹⁵

8.6. Algors

An algor, or algorithmic functor, is an object that embeds an algorithm, or simply an algorithm with

state.

The standard C++ library provides an <algorithm> header, which includes only functions. So it's natural to identify function and algorithms, but it need not be the case.

The algor object implements a simple function-like interface—typically operator () —for the execution of the algorithm, but its state grants faster repeated executions.

The simplest case where an algor is useful is buffered memory allocation.

std::stable_sort may require the allocation of a temporary buffer that's necessarily released when the function returns. Usually this is not an issue, since time spent in (a single) memory allocation is dominated by the execution of the algorithm itself. A small input will cause a small memory request, which is "fast" (operating systems tend to favor small allocations). A large input will cause a "slow" memory request, but this extra time will be unnoticed, since the algorithm will need much more time to run.

However, there are situations where a single buffer would suffice for many requests. When stable-sorting many vectors of similar length, you can save allocation/deallocation time if you maintain the buffer in an object:

```
template <typename T>
class stable sort algor
{
  buffer type buffer ;
public:
  template <RandomAccessIterator>
  void operator()(RandomAccessIterator begin, RandomAccessIterator
end)
  {
     // ensure that buffer is large enough
     // if not, reallocate
     // then perform the stable sort
  }
  ~stable sort algor()
     // release buffer
};
```

To sum up, the simplest algor is just a sort of functor with state (in the last case, a temporary buffer), but algors may have a richer interface that goes beyond functors.

As a rule, algors are not copied or assigned. They are constructed and reused (say, in a loop) or used as unnamed temporaries for a single execution. You therefore don't need to worry about efficiency, only about safety. If buffer_type cannot be safely copied (if it's a pointer), you explicitly disable all the dangerous member functions, making them private or public do-nothing operations. If buffer_type is a value type (for example, vector<T>), you let the compiler generate safe, possibly inefficient, operators.

Another useful kind of algor is a self-accumulator that holds multiple results at once. There's no

```
buffer involved (see Section 8.4).
template <typename T>
class accumulator
  T max ;
  T min ;
  T sum ;
  // ...
public:
  accumulator()
  : sum_(0) // ...
  {
  }
  template <typename iterator t>
  accumulator<T>& operator()(iterator t begin, iterator t end)
  {
     for (;begin != end; ++begin)
        sum += *begin;
        // ...
     }
     return *this;
  T max() const { return max_; }
  T min() const { return min_; }
  T sum() const { return sum ; }
  // ...
};
int main()
  double data[] = \{3,4,5\};
  // single invocation
  double SUM = accumulator<double>() (data, data+3).sum();
  // multiple results are needed
  accumulator<double> A;
  A(data, data+3);
  std::cout << "Range: " << A.max()-A.min();</pre>
```

An *interactive algor* has an interface that allows the caller to run the algorithm step-by-step. Suppose for example you have to compute the square root to some level of precision:

```
template <typename scalar t>
class interactive square root
  scalar t x ;
  scalar t y ;
  scalar t error ;
public:
  interactive square root(scalar t x)
   : X (X)
  {
     iterate();
  void iterate()
     // precondition:
     // y_ is some kind of approximate solution for y^2=x
     // error is |y^2-x|
     // now compute a better approximation
  }
  scalar t error() const
     return error ;
  operator scalar t() const
     return y ;
};
  It's the user who drives the algorithm:
int main()
  interactive square root<double> ISR(3.14);
  while (ISR.error()>0.00001)
     ISR.iterate();
  double result = ISR;
```

An algor of this kind usually takes all its parameters from the constructor. A common use-case is an algorithm that produces *a set of solutions*. After execution, a member

function permits the user to "visit" all the solutions in some order. ¹⁶ These algors might do all the work in the constructor:

```
template <typename string t>
class search a substring
  const string t& text;
  std::vector<size t> position ;
public:
  search a substring(const string t& TEXT, const string t&
PATTERN)
  : text (TEXT)
     // search immediately every occurrence of PATTERN in TEXT
     // store all the positions in position
  }
  bool no match() const { return position .empty(); }
  // the simplest visitation technique
  // is... exposing iterators
  typedef std::vector<size t>::const iterator position iterator;
  position iterator begin() const
  {
     return position .begin();
  }
  position iterator end() const
     return position .end();
```

In the case of substring matching, the iterator will likely visit the matches from the first to the last. In a numerical minimization problem, the solutions may be N points where the function has the minimum value found so far.

};

A more complex visitor-accepting interface could accept two *output* iterators, where the algor would write its solutions. You could build a "custom view" on the solutions according to the iterator value_type. For example, an algor that internally computes pairs (Xj, Yj) may emit just the first component or the entire pair (a simplified example follows):

```
class numerical_minimizer
{
   std::function<double (double)> F;
   std::vector<double> X_; // all the points where F has minima
```

```
public:
   // ...
   template <typename out t>
   out t visit(out_t beg, out_t end) const
      typedef typename std::iterator traits<out t>::value type>
val t;
      int i=0;
      while (beg != end)
        *beg++ = build result(i++, instance of<val t>());
      return beg;
   }
private:
   template <typename T>
   double build result(int i, instance of<T>) const
      return X [i];
   using std::pair;
   template <typename T>
   pair<double, double> build result(int i, instance of<pair<T,T>>)
const
      return std::make pair(X [i], F(X [i]));
};
```

8.7. Forwarding and Reference Wrappers

It's a common idiom for a class template to hold a member of a generic type, to which the class dispatches execution.

```
template <typename T>
class test
{
    T functor_;
public:
```

```
typename T::value_type operator()(double x) const
{
    return functor_(x); // call forward
}
```

Since the exact type of the member is not known, you may have to implement several overloads of test::operator(). Since this is a template, this is not a problem, because what's actually needed will be instantiated and the rest is ignored.

Invoking the wrong overload (that is, supplying too many or unsupported arguments) will cause a compiler error. However, note that arguments are forwarded by value, so you can modify the prototypes:

But if T requires an argument by non-const reference, the code will not compile.

To understand the severity of the problem, consider a slightly different example, where you *construct* a member with an unspecified number of parameters.

The STL guidelines suggest writing a single constructor for class test, which takes (possibly) a previously constructed object of type T:

```
test(const T& data = T())
: member_(data)
{
}
```

This strategy is not always possible. In particular, T might have an inaccessible copy constructor or it may be a non-const reference.

In fact, let's forget the STL style for a moment and adapt the same idiom of operator () as shown previously.

```
template <typename T>
class bad test
  T member ;
public:
  template <typename X1>
  bad test(X1 arg1)
  : member (arg1)
  }
  template <typename X1, typename X2>
  bad test(X1 arg1, X2 arg2)
  : member (arg1, arg2)
};
  As written, bad test<T&> compiles, but a subtle bug arises<sup>17</sup>:
int main(int argc, char* argv[])
{
  double x = 3.14;
                                     // unfortunately, it compiles
  bad test<double&> urgh(x);
  urgh.member = 6.28;
                                      // bang!
  int i = 0;
  assert(x == 6.28);
                                      // assertion failed!
  // ...
```

The constructor of urgh is instantiated on type double, not double &, so urgh .member_refers to a temporary location in the stack of its constructor (namely, the storage space taken by arg1), whose content is a temporary copy of x.

So you modify bad test to forward arguments by const reference. At least,

```
good test<double&> will not compile (const double& cannot be converted to
double&).
template <typename T>
class good test
  T member ;
public:
  template <typename X1>
  good test(const X1& arg1)
  : member (arg1)
  }
};
  However, an additional wrapping layer can solve both problems:
template <typename T>
class reference wrapper
  T& ref;
public:
  explicit reference wrapper (T& r)
     : ref (r)
  }
  operator T& () const
  {
     return ref ;
  T* operator& () const
     return &ref ;
};
template <typename T>
inline reference wrapper<T> by ref(T& x)
  return reference wrapper<T>(x);
int main()
 double x = 3.14;
```

```
good_test<double> y0(x); // ok: x is copied into y0.member_
good_test<double&> y1(x); // compiler error!
y1.member_ = 6.28; // would be dangerous, but does not
compile
good_test<double&> y2(by_ref(x));
y2.member_ = 6.28; // ok, now x == 6.28
```

Using by_ref, the good_test<double&> constructor is instantiated on argument constructor efference wrapper<double&>&, which is then converted to double&.

Note Once again, the argument-forwarding problem is solved in C++0x with R-value references.

¹These functors have been voted into C++14 with a slightly different terminology; they are called "transparent". To the knowledge of the author, this book was the first place where the idea appeared publicly. For all the details, see http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2012/n3421.htm.

²Since most functors are stateless, and so are not affected by initialization problems, global constants can be created in header files.

³As a rule, for expressiveness, it's best to write a fully-qualified Person: : AGE rather than just AGE, so you must assume that there's a static constant of type age_t inside Person. This also allows age_t to be friend of Person. You can also consider Person::AGE() where AGE is either a member typedef for age_t or a static member function that returns a default instance of age_t.

⁴Details are described in Section 4.2.1.

⁵The careful reader will notice that the following example does pass the length of the array, even if it is always ignored.

⁶The analogous STL structures instead merely *embed* a pointer in a functor.

⁷Some compilers, including VC, won't notice the difference; however, GCC does care.

⁸This paragraph is intended exclusively as a teaching example, not as a solution for production code. In practice, this issue would be solved by promoting a compiler warning to an error, or by using modern C++ strongly-typed enumerations. However, it's instructive as an example of how a (meta) programmer can bend the C++ syntax to solve small problems.

⁹If X has type const A, wrap will deduce T=const A and pick the second overload, if you carefully implement enum_var_helper.

¹⁰There are online accumulators that *estimate* percentiles with good accuracy, though.

¹¹ You might want to look back at Chapter 5 again.

¹²For example, the maxmin algorithm has a complexity 25% lower than computing max and min in two steps.

¹³More on this in Section **9.3**.

¹⁴Traits would be somehow equivalent to stateless drivers.

¹⁵ See also http://www.gotw.ca/publications/mill18.htm.

¹⁶This has some similarity with the way std::regex works.

 17 The example is written as if all members were public.

The Opaque Type Principle

Template type names can be too complex for the user to use directly, as they may be verbose or they may require very complex syntax. So you should either publish a convenient typedef or allow users to ignore their type altogether.

Plain C is full of opaque types.

In C, a file stream is handled via a pointer to an unknown FILE structure that resides in system memory (the C runtime pre-allocates a small number of these structures). To retrieve the current position in an open file, you call fgetpos (FILE*, fpos_t), passing the file pointer and another opaque type that acts as a bookmark. You can't know or modify the current position but can restore it via a call to fsetpos (FILE*, fpos_t). From the user perspective, an instance of fpos_t is completely opaque. Since only the name is known, the type has no interface except for the default constructor, copy constructor, and assignment.

In the opaque type principle, opaqueness is related only to the *type name*, not to the interface. In other words, the object has an unspecified type and a known interface—it may be an iterator or a functor.

Being the "difficult to write" type, you don't want to store the object but should instead use it immediately, on the creation site.

9.1. Polymorphic Results

Suppose a function performs a computation that produces several results at a time.

You can pack all of them in a polymorphic result and allow the user to select what's needed.

Let's take a simplified example:

```
template <typename iterator_t >
[[???]] average(iterator_t beg, iterator_t end)
{
   typename std::iterator_traits<iterator_t>::value_type total = 0;
   total = std::accumulate(beg, end, total);
   size_t count = std::distance(beg, end);
   return total/count;
}
```

A fixed return type will destroy the partial results, which could be useful. So you can delay the aggregation of the sub-items and change the code like this:

```
template <typename T, typename I>
class opaque average result t
  T total ;
  I count ;
public:
  opaque average result t(T total, I count)
  : total (total), count (count)
  }
  // default result is the average
  operator T () const
  {
     return total /count ;
  T get total() const
     return total ;
  I get count() const
     return count ;
};
template <typename VALUE TYPE, typename iterator t >
opaque average result t<VALUE TYPE, size t> average(iterator t
beg, iterator t end)
  VALUE TYPE total = 0;
  total = std::accumulate(beg, end, total);
  size t count = std::distance(beg, end);
  return opaque average result t<VALUE TYPE, size t>(total,
count);
}
```

Now the client can use the original algorithm in many more ways:

```
std::vector<double> v;
double avg = average<double>(v.begin(), v.end());
```

```
double sum = average<double>(v.begin(), v.end()).get total();
```

Since the return type is opaque, it's not convenient to store the result, but it's easy to pass it to a function template, if needed:¹

```
<???> x = average<double>(v.begin(), v.end());

template <typename T>
void receive(T res, double& avg, double& sum)
{
   avg = res;
   sum = res.get_total();
}

std::vector<double> v;
double avg, sum;
receive(average<double>(v.begin(), v.end()), avg, sum);
```

9.2. Classic Lambda Expressions

Lambda expressions are opaque function objects created on the call site. They combine some elementary pieces with meaningful operators. The resulting functor will later replay the operator sequence on its arguments. For example, given two suitable objects of type "lambda variables" X and Y, then (X+Y) * (X-Y) will be a functor that takes two arguments and returns their sum multiplied by their difference.

It's a good exercise to build a simplified implementation and understand the underlying template techniques. These have been proposed originally by Todd Veldhuizen in his seminal article, "Expression Templates".

You can write code like this: cos(x+2.0) is an expression that returns a functor whose operator() computes cos(x+2.0) given a double x.

```
lambda_reference<const double> X;
std::find_if(..., X<5.0 && X>3.14);
std::transform(..., cos(X+2.0));
lambda_reference<double> Y;
std::for_each(..., Y+=3.14);
lambda_reference<const double, 0> ARG1;
lambda_reference<const double, 1> ARG2;
std::sort(..., (ARG1<ARG2));</pre>
```

You can make the following assumptions. Some will be removed later and some will hopefully

become clearer as you proceed:

- For clarity, T will be a friendly scalar type, double or float, so all the operators are well-defined.
- A lambda expression will receive at most K=4 arguments, *all* of which are the same type T&. In particular:
- lambda_reference<double> and lambda_reference<const double> are different, the latter being a "lambda-const reference to double".
- An expression must contain references to objects of the same type.
- For simplicity, all constants initially have type T and X+2 is considered invalid syntax, because X refers to a double and 2 is an int. Therefore, you have to write X+2.0 (you will learn how to remove this limitation later in this chapter).
- We explicitly try to write functions that look similar, so they easily can be generated with preprocessor macros, even when they are not listed here.

9.2.1. Elementary Lambda Objects

Let's rewrite the fundamental definition here: a lambda object is a functor that is generated with a special syntax (namely, assembling some placeholders with operators). The effect of the functor is to replay the same operators on its actual arguments. For example, if X is one such placeholder, then the expression X+2 produces a functor that takes one argument and returns its argument plus 2.

First, you define an *empty* static interface. Observe that T is not used at the moment, but you will shortly realize why it's necessary.

```
template <typename true_t, typename T>
class lambda
{
protected:
    ~lambda()
    {
    }

public:
    const true_t& true_this() const
    {
      return static_cast<const true_t&>(*this);
    }
};
```

The first (trivial) object is a lambda-constant. That's a functor that returns its constant result, whatever the arguments. Since in particular it's a lambda expression, you derive this from the interface:

```
template <typename T>
class lambda_const : public lambda<lambda_const<T>, T>
{
   typedef const T& R;
   T c_;
public:
   typedef T result_type;
   lambda_const(R c)
   : c_(c)
   {
   }
   result_type operator()(R = T(), R = T(), R = T())
const
   {
     return c_;
   }
};
```

Note that a lambda-constant can take zero or more arguments, but it is a function object, so the invocation must use some form of operator ().

The second object is lambda_reference<T, N>, defined as a functor that takes at least N arguments of type T& and returns the Nth. The choice of accepting T& as an argument implies that lambda reference<T> won't work on a literal:

```
lambda_reference<double> X1;
lambda_reference<const double> Y1;

X1(3.14);  // error: needs double&
Y1(3.14);  // ok: takes and returns const double&
```

The selection of a variable is not trivial. As usual, argument rotation is the preferred technique. Furthermore, since a reference is cheap, this example introduces a technique known as the *duplication of the arguments* in order to reduce the number of overloads. The last argument of operator () is "cloned" so it always passes four items.

```
template <typename T, size_t N = 0>
class lambda_reference: public lambda<lambda_reference<T, N>, T>
{
   static T& apply_k(static_value<size_t,0>, T& x1, T&, T&, T&)
   {
     return x1;
   }

template <size_t K>
   static T& apply_k(static_value<size_t,K>, T& x1, T& x2, T& x3, T&
```

```
\times 4)
   return apply k(static value<size t,K-1>(), x2, x3, x4, x1);
public:
 typedef T& result type;
 result type operator()(T& x1, T& x2, T& x3, T& x4) const
   MXT STATIC ASSERT (N<4);
   return apply k(static value<size t, N>(), x1, x2, x3, x4);
 result type operator()(T& x1, T& x2, T& x3) const
   MXT STATIC ASSERT (N<3);
   return apply k(static value<size t, N>(), x1, x2, x3, x3);
  }
 result type operator()(T& x1, T& x2) const
   MXT STATIC ASSERT (N<2);
   return apply k(static value<size t, N>(), x1, x2, x2, x2);
  }
 result type operator()(T& x1) const
   MXT STATIC ASSERT (N<1);
   return apply k(static value<size t, N>(), x1, x1, x1, x1);
```

9.2.2. Lambda Functions and Operators

};

A unary function F applied to a lambda expression is a functor that returns F applied to the result of the lambda.²

Thanks to the static interface, the implementation can treat *any* lambda expression at once. Also, lambda<X, T> can be stored in an object of type X (and the copy is cheap).

```
template <typename F, typename X, typename T>
class lambda_unary : public lambda<lambda_unary<F,X,T>, T>
{
    X x_;
    F f_;
public:
```

```
lambda unary(const lambda<X,T>& that)
   : x (that.true this())
   }
  typedef typename F::result type result type;
  result type operator()() const
     return f (x ());
   }
  result type operator()(T& x1) const
   {
     return f(x(x1));
   }
  result type operator()(T& x1, T& x2) const
   {
     return f (x (x1, x2));
  result type operator()(T& x1, T& x2, T& x3) const
     return f (x (x1, x2, x3));
  // ...
};
  The previous code builds a functor f , whose operator () is called, but you also need to plug
in global/static member functions. Thus, a small adapter is needed:
template <typename T, T (*F)(T)>
struct unary f wrapper
  typedef T result type;
  T operator()(const T& x) const { return F(x); }
};
  Next, you collect all global functions in the traits class:
template <typename T>
struct unary f library
  static T L abs(T x) { return abs(x); }
   static T L cos(T x) { return cos(x); }
  // ...
```

```
};
  And eventually you start defining functions on lambda objects:
#define LAMBDA ABS TYPE
  lambda unary<unary f wrapper<T, &unary f library<T>::L abs>, X,
T>
template <typename X, typename T>
LAMBDA ABS TYPE abs(const lambda<X, T>& x)
  return LAMBDA ABS TYPE(x);
#define LAMBDA COS TYPE
  lambda unary<unary f wrapper<T, &unary f library<T>::L cos>, X,
T>
template <typename X, typename T>
LAMBDA COS TYPE cos (const lambda<X, T>& x)
  return LAMBDA COS TYPE(x);
  This scheme applies also to unary operators, simply using a different functor.
template <typename T>
struct lambda unary minus
 typedef T result type;
 result type operator()(const T& x) const { return -x; }
};
#define LAMBDA U MINUS TYPE lambda unary<lambda unary minus<T>,
X, T>
template <typename X, typename T>
```

The more features you add, the more complex the return types become, but these are completely hidden from the user.

LAMBDA U MINUS TYPE operator-(const lambda<X, T>& x)

return LAMBDA U MINUS TYPE(x);

{

}

A binary operation, say +, can be defined similarly: (lambda<X1, T> + lambda<X2, T>) is a functor that distributes its arguments to both its addends. So, analogous to the unary case, you will define a specific object to deal with the binary operators, namely lambda binary<X1, F,

X2, T>. In particular *mixed* binary operations, such as lambda< X1, T> + T, are a special case, handled with a promotion of T to lambda constT. template <typename X1, typename F, typename X2, typename T> class lambda binary : public lambda< lambda binary<X1,F,X2,T>, T > X1 x1 ; X2 x2 ; F f ; public: lambda binary(const lambda<X1,T>& x1, const lambda<X2,T>& x2) : x1 (x1.true this()), x2 (x2.true this()) } typedef typename F::result type result type; result type operator()() const return $f_(x1_(), x2_());$ result type operator()(T& x1) const return f (x1 (x1), x2 (x1));result type operator()(T& x1, T& x2) const return f (x1 (x1, x2), x2 (x1, x2));result type operator()(T& x1, T& x2, T& x3) const return f (x1 (x1, x2, x3), x2 (x1, x2, x3));

In this implementation, logical operators will not use short circuit. If T were int, the lambda object X>0 && (1/X)<5 will crash on a division by zero, while the analogous C++ statement returns false.

// ...

};

Arithmetic operators like + can be written as $f_{(x1_{(...)}, x2_{(...)})}$ as previously, but this is incorrect for && and ||, whose workflow is more complex:

}

In the discussion that follows, somewhat sacrificing correctness for clarity, we treat all operators as normal binary predicates, and we leave writing partial specializations of lambda_binary for logical operators from the pseudo-code above as an exercise.

```
Now you define "concrete" binary functions:
template <typename T, T (*f)(T, T)>
struct binary_f_wrapper
  typedef T result type;
  T operator()(const T& x, const T& y) const { return f(x,y); }
};
template <typename T>
struct binary f library
  static T L atan2(T x, T y) { return atan2(x, y); }
  // ...
};
#define ATAN2 T(X1,
X2)
 lambda binary<X1,
             binary f wrapper<T, &binary f library<T>::L atan2>, \
             X2, T>
template <typename X1, typename X2, typename T>
ATAN2 T(X1, X2) atan2(const lambda<X1,T>& L, const lambda<X2,T>&
R)
  return ATAN2 T(X1, X2) (L, R);
template <typename X1, typename T>
ATAN2 T(X1, lambda const<T>) atan2(const lambda<X1,T>& L, const T&
R)
{
  return atan2(L, lambda const<T>(R));
```

```
template <typename T, typename X2>
ATAN2 T(lambda const<T>, X2) atan2(const T& L, const lambda<X2,T>&
R)
{
  return atan2(lambda const<T>(L), R);
  Finally, you need another extension. There are three types of operators
       • Binary predicates, with signature bool F (const T&, const T&)
       • Binary operators, with signature T F (const T&, const T&)
       • Assignments, with signature T& F(T&, const T&)
  This translates to the following C++ code:
enum lambda tag
{
  LAMBDA LOGIC TAG,
  LAMBDA ASSIGNMENT TAG,
  LAMBDA OPERATOR TAG
};
template <typename T, lambda tag TAG>
struct lambda result traits;
template <typename T>
struct lambda result traits<T, LAMBDA ASSIGNMENT TAG>
  typedef T& result type;
  typedef T& first argument type;
  typedef const T& second argument type;
};
template <typename T>
struct lambda result traits<T, LAMBDA OPERATOR TAG>
  typedef T result type;
  typedef const T& first argument type;
  typedef const T& second argument_type;
} ;
template <typename T>
struct lambda result traits<T, LAMBDA LOGIC TAG>
  typedef bool result type;
  typedef const T& first argument type;
```

```
typedef const T& second argument type;
};
  So you can write:
template <typename T>
struct lambda less
  typedef lambda result traits<T, LAMBDA LOGIC TAG> traits t;
  typedef typename traits t::result type result type;
  typedef typename traits t::first argument type arg1 t;
  typedef typename traits t::second argument type arg2 t;
  result type operator()(arg1 t x, arg2 t y) const
     return x < y;
};
template <typename T>
struct lambda plus
  typedef lambda result traits<T, LAMBDA OPERATOR TAG> traits t;
  typedef typename traits t::result type result type;
  typedef typename traits t::first argument type arg1 t;
  typedef typename traits t::second argument type arg2 t;
  result type operator()(arg1 t x, arg2 t y) const
     return x + y;
};
template <typename T>
struct lambda plus eq
  typedef lambda result traits<T, LAMBDA ASSIGNMENT TAG> traits t;
  typedef typename traits t::result type result type;
  typedef typename traits t::first argument type arg1 t;
  typedef typename traits t::second argument type arg2 t;
  result type operator()(arg1 t x, arg2 t y) const
     return x += y;
};
```

These objects have minimal differences.

Logical and standard operators are identical to any other binary function, except the return type (compare with atan2). Here is the implementation of lambda's operator<:⁴

```
#define LSS T(X1,X2)
                             lambda binary<X1, lambda less<T>, X2, T>
template <typename X1, typename X2, typename T>
LSS T(X1,X2) operator<(const lambda<X1,T>& L, const lambda<X2,T>&
R)
{
  return LSS T(X1, X2) (L, R);
}
template <typename X1, typename T>
LSS T(X1, lambda const<T>) operator<(const lambda<X1,T>& L, const
T& R)
  return L < lambda const<T>(R);
template <typename T, typename X2>
LSS T(lambda const<T>, X2) operator<(const T& L, const
lambda < X2, T > \& R)
{
  return lambda constT>(L) < R;
}
  The assignment operators do not allow the third overload, which would correspond to a lambda
expression such as (2.0 += X), a somewhat suspicious C++ statement:
```

```
#define PEQ_T(X1,X2) lambda_binary<X1, lambda_plus_eq<T>, X2,
T>

template <typename X1, typename X2, typename T>
PEQ_T(X1,X2) operator+=(const lambda<X1,T>& L, const lambda<X2,T>&
R)
{
   return PEQ_T(X1,X2) (L,R);
}

template <typename X1, typename T>
PEQ_T(X1,lambda_const<T>) operator+=(const lambda<X1,T>& L, const
T&R)
{
   return L += lambda_const<T>(R);
}
```

Here is a sample that uses all the previous code:

```
lambda reference<double, 1> VX2;
double data[] = \{5,6,4,2,-1\};
std::sort(data, data+5, (VX1<VX2));</pre>
std::for each(data, data+5, VX1 += 3.14);
std::transform(data, data+5, data, VX1 + 3.14);
std::transform(data,data+5, data, 1.0 + VX1);
std::for each(data,data+5, VX1 += cos(VX1));
  Here's a sample that deliberately produces an error, which is still human-readable:
const double cdata[] = \{5, 6, 4, 2, -1\};
// add 3.14 to all the elements of a constant array...
std::for each(cdata, cdata+5, VX1 += 3.14);
error C2664: 'double &lambda binary<X1,F,X2,T>::operator ()(T &)
const':
cannot convert parameter 1 from 'const double' to 'double &'
           with
                  X1=lambda reference<double,0>,
                   F=lambda plus eq<double>,
                   X2=lambda const<double>,
                   T=double
           Conversion loses qualifiers
           see reference to function template instantiation being
compiled
            ' Fn1 std::for each<const
double*,lambda binary<X1,F,X2,T>>( InIt, InIt, Fn1)'
           with
       Fn1=lambda binary<lambda reference<double,0x00>,lambda plus
                             X1=lambda reference<double,0>,
                             F=lambda plus eq<double>,
                             X2=lambda const<double>,
                             T=double,
                             InIt=const double *
           ]
  You would expect the following code to work correctly; instead, it does not compile. The error
```

log may be long and noisy, but it all leads to operator+. The precise error has been isolated here:

lambda reference<double, 0> VX1;

double data[] = $\{5, 6, 4, 2, -1\}$;

```
const double cdata[] = \{5, 6, 4, 2, -1\};
lambda reference < const double > C1;
std::transform(cdata,cdata+5, data, C1 + 1.0);
error: 'lambda binary<X1, lambda plus<T>, lambda const<T>, T>
operator +(const lambda<true_t,T> &,const T &)' :
template parameter 'T' is ambiguous
     could be 'double'
            'const double'
  This issue is equivalent to:
template <typename T>
struct A
};
template <typename T>
void F(A < T >, T)
A<const double> x;
double i=0;
F(x, i); // error: ambiguous call.
          // deduce T=const double from x, but T=double from i
  This is where type traits come in. You take the parameter T only from the lambda expression and
let the type of the constant be dependent. More precisely all mixed operators with an argument of type
const T& should be changed to accept typename lambda constant arg<T>::type.
template <typename T>
struct lambda constant arg
   typedef const T& type;
};
template <typename T>
struct lambda constant arg<const T>
```

The C++ Standard specifies that if a parameter can be deduced from one of the arguments, it's deduced and then substituted in the remaining ones. If the result is feasible, then the deduction is

typedef const T& type;

};

accepted as valid, so in particular in a signature like this:

```
template <typename T>
void F(A<T> x, typename lambda_constant_arg<T>::type i);
```

the only context where T is deducible is the type of x, so ambiguities cannot occur any more. In particular, it's now possible to add constants of any type convertible to T:

```
std::transform(cdata, cdata+5, data, C1 + 1);
// don't need to write C1 + 1.0
```

Note finally that these lambda expressions are not too rigorous about the number of arguments. The only explicit check occurs as a static assertion in lambda_reference.⁵

9.2.3. Refinements

Note that unary and binary operations do contain a copy of the functor representing the operation, but the functor is always default-constructed. You can add a wrapper that embeds any user functor in a lambda expression. Just modify the constructor as follows:

```
public:
    lambda_unary(const lambda<X,T>& that, F f = F())
    : x_(that.true_this()), f_(f)
    {
    }
}
```

This example uses this feature immediately to create a *functor* that takes a functor-on-T and returns a functor-on-lambda:

```
int main()
{
    MyFunctor F;
    lambda_reference<double> X;

    std::transform(data, data+n, data, lambda_wrap[F](3*X+14)); //
    = F(3*X+14)
}

lambda_wrap is a global instance of lambda_wrap_t<void> whose operator[]
```

absorbs a suitable user functor. The choice of [] instead of () gives extra visual clarity, since it avoids confusion with function arguments.

```
template <typename F = void>
class lambda wrap t
  F f _;
public:
  lambda wrap t(F f)
  : f (f)
  template <typename X, typename T>
  lambda unary<F, X, T> operator()(const lambda<X, T>& x) const
     return lambda unary<F, X, T>(x, f);
} ;
template <>
class lambda wrap t<void>
public:
  lambda wrap t(int = 0)
  }
  template <typename F>
  lambda wrap t<F> operator[](F f) const
     return f;
};
const lambda wrap t<void> lambda wrap = 0;
  This is used as in:
struct MyF
  typedef double result type;
  result type operator()(const double& x) const
     return 7*x - 2;
};
```

The same technique can be extended even further to implement the ternary operator (which cannot be overloaded) and the hypothetical syntax could be:

```
if_[CONDITION].then_[X1].else_[X2]
```

The dot that links the statements together shows clearly that the return type of if_[C] is an object whose member then has another operator[], and so on.

9.2.4. Argument and Result Deduction

Loosely speaking, a composite lambda object $G := F(\lambda)$ takes an argument x and returns $F(\lambda(x))$. The *argument type* of G is the argument type of X and the *result type* of Y is the result type of Y.

Up to now, we avoided the problem of defining these types, because they were either fixed or explicitly given.

- The scalar type T in the lambda interface acts as the argument of its operator(). Whenever a function is applied to lambda<X, T>, T is borrowed and plugged in the result, which is say lambda<Y, T>.
- The return type of lambda's operator () instead may vary, so it's published as result_type. For example, lambda_unary<F, X, T> takes T& x from the outside and returns whatever F gives back from the call F(X(x)). F may return a reference to T or bool.

In the process, however, silent casts from bool to T may occur.

For example, the function object abs (C1<C2) takes two arguments of type double. It feeds them to less, which in turn returns bool, but this is promoted again to double before entering abs.

In general, this is the desired behavior:

operator & & can be implemented as a clone of operator <; however, & & would take two Ts, not two bools. In simple cases, this will just work, but in general you'll need more flexibility.

```
(C1<C2) && (C2>C1); // operator&& will promote two bools to double, then return bool
```

You should prescribe only the arguments of lambda_reference and let every lambda object borrow both arguments and results correctly. lambda_reference is in fact the only user-visible object and its type parameter is sufficient to determine the whole functor.

This change also allows you to remove T from the lambda interface:⁶ template <typename X> class lambda protected: ~lambda() } public: const X& true this() const return static cast<const X&>(*this); **}**; template <typename T, size t N = 0> class lambda reference : public lambda< lambda reference<T, N> > public: typedef T& result type; typedef T& argument type; result type operator()(argument type x1) const MXT STATIC ASSERT (N<1); return apply k(static value<size t, N>(), x1, x1, x1, x1); } // ... }; You are going to replace the usage of T in every "wrapping" lambda class with a (meta) function of the result type of the inner object: template <typename F, typename X> class lambda unary : public lambda< lambda unary<F,X> > X x ; F f; public: typedef typename F::result type result type; typedef typename X::argument type argument type;

// ...

};

However, while T is a plain type (maybe const qualified, but never a reference), argument type will often be a reference. So you need a metafunction to remove any qualifier:

```
template <typename T>
struct plain
  typedef T type;
};
template <typename T>
struct plain<T&> : plain<T>
};
template <typename T>
struct plain<const T> : plain<T>
};
template <typename T>
class lambda const : public lambda< lambda const<T> >
  typedef typename plain<T>::type P;
  P c ;
public:
  typedef P result type;
  typedef const P& argument type;
  // ...
};
```

The nasty issue lies in the binary operators, where you have two lambdas, X1 and X2.

Whose argument_type should you borrow? It's easy to see that both types must be inspected, because some deduction must be performed.

For example, if X is a lambda non-const reference, it needs T&. A lambda-constant needs const T&. The expression (X+1.0) is a functor that takes an argument and passes it to both a lambda reference and a lambda-constant, so this should be T&. In general, you need a commutative metafunction that is able to "deduce" a feasible common argument type.

```
template <typename X1, typename F, typename X2>
class lambda_binary : public lambda< lambda_binary<X1,F,X2> >
{
   X1 x1_;
   X2 x2_;
   F f_;
public:
```

```
typedef typename F::result_type result_type;
typedef typename
    deduce_argument<typename X1::argument_type, typename
X2::argument_type>::type
    argument_type;

// ...
};
```

The problem of combining two arbitrary functionals is even deeper. First, the elimination of T makes all the return types more complex. For example, now the lambda_plus object will have to take care of the addition not of two Ts, but of *any* two different results coming from any different lambdas:

```
// before
lambda binary<X1, lambda plus<T>, X2, T>
// after
lambda binary<X1, lambda plus<typename X1::result type, typename
X2::result_type>, X2>
  Furthermore, the return type of "a generic addition" is not known:<sup>7</sup>
template <typename T1, typename T2>
struct lambda plus
  typedef const typename plain<T1>::type& arg1 t; // not
a problem
  a problem
  typedef [[???]] result type;
  result type operator()(arg1 t x, arg2 t y) const
    return x + y;
};
```

So you need another metafunction "deduce result" that takes arg1_t and arg2_t and gives back a suitable type.

Luckily, this issue is solvable by TMP techniques under reasonable assumptions, because you have only a few degrees of freedom. Involved types are T (deduced from lambda_reference and unique in the whole template expression), T&, const T&, and bool.

9.2.5. Deducing Argument Type

You'll now look for a metafunction F that deduces the common argument type. F should satisfy:

• Symmetry: F<T1, T2> := F<T2, T1>

template <typename T>

- The strongest requirement prevails: F < T &, ...> = T &
- const T& and T have the same behavior: F<const T&, ...> = F<T, ...>

Meta-arguments of F are argument types of other lambda objects:

```
F<typename X1::argument_type, typename X2::argument_type>
```

Eventually it suffices that F returns either T& or const T&. The simplest implementation is to reduce both arguments to references. If they have the same underlying type, you should pick the strongest; otherwise, the compiler will give an error:

```
struct as reference
  typedef const T& type;
};
template <typename T>
struct as reference<T&>
  typedef T& type;
};
template <typename T>
struct as_reference<const T&> : as reference<T>
};
template <typename T1, typename T2>
struct deduce argument
: deduce argument<typename as reference<T1>::type, typename
as reference<T2>::type>
};
template <typename T>
struct deduce argument<T&, T&>
  typedef T& type;
};
template <typename T>
struct deduce argument<T&, const T&>
  typedef T& type;
};
```

```
template <typename T>
struct deduce_argument<const T&, T&>
{
   typedef T& type;
};
```

Observe that the specialization $deduce_argument<T\&$, T&>will be used also when T is a const type.

9.2.6. Deducing Result Type

struct deduce result<bool, T>

You can use a similar methodology to write code that deduces the result type. Namely, you will break down the list of cases you want to cover and implement additional metafunctions as needed. First, notice that the expected result of a function call is *never* a reference, so you must start ensuring that at the call location no references are passed:

```
template <typename T1, typename T2>
struct lambda plus
  typedef const typename plain<T1>::type& arg1 t;
  typedef const typename plain<T2>::type& arg2 t;
  typedef
     typename deduce result<typename plain<T1>::type, typename
plain<T2>::type>::type
     result type;
  result type operator()(arg1 t x, arg2 t y) const
     return x + y;
};
  This time you need four specializations:
template <typename T1, typename T2>
struct deduce result;
template <typename T>
struct deduce result<T, bool>
  typedef T type;
};
template <typename T>
```

```
{
  typedef T type;
};

template <typename T>
struct deduce_result<T, T>
{
  typedef T type;
};

template <>
struct deduce_result<bool, bool>
{
  typedef bool type;
};
```

The last specialization is necessary; otherwise, <bool, bool> would match any of the three (with T=bool), so it would be ambiguous.

9.2.7. Static Cast

The limitations of a result/argument deduction may lead to some inconsistency. While a classic addition bool+bool has type int, the addition of Boolean lambda objects returns bool:

```
lambda_reference<const double,0> C1;
lambda_reference<const double,1> C2;

((C1<C2) + (C2<C1))(x, y); // it returns bool</pre>
```

Both (C1<C2) and (C2<C1) have "function signature" bool (const double &, const double &) and so lambda_plus will be instantiated on <bool, bool>. By hypothesis, when arguments are equal, deduce_result<X, X> gives X.

The only way to solve similar issues is a lambda-cast operator. Luckily, it's easy to reproduce the syntax of static cast using a non-deducible template parameter:

```
template <typename T1, typename T2>
struct lambda_cast_t
{
   typedef T2 result_type;
   result_type operator()(const T1& x) const
   {
      return x;
   }
};
#define LAMBDA_CAST_T(T,X) \
```

```
lambda_unary<lambda_cast_t<typename X::result_type, T>, X>

template <typename T, typename X>
LAMBDA_CAST_T(T,X) lambda_cast(const lambda<X>& x)
{
   return x;
}

(lambda_cast<double>(C1<C2)+lambda_cast<double>(C1<C2)) (3.14, 6.28);
// now returns 2.0</pre>
```

9.2.8. Arrays

Todd Veldhuizen pioneered the application of "template expressions" to fast operation on arrays, in order to minimize the use of temporaries.⁸

```
valarray<double> A1 = ...;
valarray<double> A2 = ...;
valarray<double> A3 = 7*A1-4*A2+1;
```

Naive operators will, in general, produce more "copies" of the objects than necessary. The subexpression 7*A1 will return a temporary array, where each element is seven times the corresponding entry in A1; 4*A2 will return another temporary, and so on.

Instead, you can use a lambda-like expression:

```
template <typename X, typename T>
class valarray_interface
{
    // X is the true valarray and T is the scalar
    // ...

public:
    // interface to get the i-th component

T get(size_t i) const
    {
      return true_this().get(i);
    }

size_t size() const
    {
      return true_this().size();
    }

operator valarray<T>() const
```

```
{
    valarray<T> result(size());
    for (size_t i=0; i<size(); ++i)
        result[i] = get(i);

    return result;
}
};</pre>
```

The interface can be cast to a real valarray. This cast triggers the creation of *one* temporary object, which is filled componentwise (which is the most efficient way).

The product valarray<T> * T returns a valarray_binary_op< valarray<T>, std::multiplies<T>, scalar_wrapper<T>, and T>. This object contains a const reference to the original valarray.

```
template <typename VA1, typename F, typename VA2, typename T>
class valarray_binary_op
: public valarray_interface< valarray_binary_op<VA1,F,VA2,T> >
{
   const VA1& va1_;
   const VA2& va2_;
   F op_;

public:
   // ...

   T get(size_t i) const
   {
     return op_(va1_.get(i), va2_.get(i));
   }
};
```

The key optimization for successfully using expression templates with complex objects, such as arrays, is carefully using const references:

```
const VA1& va1_;
const VA2& va2_;
```

A const reference is generally fine, since it binds to temporaries, but it will not prevent the referenced object from dying.

For example, (A*7) +B will produce one temporary (A*7), and another object that has a const reference to it, and a const reference to B. Since A*7 is alive "just in that line of code", if one could store the expression and evaluate it later, it would crash the program.

You may actually want to use traits to determine a suitable storage type. If VA1 is valarray<T>, then it's convenient to use const VA1&. If VA1 is simply a scalar, const VA1 is safer.

To sum up, the line

```
valarray<double> A3 = A1*7;
```

will magically trigger the componentwise evaluation of the template expression on the right, using either a cast operator in the interface, or—even better—a dedicated template constructor/assignment in valarray<T>.9

The cast operator is not easy to remove. Since A1*7 is expected to be a valarray, it might even be used as a valarray, say writing (A1*7) [3] or even (A1*7) . resize (n). This implies that valarray and valarray interface should be very similar, when feasible.

Another advantage of the static interface approach is that many different objects can behave as a fake valarray. As an equivalent of lambda_const, you can let a scalar c act as the array [c, c, ..., c]:

```
template <typename T>
class scalar_wrapper
: public valarray_interface< scalar_wrapper<T> >
{
    T c_;
    size_t size_;

public:
    scalar_wrapper(T c, size_t size)
    : c_(c), size_(size)
    {
    }

    T get(size_t i) const
    {
        return c_;
    }
};
```

9.3. Creative Syntax

This section is devoted to exploiting template syntax tricks, such as operator overloads, to express concepts that differ from the standard meaning.

Some operators convey a natural associativity; the simplest examples are sequences connected with +, <<, and comma:

```
std:string s = "hello";
std:string r = s + ' ' + "world" + '!';
std::ofstrean o("hello.txt");
```

```
o << s << ' ' << "world" << '!';
int a = 1,2,3,4,5,6,7;
```

The user expects these operators to be able to form chains of arbitrary length. Additionally, operator[] and operator() can sometimes have a similar meaning; in particular, the former should be used when the length of the chain is *fixed*:

You can exploit this syntax by writing operators that consume the first argument and return *something* that can handle the remaining chain. Consider the line:

```
std::cout << a << b << c;
```

This expression has the form: F(F(F(cout, a), b), c), so F(cout, a) should return an object X such that there exists an overload of F that accepts X and b, and so on. In the simplest case, F(cout, a) just returns cout.

You are now going to cover this argument in full detail.

9.3.1. Argument Chains with () and []

Sometimes operator () is used to form chains, starting from a function object.

Let's analyze some hypothetical code:

Given that f is a function taking N arguments, you can guess the following facts:

- bind_to(f) returns an object with two different operator().
- The first form takes an expression whose syntax is argument<K>(x) and returns a functor that fixes x as the Kth argument of f. This first form can be invoked repeatedly to fix several arguments in the same statement.
- The second operator () takes all the remaining arguments at a time and evaluates the function.

```
\frac{\left| \frac{\text{bind_to}(f)}{\text{nothing fixed}} (\text{argument} < 0 > (17)) (\text{argument} < 2 > ('p')) \frac{(3.14, "hello")}{\text{all the remaining arguments}} \right|
\frac{\text{first argument fixed}}{\text{f(17,*,*,*)}}
first and third arguments fixed f(17,*,*,'p',*)
```

Another paradigmatic example is a function that needs several objects (usually functors or accessors, but there's no formal requirement), which you don't want to mix with the other arguments, because either:

- There are too many: $F(\ldots, x1, x2, x3, x4\ldots)$.
- They cannot be sorted by "decreasing probability of having the default value changed" and the caller may have to put arbitrary values in unspecified positions. F(..., X1 x1 = X1(), X2 x2 = X2()...) may need to be invoked as F(..., X1(), X2(), ..., x7, X8(), ...).
- Each object is associated with a distinct template parameter, say X1, X2..., so a function call with two arguments swapped by mistake will likely compile. 10

To illustrate the situation, let's pick an algorithm that needs *three* objects: a less-comparator, a unary predicate, and a logger:

```
template <typename iterator_t, typename less_t, typename pred_t,
typename logger_t>
void MyFunc(iterator_t b, iterator_t e, less_t less, pred_t p,
logger_t& out)
{
```

```
std::sort(b, e, less);
   iterator t i = std::find if(b, e, p);
   if (i != e)
     out << *i;
  However, all these arguments would have a default type, namely std::ostream as logger (and
std::cout as a default value for out) and the following two types:
struct basic comparator
  template <typename T>
  bool operator()(const T& lhs, const T& rhs) const
   { return lhs < rhs; }
};
struct accept first
  template <typename T>
  bool operator()(const T&) const { return true; }
};
  You might often want to change one of those, maybe the last. However, it's difficult to provide
overloads, because arguments cannot be distinguished on their type:
template <typename iterator t, typename less t >
void MyFunc(iterator t b, iterator t e, less t less)
{ . . . }
template <typename iterator t, typename logger t>
void MyFunc(iterator t b, iterator t e, logger t& out)
{ . . . }
// ambiguous: these functions will generate errors, if given
a named variable as 3rd argument
  So you use the argument pack technique. First, you tag the arguments.
enum { LESS, UNARY P, LOGGER };
template <size t CODE, typename T = void>
struct argument
  T arg;
  argument (const T& that)
   : arg(that)
   {
```

}

};

```
template <size t CODE>
struct argument<CODE, void>
  argument(int = 0)
  template <typename T>
  argument<CODE, T> operator=(const T& that) const
     return that;
  argument<CODE, std::ostream&> operator=(std::ostream& that)
const
     return that;
};
  Then you provide named global constants:
const argument<LESS> comparator = 0;
const argument<UNARY P> acceptance = 0;
const argument<LOGGER> logger = 0;
template <typename T1, typename T2, typename T3>
struct argument pack
  T1 first;
  T2 second;
  T3 third;
  argument pack(int = 0)
   {
  }
  argument pack(T1 a1, T2 a2, T3 a3)
   : first(a1), second(a2), third(a3)
  argument pack::operator[] takes an argument < N, T > and replaces its Nth template
argument with T:
       template <typename T>
       argument pack<T, T2, T3> operator[] (const argument<0, T>& X) const
```

```
return argument pack<T, T2, T3>(x.arg, second,
third);
       template <typename T>
       argument pack<T1, T, T3> operator[](const argument<1, T>& x) const
       {
          return argument pack<T1, T, T3>(first, x.arg, third);
       }
       template <typename T>
       argument pack<T1, T2, T> operator[](const argument<2, T>& x) const
              return argument pack<T1, T2, T>(first, second,
x.arg);
};
  This code introduces a global constant named where and overloads the original function twice
(regardless of the actual number of parameters):
typedef argument pack<basic comparator, accept first,
std::ostream&> pack t;
// note: a global variable called "where"
static const pack t where (basic comparator(), accept first(),
std::cout);
template <typename iterator t, typename T1, typename T2, typename
T3>
void MyFunc(iterator t b, iterator t e, const
argument pack<T1,T2,T3> a)
  return MyFunc(b, e, a.first, a.second, a.third);
template <typename iterator t >
void MyFunc(iterator t b, iterator t e)
{
  return MyFunc(b, e, where);
}
  So now it's possible to write:
MyFunc(v.begin(), v.end(), where[logger=std::clog]);
MyFunc(v.begin(), v.end(), where[logger=std::cerr]
[comparator=greater<int>()]);
  logger is a constant of type argument<2, void>, which gets upgraded to
```

argument<2, std::ostream&>. This instance replaces the third template parameter of pack_t with std::ostream& and the value of pack_t::third with a reference to std::cerr.

Observe that the code shown in this section is not generic, but it's strongly tied to the specific function call. However, complex functions that require argument packs should generally be just a few per project.

9.4. The Growing Object Concept

Let's start with an example. String sum has an expected cost of a memory reallocation¹¹:

```
template <typename T>
std::string operator+(std::string s, const T& x)
{
    // estimate the length of x when converted to string;
    // ensure s.capacity() is large enough;
    // append a representation of x to the end of s;
    return s;
}
```

If there are multiple sums on the same line, evidently, the compiler knows the sequence of arguments:

```
std::string s = "hello";
std::string r = s + ' ' + "world!";

// repeated invocation of operator+ with arguments: char, const char*
// may cause multiple memory allocations
```

So you will want to:

- Collect all the arguments at once and sum their lengths
- Execute a single memory allocation
- Traverse the sequence of arguments again and concatenate them

The *growing object* is a pattern that allows traversing a C++ expression before execution. The idea of the technique is to inject in the expression a proxy with special operators that "absorb" all the subsequent arguments.

The *proxy* is a temporary agglomerate object whose operators make it "grow" including references to their arguments. Finally, when growth is complete, the object can process *all* arguments at once and transform them into the desired result.

Thus in the previous example, s+' ' is not a string, but a proxy that contains a reference to s and a char. This object grows when "world" is added, so s+' '+"world" contains also a

const char*.

Informally, a growing object is implemented as a pair containing the previous state of the object and some new tiny data (say, a reference). Additionally, there are three possible variants of "pair":

- A class with two members: a reference to the previous growing object and a tiny object
- A class with two members: a copy of the previous growing object and a tiny object
- A class that derives from the previous growing object, with a tiny object as the only member

In pseudo-template notation, the three different models can be written:

```
template <...>
class G1<N>
{
    const G1<N-1>& prev_;
    T& data_;
};

template <...>
class G2<N>
{
    G2<N-1> prev_;
    T& data_;
};

template <...>
class G3<N> : public G3<N-1>
{
    T& data_;
};
```

The first is the fastest to build, because augmenting a temporary object G1 with new data involves no copy, but the lifetime of G1 is the shortest possible. The other types have similar complexity, since their construction involves copying Gj < N-1 > anyway, but they slightly differ in the natural behavior.

The great advantage of G1 is that both constructors and destructors run exactly once and in order. Instead, to create a G2<N>, you must produce two copies of G2<N-1>, three copies of G2<N-2>..., K+1 copies of G2<N-K>..., and so on.

This is especially important because you might need G<N> to run some code when the growth is complete and the destructor of G<N> would be one of the options.

Any of these Gj contains references, so for example no growing object can be thrown.

Furthermore, there are some known recursive patterns in computing the result:

• *Inward link*: G<N> either computes the result directly, or it delegates G<N-1>, passing information "inward":

• Outward link: G<N> asks recursively for a result from G<N-1> and post-processes it.

```
result do_it()
{
    result temp = prev_.do_it();
    return modify(temp);
}
```

• *Direct access*: G<N> computes J and asks G<J> for a result. This pattern has a different implementation for inheritance-based growing objects.

```
template <...>
class G1<N>
{
    result do_it_myself(static_value<int, 0>)
    {
        // really do it
    }

    template <int K>
    result do_it_myself(static_value<int, K>)
    {
        return

prev_.do_it_myself(static_value<int, K-1>());
    }
}
```

```
public:
       result do it()
              static const int J = [...];
              return do it myself(static value<int, J>
());
       }
};
template <...>
class G3<N> : G3<N-1>
       result do it myself()
            // ...
public:
       result do it()
              static const int J = ...;
              return static cast<growing<J>&>
(*this).do it myself();
};
```

9.4.1. String Concatenation

You implement the first growing object with a sequence of *agglomerates* (see Section 3.6.8).

Since objects involved in a single statement live at least until the end of the expression, you can think of an agglomeration of const references. The expression (string+T1) +T2 should not return a string, but rather a structure containing references to the arguments (or copies, if they are small).¹²

```
template <typename T1, typename T2>
class agglomerate;

template <typename T>
agglomerate<string, const T&> operator+(const string&, const T&);

template <typename T1, typename T2, typename T>
agglomerate<agglomerate<T1, T2>, const T&>
    operator+(const agglomerate<T1, T2>, const T&);
```

So the sum in the prototype example below would return agglomerate< agglomerate<string, char>, const char*>:

```
std::string s = "hello";
std::string r = s + ' ' + "world!";
```

Eventually, all the work is done by a cast operator, which converts agglomerate to string:

- Sum the lengths of this->first and this->second (first is another agglomerate or a string, so both have a size () function; second is a reference to the new argument).
- Allocate a string of the right size.

// using namespace std;

• Append all the objects to the end of the string, knowing that internally no reallocation will occur.

Note that the agglomerates are built in reverse order, with respect to arguments; that is, the object that executes the conversion holds the last argument. So, it has to dump its agglomerate member *before* its argument member.

```
template <typename T, bool SMALL = (sizeof(T) <= sizeof(void*))>
struct storage traits;
template <typename T>
struct storage traits<T, true>
  typedef const T type;
};
template <typename T>
struct storage traits<T, false>
{
  typedef const T& type;
};
// assume that T1 is string or another agglomerate
// and T2 is one of: char, const char*, std::string
template <typename T1, typename T2>
class agglomerate
{
  T1 first;
  typename storage traits<T2>::type second;
  void write(string& result) const
  {
     // member selection based on the type of 'first'
     write (result, &first);
  }
```

```
template <typename T>
  void write(string& result, const T*) const
     // if we get here, T is an agglomerate, so write recursively:
     // mind the order of functions
     first.write(result);
     result += this->second;
  }
  void write(string& result, const string*) const
     // recursion terminator:
     // 'first' is a string, the head of the chain of arguments
     result = first;
  }
  size t size()
     return first.size() + estimate length(this->second);
  static size t estimate length(char)
     return 1;
  static size t estimate length(const char* const x)
  {
     return strlen(x);
  }
  static size t estimate length(const string& s)
  {
     return s.size();
  }
public:
  operator string() const
     string result;
     result.reserve(size());
     write (result);
     return result; // NVRO
};
```

The first enhancement allows accumulating information in a *single pass* through the chain:

```
void write(string& result, size t length = 0) const
  write(result, &first, length + estimate length(this->second));
template <typename T>
void write(string& result, const T*, size t length) const
  first.write(result, length);
  result += this->second;
}
void write(string& result, const string*, size t length) const
  result.reserve(length);
  result = first;
}
operator string() const
  string result;
  write (result);
  return result;
std::string s = "hello";
std::string r = s + ' ' + "world!";
```

In classic C++, each call to string::operator+ returns a different temporary object, which is simply copied. So the initial example produces two intermediate strings: namely t_1 ="hello" and t_2 ="hello world!". Since each temporary involves a copy, this has quadratic complexity.

With C++0x language extensions, std::string is a moveable object. In other words, its operators will detect when an argument is a temporary object and allow you to steal or reuse its resources. So the previous code might actually call two different sum operators. The first produces a temporary anyway (because you are not allowed to steal from local variable 's'); the second detects the temporary and reuses its memory.

Conceptually, the implementation could look like this:

```
string operator+(const string& s, char c)
{
    string result(s);
    return result += c;
}
```

```
string operator+(string&& tmp, const char* c)
{
    string result;
    result.swap(tmp);
    return result += c;
}
```

Where the notation string&& denotes a reference to temporary. Even more simply:

```
string operator+(string s, char c)
{
    return s += c;
}

string operator+(string&& tmp, const char* c)
{
    return tmp += c;
}
```

In other words, C++0x string sum is conceptually similar to: 14

```
std::string s = "hello";
std::string r = s;
r += ' ';
r += "world!";
```

But a growing object performs even better, being equivalent to:

```
std::string s = "hello";
std::string r;
r.reserve(s.size()+1+strlen("world!");
r += s;
r += ' ';
r += "world!";
```

So C++0x extensions alone will not achieve a better performance than a growing object.

9.4.2. Mutable Growing Objects

```
A growing object may be used to provide enhanced assertions: 15
std::string s1, s2;
SMART ASSERT(s1.empty() && s2.empty())(s1)(s2);
Assertion failed in matrix.cpp: 879412:
Expression: 's1.empty() && s2.empty()'
Values: s1 = "Wake up, Neo"
       s2 = "It's time to reload."
  This code may be implemented with a plain chainable operator ():
class console assert
   std::ostream& out ;
public:
   console assert(const char*, std::ostream& out);
  console assert& operator()(const std::string& s) const
   {
     out << "Value = " << s << std::endl;
     return *this;
  console assert& operator()(int i) const;
  console assert& operator()(double x) const;
   // ...
};
#define SMART ASSERT(expr) \
     if (expr) {} else console assert(#expr, std::cerr)
  This macro starts an argument chain using operator (), and since it's not a growing object,
arguments must be used immediately. But you could have a more intricate "lazy" approach: 16
template <typename T1, typename T2>
class console assert
  const T1& ref;
  const T2& next;
  mutable bool run ;
public:
   console assert (const T1& r, const T2& n)
   : ref (r), next (n), run (false)
   { }
```

```
std::ostream& print() const
     std::ostream& out = next .print();
     out << "Value = " << ref << std::endl;
     run = true;
     return out;
  }
  template <typename X>
  console assert<X, console assert<T1, T2> >
     operator()(const X& x) const
     return console assert<X, console assert<T1, T2> >(x, *this);
  ~console assert()
     if (!run )
       print();
};
template < >
class console assert<void, void>
 std::ostream& out ;
public:
 console assert(const char* msg, std::ostream& out)
  : out (out << "Assertion failed: " << msg << std::endl)</pre>
  {
 std::ostream& print() const
  {
    return out ;
 template <typename X>
 console assert<X, console assert<void, void> >
    operator()(const X& x) const
 {
    return console assert<X, console assert<void, void> >(x,
*this);
};
#define SMART ASSERT(expr) \
     if (expr) {} else console assert<void, void>(#expr,
```

```
std::cerr)
```

The previous sample shows that it's possible to modify the growing object from inside, stealing or passing resources through the members.

In particular, a step-by-step code expansion yields the following:

```
SMART_ASSERT(sl.empty() && s2.empty())(s1)(s2);

if (sl.empty() && s2.empty())
    {}

else
    console_assert<void, void>("sl.empty() && s2.empty()",
    std::cerr)(s1)(s2);
// constructor of console_assert<void, void>
```

If assertion is false, three nested temporaries are created:

```
T0 console_assert<void,void>
T1 console_assert<string,console_assert<void,void>>
T2 console assert<string,console assert<string,console assert<void,void>>>
```

- T2 is created and immediately destroyed. Since run_ is false, it invokes print.
- print calls next_.print.
 - T0 passes its stream to T1.
 - T1 prints its message, sets run_=true, and passes the stream up to T2.
- T2 prints its message and dies.
- T1 is destroyed, but since run is true, it stays silent.
- T0 is destroyed.

A specialization such as console_assert<void, void> is called the *chain starter*. Its interface might be significantly different from the general template.

The interface of the growing object usually does not depend on the number of arguments that were glued together.¹⁷

9.4.3. More Growing Objects

Generalizing the pattern, given a type container C, you implement a generic agglomerate chain<traits, C>. The only public constructor lies in the chain starter and the user can sum any chain and an argument.

For simplicity, the decision about how to store arguments in the chain (by copy or by reference) is

```
given by a global policy: 18
template <typename T>
struct storage traits
  typedef const T& type;
};
template <typename T>
struct storage traits<T*>
  typedef T* type;
};
template <typename T>
struct storage traits<const T*>
  typedef const T* type;
};
template <>
struct storage traits<char>
  typedef char type;
};
```

During the "agglomeration," a chain of length N+1 is generated by a chain of length N and a new argument. The new chain stores both and combines some new piece of information (for example, it sums the estimated length of the old chain with the expected length of the new argument).

Eventually, this piece of information is sent to a "target object," which then receives all the arguments in some order.

Since all these actions are parametric, you can combine them in a traits class:

- update collects information from arguments, one at a time.
- dispatch sends the cumulative information to the target object.
- transmit sends actual arguments to the target object.

```
struct chain_traits
{
   static const bool FORWARD = true;

   struct information_type
   {
       // ...
   };

   typedef information_type& reference;
```

```
typedef ... target_type;

template <typename ARGUMENT_T>
   static void update(const ARGUMENT_T&, information_type&);

static void dispatch(target_type&, const information_type&);

template <typename ARGUMENT_T>
   static void transmit(target_type&, const ARGUMENT_T&);
};
```

Update will be called automatically during object growth; dispatch and transmit will be called lazily if the chain is cast to or injected in a target type.

First you implement the empty chain.

Analogous to the stream reference in Section 9.4.2 above, this class will store just the common information. Additional layers of the growing object will refer to it using traits::reference.

```
template <typename traits t, typename C = empty>
class chain;
template <typename traits t>
class chain<traits t, empty>
  template <typename ANY1, typename ANY2>
  friend class chain;
  typedef typename traits t::information type information type;
  typedef typename traits t::target type target type;
  information type info;
  void dump(target type&) const
public:
  explicit chain(const information type& i = information type())
  : info (i)
#define PLUS T \
      chain<traits t, typename push_front<empty, T>::type>
  template <typename T>
  PLUS T operator+(const T& x) const
     return PLUS T(x, *this);
```

```
const chain& operator >> (target_type& x) const
{
    x = target_type();
    return *this;
}

operator target_type() const
{
    return target_type();
}
```

A nonempty chain instead contains:

- A data member of type front<C>, stored as a storage_traits<local_t>::type.
- A chain of type chain<pop_front<C>>, stored by const reference. Since C is nonempty, you can safely pop_front it.
- A reference to the information object. Storage of information_type is traits-dependent. It may be a copy (when traits_t::reference and traits t::information type are the same) or a true reference.

The private constructor invoked by operator+ first copies the information carried by the tail chain, then it updates it with the new argument.

```
template <typename traits_t, typename C>
class chain
{
  template <typename ANY1, typename ANY2>
  friend class chain;

  typedef typename traits_t::target_type target_type;
  typedef typename front<C>::type local_t;
  typedef chain<traits_t, typename pop_front<C>::type> tail_t;

  typename storage_traits<local_t>::type obj_;
  typename traits_t::reference info_;
  const tail_t& tail_;

  void dump(target_type& x) const
  {
    tail_.dump(x);
    traits_t::transmit(x, obj_);
  }
}
```

```
chain (const local t& x, const tail t& t)
  : obj_(x), tail_(t), info_(t.info_)
     traits t::update(x, info);
public:
  template <typename T>
  chain<traits t, typename push front<C, T>::type> operator+(const
T& x) const
  {
     typedef
       chain<traits t, typename push front<C, T>::type> result t;
     return result t(x, *this);
  }
  const chain& operator >> (target type& x) const
     traits t::dispatch(x, info);
     dump(x);
     return *this;
  }
  operator target type() const
     target type x;
     *this >> x;
     return x;
};
```

The private dump member function is responsible for transmitting recursively all arguments to the target. Note that you can make the traversal parametric and reverse it with a simple Boolean:

```
void dump(target_type& x) const
{
  if (traits_t::FORWARD)
  {
    tail_.dump(x);
    traits_t::transmit(x, obj_);
  }
  else
  {
    traits_t::transmit(x, obj_);
    tail_.dump(x);
  }
}
```

Finally, you show an outline of the traits class for string concatenation:

```
struct string chain traits
  static const bool FORWARD = true;
  typedef size t information type;
  typedef size t reference;
  typedef std::string target type;
  template <typename ARGUMENT T>
  static void update(const ARGUMENT T& x, information type& s)
  {
    s += estimate length(x);
  }
  static void dispatch(target type& x, const information type s)
     x.reserve(x.size()+s);
  }
  template <typename ARGUMENT T>
  static void transmit(target_type& x, const ARGUMENT_T& y)
    x += y;
};
typedef chain<string chain traits> begin chain;
std::string q = "lo ";
std::string s = (begin chain() + "hel" + q + 'w' + "orld!");
```

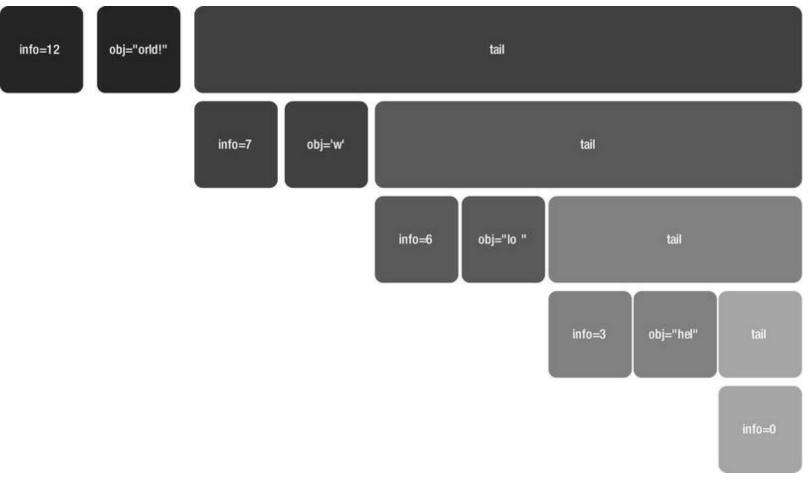


Figure 9-1. Chain diagram. Objects are constructed from bottom to top

- Since you are not allowed to modify std::string, you have to start the chain explicitly with a default-constructed object.
- Code that runs only once before a chain starts can be put in information_type constructor. Then you can begin the chain with begin chain (argument).
- The *storage policy* is a place where custom code may be transparently plugged in to perform conversions. For example, to speed up the int-to-string conversion, you could write:

```
template <>
struct storage_traits<int>
{
   class type
   {
      char data_[2+sizeof(int)*5/2];19

   public:
      type(const int i)
      {
            // perform the conversion here
            _itoa(i, data_, 10);
      }
}
```

```
operator const char* () const
{
    return data_;
}
};
```

9.4.4. Chain Destruction

It may be possible to write custom code in the *chain destructor*.

Since you have only one copy of each chain (they are linked by const references), chain pieces are constructed in order from the first argument to the last. They will be destroyed in the reverse order. You have an opportunity to execute some finalization action at the end of the statement.

```
~chain()
{
    traits_t::finalize(obj_, info_);
}

std::string s;
std::string q = "lo ";
(begin_chain() + "hel" + q + 'w' + "orld!") >> s;
```

Then the leftmost object will append "hello world!" to s with at most a single reallocation. Finally, the destructors will run finalize in reverse order (from left to right).

If chains are stored by value, the order of destruction is fixed (first the object, then its members). But there will be multiple copies of each sub-chain (namely, all the temporaries returned by operator+). Evidently, if C1 holds a copy of C0 and C2 holds a copy of C1, there are three copies of C0 and so, without some additional work, you will not know which sub-chain is being destroyed.

9.4.5. Variations of the Growing Object

If you have to add growing objects to a read-only class (as std::string should be), instead of inserting manually a chain starter, you can:

- Replace the chain starter with a global function that processes the first argument (this is equivalent to promoting the empty chain's operator+ to a function).
- Switch to operator() for concatenation (this makes the bracket syntax uniform).

```
template <typename traits_t, typename T>
chain<traits_t,typename push_front<empty,T>::type> concatenate(const T&
x)
```

```
{
   typedef chain<traits_t, typename push_front<empty, T>::type>
result_t;
   return result_t(x, chain<traits_t>());
}

std::string s = concatenate("hello")(' ')("world");
```

Another variation involves the extraction of the result. Sometimes the cast operator is not desirable. You may decide to replace both = and + with the stream insertion syntax, so you'd write:

```
std::string s;
s << begin_chain() << "hello" << ' ' << "world";</pre>
```

This is feasible, but it requires some trick to break the associativity, because the language rules will make the compiler execute:

In the old approach, the result was the last piece of information; now it's the first. So you have to modify the chain and carry it around. You store a pointer to the result in the empty chain so that it can be read only once. The unusual operator<< fills this pointer and then returns its *second* argument, not the first; this is the associativity-breaker.

This section shows only briefly the differences from the previous implementation:

```
template <typename traits_t, typename C = empty>
class chain;

template <typename traits_t>
class chain<traits_t, empty>
{
    // ...
    mutable target_type* result_;

public:
    // ...

const chain& bind_to(target_type& x) const
    {
      result_ = &x;
      return *this;
    }

    target_type* release_target() const
```

```
{
     target type* const t = result ;
     result = 0;
     return t;
};
template <typename traits t>
const chain<traits t>& operator<<(typename traits t::target type&</pre>
                             const chain<traits t>& c)
{
  return c.bind to(&x);
template <typename traits t, typename C>
class chain
 // ...
 target type* release target() const
   return tail .release target();
public:
 template <typename T>
 chain<traits t, typename push front<C,T>::type> operator<<(const
T& x) const
   typedef chain<traits t, typename push front<C, T>::type>
result t;
   return result t(x, *this);
  }
 ~chain()
    if (target type* t = release target())
      dump(*t);
```

The last object in the chain will be destroyed first, and it will be the only one to succeed in release target.

9.5. Streams

};

As introduced in the previous section, the stream insertion syntax is one of the most uniform, so it's visually clear yet flexible and open to customizations.

9.5.1. Custom Manipulators and Stream Insertion

Say you want to print a bitstring (see Section 5.2.3) in the C++ way, via stream insertion. A bitstring implements many static interfaces at the same time.

```
class bitstring
: public pseudo_array<bitstring, bit_tag>
, public pseudo_array<bitstring, nibble_tag>
, public pseudo_array<bitstring, byte_tag>
{ ... };
```

How do you decide which of the interfaces should send its data to the stream? In other words, how can you elegantly select between bit-wise, byte-wise, and nibble-wise printing?

Recall that a manipulator is an object that flows in the stream, takes the stream object, and modifies its state:²⁰

```
using namespace std;
ostream& flush(ostream& o)
{
    // flush the stream, then...
    return o;
}

// a manipulator is a function pointer that takes and returns
a stream by reference
typedef ostream& (*manip_t) (ostream&);
ostream& operator<<(ostream& o, manip_t manip)
{
    manip(o);
    return o;
}

cout << flush << "Hello World!";</pre>
```

Note that, while some objects modify the state of the stream permanently, in general the effect of a manipulator insertion is lost after the next insertion. In the previous code, cout will need reflushing after the insertion of the string.

However, nothing prevents the manipulator from returning an entirely different stream. Being part of a subexpression, the original stream is surely alive, so it can be wrapped in a shell that intercepts any further call to operator<<.

```
class autoflush_t
```

```
{
   ostream& ref;
public:
   autoflush t(ostream& r)
   : ref(r)
   { }
   template <typename T>
   autoflush t& operator<<(const T& x)</pre>
   {
      ref << x << flush;
      return *this;
   operator ostream& () const
      return ref;
};
autoflush t* autoflush() { return 0; }
inline autoflush t operator << (ostream @ out, autoflush t* (*)())
   return autoflush t(out);
cout << autoflush << "Hello" << ' ' << "World";</pre>
   All insertions after autoflush are actually calls to autoflush t::operator<<, not to
std::ostream.
   Note also that the code generates a unique signature for the manipulator with the proxy itself.
   A stream proxy need not be persistent. It may implement its own special insertion and a generic
operator that "unwraps" the stream again if the next object is not what's expected.
   Suppose you have a special formatter for double:
class proxy
   ostream& os ;
public:
   explicit proxy(ostream& os)
   : os (os)
   ostream& operator << (const double x) const
```

```
{
     // do the actual work here, finally clear the effect of the
     // manipulator, unwrapping the stream
     return os ;
  }
  // the default insertion simply reveals the enclosed stream
  template <typename T>
  ostream& operator<<(const T& x) const
     return os << x;
};
proxy* special numeric() { return 0; }
inline proxy operator << (ostream @ os, proxy* (*)())
  return proxy(os);
}
cout
  << special numeric << 3.14 // ok, will format a double
  << special numeric << "hello"; // ok, the manipulator has no
effect
```

If instead the template operator<< is omitted, a double will be *required* after the manipulator. To sum up, by altering the return type of operator<<, you can write manipulators that:

- Affect only the next insertion, as long as an instance of X is inserted; otherwise, they are ignored.
- Affect only the next insertion and require the insertion of X immediately thereafter; otherwise, there's a compiler error.
- Affect all the next insertions until the end of the subexpression.
- Affect all the next insertions until X is inserted:

```
template <typename any_t>
proxy_dumper& operator<<(const any_t& x) const
{
  os_ << x;
  return *this;
}</pre>
```

This is exactly the solution you need for bitstring; treating the static interface type tags as manipulators. Insertion returns a template proxy that formats the next bitstring according to the (statically known) type tag, using a suitable function from the static interface itself.

```
using std::ostream;
template <typename digit t>
class bistring_stream_proxy
  ostream& os ;
public:
  bistring stream proxy(ostream& os)
  : os (os)
  ostream& operator<<(const pseudo array<bitstring, digit t>& b)
const
     b.dump(os);
     return os ;
  }
  template <typename any t>
  ostream& operator << (const any t& x) const
     return os << x;
};
inline bistring stream proxy<bit t> operator<<(ostream& o, bit t)</pre>
  return bistring stream proxy<bit t>(o);
inline bistring stream proxy<octet t> operator<<(ostream& o,
octet t)
  return bistring stream proxy<octet t>(o);
inline bistring stream proxy<nibble t> operator<< (ostream& o,
nibble t)
  return bistring stream proxy<nibble t>(o);
```

9.5.2. Range Insertion with a Growing Object

Another exercise is the insertion of a range into a stream. You need a custom item to start a chain:

```
cout << range << begin << end;</pre>
```

The first proxy (returned by std::cout << range) takes an iterator and grows (see the previous section). The insertion of a second iterator of the same kind triggers the full dump:

```
template <typename iterator t = void*>
class range t
  std::ostream& ref ;
  iterator t begin ;
public:
  explicit range t(std::ostream& ref)
  : ref (ref), begin ()
  }
  range t(range t<> r, iterator t i)
     : ref (r.ref ), begin (i)
  {
  }
  std::ostream& operator<<(iterator t end)</pre>
     while (begin != end)
        ref << *(begin_++);
     return ref ;
  }
  std::ostream& operator<<(size t count)</pre>
     while (count--)
        ref << *(begin ++);
     return ref ;
  }
};
range t<>* range() { return 0; }
inline range t<> operator<<(std::ostream& os, range<>* (*)())
  return range t<>(os);
template <typename iterator t>
inline range t<iterator t> operator<<(range t<> r, iterator t
begin)
```

```
{
  return range t<iterator t>(r, begin);
  The range proxy accepts a range represented either by [begin...end) or by [begin, N).
In theory, it's possible to specialize even more:
template <typename iterator t = void*>
class range t
private:
 // ...
 void insert(iterator t end, std::random access iterator tag)
    // faster algorithm here
public:
 // ...
  std::ostream& operator<<(iterator t end)</pre>
    insert (end, typename
iterator traits<iterator t>::iterator category());
    return ref ;
};
```

9.6. Comma Chains

The comma operator is sometimes overloaded together with assignment to get some form of lazy/iterative initialization. This mimics the common C array initialization syntax:

```
int data[] = { 1,2,3 };

// equivalent to:
// data[0] = 1; data[1] = 2; data[2] = 3
```

Because of standard associativity rules, regardless of its meaning, an expression like this:

```
A = x, y, z;
is compiled as
(((A = x), y), z);
```

where each comma is actually a binary operator, so actually

```
((A.operator=(x)).operator,(y)).operator,(z)
```

Note the difference between this syntax and the *growing object*. The latter associates all the items on the right side of assignment, left to right:

```
A = ((x+y)+z);
```

Here, you have the opportunity to modify A iteratively, because the part of the expression containing A is the first to be evaluated:

- Define a proxy object P<A>, which contains a reference to A.
- Define P<A>:: operator so that it takes an argument x. It combines A and x and returns *this (which is the proxy itself).
- Define A::operator=(x) as return P<A>(*this), x.

Suppose you have a wrapper for a C array:

```
template <typename T, size_t N>
struct array
{
   T data[N];
};
```

Being a struct with public members, such an object can be initialized with the curly bracket syntax:

```
array<double, 4 > a = \{ 1, 2, 3, 4 \};
```

However, you cannot do the same on an existing object:²¹

```
array<double, 4> a = { 1,2,3,4 };

// ok, but now assign {5,6,7,8} to a...

const array<double, 4> b = { 5,6,7,8 };

a = b;

// is there anything better?
```

Let the assignment return a proxy with a special comma operator:

```
template <typename T, size_t N>
struct array
{
   T data[N];
private:
```

```
template <size_t J>
class array_initializer
{
   array<T, N>* const pointer_;
   friend struct array<T, N>;
   template <size_t K>
   friend class array_initializer;
   array_initializer(array<T, N>* const p, const T& x)
   : pointer_(p)
   {
     MXT_ASSERT(J<N);
     pointer_->data[J] = x;
}
```

The proxy, being the result of operator=, is conceptually equivalent to a reference, so it's quite natural to forbid copy and assignment by declaring a member const (as in this case) or reference.

For convenience, the proxy is an inner class of array and its constructor is private; array itself and all proxies are friends. Note that the constructor performs a (safe) assignment.

The proxy has a public comma operator that constructs another proxy, moving the index to the next position. Since the user expects the expression A = x to return a reference to A, you can also add a conversion operator:

```
class array_initializer
{
    // ...
public:
    array_initializer<J+1> operator, (const T& x)
    {
        return array_initializer<J+1>(pointer_, x);
    }
    operator array<T, N>& ()
    {
        return *pointer_;
    }
    }; // end of nested class
    Finally, the array assignment just constructs the first proxy:

public:
    array initializer<0> operator=(const T& x)
```

return array initializer<0>(this, x);

```
} ;
```

Note that the fragment:

```
array<int, 4> A;
A = 15, 25, 35, 45;
is roughly equivalent to:
((((A = 15), 25), 35), 45);
```

```
where, as mentioned, each comma is an operator. This expression, after array::operator=, expands at compile time to:
```

```
(((array_initializer<0>(A, 15), 25), 35), 45);
```

The construction of array_initializer<0> sets A[0]=15, then the array_initializer comma operator constructs another initializer that assigns A[1], and so on.

To build a temporary array_initializer<I>, you have to store a const pointer in a temporary on the stack, so the whole process is somehow equivalent to:

```
array<int,4>* const P1 = &A;
P1->data[0] = 15;
array<int,4>* const P2 = P1;
P2->data[1] = 25;
array<int,4>* const P3 = P2;
P3->data[2] = 35;
array<int,4>* const P4 = P3;
P4->data[3] = 45;
```

If the compiler can propagate the information that all assignments involve A, the code is equivalent to a hand-written initialization. All const modifiers are simply hints for the compiler to make its analysis easier.

Comma chains often exploit another language property: destruction of temporary proxy objects. In general, the problem can be formulated as: *how can a proxy know if it is the last one?*

In the previous example, you might like:

```
array<int, 4> A;
A = 15, 25;
// equivalent to {15, 25, 0, 0}

but also

array<int, 4> A;
A = 15;
// equivalent to {15, 15, 15, 15} not to
{15, 0, 0, 0}
```

The expression compiles as array initializer<0>(&A, 15).operator, (25); this

```
returns array initializer<1>(&A,25).
```

The only way a proxy can transmit information to the next is via the comma operator. The object can keep track of invocation and its destructor can execute the corresponding action:

```
template <size t J>
class array initializer
{
  array<T, N>* pointer; // <-- non-const
public:
  array initializer(array<T, N>* const p, const T& x)
   : pointer (p)
     MXT ASSERT (J<N);
     p->data[J] = x;
  }
  array initializer<J+1> operator, (const T& x)
     array<T, N>* const p = pointer;
     pointer = 0; // <-- prevent method re-execution</pre>
     return array initializer<J+1>(p, x);
  }
  ~array initializer()
     // if operator, has not been invoked
     // then this is the last proxy in chain
     if (pointer )
        if (J == 0)
          std::fill n(pointer ->data+1, N-1, pointer ->data[0]);
        else
          std::fill n(pointer \rightarrow data+(J+1), N-(J+1), T());
     }
};
```

Altering the semantics of destructors, in general, is risky. Here, however, you can assume that these objects should not be stored or duplicated, and the implementation enforces that so that (non-malicious) users cannot artificially prolong the life of these proxies.²²

- Put the proxy in a private section of array, so its name is inaccessible to the user.
- Declare all dangerous operators as non-const, so if a proxy is passed to a function by const reference, they cannot be invoked. A non-const reference

must refer to a non-temporary variable, which is unlikely.

• Forbid copy construction.

While it is possible to perform an illegal operation, it *really* requires malicious code:

```
template <typename T>
T& tamper(const T& x)
{
   T& r = const_cast<T&>(x);
   r, 6.28;
   return r;
}
array<double, 10> A;
array<double, 10> B = tamper(A = 3.14);
```

- The argument of tamper is const T&, which can bind to any temporary. Thus it defeats the name-hiding protection.
- const_cast removes the const protection and makes the comma operator callable.
- r.operator, (6.28) as a side effect sets r.pointer = 0.
- The returned reference is still alive when the compiler is going to construct B, but the conversion operator dereferences the null pointer.

Observe that a function like tamper looks harmless and may compile for every T.

9.7. Simulating an Infix

Let's analyze the following fragment:

```
double pi = compute_PI();
assert(pi IS_ABOUT 3.14);
```

We will not solve the problem of comparing floats, but this paragraph will give you an idea of simulating new infixes. If an expression contains operators of different priority, you can take control of the right part before you execute the left part (or vice versa). For example, IS_ABOUT may be a macro that expands to:

```
assert(pi == something() + 3.14);
```

SOMETHING:: operator+ runs first, so you immediately capture 3.14. Then a suitable operator== takes care of the left side.

Here's some code that will do:

```
template <typename float t>
class about t
   float t value ;
public:
   about t(const float t value)
   : value (value)
   }
  bool operator == (const float t x) const
     const float t delta = std::abs(value - x);
     return delta < std::numeric limits<float t>::epsilon();
};
template <typename float t>
inline bool operator == (const float t x, const about t < float t > a)
  return a == x;
struct about creator t
  template <typename float t>
  inline about t<float t> operator+(const float t f) const
     return about t<float t>(f);
};
#define IS ABOUT
                            == about creator t() +
  Obviously, the role of + and == can be reversed, so as to read the left side first.
  Note also that if all these objects belong to a namespace, the macro should qualify
```

about creator tfully.

The curious reader may wish to investigate the following algorithm, which is given without explanations.

Two numbers X and Y are given.

- 1. If X==Y return true.
- 2. Check trivial cases when one or both numbers are infinite or NAN and return

accordingly.

- 3. Pick epsilon from std::numeric_limits.
- 4. Let D := |X-Y|.
- 5. Let R := $\max(|X|, |Y|)$.
- 6. Return R<epsilon OR D<(R*epsilon).
- A variant of the algorithm also tests D<epsilon.

¹See also Section 12.3.

²In symbols, $(F(\lambda))(x) := F(\lambda(x))$, where x may be a tuple.

³In symbols again, $(\lambda 1 + \lambda 2)$ (x) := $\lambda 1$ (x) + $\lambda 2$ (x), where x may be a tuple.

⁴We don't explicitly list the code for operator+, which is almost identical, but we will use it freely in the rest of the section.

⁵This can be fixed, by storing a static constant named min_number_of_arguments in every lambda implementation. Atomic lambdas, such as lambda_reference, will define it directly and derived lambdas will take the maximum from their nested types. Finally, this constant may be used for static assertions. We leave this as an exercise.

⁶The rest of the chapter assumes that all the code presented up to now has been updated.

⁷On the other hand, assignment and logical operators have a deducible return type. The former returns its first argument (a non-const reference); the latter returns bool. The new keyword decltype of C++0x would allow deducing this type automatically.

⁸The name valarray is used on purpose, to remark that these techniques fit std::valarray.

⁹In other words, a constructor that takes const valarray_interface<X, T>&. The details should follow easily and are left to the reader. The cast operator is required if operators return a type that you cannot modify directly (such as std::string).

¹⁰See "price" and "quality" in the knapsack examples in Section 6.2.1.

¹¹Note that s is passed by value. According to the NVRO (Named Value Return Optimization), if there is only one return statement and the result is a named variable, the compiler can usually elide the copy, directly constructing the result on the caller's stack.

¹²See Section 9.2.8.

¹³The example is obviously fictitious, as you cannot really add the operator to std::string.

¹⁴See also Scott Meyers, "Effective Modern C++," Item 29: "Assume that move operations are not present, not cheap, and not used."

¹⁵See the article on assertions by Alexandrescu and Torjo, which is also the source of the first sample in this paragraph: .

¹⁶For extra clarity, we omitted the information collection phase (the estimation of the string length) from this example. In fact std::ostream does not need to be managed.

¹⁷ This is not obvious at all. In the last example of Section 9.3, we considered an agglomerate, namely bind_to(f) (argument)... (argument), whose syntax depends on the length of the chain. In fact, binding one argument of four yields a functor that takes 4-1=3 free arguments, and so on.

¹⁸This allows you to present simplified code. You can easily add a storage policy as a template parameter.

 $^{^{19}}$ If n is an integer of type <code>int_t</code>, the number of digits in base 10 for n is <code>ceil(log10(n+1))</code>. Assuming that a byte contains eight bits and that <code>sizeof(int_t)</code> is even, the largest integer is <code>256^sizeof(int_t)-1</code>. When you put this in place of n, you'll obtain a maximum number of <code>ceil(log10(256)*sizeof(int_t)) ~ (5/2)*sizeof(int_t)</code> digits. You would add 1

for sign and 1 for the terminator.

- ²⁰See Section 1.4.7 on manipulators.
- ²¹C++0x language extensions allow you to initialize some objects (including std::array) with a list in curly brackets. For more details, refer to http://en.cppreference.com/w/cpp/utility/initializer list.
- ²²Exception safety may be a dependent issue. If the destructor of a proxy performs non-trivial work, it could throw.

PART 3

#include #include <techniques>
#include <applications>

CHAPTER 10

Refactoring

Templates can be considered a generalization of ordinary classes and functions. Often a preexisting function or class, which is already tested, is promoted to a template, because of new software requirements; this will often save debugging time.

However, be careful before adding template parameters that correspond to implementation details, because they are going to be part of the type. Objects that do not differ significantly may not be interoperable. Consider again the example from Section 1.4.9, which is a container that violates this rule:

```
template <typename T, size_t INITIAL_CAPACITY = 0>
class special vector;
```

It makes sense to have operators test equality on any two special_vector<double>, regardless of their initial capacity.

In general, all member functions that are orthogonal to extra template parameters either need to be promoted to templates or be moved to a base class.¹

In fact, two implementations are possible:

• A template function special_vector<T, N>::operator== that takes const special vector<T, K>& for any K:

```
template <typename T, size_t N>
class special_vector
{
public:
   template <size_t K>
   bool operator==(const special_vector<T, K>&);
   // ...
};
```

• special_vector<T, N> inherits from a public special_vector_base<T>. This base class has a protected destructor and operator== (const special vector base<T>&):

```
template <typename T>
class special_vector_base
{
public:
bool operator==(const special_vector_base<T>&);

    // ...
};

template <typename T, size_t N>
class special_vector : public special_vector_base<T>
{
    // ...
};
```

template <typename T>

The latter example allows more flexibility. The base class should not be directly used, but you can expose wrappers as smart pointers/references, to allow arbitrary collections of special vectors (having the same T) without risking accidental deletion. To illustrate this, suppose you were to change the code slightly as follows:

```
class pointer to special vector;
template <typename T, size t N>
class special vector : private special vector base<T>
  // thanks to private inheritance,
  // only the friend class will be able to cast special vector to
  // its base class
  friend class pointer to special vector<T>;
};
template <typename T>
class pointer to special vector // <-- visible to users
  public:
  template <size t K>
  pointer to special vector(special vector<T,K>* b = 0)
  : ptr (b)
  { }
  // fictitious code...
  T at(size t i) const { return (*ptr )[i]; }
};
```

10.1. Backward Compatibility

A typical refactoring problem consists of modifying an existing routine so that any caller can choose either the original behavior or a variation.

To begin with a rather trivial example, assume you want to (optionally) log the square of each number, and you don't want to duplicate the code. So, you can modify the classic function template sq:

```
template <typename scalar_t>
inline scalar_t sq(const scalar_t& x)
{
    return x*x;
}

template <typename scalar_t, typename logger_t>
inline scalar_t sq(const scalar_t& x, logger_t logger)
{
    // we shall find an implementation for this...
}

struct log_to_cout
{
    template <typename scalar_t>
    void operator() (scalar_t x, scalar_t xsq) const
    {
        std::cout << "the square of " << x << " is " << xsq;
    }
};

double x = sq(3.14);
    double y = sq(6.28, log to cout());
    // not logged</pre>
```

The user will turn on the log, passing a custom functor to the two-argument version of sq. But

there are different ways to implement the new function over the old one:

• Encapsulation: Make a call to sq(scalar_t) inside sq(scalar_t, logger t). This solution's implementation risk is minimal.

```
template <typename scalar t>
inline scalar t sq(const scalar t& x)
  return x*x;
}
template <typename scalar t, typename logger t>
inline scalar t sq(const scalar t& x, logger t logger)
  const scalar t result = sq(x);
  logger(x, result);
  return result;
}
       • Interface adaptation: Transform sq(scalar t) so as to secretly call
         sq(scalar t, logger t) with a no-op logger. This is the most flexible
         solution.<sup>2</sup>
struct dont log at all
  template <typename scalar t>
  void operator()(scalar t, scalar t) const
}
template <typename scalar t, typename logger t>
inline scalar t sq(const scalar t& x, logger t logger)
  const scalar t result = x*x; // the computation is performed
  logger(x, result);
  return result;
}
template <typename scalar t>
inline scalar t sq(const scalar t& x)
{
  return sq(x, dont log at all());
```

• *Kernel macros*: Work when the core of the algorithm is extremely simple and needs to be shared between static and dynamic code.

```
#define MXT M SQ(x) ((x)*(x))
template <typename scalar t>
inline scalar t sq(const scalar t& x)
  return MXT M SQ(x);
template <typename int t, int t VALUE>
struct static sq
  static const int t result = MXT M SQ(VALUE);
};
Note The use of kernel macros will be superseded by the C++0x keyword constexpr.
  The square/logging example is trivial, but code duplication is regrettably common. In many STL
implementations, std::sort is written twice:
template <typename RandomAccessIter>
void sort(RandomAccessIter first, RandomAccessIter last);
template <class RandomAccessIter, typename Compare>
void sort (RandomAccessIter first, RandomAccessIter last,
Compare less);
  Using interface adaptation, the first version is a special case of the second:
struct weak less compare
  template <typename T1, typename T2>
  bool operator()(const T1& lhs, const T2& rhs) const
     return lhs < rhs;
};
template <typename RandomAccessIter>
void sort(RandomAccessIter first, RandomAccessIter last)
  return sort( first, last, weak less compare());
}
```

10.2. Refactoring Strategies

This section considers an example problem and exposes some different techniques.

10.2.1. Refactoring with Interfaces

A preexisting private_ptr class holds the result of a malloc in a void* and frees the memory block in the destructor:

```
class private_ptr
{
   void* mem_;

public:
   ~private_ptr() { free(mem_); }

   private_ptr() : mem_(0)
   { }

   explicit private_ptr(size_t size) : mem_(malloc(size))
   { }

   void* c_ptr() { return mem_; }

   //...
};
```

Now you need to extend the class so that it can hold a pointer, either to a malloc block or to a new object of type T.

Since private_ptr is responsible for the allocation, you could just introduce a private interface with suitable virtual functions, create a single derived (template) class, and let private_ptr make the right calls:

```
class private_ptr_interface
{
  public:
    virtual void* c_ptr() = 0;
    virtual ~private_ptr_interface() = 0;
};

template <typename T>
  class private_ptr_object : public private_ptr_interface
{
    T member_;

public:
    private_ptr_object(const T& x)
    : member_(x)
    {
}
```

```
}
  virtual void* c ptr()
     return &member ;
  virtual ~private ptr object()
};
template < >
class private ptr object<void*> : public private ptr interface
  void* member ;
public:
  private ptr object(void* x)
  : member (x)
  virtual void* c ptr()
     return member ;
  virtual ~private ptr object()
     free(member);
};
class private ptr
  private ptr interface* mem ;
public:
  ~private ptr()
     delete mem ;
  private ptr()
    mem (0)
```

```
explicit private_ptr(size_t size)
: mem_(new private_ptr_object<void*>(malloc(size)))
{
}

template <typename T>
explicit private_ptr(const T& x)
: mem_(new private_ptr_object<T>(x))
{
}

void* c_ptr()
{
   return mem_->c_ptr();
}

//...
```

Note that virtual function calls are invisible outside private ptr.³

10.2.2. Refactoring with Trampolines

The former approach uses two allocations to store a void*: one for the memory block and one for the auxiliary private ptr object. Trampolines can do better:

```
template <typename T>
struct private_ptr_traits
{
    static void del(void* ptr)
    {
        delete static_cast<T*>(ptr);
    }
};

template <typename T>
struct private_ptr_traits<T []>
{
    static void del(void* ptr)
    {
        delete [] static_cast<T*>(ptr);
    }
};

template < >
struct private_ptr_traits<void*>
{
    static void del(void* ptr)
```

```
{
     free (ptr);
};
template < >
struct private ptr traits<void>
  static void del (void*)
};
class private ptr
  typedef void (*delete t) (void*);
  delete t del ;
  void* mem ;
public:
  ~private ptr()
  {
     del (mem );
  private ptr()
  : mem (0), del (&private ptr traits<void>::del)
  }
  explicit private ptr(size t size)
     mem = malloc(size);
     del = &private ptr traits<void*>::del;
  }
  template <typename T>
  explicit private ptr(const T& x)
     mem = new T(x);
     del = &private ptr traits<T>::del;
  template <typename T>
  explicit private ptr(const T* x, size t n)
     mem = x;
```

```
del_ = &private_ptr_traits<T []>::del;
}

void* c_ptr()
{
   return mem_;
}

//...
;
```

10.2.3. Refactoring with Accessors

Suppose you have algorithms that process a sequence of simple objects:

Refactoring may be needed to allow you to process data from two independent containers:

```
std::vector<double> prices;
std::vector<time_t> dates;

// problem: we cannot call computePriceIncrease
```

You have several choices for the new algorithm I/O:

- Assume that iterators point to pair, where first is price and second is date (in other words, write end->first begin->first...). This in general is a poor style choice, as discussed previously.
- Mention explicitly begin->price and begin->date (as shown previously). The algorithm does not depend on the iterator, but the underlying type is constrained to the interface of stock price.
- Pass two disjoint ranges. The complexity of this solution may vary.

```
template <typename I1, typename I2>
```

```
double computePriceIncrease(I1 price_begin, I1 price end, I2
date begin, I2 date end)
  // the code must be robust and handle ranges of different
length, etc.

    Pass one range and two accessors.

template <typename I, typename price t, typename date t>
double computePriceIncrease(I begin, I end, price t PRICE, date t
DATE)
{
  double p = PRICE(*begin);
  time t t = DATE(*begin);
  //...
}
struct price accessor
  double operator()(const stock price& x) const
     return x.price;
} ;
struct date accessor
{
  time t operator()(const stock price& x) const
     return x.date;
};
computePriceIncrease(begin, end, price accessor(),
date accessor());
  Note that you can trick accessors into looking elsewhere, for example, in a member variable:
struct price accessor ex
{
  const std::vector<double>& v ;
  double operator()(const int x) const
     return v [x];
};
```

```
struct date accessor ex
  const std::vector<time t>& v ;
  time t operator()(const int x) const
  {
     return v [x];
};
int main()
  std::vector<double> prices;
  std::vector<time t> dates;
  // ...
  assert(prices.size() == dates.size());
  std::vector<int> index(prices.size());
  for (int i=0; i<prices.size(); ++i)</pre>
     index[i] = i;
  price accessor ex PRICE = { prices };
  date accessor ex DATE = { dates };
  computePriceIncrease(index.begin(), index.end(), PRICE, DATE);
```

Accessors may carry around references to an external container, so they pick an element deduced from the actual argument. In some special cases, you can use pointers to avoid creating a container of indices. This approach, however, should be used with extreme care.

```
// warning: this code is fragile:
// changing a reference to a copy may introduce subtle bugs
struct price_accessor_ex
{
    double operator() (const double& x) const
    {
       return x;
    }
};
struct date_accessor_ex
{
    const double* first_price_;
    size_t length_;
    const time_t* first_date_;
```

```
time_t operator()(const double& x) const
{
    if ((&x >= first_price_) && (&x < first_price_+length_))
        return first_date_[&x - first_price_];
    else
        throw std::runtime_error("invalid reference");
};
int main()
{
    price_accessor_ex PRICE;
    date_accessor_ex DATE = { &prices.front(), prices.size(), &dates.front() };
    computePriceIncrease(prices.begin(), prices.end(), PRICE, DATE);
}</pre>
```

The algorithm takes a reference to a price and deduces the corresponding date accordingly.

10.3. Placeholders

Every C++ object can execute some actions. Empty objects, such as instance_of, can execute meta-actions, such as declare their type and "bind" their type to a template parameter or to a specific function overload.

Sometimes the job of TMP is to *prevent* work from being done, by replacing an object with a similar empty object and an action with a corresponding meta-action.

Type P < T > is called a *placeholder for* T if P < T > is a class whose public interface satisfies the same pre- and post-conditions as T, but has the least possible runtime cost. In the most favorable case, it does nothing at all.

10.3.1. Switch-Off

The switch-off is an algorithm-refactoring technique that allows you to selectively "turn off" some features without rewriting or duplicating functions. The name comes from the paradigmatic situation where a function takes an object by reference, which is "triggered" during execution, and eventually returns an independent result, which is a by-product of the execution. The object may be a container that receives information during the execution or a synchronization object.

```
void say_hello_world_in(std::ostream& out)
{
  out << "hello world";
}</pre>
```

```
double read_from_database(mutex& s)
{
    // acquire the mutex, return a value from the DB, and release the mutex
}
```

A quick and elegant way to get a different result with minimal code rework is to supply a hollow object with a reduced interface and that, in particular, does not need any dynamic storage. Step by step:

• Rename the original function and promote the parameter to a template type:

```
template <typename T>
void basic_say_hello_world_in(T& o)
```

• Add an overload that restores the original behavior:

```
inline void say_hello_world_in(std::stream& o)
{
   return basic_say_hello_world_in(o);
}
```

• Finally, provide an object that "neutralizes" most of the effort:

```
struct null_ostream
{
   template <typename T>
   null_ostream& operator<<(const T&)
   {
      return *this;
   }
};
inline void say_hello_world_in()
{
   null_stream ns;
   basic_say_hello_world_in(ns);
}</pre>
```

The switch-off idiom requires exact knowledge of the (subset of the) object's interface used in the main algorithm.

When you're designing a custom container, it may occasionally be useful to add an extra template parameter to enable a *hollow-mode*. You take the original class and promote it to a template:

```
template <bool SWITCH_ON =
true>
class spinlock;
```

```
template < >
                                                        class spinlock<true>
                                                        typedef void* ptr t;
                                                        typedef volatile ptr t
                                                        vptr t;
class spinlock
                                                        public:
                                                        spinlock(vptr t* const);
typedef void* ptr_t;
                                                        bool try_acquire();
typedef volatile ptr t vptr t;
                                                        bool acquire();
public:
                                                        // ...
spinlock(vptr_t*const);
                                                        };
bool try acquire();
                                                        template < >
bool acquire();
                                                        class spinlock<false>
// ...
                                                        // hollow implementation
};
                                                        spinlock(void*)
                                                        { }
                                                        bool try acquire()
                                                        { return true; }
                                                        bool acquire()
                                                        { return true; }
                                                        //...
                                                        };
```

Had the class been a template, you would need to add one more Boolean parameter.

Of course, the crucial point of the duplication of the interface is the set of cautious, but meaningful, default answers of the hollow class, provided that such duplication is possible (see below for a counter-example). This also allows you to identify the minimal interface for an object to be considered "valid". The interface of an object is defined by its usage.

Finally, you can restrict the program to spinlocks (which may be "on" or "off"):

```
template <typename ..., bool IS_LOCKING_REQURED>
void run_simulation(..., spinlock<IS_LOCKING_REQURED>& spin)
{
   if (spin.acquire())
   {
      //...
}
```

Or to objects of an unspecified type, whose interface is implicitly assumed compatible with spinlock:

```
template <typename ..., typename lock_t>
```

```
void run_simulation(..., lock_t& lock)
{
   if (lock.acquire())
   {
      //...
}
```

Either choice is valid, but there are situations where one is preferred (see Section 5.2 for more details).

Another application is *twin reduction*. There are algorithms that manipulate one or two items at a time and execute the same actions simultaneously on both. To avoid duplication, you want a single implementation of the algorithm that accepts one or two arguments.

Prototype examples are sorting two "synchronized" arrays and matrix row reduction. This algorithm, due to Gauss, performs a sequence of elementary operations on a matrix M and turns it into a diagonal (or triangular) form. If the same operations are applied in parallel on an identity matrix, it also obtains the inverse of M.⁴

So you can write a general-purpose function that *always takes two* matrices of different static type and treats them as identical:

```
template <typename scalar_t>
class matrix
{
public:
   typedef scalar_t value_type;
   size_t rows() const;
   size_t cols() const;
   void swap_rows(const size_t i, const size_t j);
   value_type& operator()(size_t i, size_t j);
```

```
value_type operator()(size_t i, size_t j) const;
};
```

It's not possible to extend it following the hollow-mode idiom, because there's no satisfactory default answer for functions returning a reference:⁶

```
template <typename scalar_t, bool NO_STORAGE = false>
class matrix;

template <typename scalar_t>
class matrix<scalar_t, false>
{
    /* put the usual implementation here */
};

template <typename scalar_t>
class matrix<scalar_t, true>
{
public:
    value_type& operator()(size_t i, size_t j)
    {
        return /* what? */
    }
    //...
};
```

So you drop the reference entirely and move down one level. You neutralize both the container and the contained object. The twin matrix is a container defined on a *ghost scalar*; a class whose operators do nothing:

```
template <typename T>
struct ghost
{
    // all operators return *this

    ghost& operator-=(ghost)
    {
        return *this;
    }

    //...
};

template <typename T>
inline ghost operator*(T, ghost g) { return g; }

template <typename T>
inline ghost operator*(ghost g, T) { return g; }
```

```
template <typename scalar_t>
class matrix<scalar_t, true>
{
    size_t r_;
    size_t c_;

public:
    typedef ghost<scalar_t> value_type;

    size_t rows() const { return r_; }
    size_t cols() const { return c_; }

    void swap_rows(const size_t, const size_t) {}

    value_type operator()(size_t i, size_t j) {
        return value_type();
    }

    const value_type operator()(size_t i, size_t j) const {
        return value_type();
    }
};
```

ghost<T> will be a stateless class such that every operation is a no-op. In particular, the line twin (j, k) -= pivot*twin (i, k) translates into a sequence of do-nothing function calls. Some more detail on this point is needed.

10.3.2. The Ghost

There are no truly satisfactory ways to write ghost scalars. Most implementations are semi-correct but they can have nasty side effects:

- Ghosts are likely to haunt your namespaces if they are not properly constrained. Since their interfaces should support virtually all C++ operators, you will probably need to write some global operators, and you want to be sure these will appear only when necessary.
- The main purpose of ghosts is to prevent work from being done. If G is a ghost, then G*3+7 should compile and do nothing. It's very easy to obtain an implementation that compiles, but erroneously does some work—say, because G is converted to integer 0.

A ghost should be a class template that mimics its template parameter \mathbb{T} and it resides in a different namespace. You can assume for simplicity that \mathbb{T} is a built-in numeric type, so you can implement all possible operators.

```
template <typename T>
struct ghost
{
   ghost(T) {}
   ghost() {}
   //...
};
```

#define mxt GHOST ASSIGNMENT(OP)

For coherence, comparison operators return a result compatible with the fact that ghost is monostate (all ghosts are equivalent), so operator< is always false and operator== is always true.

As a rule, most arithmetic operators can be defined with suitable macros: ⁷

```
ghost& operator OP##= (const ghost) { return *this; }
#define mxt GHOST UNARY(OP)
     ghost operator OP() const { return *this; }
#define mxt GHOST INCREMENT(OP)
     qhost& operator OP () { return *this; }
     const ghost operator OP (int) { return *this; }
template <typename T>
struct qhost
  ghost(const T&) { }
  qhost() {}
                              // defines pre- and post-increment
  mxt GHOST INCREMENT(++);
  mxt GHOST INCREMENT(--);
                            // defines operator+=
  mxt GHOST ASSIGNMENT(+);
  mxt GHOST ASSIGNMENT(-);
  // ...
  mxt GHOST UNARY(+);
  mxt GHOST UNARY(-);
  //...
};
```

For the arithmetic/comparison operators, you need to investigate these possibilities:

- 1. Member operators with argument ghost<T>.
- 2. Member operators with argument T.
- 3. Template member operators with argument const X&, where X is an independent template parameter.

4. Nonmember operators, such as

Each choice has some problems.

- 1. Member operators will perform argument promotion on the right side, but template global operators require a perfect match for argument deduction. With member operators ghost<T>::operator+(ghost<T>) const, any sum of the form ghost<T> + X will succeed whenever it's possible to build a temporary ghost<T> from X (since a ghost constructor is not explicit). However, X + ghost<T> will not compile.
- 2. The problem is mostly evident when T is a numeric type (say, double) and X is a literal zero. A member operator+ will take care of ghost<double> + 0, since 0 (int) → 0.0 (double) → ghost<double>, but 0 + ghost<double> must be handled by a global operator whose signature cannot be too strict, as 0 is not a double.
- 3. This implies that in this case, only variant #4 is feasible, because no other operator would match exactly (int, ghost<double>).
- 4. However, you want operators to match as many types as possible, but not more. While you should be able to write int + ghost<double>, you don't want to accept *anything*.

As a rule, the global operator should delegate the execution to a member function:

```
template <typename T1, typename T2>
inline ghost<T2> operator+ (T1 x, const ghost<T2> y)
{
  return y + x;
```

}

y + x is indeed a call to any member operator+, so you can pass the responsibility for accepting T1 as argument to the ghost's own interface (the compiler will try any overloaded operator+).

A conversion operator is necessary to make assignments legal:

Conversely, with the conversion operator and a bad implementation of operators, innocuous code will suddenly become ambiguous:

```
ghost<double> g;
g + 3.14;
```

For example, there may be an ambiguity between:

- Promotion of 3.14 to ghost<double>, followed by ghost<double>::operator+(ghost<double>).
- Conversion of g to double, followed by an ordinary sum.

Since both paths have equal rank, the compiler will give up. In different situations, the conversion will be unexpectedly called:

```
ghost<double> g = 3.14;double x = 3*q + 7;
```

This code should be translated by the compiler into this sequence:

```
double x = (double) (operator*(3, g).operator*(ghost<double>(7)));
```

If the global operator* cannot be called for any reason (say, it expects double, ghost<double>, so it won't match), the code is still valid, but it silently executes something different:

```
double x = 3*(double)(g) + 7;
```

This costs two floating-point operations at runtime, so it defeats the ghost purpose. Summing up, in the best implementation:

• The ghost constructor is strongly typed, so it needs one argument convertible to T.

- You need both member and non-member operators:
 - Member operators that will accept any argument (any type X) and will check X with a static assertion (using the constructor itself).
 - Non-member operators that will blindly delegate anything to member functions.

What's described here is an implementation without making use of macros. Anyway, functions generated by the same preprocessor directive have been grouped:

```
\#define mxt GHOST GUARD(x) sizeof(ghost<T>(x))
template <typename T>
struct ghost
{
  ghost(const T&) {}
  ghost() {}
  operator T() const
    return T();
  }
  ghost& operator++ () { return *this; }
  const ghost operator++ (int) { return *this; }
  ghost& operator-- () { return *this; }
  const ghost operator-- (int) { return *this; }
  template <typename X> ghost& operator+= (const X& x)
  { mxt GHOST GUARD(x); return *this; }
  template <typename X> ghost& operator-= (const X& x)
  { mxt GHOST GUARD(x); return *this; }
  template <typename X> ghost operator+ (const X& x) const
  { mxt_GHOST_GUARD(x); return *this; }
  template <typename X> ghost operator- (const X& x) const
  { mxt GHOST GUARD(x); return *this; }
  template <typename X> bool operator== (const X& x) const
  { mxt GHOST GUARD(x); return true; }
  template <typename X> bool operator!= (const X& x) const
  { mxt GHOST GUARD(x); return false; }
  ghost operator+() const { return *this; }
```

```
ghost operator-() const { return *this; }
};

template <typename X, typename Y>
ghost<Y> operator+ (const X& x, const ghost<Y> y) { return y + x;
}

template <typename X, typename Y>
ghost<Y> operator- (const X& x, const ghost<Y> y) { return -(y - x); }

template <typename X, typename Y>
bool operator== (const X& x, const ghost<Y> y) { return y == x; }

template <typename X, typename Y>
bool operator!= (const X& x, const ghost<Y> y) { return y != x; }
```

¹A similar debate was raised about STL allocators. The notion of "equality of two containers of the same kind" obviously requires the element sequences to be equal, but it's unclear whether this is also sufficient.

²While encapsulation conveys to the user a "sense of overhead," interface adaptation suggests that the new sq is much better and can be used freely.

³In other words, callers of the code do not have to worry about inheritance. They can pass any T and the class will wrap it silently and automatically. This idea was developed further in a talk by Sean Parent and is freely downloadable from this link:

http://channel9.msdn.com/Events/GoingNative/2013/Inheritance-Is-The-Base-Class-of-Evil.

⁴A non-mathematically inclined reader may want to consider an analogous case: software that executes a series of actions and at the same time records a list of "undo" steps.

⁵The interface of a data structure is frequently remodeled for ease of algorithms. This lesson was one of the milestones of the STL design.

⁶As a rule, hollow containers own no memory. You could object that here you could use a single scalar_t data member and return a reference to the same object for any pair of indices, but this strategy would consume a lot of CPU runtime, overwriting the same memory location for no purpose.

^{&#}x27;Mind the use of token concatenation ##. You might be tempted to write operator ## OP to join operator and +, but this is illegal, because in C++, operator and + are two different tokens. On the other hand, ## is required between + and = to generate operator +=, so you need to write operator OP ## =.

⁸The user-defined constructor that converts T to ghost<T> is considered only after template argument deduction. Note that the constructor here is not even explicit. See [2], Section B.2.

⁹Hint: always leave a breakpoint in the conversion operator.

CHAPTER 11

Debugging Templates

As TMP code induces the compiler to perform calculations, it's virtually impossible to follow it step by step. However, there are some techniques that can help. This chapter in fact contains a mix of pieces of advice and debugging strategies.

11.1. Identify Types

Modern debuggers will always show the exact type of variables when the program is stopped. Moreover, a lot of information about types is visible in the call stack, where (member) functions usually are displayed with their full list of template arguments. However, you'll often need to inspect intermediate results and return types.

The following function helps:

```
template <typename T>
void identify(T, const char* msg = 0)
{
   std::cout << (msg ? msg : "") << typeid(T).name() << std::endl;
}</pre>
```

Remember that type_info::name gives no guarantees about the readability of the returned string.¹ Using a free function to return void makes it easy to switch between debug and optimized builds, as the code can simply use a preprocessor directive to replace the function, say, with an empty macro. However, this approach does not work when you need to identify a class member, such as when you're debugging lambda expressions. (See Section 9.2). You may want to check if the return type has been correctly deduced; the best solution is to add a small, public data member:

```
template <typename X1, typename F, typename X2>
class lambda_binary : public lambda< lambda_binary<X1,F,X2> >
{
    // ...
    typedef typename
    deduce_argument
```

```
typename X1::argument_type,
    typename X2::argument_type
>::type
argument_type;

#ifdef MXT_DEBUG
    instance_of<result_type> RESULT_;
#endif

result_type operator()(argument_type x1, argument_type x2) const
{
    identify(RESULT_);
    return f_(x1_(x1, x2), x2_(x1, x2));
};
```

Adding a data member is especially useful because interactive debuggers allow you to inspect objects in memory and display their exact type.

In general, whenever a metafunction compiles but gives the wrong results, add members of type <code>instance_of</code> and <code>static_value</code> to inspect the intermediate steps of the computation, then create *a local instance* of the metafunction on the stack.

```
template <size_t N>
struct fibonacci
{
    static const size_t value = fibonacci<N-1>::value + fibonacci<N-
2>::value;
    static_value<size_t, value> value_;
    fibonacci<N-1> prev1_;
    fibonacci<N-2> prev2_;
};
int main()
{
    fibonacci<12> F;
}
```

Then look at F in the debugger. You can inspect the constants from their type.²

11.1.1. Trapping Types

Sometimes in large projects, an erroneous pattern is detected. When this happens, you need to list all the code lines where the bad pattern is used. You can use templates to create *function traps* that do not compile and inject them in the error pattern, so that the compiler log will point to all the lines you are looking for.

Suppose for a moment that you discover that a std::string is passed to printf and you

suspect this happens several times in the project.

```
std::string name = "John Wayne";
printf("Hello %s", name);  // should be: name.c_str()
class Foo{};
printf("I am %s", Foo());
```

Brute-force iteration through all occurrences of printf would take too much time, so you can instead add some trap code in a common included file. Note that you have to write a static assertion that is always false, but depends on an unspecified parameter T. In the following code,

```
MXT ASSERT is a static assertion:
template <typename T>
void validate(T, void*)
}
template <typename T>
void validate(T, std::string*)
  MXT ASSERT(sizeof(T) == 0); // if this triggers, someone is
passing
                         // std::string to printf!
template <typename T>
void validate(T x)
  validate(x, &x);
template <typename T1>
void printf trap(const char* s, T1 a)
  validate(a);
}
template <typename T1, typename T2>
void printf trap(const char* s, T1 a, T2 b)
  validate(b);
  printf trap(s, a);
template <typename T1, typename T2, typename T3>
void printf trap(const char* s, T1 a, T2 b, T3 c)
  validate(c);
```

```
printf_trap(s, a, b);
}
// ...
#define printf printf_trap
```

This trap code will cause a compiler error every time a string is passed to printf.

It's important to be able to mention std::string (in validate), so the previous file must include <string>. But if you are testing a user class, this might not be feasible (including project headers that might cause loops), so you simply replace the explicit validation test with a generic SFINAE static assertion:

```
template <typename T>
void validate(T, void*)
{
    MXT_ASSERT(!is_class<T>::value); // don't pass classes to printf;
}
```

11.1.2. Incomplete Types

Class templates might not require that T be a complete type. This requirement is usually not explicit, and it depends on the internal template implementation details.

STL containers, such as vector, list, and set, can be implemented so as to accept incomplete types, because they allocate storage dynamically. A necessary and sufficient condition to decide if \mathbb{T} may be incomplete is to put in a class a container of itself.

```
struct S1
{
   double x;
   std::vector<S1> v;
};

struct S2
{
   double x;
   std::list<S2> l;
};
```

In particular, an *allocator* should not assume that \mathbb{T} is complete; otherwise, it might be incompatible with standard containers.

A static assertion is easily obtained by just asking the compiler the size of a type:

```
template <typename T>
struct must_be_complete
{
```

```
static const size t value = sizeof(T);
};
struct S3
   double x;
  must be complete<S3> m;
};
test.cpp: error C2027: use of undefined type 'S3'
   This technique is used to implement safe deletion. A pointer to an incomplete type may be
deleted, but this causes undefined behavior (in the best case, T's destructor won't be executed).
template <typename T>
void safe delete(T* p)
    typedef T must be complete;
    sizeof(must be complete);
    delete x;
}
   Determining if a template will get a complete type as an argument may not be easy.
   Standard allocators have a rebind member that allows any allocator<T> to create
allocator<X>, and different implementations will take advantage of the feature to construct their
own private data structures. A container, say std::list<T>, may need
allocator<node<T>> and this class may be incomplete.
template <typename T>
class allocator
  typedef T* pointer;
   template <typename other t>
   struct rebind
      typedef allocator<other t> other;
   };
   // ...
};
template <typename T, typename allocator t>
struct list
   struct node;
   friend struct node;
   typedef typename allocator t::template
```

```
rebind<node>::other::pointer node pointer;
   // the line above uses allocator<node> when node is still
incomplete
   struct node
   {
      node(node pointer ptr)
   };
   // ...
};
   To compile the node constructor, node pointer is needed. So the compiler looks at
allocator ::rebind<node>::other, which is in fact allocator<node>.
   Suppose you now have an efficient class that manages memory blocks of fixed length N:
template <size t N>
class pool;
   To wrap it correctly in a generic stateless allocator, you may be tempted to write:
template <typename T>
class pool allocator
   static pool<sizeof(T)>& get storage();
   // ...
};
   But in this case, the presence of sizeof (T) at class level requires T to be complete. Instead,
you switch to a lazy instantiation scheme with a template member function:
template <typename T>
class pool allocator
   template <typename X>
   static pool<sizeof(X)>& get storage()
      static pool\langle \text{sizeof}(X) \rangle^* p = new pool\langle \text{sizeof}(X) \rangle^*;
      return *p;
   }
   // ...
   void deallocate(pointer ptr, size type)
      get storage<T>().release(ptr);
```

```
};
```

Now, at class level, sizeof(T) is never mentioned.

Note As mentioned in Section 10.14 of [7], there's a difference between stack and heap allocation:

```
static T& get1()
{
   static T x;
   return x;
}
static T& get2()
{
   static T& x = *new T;
   return x;
}
```

The former will destroy x at some unspecified moment at the end of the program, while the latter never destroys x.

So, if T:: T () releases a resource, say a mutex, the first version is the right one. However, if the destructor of another global object invokes get1 (), it might be that x has already been destroyed (a problem known as "static initialization order fiasco").

11.1.3. Tag Global Variables

A non-type template parameter can be an arbitrary pointer to an object having an external linkage. The limitation is that this pointer cannot be dereferenced at compile time:

```
template <int* P>
struct arg
{
    arg()
    {
       myMember = *P; // dereference at runtime
    }
    int myMember;
};
extern int I;
int I = 9;
arg<&I> A;
```

It would be illegal instead to write:

template <int* P>

```
struct arg : static value<int, *P> // dereference at compile time
  You can use pointers to associate some metadata to global constants:
// metadata.hpp
template <typename T, T* global>
struct metadata
 static const char* name;
};
#define DECLARE CPP GLOBAL(TYPE, NAME)
  TYPE NAME;
  template <> const char* metadata<TYPE, &NAME>::name = #NAME
// main.cpp
#include "metadata.hpp"
DECLARE CPP GLOBAL (double, xyz);
int main()
   printf(metadata<double, &xyz>::name); // prints "xyz"
```

11.2. Integer Computing

This section quickly reviews some problems that static integer computations may cause.

11.2.1. Signed and Unsigned Types

Common issues may arise from the differences between T(-1), -T(1), T()-1, and -T() when T is an integer type.

- If T is unsigned and large, they are all identical.
- If \mathbb{T} is signed, the first three are identical.
- If T is unsigned and small, the second and third expressions may give unexpected results.

Let's borrow a function from the implementation of is signed integer (see Section 4.3.2).

```
template <typename T>
static selector<(T(0) > T(-1))> decide_signed(static_value<T, 0>*);

Replace T(-1) with -T(1) and suddenly two regression tests fail. (But which ones?)

bool t01 = (!is_signed_integer<unsigned char>::value);
bool t02 = (!is_signed_integer<unsigned int>::value);
bool t03 = (!is_signed_integer<unsigned long long>::value);
bool t04 = (!is_signed_integer<unsigned long>::value);
bool t05 = (!is_signed_integer<unsigned short>::value);
bool t11 = (is_signed_integer<char>::value);
bool t12 = (is_signed_integer<int>::value);
bool t13 = (is_signed_integer<long long>::value);
bool t14 = (is_signed_integer<long long>::value);
bool t15 = (is_signed_integer<long>::value);
```

The reason for failure is that the "unary minus" operator promotes small unsigned integers to int, so -T(1) is int and the whole comparison is shifted into the int domain, where 0 > -1 is true. To see this, execute the following:

```
unsigned short u = 1;
identify(-u);
```

11.2.2. References to Numeric Constants

As a rule, don't pass static constants to functions directly:

```
struct MyStruct
{
   static const int value = 314;
}
int main()
{
   double myarray[MyStruct::value];
   std::fill_n(myarray, MyStruct::value, 3.14); // not recommended
}
```

If fill_n takes the second argument by const reference, this code may fail *linking*. Taking the address of the constant requires the constant to be redeclared in the .cpp file (as is the case for any other static member). In TMP, this is rarely the case.

As a cheap workaround, you can build a temporary integer and initialize it with the constant:

```
// not guaranteed by the standard, but usually ok
std::fill n(myarray, int(MyStruct::value), 3.14);
```

For extreme portability, especially for enumerations and bool, you can build a function on the

```
template <bool B> struct converter;

template <> struct converter<true>
{ static bool get() { return true; } };

template <> struct converter<false>
{ static bool get() { return false; } };

// instead of: DoSomethingIf(MyStruct::value);
DoSomethingIf(converter<MyStruct::value>::get());
```

11.3. Common Workarounds

11.3.1. Debugging SFINAE

fly:

A common "cut and paste" error is the addition of a useless non-deducible template parameter to a function. Sometimes, the compiler will complain, but if the function is overloaded, the SFINAE principle will silently exclude it from overload resolution, which will generally lead to subtle errors:

```
template <typename X, size_t N>
static YES<[condition on X]> test(X*);
static NO test(...);
```

In this fragment, N cannot be deduced, thus the second test function will always be selected.

11.3.2. Trampolines

Compiler limitations may affect trampolines. In classic C++, local classes have some limitations (they cannot bind to template parameters). They may cause spurious compiler and linker errors:

```
// call local::myFunc(&m);
};
  The workaround is to move most of template code outside of the local class:
template <typename T>
struct MyStruct
  template <typename X>
  static T* MyFunc(const X& m)
     // do the work here
  template <typename X>
  void DoSomething(const X& m)
     struct local
     {
        static T* MyFunc(const void* p)
           // put nothing here, just a cast
           return MyStruct<T>::MyFunc(*static cast<const X*>(p));
     };
     // ...
```

11.3.3. Compiler Bugs

};

};

Compiler bugs are rare, but they do occur, especially within template metaprogramming. They usually produce obscure diagnostics.³

```
error C2365: 'function-parameter': redefinition; previous definition was a 'template parameter'. see declaration of 'function-parameter'
```

Compilers get confused by templates when:

- They cannot deduce that an expression is a type.
- They don't perform automatic conversion correctly, or in the right order, so they

emit incorrect diagnostics.

• Some language keywords may not work correctly in a static context.

Here is an example of this last statement. sizeof will usually complain if an expression is invalid. Here is what happens when you try to dereference a double:

```
int main()
{
    sizeof(**static_cast<double*>(0));
}
error: illegal indirection
```

The same test may fail to trigger SFINAE correctly. The following code used to print "Hello" with an old version of a popular compiler:⁴

```
template <size_t N>
struct dummy
{
};

template <typename X>
dummy<sizeof(**static_cast<X*>(0))>* test(X*)
{
   printf("Hello");
   return 0;
}

char test(...)
{
   return 0;
}

int main()
{
   double x;
   test(&x);
}
```

The next example is due to implicit conversions:

```
double a[1];
double b[1];
double (&c)[1] = true ? a : b;
error: 'initializing' : cannot convert from 'double *' to 'double
(&)[1]'
         A reference that is not to 'const' cannot be bound to a non-
```

Thus you can see that the compiler is erroneously converting the array to pointer in the ternary operator. However, the *bug might not trigger* inside a template function:

```
template <typename T>
void f()
{
    T a;
    T b;
    T& c = true ? a : b;
}
f<double [1]>();
```

Ensuring *portability* is a non-trivial development effort. An informal definition of portability is, "code that works in multiple platforms, potentially adapting to the platform itself (with preprocessor directives, and so on)". Code that is *standard conformant* will work everywhere, without changes (given a bug-free compiler). In practice, portability is a combination of both standard conformant code and code that works around some specific compiler limitations/bugs. Some compilers have subtle non-standard behavior; they may have extensions (for example, they may silently allow creating variable-length arrays on the stack), they may tolerate minor syntax errors (such as this-> or the use of ::template), and even some ambiguities (for example, static casts of objects with multiple bases). However, aiming for standard conformance is extremely important, because it guarantees that if a piece of (metaprogramming) code works, it will continue working even with future versions of the same compiler.

If code that looks correct does not compile, it may help to:

- Simplify a complex type introducing extra typedefs or vice versa.
- Promote a function to template or vice versa.
- Test a different compiler if the code cannot be changed further.

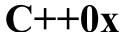
¹See http://en.cppreference.com/w/cpp/types/type_info/name.

²Additionally, there exist interactive meta-debuggers. Meta-debuggers use their own compiler under the hood, so their output might differ from what is observed in the actual binary, but they are extremely valuable when investigating a metafunction that does *not* compile. One can be found here: http://metashell.readthedocs.org/en/latest/

³Note that all the examples in this section rely on bugs of a specific version of some popular C++ compilers (which we don't mention), so hopefully, they won't be reproducible. However, they are good examples of what could go wrong.

⁴decltype may suffer from similar issues.

CHAPTER 12



"I note that every C++0x feature has been implemented by someone somewhere."

Bjarne Stroustrup

We conventionally call "classic C++" the language in its final revision in 2003, as opposed to "modern C++" (also informally known as C++0x), introduced in 2011 and subsequently refined in 2014. The set of changes was huge, but the new rules in general were written to ease TMP and make the code less verbose. Additionally, compilers come with a new arsenal of standard classes, containers, language tools (like std::bind), and traits that expose meta-information previously known only to the compiler. 1

The simplest example is the metafunction std::has trivial destructor<T>.

It's not possible to detect if a type has a trivial destructor by language only. The best default implementation in classic C++ would be "return false unless T is a native type".²

This chapter briefly scratches the surface of a huge topic, so don't consider this chapter a complete reference. Some of the descriptions are slightly simplified, for the benefit of extra clarity.

12.1. Type Traits

Compilers already offer a complete set of metafunctions:

```
#include <type traits>
```

This will bring some metafunctions in namespace std or std::tr1 (depending on the compiler and the standard library).³

In particular, some metafunctions that were described in this book are present in C++0x, with a different name. Some examples are listed in the following table.⁴

This Book	C++0x Equivalent	
static_value	std::integral constant	
only_if	std::enable_if	
typeif	std::conditional	

12.2. Decltype

Similarly to sizeof, decltype resolves to the type of the C++ expression given in brackets (without evaluating it at runtime), and you can put it wherever a type is required:

```
int a;
double b;
decltype(a+b) x = 0;  // x is double
```

decltype can have a positive impact on SFINAE. The following metafunction detects correctly a swap member function, testing the expression x. swap (x), where x is a non-constant reference to x.

Since swap usually returns void, you use pointer-to-decltype for types that pass the test, and a non-pointer class for the rest. Then you cast this to yes/no as usual:

```
#define REF_TO_X (*static_cast<X*>(0))
struct dummy {};
template <typename T>
struct has_swap
{
   template <typename X>
   static decltype(REF_TO_X.swap(REF_TO_X))* test(X*);
   static dummy test(...);
   template <typename X>
   static yes_type cast(X*);
   static no_type cast(dummy);
   static const bool value = sizeof( cast(test((T*)0)))==sizeof(yes_type);
};
```

Additionally, the C++11 header <utility> adds a new function equivalent to the macro REF TO X.

In a SFINAE-expression, you may mention a member function call (the previous example reads "the result of REF_TO_X.swap (REF_TO_X)"), so you need an instance of T. However, you cannot simply call a constructor, say as T(), because T may not have a public default constructor. A workaround is to produce a fake reference, such as REF_TO_X, as the expression is not evaluated anyway. But in C++11 you can just use the expression std::declval<T>(). This is safer because, as opposed to macros, it will work only in an unevaluated context.

12.3. Auto

The keyword auto has a new meaning since C++11. It is used to declare a local variable that needs to be initialized immediately. The initialization object is used to deduce the actual type of the variable, exactly as it happens for template parameters:

```
auto i = 0;
```

The actual type of i is the same as the template deduced from the call f(0), where f would be (pseudo-code):

```
template <typename auto>
void f(auto i);
```

auto will always resolve to a value type. In fact, its intended use is to store results coming from a function, without explicitly mentioning their type (think auto i = myMap.begin()). If the user really wants a reference, auto can be explicitly qualified (as any template parameter):

```
const auto& i = cos(0.0);
```

auto will resolve to double, because that's what would happen when calling q(cos(0.0)), with

```
template <typename auto>
void g(const auto& i);
```

Remember that a generic template parameter will not match a reference:

```
int& get_ref();
template <typename T>
void f(T x);
f(get_ref());  // T = int, not reference-to-int
```

On the other hand, decltype returns the exact static type of an expression, as defined:⁵

decltype has a few rules for handling references:

- decltype (variable) or decltype (class member) result in the same declared type as the operand; if x is a double in current scope, decltype (x) is deduced to be double, not double &.
- decltype (function call) is the type of result returned by the function.
- If none of the previous rules is true and if the expression is an lvalue of type T, the result is T&; otherwise, it's T.

In particular, some "bizarre-looking" expressions like decltype (*&x), decltype ((x)), or decltype (true ? x : x) will yield double & because none of the operands is a plain variable, so the third rule prevails.

12.4. Lambdas

}

Lambda expressions ("lambdas" for short) provide a concise way to create function objects on the fly. They are not a new language feature, but rather a new syntax:

```
[](int i) { return i<7; }
[](double x, double y) { return x>y; }
```

std::partition(begin, end, [](int i) { return i<7; });</pre>

Each line represents an *instance* of an object of type "functor" (called *closure*) taking one or more arguments and returning decltype (return statement). So you can pass this object to an algorithm:

```
std::sort(begin, end, [](double x, double y) { return x>y; });
   This is equivalent to the more verbose:
struct LessThan7
{
   bool operator()(int i) const
   {
      return i<7;
   }
};
int main()
{
   std::vector<int> v;
   std::partition(v.begin(), v.end(), LessThan7());
```

The obvious advantages are more clarity (the line that executes the partitioning becomes self-contained) and omission of irrelevant information (as you don't need to find a meaningful name for the functor, nor for its parameters).

The brackets [] are called *lambda introducers*, and they can be used to list local variables that you want "captured," which means added to the functor as members. In the example that follows, the closure gets a copy of \mathbb{N} (the introducer [&N] would pass a reference).

```
int N = 7; std::partition(v.begin(), v.end(), [N](int i) { return i<N; });
```

Again, this lambda is equivalent to the more verbose:

```
class LessThanN
{
  private:
    int N_;

public:
  LessThanN(int N)
  : N_(N)
  {}

  bool operator()(int i) const
  {
    return i<N;
  }
};</pre>
```

[](int i) -> bool { ... }

There are some more syntax details. You can specify the return type explicitly after the argument list. This is indeed useful when you want to return a reference (by default, the return type is an rvalue).

```
Closures can be stored using auto:

auto F = [] (double x, double y) { return cos(x*y); }
```

Finally, a lambda created inside a member function is allowed to capture this; the lambda function call operator will be able to access anything that was available in the original context. In practice, the code of the lambda body works *as if* it were written directly in the place it's declared.

```
read
               };
      return L();
};
```

The following example (due to Stephan T. Lavavej) shows that lambdas can interact with template parameters. Here a lambda is used to perform the logical negation of an unspecified unary predicate.

```
template <typename T, typename Predicate>
void keep if(std::vector<T>& v, Predicate pred)
   auto notpred = [&pred](const T& t) { return !pred(t); };
   v.erase(remove if(v.begin(), v.end(), notpred), v.end());
```

12.5. Initializers

template <typename X>

If a function has a long return type, you may be forced to write it twice—both in the function signature and when building the result. This redundancy is likely to cause maintenance and refactoring problems. Consider the following example from 9.4.2:

```
console assert<X, console assert<T1, T2> > operator() (const X& x) const
  return console assert<X, console assert<T1, T2> >(x, *this);
```

In classic TMP, this is avoided with non-explicit single-argument constructors (when feasible):

```
template <typename T1, typename T2>
class console assert
  public:
     console assert(int = 0) {}
};
template <typename X>
console_assert<X, console_assert<T1, T2> > operator() (const X& x) const
  return 0; // much simpler, but we cannot pass parameters...
```

In C++0x, a new language feature called *braced initializer list* allows you to build an object using curly brackets and (in some cases) to omit the type name:

```
std::pair<const char*, double> f()
{
   return { "hello", 3.14 };
}

template <typename X>
console_assert<X, console_assert<T1, T2> > operator() (const X& x)
const
{
   return { x, *this };
}
```

The compiler will match the items in the initializer list against the arguments of all constructors and pick the best, according to the overload resolution rules.

12.6. Template Typedefs

C++0x extends the traditional typedef syntax with a new using statement:

```
typedef T MyType; // old syntax
using MyType = T; // new syntax
```

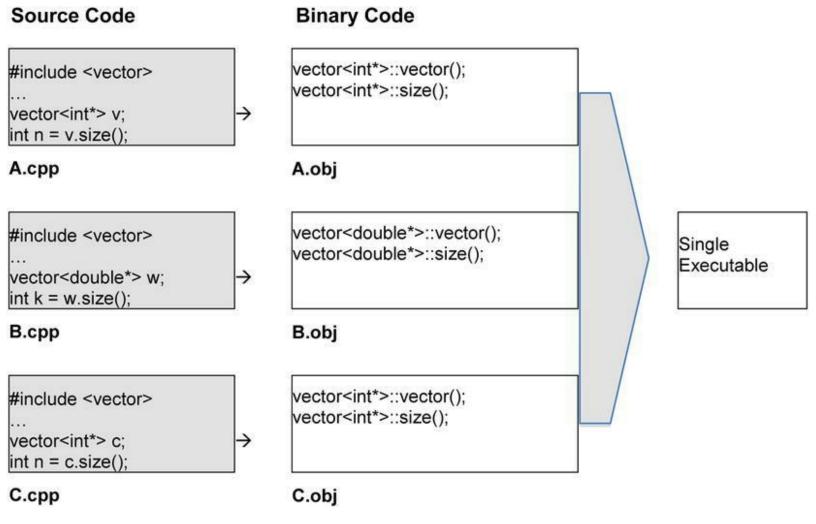
However, the new syntax is also valid with templates:

```
template <typename T>
using MyType = std::map<T, double>; // declares MyType<T>
MyType<string> m; // std::map<string, double>
```

12.7. Extern Template

12.7.1. Linking Templates

In classic C++, the compiler needs to see the entire body of the function/class template, to be able to generate template instantiations. The default behavior is to generate only member functions that are actually used in the translation unit, so roughly speaking, every .cpp file that uses a template class will produce a copy of the code in the corresponding binary object. Finally, the linker will collect all the binary objects and produce a single executable, usually identifying and removing duplicates correctly.



In ordinary code, symbols cannot be defined twice, but template-generated code is marked as "de-duplicable," and the linker in the final step will remove both C++ duplicates (like vector < int > : size(), which was generated twice) and machine-code duplicates. It may detect that all vector < T * > produce the same assembly for every T, so the final executable will contain just *one* copy of each member function.

However, this happens because the vector header contains all the relevant code. Let's write a template class as if it were a plain class (remember that as a rule, this is incorrect).

```
// xyz.h

template <typename T>
class XYZ
{
   public:
       int size() const;
};

// xyz.cpp

template <typename T>
int XYZ<T>::size() const
{
   return 7;
};
```

Now any translation unit that includes xyz. h (and links against xyz. cpp) will be able to compile correctly any code, including:

```
// main.cpp
#include <xyz.h>
int main()
{
    XYZ<int> x;
    return x.size();
}
```

However, the program won't link, because in the translation unit main.cpp the compiler does not see the relevant template bodies. On the other hand, XYZ can be fully used *inside* xyz.cpp:

Now, as a side effect, the binary object xyz.obj will contain the binary code for the relevant member functions that are used (namely, the constructor XYZ::XYZ() and XYZ::size). This implies that main.cpp will now link correctly!

The compiler will verify that main.cpp is syntactically correct. Since it's unable to produce the code in-place, it will mark the symbols as "missing," but the linker will eventually find and borrow them from xyz.cpp.

Needless to say, this works because both files are using XYZ<same type> and the same member functions.

The standard offers a way to force instantiation of a template *and all its member functions*, in a translation unit. This is called *explicit instantiation*.

```
template class XYZ<int>;
```

Namespaces and functions can be used:

```
// assume that we included <vector>
template class std::vector<int>;
// assume that we included this template function:
```

```
// template <typename T>
// void f(T x)

template void f<int>(int x);
```

A possible use is to limit the set of types that the user can plug in to a template:

```
// xyz.cpp

template <typename T>
int XYZ<T>::size() const
{
   return ...;
};

// these are the only types that the user will be able to plug in
// XYZ<T>. otherwise the program won't link.

template class XYZ<int>;
template class XYZ<double>;
template class XYZ<char>;
```

Now this translation unit will contain the binary code for all member functions of XYZ, so they can be correctly "exported" to other units when assembling the final executable.

12.7.2. Extern Template

In C++0x (and as an extension in many classic C++ compilers), it's possible to *prevent* the compiler from instantiating a template automatically and force a behavior like the one described in the last section.

```
extern template class XYZ<int>;
```

This forces the template class to link like an ordinary class (so in particular, inlining is still possible), and it may save compilation time.

According to the C++ standard, this syntax prevents implicit instantiation, but *not* explicit instantiation. So you can in principle put a single extern template declaration in an .hpp file (after the template code), and a single explicit instantiation in a .cpp file.⁷

12.9. Variadic Templates

Since C++11, the list of template arguments can have a variable length:

The ellipsis (...) to the *left* of T declares that T can match a (possibly empty) list of parameters. T is indeed called a *template parameter pack*. On the other hand, an ellipsis to the *right* of an expression involving a parameter pack name expands it (put simply, it clones the expression for every type in the pack):

You can use pattern matching to "iterate" over a parameter pack:

```
0 arguments
template <typename HEAD, typename... TAIL>
void doSomething(HEAD h, TAIL... tail) // will match 1 or more
arguments
    std::cout << h << std::endl;</pre>
   doSomething(tail...);
}
  As an exercise, take a look at this metafunction count<T, A...>, which counts how many
times a type T appears in a pack A:
template <typename T, typename... A>
struct count;
template <typename T, typename... A>
struct count<T, T, A...>
   static const int value = 1 + count<T, A...>::value;
};
template <typename T, typename T2, typename... A>
struct count<T, T2, A...> : count<T, A...>
{ } ;
template <typename T>
struct count<T> : std::integral constant<int, 0>
{ } ;
  The ellipsis can trigger more than one expansion at the same time. Suppose for example that you
want to check that no type in a pack was repeated twice:
template <typename T, typename... A>
int assert()
{
   static assert(count<T, A...>::value <= 1, "error");
   return 0;
}
template <typename... N>
void expand all(N...)
template <typename... A>
void no duplicates(A... a)
```

```
expand_all(assert<A, A...>()...); // double expansion
}
```

This double expansion will call:

```
expand all(assert<A_1, (all A)>(), assert<A_2, (all A)>(), ...)
```

expand_all gets any number of arguments of any type and ignores them entirely. This is necessary to trigger the expansion of the parameter pack. In practice, all assert<...> functions will either fail to compile or return 0, so no_duplicates will be easily inlined and produce almost no code.

As the situation is evolving quickly, refer to online documentation. It's not easy to find a comparison table that is simultaneously complete and up-to-date, but at the time of this writing, good references are

http://wiki.apache.org/stdcxx/C++0xCompilerSupport and http://cpprocks.com/c11-compiler-support-shootout-visual-studio-gcc-clang-intel/.

²As a rule, however, it's acceptable for metafunctions to return a "suboptimal" value. If a class destructor is known to be trivial, then the code may be optimized. A drastic assumption like "no destructor is trivial" will probably make the program slower, but it shouldn't make it wrong.

³They are described in the freely downloadable "Draft Technical Report on C++ Library Extensions" (http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2005/n1836.pdf).

 $^{^4}$ A list of metafunctions that ship with C++11-compliant compilers can be found here: http://en.cppreference.com/w/cpp/header/type traits.

⁵For a detailed explanation of the differences between auto and decltype, see [17].

⁶The compiler will find the appropriate function with the standard overload resolution rules, as if it were a normal call.

⁷Older compilers need not respect this behavior when both directives are present, so some use of the preprocessor may be required.

APPENDIX A

Exercises

A.1. Exercises

In all the following problems, the reader should assume that a single writable file with some template code is given; more files can be added, and the rest of the project is read-only.

A.1.1. Extension

A function template is given:

```
template <typename T>
void f(T x)
{
   printf("hello T");
}
```

- Add another overload that is to be called for every class that derives from BASE and prints "hello BASE-or-derived"
- Ensure that your solution is robust. Change the return type of f to int and see if your solution still holds
- Ensure that your solution is robust. Add a plain function say int f(double x) in the same file and see if compilation fails
- Think of an alternative solution that minimizes the changes to the existing code.

A.1.2. Integer

The following code:

```
template <typename T>
uint32 f(T x) { ... }
```

```
// ... printf("%x", f(a));
```

is emitting a warning: return type of f is incompatible with %x.

What kind of investigation would you perform?

A.1.3. Date Format

Following Section 7.7.1, implement a constant generator with an even more natural syntax, such as:

A.1.4. Specialization

A template class is given:

```
template <typename T>
class X
{ /* very long implementation... */ };
```

Modify X so that X<double> has precisely one additional data member (say, int) and one extra member function. Perform minimal changes to existing code (so if the source file is under version control software, the differences are self-explanatory).

A.1.5. Bit Counting

The code below:

```
template <> struct nb<2> { static const size_t value = 1; };
template <> struct nb<3> { static const size_t value = 2; };
template <> struct nb<4> { static const size_t value = 1; };
template <> struct nb<5> { static const size_t value = 2; };
template <> struct nb<6> { static const size_t value = 2; };
template <> struct nb<6> { static const size_t value = 2; };
```

- Is completely correct and it shows a new technique not seen previously in this book (or in any other equivalent book)
- Has a trivial bug, but the technique is comparable to Section 3.6.6 and thereafter
- Has at least one nontrivial bug, which cannot be easily fixed

A.1.6. Prime Numbers

As an exercise for debugging techniques, we present an example of a non-trivial metafunctionis prime<N>::value.

The reader is expected to be able to understand the code, at least in principle, even if some of the algorithm details are not known.

```
#define mxt EXPLICIT VALUE(CLASS, TPAR, VALUE)
template <> struct CLASS<TPAR> { static const size t value
= VALUE; }
template <size t N>
struct wheel prime;
mxt EXPLICIT VALUE(wheel prime, 0, 7);
mxt EXPLICIT VALUE (wheel prime, 1, 11);
mxt EXPLICIT VALUE (wheel prime, 2, 13);
mxt EXPLICIT VALUE (wheel prime, 3, 17);
mxt EXPLICIT VALUE (wheel prime, 4, 19);
mxt EXPLICIT VALUE (wheel prime, 5, 23);
mxt EXPLICIT VALUE (wheel prime, 6, 29);
mxt EXPLICIT VALUE(wheel prime, 7, 31);
template <size t A>
struct nth tentative prime
  static const size t value
     = 30*((A-3)/8) + wheel prime<(A-3) % 8>::value;
} ;
mxt EXPLICIT VALUE(nth tentative prime, 0, 2);
mxt EXPLICIT VALUE (nth tentative prime, 1, 3);
mxt EXPLICIT VALUE (nth tentative prime, 2, 5);
```

```
template
  size t A,
  size t N,
  size t K = nth tentative prime<N>::value,
  size t M = (A \% K)
>
struct is prime helper
  static const bool EXIT = (A < MXT M SQ(K));
  static const size t next A = (EXIT ? 0 : A);
  static const size t next N = (EXIT ? 1 : N+1);
};
template <size t A, size t N, size t K>
struct is_prime_helper<A, N, K, 0>
  static const size t next A = 0;
  static const size t next N = 0;
};
template \langle \text{size t A}, \text{ size t N} = 0 \rangle
struct is prime
: is prime<is prime helper<A, N>::next A,
          is prime helper<A, N>::next N>
};
template <> struct is prime<0,0> { static const bool value
= false; };
template <> struct is prime<0,1> { static const bool value
= true;
         } ;
template <> struct is prime<1,0> { static const bool value
= true; };
template <> struct is prime<2,0> { static const bool value
```

A.1.7. Typeinfo without RTTI

= true; };

The typeinfo wrapper in Section 5.3.2 relies on the compiler to generate a runtime identifier for different types. If this is not available, then a different implementation can be used (at least in some cases):

- Create a traits class TI<T> having a single static member function T f() that returns T()
- Use a reinterpret_cast and convert &TI<T>::f to void (*)()

- Use this latter pointer as index in a std::map
- Prove that this works (Hint: step #1 is necessary because of ICF, see page 354; for step #3, See Section 20.3.3 of the Standard)
- Note that pointer-based type identifiers work with static types, while typeinfo uses dynamic types, so this technique in general is weaker.

A.1.8. Hints and Partial Solutions

We give a solution to Exercise #1, because of its practical importance.

Obviously overload alone doesn't work. We can select BASE but a DERIVED will prefer the template function (with T=DERIVED it's an exact match).

```
template <typename T>
void f(T x)
{
  printf("hello T");
}

void f(BASE& x)
{
  printf("hello BASE");
}
```

Instead we introduce another layer:

```
template <typename T>
void f(T x)
{
    g(&x, &x);
}

template <typename T>
void g(T* p, void*)
{
    printf("hello T");
}

template <typename T>
void g(T* p, BASE*)
{
    printf("hello BASE-OR-DERIVED");
}
```

Conversion of T* to BASE* is preferred.

Note that the same technique solves also another problem; when invoking a member function on the argument, we can prevent virtual calls:

```
template <typename T>
void g(T* p, void*)
  printf("Not a BASE. No call was made.");
template <typename T>
void g(T* p, BASE* b)
  b->doit(); // always virtual, if BASE::doit is virtual
  p->doit();  // may be virtual or not
  p->T::doit(); // always non-virtual
  Observe that in principle T may hide BASE::doit, so the second call won't be virtual:
class BASE
{
public:
  virtual void doit();
};
class D1 : public BASE
public:
  void doit(int i = 0);
};
class D2 : public D1
{
public:
   virtual void doit();
};
```

APPENDIX B

Bibliography

- [1] Alexandrescu A., "Modern C++ Design", Addison-Wesley
- [2] Vandevoorde, D. and Josuttis, N., "C++ Templates: The Complete Guide", Addison-Wesley
- [3] Abrahams D. and Gurtovoy A., "C++ Template Metaprogramming", Addison-Wesley
- [4] Sutter H., "Exceptional C++ Style", Addison-Wesley
- [5] Wilson, "Imperfect C++", Addison-Wesley
- [6] Austern M., "Generic Programming and the STL", Addison-Wesley
- [7] Cline M., "C++ FAQ (lite)", http://www.cs.rit.edu/~mjh/docs/c++-faq/
- [8] Meyers S., "Effective STL"
- [9] Coplien, J., "Curiously Recurring Template Patterns", C++ Report, February 1995, pp. 24-27.
- [10] Stroustrup, B., "Design and Evolution of C++", Addison-Wesley, Reading, MA, 1993.
- [11] Barton, J.J. and Nackman L.R., "Scientific and Engineering C++", Addison-Wesley, Reading, MA, 1994.
- [12] Veldhuizen, T. "Expression Templates", C++ Report, June 1995, reprinted in C++ *Gems*, edited by Stanley Lippman.
- [13] Myers Nathan C., "A New and Useful Template Technique: Traits", C++ Report, June 1995, http://www.cantrip.org/traits.html
- [14] C++ 0x (C++11) Final Committee Draft: C++11 - ISO/IEC 14882:2011: \$60 from ansi.org
- [15] Meyers S., Effective Modern C++, O'Reilly 2014
- [16] Meucci, A. "Risk and Asset Allocation", Springer 2005

Index

```
Accessors
Accumulation
    accessor
    binary operations
    collect fuction
    elementary operations
    global constant objects
    global helper function
    i-th operation
    multi-layered accumulators
    operator*
    operator+
    op_void
    output
    recursion-stopping specialization
    runtime error
    static recursion
    template parameters
    template rotation
    types
Agglomeration
Algorithms
    accessors
    algebraic requirements
    algorithm I/O
    Barton-Nackman trick
    classification
    *i returns
    iterator (see Iterators)
    mimesis
    properties
    range begin...end process
    receipts
    reordering algorithms
    selective copying algorithms
    set partition
    swap-based/copy-based
Angle brackets
Argument pack technique
Argument rotation
Artistic Style (AStyle)
```

B

Barton-Nackman trick
Base copy constructor

```
Bitstring
Body level
Boolean type
Buffer type algor
C++0x
    auto
    decltype
    extern template
    initializers
    lambda expressions
    linking templates
    metafunctions
    typedef syntax
    variadic templates
Chain destructor
Chain starter
Classic patterns
    action(range)
    action(value)
    boolean type
    bool T::empty() const
    default and value initialization
    literal zero
    manipulators
    operators position
    ptrdiff t
    secret inheritance
    size t
    void T::clear()
    void T::property(X)
    void T::swap(T&)
    X T::base() const
    X T::get() const
    X T::property() const
Class level
Class template
Code generators
    double checked stop
        arithmetic/floating point operations
        integrize template
        loop unrolling
        vector-sum
    enumeration types
    If-Less code (see If-Less code)
    Nth Minimum
    static (see Static code generators)
    static and dynamic hashing (see Static and dynamic hashing)
    template factory pattern
Code safety
Comma chains
CompareAndSwap
Compile-time polymorphism
Conversion functions
```

Bit counting

```
C++ templates
    angle brackets
    class template
    compiler error
    compile-time constants
    division by zero
    errors
    explicit specialization
    function template
    function types and function pointers
    integer literals
    integer operation
    metafunction
    non-standard language operators
    non-template base classes
    non-type parameters
    parameter list
    position
    sizeof
    sq<double>
    static constants
    template
        arguments
        member functions
        parameters
    typename
    type parameters
    types and non-types integers and pointers
    universal constructors
Curiously recurring template pattern (CRTP)
    D
Date format
Debugging templates
    compiler bugs
    integer computations
        references to numeric constants
        signed and unsigned types
    portability
    SFINAE principle
    trampolines
    types
        incomplete types
        tag global variables
```

E

Drivers

Enum conversion Exponential unrolling Extension

Discriminated unions

trapping types type_info::name

```
fgetpos
Forwarding/reference wrappers
fsetpos
Function pointers
Function template
Function types
Functors
    accumulation (see Accumulation)
    buffer_type algor
    composition tools
        CRTP
        helper function
        lambda predicate
        person/age relationship
        template parameter
        template-template parameter
    drivers
    forwarding/reference wrappers
    inner template
        conversion functions
        members conversion (see Members conversion)
    interactive algor
    self-accumulator algor
    static member function
    strong/weak
get ptr functions
    H
has abs method
Heuristic algorithms
Hints and partial solutions
    I, J, K
Identical code folding (ICF)
If-Less code
    enum conversion
    self-modifying function
    smart constants
Implicit promotion techniques
Inline function
Inner class templates
Input iterator
Integer
Interactive algor
Interfaces
    code unification
    compile-time polymorphism
    copy constructor
    dereferenceable entity
```

```
emptiness
    global functions
    input iterator and output iterator
    member functions
    minimal set of concepts
    object returned by operator
    static
        ambiguous inheritance diagram
        base class
        BASE<T> object
        bitstring class
        clone of metafunction
        common errors
        compile-time steps
        CRTP
        declaration
        derived class
        grouping ideas
        improved inheritance diagram
        macro
        member selection
        memberspace problem
        PrintSomeNumber function
        random algorithms
        reference-to-base
        static cast
        summable<...> interface
        template
        "virtual" callback mechanism
    type hiding
        *begin parameter
        boundary crossing with trampolines
        generic t
        iterator handling
        metafunctions
        option_map container
        option parser container
        parameterless options
        storing object
        trampolines
        typeinfo wrapper
        variant objects
        virtual function
    variant
        opaque object
        parameter deletion with virtual calls
        with visitors
    wrapping
        containers
        references
Is PRIME<N> value
Iterators
    definition
    identification
    iterator t
    metafunction
    requirements
    set partition
```

```
sort algorithm
    value type
    wrapping
        const iterators
        expander
        fake pairs
        iterator_expander
        multiplier_iterator
        operator->
        pair iterator
    T
Lambda-constant
```

Lambda expressions argument and result deduction arithmetic operators arrays assignments binary operation binary operators binary predicates concrete binary functions cos(X+2.0)deducing argument type deducing result type elementary lambda object enum lambda tag error log global/static member functions lambda-const reference logical and standard operators refinements static cast unary function unary operators lambda reference Lambda semantics lambda wrap Less and NaN Literal zero Logger

M

Macro expansion rules Members conversion enum const get < X > (0)nested class template pointer-to-member **PTR** wrap Memberspace problem Metafunctions Mimesis interface Mimesis techniques

Multi-layered accumulators Multi-layer template factory

N

Namespace level Non-inline member functions Non-mutating function Non-template base classes Nth Minimum

```
Opaque type principle
    bitstring
    comma
    comma chains
    growing object concept
        chain destruction
        chain traits
        console assert
        definition
        direct access
        inward link
        outward link
        proxy
        pseudo-template notation
        SMART ASSERT
        string concatenation
        string sum
        variations
    infix simulation
    lambda expressions (see Lambda expressions)
    operator()
    operator[]
    polymorphic result
    range insertions
operator()
Output iterator
Overload resolution
    concept traits
    fake_incrementable
    function pointers
        argument dominance
        erase
        swap
    groups
        auxiliary functions
        clustered implementation
        companion global function template
        default template implementation
        definition
        "double-layer" template
        has abs method
        implicit argument promotion
        maths < double >
        non-template function
```

```
operator%
    runtime decay
    single action
    template struct
    typedefs
has conversion<L,R>::L2R
implicit promotion techniques
merging traits
    alternative blocks
    binary relation traits
    code rearrangement
    derivation and hidden template parameter
    flags idiom
    "native" operators
    static highest bit<N>::value
namespace
platform-specific traits
SFINAE
    advantage
    A<int>
    limitations
   multiple decisions
    Only If
    partial specializations
    pointer
   returned functors
    sizeof
    software updates
    workarounds
string function
    constant-time
    const char*
    semi-opaque interface
    std::string
   strlen
```

■ P, Q

Partial specialization

Placeholder
Pointer paradox
Pointer-to-function type
Preprocessor
include guards
macro expansion rules
PrintSomeNumber function
Push back

\mathbf{R}

Refactoring
 accessors
 compatibility
 encapsulation
 interface adaptation
 kernel macros
 template sq

```
ghost
interfaces
placeholder
switch-off
template parameters
trampolines
Reordering minmax algorithm
```

S

```
Secret inheritance
Self-accumulator algor
SFINAE principle
Skeleton technique
Small object toolkit
    hollow types
        constraints
        instance of
        selector
        static value
    static assertions
        assert legal
        Boolean assertions
        overloaded operators
        pointers
    tagging techniques
        function pointers
        inheritance
        overload mechanisms
        program compiles
        runtime argument
        similar signature
        static useless argument
        tag iteration
        tags type
sorted_vector
Specialization
sq function
Static and dynamic hashing
    algorithm I/O
    ambiguous overload
    changing case
    character set functions
    classic code
    macros
    mimesis interface
    mimesis techniques
    template rotation
Static code generators
    static pow
    static_raise
    strongly typed template
    struct deduce
Static interfaces
    ambiguous inheritance diagram
    base class
    BASE<T> object
```

```
bitstring class
    clone of metafunction
    common errors
    compile-time steps
    CRTP
    declaration
    derived class
    grouping ideas
    improved inheritance diagram
    macro
    member selection
    memberspace problem
    PrintSomeNumber function
    random algorithms
    reference-to-base
    static cast
    summable<...> interface
    template
Static member function
Static programming
    compilation complexity
        auxiliary class
        full recursion
        linear complexity
        low-complexity
        MXT_M_SQ
        template instances
    hidden template parameters
        companion technique
        disambiguation
        primary template
        static recursion
    metaprogramming idioms
        header files
        random elements
        static short circuit
        struct
        variables
    preprocessor
    traits (see Traits)
    type containers
        agglomerates
        arrays
        conversions
        depth
        empty/void
        find
        front and back
        metafunctors
        push and pop
        returning error
        template rotation
        typeat
        typelist
        typepair
STL containers
Strong/weak functors
Style
```

```
AStyle plugin
    comments
    elements
    generality
    macros
        code appearance
        constexpr
        infix _M_
        integer functions
        intercepting
        lowercase prefix mxt
        macro directive
        MXT
        namespace directives
    metafunctions
    multidimensional_vector::empty
    namespaces and using declarations
    naming convention
    symbols
    template parameters
Substitution failure is not an error (SFINAE)
    advantage
    A<int>
    limitations
    multiple decisions
    Only If
    partial specializations
    pointer
    returned functors
    sizeof
    software updates
    workarounds
Switch-off
    T
Tamper
Templates
    argument deduction
        automatic deduction
        built-in C++ cast operator
        cast and function template invocation
        class template
        disambiguation
        dummy argument
        function templates
        non-deduced parameters
        non-deducible
        non-type arguments
        obscure error messages
        overload resolution
        pattern matching
        template-dependent
    C++ (see C++ templates)
    class and typename
    classic patterns (see Classic patterns)
    code safety
```

```
compiler assumptions
        argument names
        bugs
        catalog
        dirty code modifications
        empty destructor addition
        error messages
        inline function
        language features
        language-neutral idioms
        memory-helping tricks
        non-standard behavior
        standard-conforming compiler
        warning, default level
    function template
    global object
    inner class
    parameter
    position
    preprocessor (see Preprocessor)
    purpose of TMP
    rotation
    self-adapting
    specialization
        class templates
        compiler error
        forward declaration
        full specialization
        global function template
        namespace level
        non-deducible types
        and overload
        partial specialization
    squaring numeric type
    style conventions (see Style)
    template name
    template-template parameters
    transforming sequence
    type
    visual equivalence
Tentative interface
Traits
    concept
    definition
    ios char traits
    largest unsigned integer
    merging
        alternative blocks
        binary relation traits
        code rearrangement
        derivation and hidden template parameter
        flags idiom
        "native" operators
        static highest bit<N>::value
    platform-specific
    streambuf template
    string function
```

```
constant-time
        const char*
        semi-opaque interface
        std::string
        strlen
    type dismantling
    type traits
        argument_type
        assignable
        base class
        codes implementation
        const objects
        definition
        DER
        immutable
        inheritance
        intrusive const-ness
        operator=
        specializations
        static constants
Trampolines
    boundary crossing with
    virtual function
typeinfo wrapper
typename
```

U

Universal assignments
Universal constructors
Universal copy constructors
Using-namespace declarations

■ V, W, X, Y, Z

Variadic templates Virtual function table pointer Virtual inheritance